



- Expert Verified, Online, **Free**.

⚙ Custom View Settings

Topic 1 - Single Topic



Every employee of your company has a Google account. Your operational team needs to manage a large number of instances on Compute Engine. Each member of this team needs only administrative access to the servers. Your security team wants to ensure that the deployment of credentials is operationally efficient and must be able to determine who accessed a given instance. What should you do?



- A. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key in the metadata of each instance.
- B. Ask each member of the team to generate a new SSH key pair and to send you their public key. Use a configuration management tool to deploy those keys on each instance.
- C. Ask each member of the team to generate a new SSH key pair and to add the public key to their Google account. Grant the `compute.osAdminLogin` role to the Google group corresponding to this team.
- D. Generate a new SSH key pair. Give the private key to each member of your team. Configure the public key as a project-wide public SSH key in your Cloud Platform project and allow project-wide public SSH keys on each instance.



**Correct Answer:** D



Reference:



<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>



-   **dan80** Highly Voted 1 year, 1 month ago



C is correct - <https://cloud.google.com/compute/docs/instances/managing-instance-access>  
upvoted 20 times
-   **adedj99** 9 months, 1 week ago



We recommend collecting users with the same responsibilities into groups and assigning IAM roles to the groups rather than to individual users. For example, you can create a "data scientist" group and assign appropriate roles to enable interaction with BigQuery and Cloud Storage. When a new data scientist joins your team, you can simply add them to the group and they will inherit the defined permissions. You can create and manage groups through the Admin Console.  
upvoted 2 times
-   **zakhili** Highly Voted 1 year, 1 month ago



Send private key to users is not safe, i think it's C  
upvoted 9 times
-   **shayanahmed** Most Recent 3 weeks, 4 days ago



All these questions are still valid. Yesterday i cleared my paper. I have all correct answers as well (shayan18@live.com)  
upvoted 3 times
-   **vmart** 1 month, 1 week ago

c is correct  
upvoted 1 times
-   **KiranChinna** 1 month, 2 weeks ago

C is correct  
upvoted 1 times
-   **Finger41** 2 months, 2 weeks ago

C - Need compute.osAdminLogin role to perform this role.  
upvoted 1 times
-   **mcaromit** 2 months, 2 weeks ago

C is correct  
upvoted 2 times
-   **viswanand** 2 months, 3 weeks ago

C is correct as it is recommended approach from google to use groups  
upvoted 1 times
-   **kopper2019** 3 months, 2 weeks ago


Answers QUESTION 179  
A. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN.  
2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.  
B. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN.  
2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.  
C. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN.  
2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for

the DMZ.

D. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN.

2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.

upvoted 4 times

  **Ash543211** 2 months, 1 week ago

Hi, Can you send all the remaining question to me .

upvoted 1 times

  **mdc** 3 months, 1 week ago

what is the answer to this? is it C?

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Hey Kopper2019 , can you post all the questions to me?

upvoted 1 times

  **AmineHM** 3 months, 1 week ago

Do you have another questions ?

upvoted 2 times

  **kopper2019** 3 months, 2 weeks ago

Answer is C, since OsLogin is needed

upvoted 1 times

  **kopper2019** 3 months, 2 weeks ago

QUESTION 177

You are creating an application that will run on Google Kubernetes Engine. You have identified MongoDB as the most suitable database system for your application and want to deploy a managed MongoDB environment that provides a support SLA. What should you do?

A. Create a Cloud Bigtable cluster, and use the HBase API.

B. Deploy MongoDB Atlas from the Google Cloud Marketplace.

C. Download a MongoDB installation package, and run it on Compute Engine instances.



D. Download a MongoDB installation package, and run it on a Managed Instance Group

upvoted 2 times

  **mathurshagun17** 1 month, 3 weeks ago

B. Deploy from Marketplace.

upvoted 1 times

  **Ashii** 2 months, 1 week ago

Mongo is available in Marketplace, its B

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago


Simple, it's D.

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Sorry, its B.

upvoted 4 times

  **kopper2019** 3 months, 2 weeks ago

QUESTION 176

You will have several applications running on different Compute Engine instances in the same project. You want to specify at a more granular level the service account each instance uses when calling Google Cloud APIs. What should you do?

A. When creating the instances, specify a Service Account for each instance.

B. When creating the instances, assign the name of each Service Account as instance metadata.

C. After starting the instances, use gcloud compute instances update to specify a Service Account for each instance.

D. After starting the instances, use gcloud compute instances update to assign the name of the relevant Service Account as instance metadata.

upvoted 3 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

It's A, as In the Identity and API Access section, choose the service account you want to use from the drop-down list.

upvoted 5 times

  **kopper2019** 3 months, 2 weeks ago

QUESTION 175

You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

A. Create and deploy a Custom Resource Definition per microservice.

B. Create and deploy a Docker Compose File.

- C. Create and deploy a Job per microservice.  
 D. Create and deploy a Deployment per microservice.  
 upvoted 4 times

  **meh009** 3 months, 1 week ago

D:  
<https://cloud.google.com/architecture/scaling-commerce-workloads-architecture>  
 upvoted 2 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

When you use GKE, each microservice has a separate development and deployment lifecycle. Each microservice is packaged as a Docker container. You deploy those containers as a Kubernetes Pod and Service using one of these ways:

imperative commands (using kubectl)  
 declarative specifications (Kubernetes Deployments)  
 packaged as Helm charts

So, we can stick with D. but It can be B too as docker-compose is a tool for defining and running multi-container Docker applications. With Compose, you use a YAML file to configure your application's services. Then, with a single command, you create and start all the services from your configuration.

upvoted 2 times

  **EricJason** 2 months, 2 weeks ago

It's all about K8s, so just forget the docker-compose  
 upvoted 1 times

  **kopper2019** 3 months, 2 weeks ago

QUESTION 174



You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.  
 B. Provision Compute Engine instances with GPUs attached.  
 C. Provision Compute Engine instances with local SSDs attached.  
 D. Provision Compute Engine instances with M1 machine type.

upvoted 3 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

this is D, as M1 machine types are specifically for in-memory databases and operations.  
 upvoted 5 times



  **kopper2019** 3 months, 2 weeks ago

NEW Q, 173

You are developing a financial trading application that will be used globally. Data is stored and queried using a relational structure, and clients from all over the world should get the exact identical state of the data. The application will be deployed in multiple regions to provide the lowest latency to end users. You need to select a storage option for the application data while minimizing latency. What should you do?

- A. Use Cloud Bigtable for data storage.  
 B. Use Cloud SQL for data storage.  
 C. Use Cloud Spanner for data storage.  
 D. Use Firestore for data storage.

upvoted 2 times

  **dude11** 3 months, 2 weeks ago

C. Cloud spanner (it says relational, all over the world)  
 upvoted 3 times

  **kranti1020** 3 months, 3 weeks ago

Its not advisable to share keys to everyone and best answer is C  
 upvoted 1 times

  **[Removed]** 4 months, 1 week ago

C is correct. Ask each member of the team to generate a new SSH key pair and to add the public key to their Google account. Grant the "compute.osAdminLogin" role to the Google group corresponding to this team.

upvoted 1 times

You need to create a custom VPC with a single subnet. The subnet's range must be as large as possible. Which range should you use?

- A. 0.0.0.0/0
- B. 10.0.0.0/8
- C. 172.16.0.0/12
- D. 192.168.0.0/16

Correct Answer: A

- CarlS

Highly Voted

1 year, 3 months ago

I got this question in one the trainings I did in udemy: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?referralCode=810D02D4A159FC3E36CC>

B is correct. Pay attention to the question, is talking about custom VPC subnet and is not mentioning you will use automatic subnet mode creation. If you set subnet to custom, the minimum size is /8.

upvoted 26 times
- passnow

Highly Voted

1 year, 4 months ago

B is legit as fuck!

upvoted 12 times
- hems4all

8 months, 4 weeks ago

B is correct

Use 10.0.0.0/8 CIDR range. is the right answer.

The private network range is defined by IETF (Ref: <https://tools.ietf.org/html/rfc1918>) and adhered to by all cloud providers. The supported internal IP Address ranges are

1. 24-bit block 10.0.0.0/8 (16777216 IP Addresses)

2. 20-bit block 172.16.0.0/12 (1048576 IP Addresses)

3. 16-bit block 192.168.0.0/16 (65536 IP Addresses)

10.0.0.0/8 gives you the most extensive range - 16777216 IP Addresses.

upvoted 2 times
- Rafael991

Most Recent

2 weeks, 5 days ago

B is correct

upvoted 1 times
- kashi123

2 weeks, 6 days ago

B is correct

upvoted 1 times
- tanito83

1 month, 3 weeks ago

The correct answer is B. Please, modify it.

upvoted 2 times
- Finger41

2 months, 2 weeks ago

Answer is B, range is /8 to /32.

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct

upvoted 1 times
- viswanand

2 months, 3 weeks ago

Option A gives invalid even though it allow to provision but it fails eventually stating 0.0.0.0/0 invalid

upvoted 1 times
- SaPa2021

3 months ago

B is correct answer ..10.0.0.0/8 CIDR range can be maximum in VPC subnet.

upvoted 1 times

  **meh009** 3 months, 1 week ago

B: Largest Private IP Range defined by IETF

upvoted 1 times

  **kopper2019** 3 months, 2 weeks ago

B, [https://cloud.google.com/vpc/docs/vpc#manually\\_created\\_subnet\\_ip\\_ranges](https://cloud.google.com/vpc/docs/vpc#manually_created_subnet_ip_ranges)

upvoted 1 times

  **kranti1020** 3 months, 3 weeks ago

0.0.0.0/0 is not custom VPC and hence answer is B

upvoted 2 times

  **Pawar\_Shah** 4 months ago

the correct answer is b.

valid ranges :



10.0.0.0/8

172.16.0.0/12

192.168.0.0/16



so 10.0.0.0/8 is have most ip ranges.

upvoted 2 times

  **[Removed]** 4 months, 1 week ago

B is correct. 10.0.0.0/8

upvoted 1 times

  **rgalfaro** 4 months, 1 week ago

The correct answer is B. It can be checked under <https://cloud.google.com/vpc/docs/vpc>, look for the "valid ranges" section, it's all in there.

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

Answer is B

upvoted 1 times

  **JackGlemins** 5 months ago

B is right: [https://cloud.google.com/vpc/docs/vpc#manually\\_created\\_subnet\\_ip\\_ranges](https://cloud.google.com/vpc/docs/vpc#manually_created_subnet_ip_ranges)

upvoted 1 times

Question #3

Topic 1

You want to select and configure a cost-effective solution for relational data on Google Cloud Platform. You are working with a small set of operational data in one geographic location. You need to support point-in-time recovery. What should you do?

- A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.
- B. Select Cloud SQL (MySQL). Select the create failover replicas option.
- C. Select Cloud Spanner. Set up your instance with 2 nodes.
- D. Select Cloud Spanner. Set up your instance as multi-regional.

**Correct Answer: A**

Reference:

<https://cloud.google.com/sql/docs/mysql/backup-recovery/restore>

- YashBindlish**

Highly Voted

1 year, 2 months ago

A is Correct. You must enable binary logging to use point-in-time recovery. Enabling binary logging causes a slight reduction in write performance.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 19 times
- Bharathy**

Highly Voted

1 year, 4 months ago

A is correct, as Binary Logging enables Point in Recovery in Cloud SQL

upvoted 9 times
- Israel**

Most Recent

4 days, 8 hours ago

Before you begin

Before completing this task, you must have:

Binary logging and backups enabled for the instance, with continuous binary logs since the last backup before the event you want to recover from. For more information, see Enabling binary logging.

upvoted 1 times
- YAS007**

1 week, 2 days ago

A is correct :

<https://cloud.google.com/sql/docs/mysql/backup-recovery/pitr>

upvoted 1 times
- curious\_kitty9**

2 months, 1 week ago

Your Company is planning to migrate all Java web applications to Google App Engine. However, you still want to continue using your on-premise database. How can you set up the app engine to communicate with your on-premise database while minimizing effort?

A. Setup the application using App Engine Flexible environment with Cloud Router to connect to an on-premise database.

B. Setup the application using App Engine Standard environment with Cloud VPN to connect to an on-premise database.

C. Setup the application using App Engine Standard environment with Cloud Router to connect to an on-premise database.

D.Setup the application using App Engine Flexible environment with Cloud VPN to connect to an on-premise database.

upvoted 1 times
- nicodishanth**

2 months, 1 week ago

D Is the right one

upvoted 2 times
- Finger41**

2 months, 2 weeks ago

A is correct. need to enable binary logging.

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

A is correct

upvoted 1 times
- Pawar\_Shah**

4 months ago

Here the solution should be cost-effective so mysql and binary logging enable option is for the recovery point objective. A is correct.

upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.

upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

A is correct

upvoted 1 times

  **ozhai** 4 months, 3 weeks ago

A is correct.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago



A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.

upvoted 2 times

  **INASR** 6 months ago



A is Correct. You must enable binary logging to use point-in-time recovery

upvoted 1 times

  **swatititame** 9 months, 1 week ago

A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.

upvoted 1 times

  **glam** 9 months, 3 weeks ago

A. Select Cloud SQL (MySQL). Verify that the enable binary logging option is selected.

upvoted 4 times

  **dragosup** 1 year ago

A is correct

upvoted 1 times

  **qafro** 1 year ago

A is my solution

upvoted 2 times



## Question #4

## Topic 1

You want to configure autohealing for network load balancing for a group of Compute Engine instances that run in multiple zones, using the fewest possible steps.

You need to configure re-creation of VMs if they are unresponsive after 3 attempts of 10 seconds each. What should you do?

- A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)
- B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.
- C. Create a managed instance group. Set the Autohealing health check to healthy (HTTP)**
- D. Create a managed instance group. Verify that the autoscaling setting is on.

**Correct Answer: D**

 **ReyBan** Highly Voted 1 year, 2 months ago

C, Agreed

reference : <https://cloud.google.com/compute/docs/tutorials/high-availability-autohealing>

Pro Tip: Use separate health checks for load balancing and for autohealing. Health checks for load balancing detect unresponsive instances and direct traffic away from them. Health checks for autohealing detect and recreate failed instances, so they should be less aggressive than load balancing health checks. Using the same health check for these services would remove the distinction between unresponsive instances and failed instances, causing unnecessary latency and unavailability for your users.

upvoted 33 times

 **bryanchew** Highly Voted 1 year, 3 months ago

A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy(HTTP)

This is a possible answer. This answer assumes that the existing backend is configured correctly.

B. Create an HTTP load balancer with a backend configuration that references an existing instance group. Define a balancing mode and set the maximum RPS to 10.

This is a possible answer. This answer assumes that the existing backend is configured correctly. This answer adds an additional step over answer A.

C. Create a managed instance group. Set the Autohealing health check to healthy(HTTP)

This is only a partial solution. The default configuration is auto scaling enabled. You still need to create the HTTP Load Balancer.

D. Create a managed instance group. Verify that the auto scaling setting is on.

This is only a partial solution. Creating a Managed Instance Group with Auto Scaling is required, but you still need to create the HTTP Load Balancer.

Therefore the best answer is A in my opinion.

upvoted 16 times

 **pYWORLD** 2 days, 2 hours ago

I hope you saw that in the questions is stated: "network load balancing"

upvoted 1 times

 **DickDastardly** 4 months, 3 weeks ago

It can't be A as a load balancer does not re-create unhealthy instances, as per the requirement.

Has to be C

<https://cloud.google.com/compute/docs/instance-groups>

upvoted 4 times

 **koniec** 8 months, 3 weeks ago

It's A.

Managed group already exists so create a LB with health checks.

If you go for C, you will have to create a LB anyway so it's more steps to achieve the goal

upvoted 1 times

 **tavva\_prudhvi** 3 months, 3 weeks ago

[https://www.youtube.com/watch?v=dT7xDEtALPQ&list=PLIivdWyY5sqlij\\_cgINUHZDMnGjVx3rxi&index=36](https://www.youtube.com/watch?v=dT7xDEtALPQ&list=PLIivdWyY5sqlij_cgINUHZDMnGjVx3rxi&index=36)

step-1: go to the instance group

step-2: click edit

step-3: scroll down you will see auto-healing off by default change to ON

step-4: create a health check saying 10 seconds as CHECK INTERVAL and UNHEALTHY THRESHOLD=3

upvoted 2 times

  **shubjha** Most Recent 2 weeks, 6 days ago

C is the right answer, While creating the MIG, we need to set Auto-healing setup with HEALTH CHECK policy.

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

C is the correct option as autohealing would recreate an unhealthy instance unlike a health-check configured with a lb

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

C is the correct option as using a health-check with a lb (option A) would only probe if the backend instances are receiving traffic, but won't recreate unhealthy instances which is a specific ask in the question.

upvoted 1 times

  **viswanand** 2 months, 3 weeks ago

It has to be C only

upvoted 1 times

  **ri\_unhou119** 2 months, 3 weeks ago

C

<https://cloud.google.com/compute/docs/instance-groups>

upvoted 1 times

  **ozeross** 4 months ago

Correct answer is C. ltes explained in the Compute engine section on Google documentation

upvoted 1 times

  **tavva\_prudhvi** 4 months, 1 week ago

I think it's C,

To configure the recreation of VMs, you need Autohealing. So not B and D.

A: Load balancing health checks help direct traffic away from non-responsive instances and toward healthy instances; these health checks do not cause Compute Engine to recreate instances.

C: Application-based autohealing improves application availability by relying on a health checking signal that detects application-specific issues such as freezing, crashing, or overloading. If a health check determines that an application has failed on an instance, the group automatically recreates that instance.

So the answer is C.

upvoted 1 times

  **tavva\_prudhvi** 4 months ago

Changing my answer to A. This link states that LB detect and automatically removes unhealthy VM instances using health checks. Instances that become healthy again are automatically re-added.

<https://cloud.google.com/compute/docs/load-balancing-and-autoscaling>

upvoted 1 times

  **tavva\_prudhvi** 3 months, 3 weeks ago

[https://www.youtube.com/watch?v=dT7xDEtALPQ&list=PLlivdWyY5sqlij\\_cgINUHZDMnGjVx3rxi&index=36](https://www.youtube.com/watch?v=dT7xDEtALPQ&list=PLlivdWyY5sqlij_cgINUHZDMnGjVx3rxi&index=36)

I request everyone to watch this video, as in the question we have a few conditions like,

the re-creation of VMs if they are unresponsive after 3 attempts of 10 seconds each. we can configure in the instanced group by setting the health check on!!!!


I'm sorry if I misled anyone the correct answer is C.

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)

upvoted 1 times

  **pondai** 4 months, 1 week ago


Load Balancer's "health check" just check instance health ,if it not health then load balancer doesn't transmission to the bad instance. No help to recreate instance ,so I vote C.

upvoted 1 times

  **neerajgoyal** 4 months, 2 weeks ago

can someone confirm the answer should be A or C if need to choose only one option? I am writing my GCP exam tomorrow.

upvoted 1 times

  **Fidget\_** 4 months, 3 weeks ago

<https://cloud.google.com/compute/docs/instance-groups#autohealing>

"Load balancing health checks help direct traffic away from non-responsive instances and toward healthy instances; these health checks do not cause Compute Engine to recreate instances. On the other hand, managed instance group health checks proactively signal to delete and recreate instances that become UNHEALTHY"

upvoted 2 times

🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago

Answer is A. Create an HTTP load balancer with a backend configuration that references an existing instance group. Set the health check to healthy (HTTP)

upvoted 1 times

🗨️ 👤 **DickDastardly** 4 months, 3 weeks ago

From <https://cloud.google.com/compute/docs/instance-groups>

Health checking

The health checks used to monitor MIGs are similar to the health checks used for load balancing, with some differences in behavior. Load balancing health checks help direct traffic away from non-responsive instances and toward healthy instances; these health checks do not cause Compute Engine to recreate instances. On the other hand, managed instance group health checks proactively signal to delete and recreate instances that become UNHEALTHY

Health checks for load balancers do NOT re-create the instance, as per the requirement, they merely direct traffic away from the sick instance.

A does not solve this requirement therefore it has to be C

upvoted 2 times

🗨️ 👤 **navsaru** 4 months, 3 weeks ago

C is correct answer

upvoted 1 times

🗨️ 👤 **ShakthiGCP** 5 months ago

Answer is A: Don't get confused with other things.. just verified.

upvoted 1 times

🗨️ 👤 **ShakthiGCP** 4 months, 3 weeks ago

I was wrong. A & B are wrong because it should be NetworkLoad Balancer and not HTTP Load balancer. Ans: C is correct .

upvoted 3 times

You are using multiple configurations for gcloud. You want to review the configured Kubernetes Engine cluster of an inactive configuration using the fewest possible steps. What should you do?

- A. Use gcloud config configurations describe to review the output.
- B. Use gcloud config configurations activate and gcloud config list to review the output.
- C. Use kubectl config get-contexts to review the output.
- D. Use kubectl config use-context and kubectl config view to review the output.

Correct Answer: D

Reference:

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

- zukko78

Highly Voted

1 year, 2 months ago

D is correct

upvoted 21 times
- nhusain

2 months, 2 weeks ago

<https://medium.com/google-cloud/kubernetes-engine-kubectl-config-b6270d2b656c>

explains it well

upvoted 3 times
- poogcp

Highly Voted

1 year, 1 month ago

C is correct , Use kubectl config get-contexts to review the output : shows the clusters and the configurations and based on the output we can identify the inactive configurations

upvoted 16 times
- Gurnoor

1 year, 1 month ago

This is wrong get-contexts does not show clusters it only shows contexts.

upvoted 7 times
- jilly

1 year, 1 month ago

True .

Will give only below results

kubectl config get-contexts

CURRENT NAME CLUSTER AUTHINFO NAMESPACE

\* white white dazwilkin

black black dazwilkin

upvoted 1 times
- sidharthwader

Most Recent

4 weeks, 1 day ago

kubectl config get-contexts is used to view the clusters which we have access. Here we can see name of the cluster and permission lev we have in the cluster.

Kubectl config view will list the configs of all clusters you have access.

So its D

upvoted 2 times
- Ashii

2 months, 1 week ago

D is the correct answers, single command will not fetch the full answer

upvoted 2 times
- Agraved

2 months, 2 weeks ago

C is correct, as the point is to find "inactive configuration using the fewest possible".

upvoted 1 times
- Finger41

2 months, 2 weeks ago

The answer is D

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

D is correct

upvoted 1 times
- lxgywil

2 months, 3 weeks ago

AFAIU, the question is saying "You want to REVIEW the configured Kubernetes Engine cluster...".

Using the `kubectl config get-contexts` command is not enough for that as it only lists them with the following context: NAME, CLUSTER, AUTHINFO, NAMESPACE.

To actually SELECT a cluster and REVIEW its details, you will need to run `kubectl config use-context` and `kubectl config view`.

upvoted 2 times

🗨️ 👤 **[Removed]** 4 months, 1 week ago

C is correct. Use kubectl config get-contexts to review the output.

upvoted 1 times

🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago

D is correct

upvoted 1 times

🗨️ 👤 **Hi2ALL** 4 months, 4 weeks ago

"D" Answer is correct

upvoted 1 times

🗨️ 👤 **ShakthiGCP** 5 months ago

No Confusion. It is " D"

upvoted 1 times

🗨️ 👤 **Calistus** 4 months, 4 weeks ago

I agree, "D" is the answer. The question made mention of using multiple configurations, so the "use context .." part of the code help you switch to the inactive instance you want to query, and use the second part to view the details

upvoted 1 times

🗨️ 👤 **Vinod87** 5 months ago

D is the correct answer to see the configuration details for the particular cluster we need to use-context

upvoted 2 times

🗨️ 👤 **Sarin** 5 months, 1 week ago

D assuming that we know the inactive configuration so that we can use-context

upvoted 2 times

🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago

FYI  
List authenticated contexts = kubectl config get-contexts  
Switch context = kubectl config use-context

<https://github.com/dennyzhang/cheatsheet-kubernetes-A4>

upvoted 1 times

🗨️ 👤 **mdf\_pass** 5 months, 1 week ago

Answer is D. C is not correct, kubectl config get-contexts # display list of contexts

upvoted 1 times

🗨️ 👤 **victory108** 5 months, 3 weeks ago

D - Use kubectl config use-context and kubectl config view to review the output.

upvoted 1 times

🗨️ 👤 **victory108** 4 months, 3 weeks ago

Tricky one C is correct. Use kubectl config get-contexts to review the output.

upvoted 1 times

Your company uses Cloud Storage to store application backup files for disaster recovery purposes. You want to follow Google's recommended practices. Which storage option should you use?

- A. Multi-Regional Storage
- B. Regional Storage
- C. Nearline Storage
- D. Coldline Storage

**Correct Answer:** D

Reference:

<https://cloud.google.com/storage/docs/storage-classes#nearline>

- poogcp

Highly Voted

1 year, 1 month ago

Best Answer is " Archive Storage "

<https://cloud.google.com/storage/docs/storage-classes>

But as per the given option next best solution is " Coldline Storage"

upvoted 17 times
- Mutune

5 months ago

Perfectly stated

upvoted 2 times
- zukko78

Highly Voted

1 year, 2 months ago

D is correct,

Coldline Storage COLDLINE 90 days

99.95% in multi-regions and dual-regions

99.9% in regions

upvoted 7 times
- Finger41

Most Recent

2 months, 2 weeks ago

Unless a withdraw frequency is mentioned, the recommended practice by google is to use "Archive Storage". However as this is not an available option, then "Coldline Storage" is the next best solution.

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

D is correct

upvoted 1 times
- kopper2019

3 months, 2 weeks ago

Archive Storage, under uncertainty all people hopes to never use DR site so assuming that Archive is the cheapest but it can be restored in a fast way so DR with great RTO

upvoted 1 times
- kopper2019

3 months, 2 weeks ago

I meant ColdLine, D

upvoted 1 times
- [Removed]

4 months, 1 week ago

D is correct. Coldline Storage

upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

D is correct

upvoted 2 times
- Hi2ALL

4 months, 4 weeks ago



















D is correct

Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive Storage, yet still be available if you need it.

Disaster recovery - In the event of a disaster recovery event, recovery time is key. Cloud Storage provides low latency access to data stored as Archive Storage.

upvoted 1 times
- ShakthiGCP

5 months, 1 week ago

- Verified Recently. Try to create a new Bucket in GCP and it will give options, Standard/Nearline/ColdLine/Archive. And under Coldline - google mentioned "Best for disaster recovery and data accessed less than once a quarter" .So Answer is - ColdLine and Not 'Archive' if the option is given in the future.  
upvoted 2 times
-   **Mannu12** 2 months, 1 week ago  
what if the data is accessed once in 10years  
upvoted 1 times
-   **andregrijp** 5 months, 4 weeks ago  
D.. is the answer but I think in the exam will show Archive and here you can find the new classes <https://cloud.google.com/storage/docs/storage-classes>  
upvoted 2 times
-   **INASR** 6 months ago  
Answer is D  
Coldline is a new Cloud Storage class designed for long-term archival and disaster recovery.  
upvoted 3 times
-   **guid1984** 6 months, 1 week ago  
from GCP docs "Coldline Storage: Very low cost. Good for data that can be stored for at least 90 days, including disaster recovery." Hence ANSWER IS D ColdLine  
upvoted 1 times
-   **vcafe** 6 months, 4 weeks ago  
To store backup files, should use Nearline storage rather than Coldline.  
upvoted 1 times
-   **magistrum** 7 months ago  
The most current answer is Archive storage...unsure if this question will get updated  
[https://cloud.google.com/storage/docs/storage-classes?\\_ga=2.100245342.-1331405417.1599668379&\\_gac=1.26969423.1607902528.CjwKCAiAINf-BRB\\_EiwA2osbxSJcakZWg1RPV4HNzaWrISGcB07JrE1B6KZ6OmAh1zgJoQiI\\_X1QQxoCKRMQAvD\\_BwE](https://cloud.google.com/storage/docs/storage-classes?_ga=2.100245342.-1331405417.1599668379&_gac=1.26969423.1607902528.CjwKCAiAINf-BRB_EiwA2osbxSJcakZWg1RPV4HNzaWrISGcB07JrE1B6KZ6OmAh1zgJoQiI_X1QQxoCKRMQAvD_BwE)  
upvoted 2 times
-   **vara3dk** 7 months, 1 week ago  
Correct Ans is Coldline Storage.  
upvoted 1 times
-   **zanaad231** 7 months, 1 week ago  
I think A is correct  
upvoted 1 times
-   **suresz** 7 months, 2 weeks ago  
Google Cloud Storage Coldline, which is primarily used for archiving and disaster recovery, offers slightly lower availability than the Multi-Regional and Regional Storage classes, as well higher latency.  
  
Ref :- <https://searchcloudcomputing.techtarget.com/definition/Google-Cloud-Storage-Coldline>  
upvoted 1 times



Several employees at your company have been creating projects with Cloud Platform and paying for it with their personal credit cards, which the company reimburses. The company wants to centralize all these projects under a single, new billing account. What should you do?

- A. Contact cloud-billing@google.com with your bank account details and request a corporate billing account for your company.
- B. Create a ticket with Google Support and wait for their call to share your credit card details over the phone.
- C. In the Google Platform Console, go to the Resource Manage and move all projects to the root Organizarion.
- D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

Correct Answer: D

Reference:



https://www.whizlabs.com/blog/google-cloud-interview-questions/

  **samcat84** Highly Voted 1 year, 2 months ago

C is incomplete. Moving projects under an organisation doesn't change their linked billing project.  
<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>

----  
Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet.

----  
D is incomplete as well, after setting the billing account in the organisation you need to link the projects to the new billing account.  
upvoted 24 times

  **sarahf** 7 months, 2 weeks ago

I agree that neither C or D is correct. I did the cert a month ago and this question was not on it. Although a similar question about how to change the payment method from your own card in your project to to the company's "card". So they might have removed this one.  
upvoted 2 times

  **ehizren** 6 months ago

What's was the answer your chose for your particular exam question?  
upvoted 3 times

  **poogcp** Highly Voted 1 year, 1 month ago

C is correct Answer. there will be 1 billing account for the organization and all projects under that organization are linked to single billing account.  
upvoted 8 times

  **hasib125** Most Recent 2 weeks, 4 days ago

I will go with C.

Read the question. The company wants to centralize all these projects under a single, new billing account. What should you do? Company will also reimburses employees billing.

In Google Cloud, you set up a Cloud Billing account and use it to define who pays for a given set of Google Cloud resources. You use Identity and Access Management (IAM) roles to control access to a Cloud Billing account. You can also move projects under one cloud billing account .  
<https://cloud.google.com/billing/docs/how-to/modify-project>. The Resource Manager (D) does not include billing merely resource management.  
upvoted 1 times

  **Larysmith** 1 month ago

The company wants to centralize all these projects under a single, new billing account.  
centralising all projects to a single account might require moving to the organisation root.  
upvoted 1 times

  **Sruthigp** 1 month, 1 week ago



What is thw correct answer? C or D. Please advise.  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago

D, 100%  
upvoted 3 times

  **Mannu12** 2 months, 1 week ago

How are you 100%sure?  
upvoted 1 times

  **Agraved** 2 months, 2 weeks ago

D is correct, as that is a mandatory step. C is not mandate.



upvoted 1 times

🗲️ 👤 **Finger41** 2 months, 2 weeks ago

Technically, its a combo of C and D.

upvoted 3 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

D is correct...<https://cloud.google.com/resource-manager/docs/project-migration> clearly clarifies changing Org won't impact billing

upvoted 1 times

🗲️ 👤 **meh009** 3 months, 1 week ago

Tricky one however the question does state that they want a NEW billing account for these projects and simply moving a project under the Org will likely just use the existing billing account already in situ.

My answer is D

upvoted 1 times

🗲️ 👤 **Mannu12** 2 months, 1 week ago

but what projects will be under the new account that should be build?

upvoted 1 times

🗲️ 👤 **wazza11** 2 months ago

You are correct, there will be none. D is an incomplete solution for it as the next step to D would be C(also not entirely), as we'd move the projects into the organization and change the billing accounts for them.

upvoted 1 times

🗲️ 👤 **civilizador** 3 months, 3 weeks ago

D obviously is the answer because it literally says in the question: " new billing account" . I can't believe so many people do not read the question

upvoted 1 times

🗲️ 👤 **tavva\_prudhvi** 3 months, 1 week ago

So, you create a billing account and set up a payment method, after that???

You have to ask all the users to map their projects to this billing account to get reimbursed, which neither C/D states but the closest one is C.

upvoted 4 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago

C is correct. In the Google Platform Console, go to the Resource Manage and move all projects to the root Organizarion.

upvoted 2 times

🗲️ 👤 **sumanshu** 4 months, 1 week ago

Correct answer is C as Google Cloud Resource Manager can help group the existing accounts under an Organization for centralized billing.

upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

C is correct. In the Google Platform Console, go to the Resource Manage and move all projects to the root Organizarion.

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

D. In the Google Cloud Platform Console, create a new billing account and set up a payment method.

upvoted 1 times

🗲️ 👤 **sumanshu** 4 months, 2 weeks ago

Only creating the new billing account is not sufficient.

upvoted 1 times

🗲️ 👤 **navsaru** 4 months, 3 weeks ago

C is correct answer

upvoted 1 times

🗲️ 👤 **JackGlemins** 5 months ago

C is wrong: <https://cloud.google.com/resource-manager/docs/migrating-projects-billing>

You can move the projects but the billing no chance:

Note: The link between projects and billing accounts is preserved, irrespective of the hierarchy. When you move your existing projects into the organization they will continue to work and be billed as they used to before the migration, even if the corresponding billing account has not been migrated yet. Similarly, if you move a billing account into the organization, all projects linked to it will continue to work even if they are still outside of the organization. There should be no server downtime or impact as a result of migration.

I think D is best answers

upvoted 2 times

You have an application that looks for its licensing server on the IP 10.0.3.21. You need to deploy the licensing server on Compute Engine. You do not want to change the configuration of the application and want the application to be able to reach the licensing server. What should you do?

- A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.
- B. Reserve the IP 10.0.3.21 as a static public IP address using gcloud and assign it to the licensing server.
- C. Use the IP 10.0.3.21 as a custom ephemeral IP address and assign it to the licensing server.
- D. Start the licensing server with an automatic ephemeral IP address, and then promote it to a static internal IP address.

Correct Answer: A

- zukko78** Highly Voted

1 year, 2 months ago

it's obvious, A

upvoted 16 times
- Khaled\_Rashwan** Highly Voted

1 year, 2 months ago

A

IP 10.0.3.21 is internal by default, and to ensure that it will be static non-changing it should be selected as static internal ip address.

upvoted 9 times
- Finger41** Most Recent

2 months, 2 weeks ago

Its A, if its in the same subnet you can use the internal IP, which you would want as a static IP. There is no point of connecting to the server via the internet.

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

A is correct

upvoted 1 times
- kopper2019**

3 months, 2 weeks ago

A is the only one

upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

upvoted 2 times
- INASR**

6 months ago

agreed it is A

IP 10.0.3.21 is internal by default, and to ensure that it will be static non-changing it should be selected as static internal ip address.

upvoted 2 times
- Gohs**

6 months, 1 week ago

It seems like IP addresses that start with 10.

Are usually internal IPs on Google cloud.

upvoted 1 times
- EABDAJA**

6 months, 1 week ago

Answer is A

upvoted 1 times
- Ozymandiox**

9 months, 1 week ago

B, it's just a good joke 10.x public, yeah, sure...

C, states that th IP it's ephemeral, so... Nope

D, does not guarantee the IP we do need

sooooo ....:-P

upvoted 5 times
- swatititame**


9 months, 1 week ago

A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.

upvoted 1 times
- nmohapatra87**

9 months, 2 weeks ago

A is correct  
upvoted 1 times

  **glam** 9 months, 3 weeks ago

A. Reserve the IP 10.0.3.21 as a static internal IP address using gcloud and assign it to the licensing server.  
upvoted 4 times

You are deploying an application to App Engine. You want the number of instances to scale based on request rate. You need at least 3 unoccupied instances at all times. Which scaling type should you use?

- A. Manual Scaling with 3 instances.
- B. Basic Scaling with min\_instances set to 3.
- C. Basic Scaling with max\_instances set to 3.
- D. Automatic Scaling with min\_idle\_instances set to 3.

**Correct Answer:** D

Reference:  
<https://cloud.google.com/appengine/docs/standard/python/how-instances-are-managed>

- zukko78**

Highly Voted

1 year, 2 months ago

D is correct.  
App Engine supports the following scaling types, which controls how and when instances are created:  
  
Automatic  
Basic  
Manual  
You specify the scaling type in your app's app.yaml.  
  
Automatic scaling  
Automatic scaling creates instances based on request rate, response latencies, and other application metrics. You can specify thresholds for each of these metrics, as well as a minimum number instances to keep running at all times.  
upvoted 28 times
- pacman\_user**

Most Recent

2 weeks, 4 days ago

D is correct.  
upvoted 2 times
- sidharthwader**

1 month ago

Obviously Auto-scaling with min no of instances as 3  
upvoted 1 times
- vmart**

1 month, 1 week ago

D is correct  
upvoted 1 times
- Finger41**

2 months, 2 weeks ago

D is correct : <https://cloud.google.com/appengine/docs/standard/go/config/appref>  
"App Engine calculates the number of instances necessary to serve your current application traffic based on scaling settings such as target\_cpu\_utilization and target\_throughput\_utilization. Setting min\_idle\_instances specifies the number of instances to run in addition to this calculated number. For example, if App Engine calculates that 5 instances are necessary to serve traffic, and min\_idle\_instances is set to 2, App Engine will run 7 instances (5, calculated based on traffic, plus 2 additional per min\_idle\_instances)."  
upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

D is correct  
upvoted 1 times
- [Removed]**

4 months, 1 week ago

D is correct. Automatic Scaling with min\_idle\_instances set to 3.  
upvoted 2 times
- pondai**

4 months, 1 week ago

D is right  
B is wronth<https://cloud.google.com/compute/docs/disks/create-snapshots?hl=zh-tw>  
upvoted 2 times
- EABDAJA**

4 months, 3 weeks ago

D is the Answer  
upvoted 1 times
- JackGlemins**

5 months ago

D is right: <https://cloud.google.com/appengine/docs/standard/go/config/appref>.  
App Engine calculates the number of instances necessary to serve your current application traffic based on scaling settings such as target\_cpu\_utilization and target\_throughput\_utilization. Setting min\_idle\_instances specifies the number of instances to run in addition to this

calculated number. For example, if App Engine calculates that 5 instances are necessary to serve traffic, and min\_idle\_instances is set to 2, App Engine will run 7 instances (5, calculated based on traffic, plus 2 additional per min\_idle\_instances).

upvoted 2 times

  **Vinod87** 5 months ago

D is correct

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

D. Automatic Scaling with min\_idle\_instances set to 3.

upvoted 2 times

  **swatititame** 9 months, 1 week ago

- D. Automatic Scaling with min\_idle\_instances set to 3.

upvoted 1 times

  **glam** 9 months, 3 weeks ago



D. Automatic Scaling with min\_idle\_instances set to 3.

upvoted 4 times

  **[Removed]** 10 months, 3 weeks ago

min\_idle\_instances is the app.yml file changes you need to make and include in your application and deploy on app engine  
<https://cloud.google.com/appengine/docs/standard/nodejs/config/appref>

upvoted 1 times

  **Xolex** 11 months, 2 weeks ago

definitely D

upvoted 1 times

  **professor** 1 year, 1 month ago

Ans is D

You need to scale based on the request rate. only possible with automatic scaling

upvoted 3 times

You have a development project with appropriate IAM roles defined. You are creating a production project and want to have the same IAM roles on the new project, using the fewest possible steps. What should you do?

- A. Use gcloud iam roles copy and specify the production project as the destination project.
- B. Use gcloud iam roles copy and specify your organization as the destination organization.
- C. In the Google Cloud Platform Console, use the "create role from role" functionality.
- D. In the Google Cloud Platform Console, use the "create role" functionality and select all applicable permissions.

**Correct Answer:** B

Reference:

<https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>

- coldpar

Highly Voted

1 year, 4 months ago

Correct Answer is A not B

upvoted 23 times
- Agents89

Highly Voted

1 year, 2 months ago

Correct answer is A

upvoted 13 times
- Rajagopalan

Most Recent

2 weeks, 1 day ago

Answer is B

gcloud iam roles copy [--dest-organization=DEST\_ORGANIZATION] [--dest-project=DEST\_PROJECT] [--destination=DESTINATION] [--source=SOURCE] [--source-organization=SOURCE\_ORGANIZATION] [--source-project=SOURCE\_PROJECT] [GLOUD\_WIDE\_FLAG ...]

upvoted 1 times
- pacman\_user

2 weeks, 4 days ago

Correct Answer is A

gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomSpannerDbAdmin --dest-project=PROJECT\_ID

upvoted 1 times
- AB15

1 month ago

A is correct. You do not want to copy the roles to all the projects under the organisation. Hence A.

upvoted 2 times
- tanito83

1 month, 3 weeks ago

The correct answer is A. Please, modify it.

upvoted 1 times
- Agraved

2 months, 2 weeks ago

A is correct, as role has to copied from Dev project to Prod, not across org.

upvoted 1 times
- Finger41

2 months, 2 weeks ago

Copy roles to a "project", not "organization" as per question. The link provided, additionally has instructions on deploying roles into projects : <https://cloud.google.com/sdk/gcloud/reference/iam/roles/copy>

gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomSpannerDbAdmin --dest-project=PROJECT\_ID

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct

upvoted 1 times
- lxgywil

2 months, 3 weeks ago

With `gcloud iam roles copy` you have to specify either the PROJECT or ORGANIZATION name for the destination custom role.

Since the question specifies that you need project-level roles, option A should be correct.


Ref: [https://https://www.examttopics.com/exams/google/associate-cloud-engineer/view/3/#cloud.google.com/sdk/gcloud/reference/iam/roles/copy](https://www.examttopics.com/exams/google/associate-cloud-engineer/view/3/#cloud.google.com/sdk/gcloud/reference/iam/roles/copy)

upvoted 1 times
- r1ck

3 months, 2 weeks ago

Yes A, b is wrong answer

upvoted 1 times

  **kopper2019** 3 months, 2 weeks ago



A, is the answer  
upvoted 1 times

  **Awot** 4 months ago



Obviously A. You're replicating roles for project not new organization.  
upvoted 2 times

  **[Removed]** 4 months, 1 week ago



A is correct. Use gcloud iam roles copy and specify the production project as the destination project.  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct  
upvoted 3 times

  **navsaru** 4 months, 3 weeks ago

Correct answer is A  
upvoted 1 times



  **Hi2ALL** 4 months, 4 weeks ago



I think B is correct answer  
EXAMPLES  
To create a copy of an existing role spanner.databaseAdmin into an organization with 1234567, run:  
  
gcloud iam roles copy --source="roles/spanner.databaseAdmin" --destination=CustomViewer --dest-organization=1234567  
upvoted 2 times



You need a dynamic way of provisioning VMs on Compute Engine. The exact specifications will be in a dedicated configuration file. You want to follow Google’s recommended practices. Which method should you use?



- A. Deployment Manager
- B. Cloud Composer
- C. Managed Instance Group
- D. Unmanaged Instance Group



**Correct Answer: C**  
Reference:  
<https://cloud.google.com/compute/docs/instances/>



-   **[Removed]** Highly Voted 4 months, 1 week ago



A is correct. Deployment Manager  
upvoted 18 times
-   **rgalfaro** Highly Voted 4 months, 1 week ago



Taking out D and B (Apache AirFlow) then we have A and C. Now, from MIG description from Google Cloud it states:  
A MIG creates each of its managed instances based on the instance template and optional stateful configuration that you specify.  
  
So, here on MIG, configuration files are optional and what's required is an Instance Template, not even mentioned on the question itself.  
  
From that logic, and assuming we have only a configuration file, Deployment Manager is the correct answer, so A.  
upvoted 7 times
-   **yite** 3 months, 4 weeks ago



Deployment Manager speaks to using a config file, I don't see that mentioned with MIG. I agree with A.  
upvoted 2 times
-   **associateclouDEXAMUSER** Most Recent 3 weeks ago



A is correct.  
Deployment manager can create, update and delete resources inside GCP by simply using a config file  
upvoted 1 times
-   **Agraved** 2 months, 2 weeks ago



Has to be "A" as it say specifications will be in a dedicated configuration file  
upvoted 1 times
-   **Finger41** 2 months, 2 weeks ago

A is correct, you deploy and configure infrastructure via Deployment Manager  
upvoted 1 times
-   **mcaromit** 2 months, 2 weeks ago

A is correct  
upvoted 1 times
-   **yvinisiupacuando** 2 months, 3 weeks ago

I vote for A  
upvoted 1 times
-   **Amu89** 2 months, 3 weeks ago

Did you wrote the exam?  
upvoted 2 times
-   **Nee26** 3 months, 4 weeks ago

The correct answer is A Deployment Manager:  
<https://cloud.google.com/deployment-manager?hl=en>  
  
Template-driven  
Templates allow the use of building blocks to create abstractions or sets of resources that are typically deployed together (e.g. an instance template, instance group, and autoscaler). These templates can be parameterized to allow them to be used over and over by changing input values to define what image to deploy, the zone in which to deploy, or how many virtual machines to deploy.  
upvoted 4 times
-   **ozeross** 4 months ago

correct answer - MIG.




upvoted 1 times

  **Luck345** 4 months, 1 week ago

I think A is answer

upvoted 2 times

  **Abhee12** 4 months, 2 weeks ago


A, question is about dynamic way to provision VM, it can be achieved by Deployment manager or by using terraform. MIG is creating multiple machine based on template by loadbalancing

upvoted 4 times

  **CloudMom** 4 months, 2 weeks ago



so which is it? A or C? i vote for A

upvoted 2 times

  **sasha0518** 4 months, 2 weeks ago

c is right

upvoted 2 times

  **lio123** 4 months, 2 weeks ago

it is C.

<https://cloud.google.com/compute/docs/instance-groups>

Managed instance groups (MIGs) let you operate apps on multiple identical VMs. You can make your workloads scalable and highly available by taking advantage of automated MIG services, including: autoscaling, autohealing, regional (multiple zone) deployment, and automatic updating.

upvoted 1 times

  **neerajgoyal** 4 months, 2 weeks ago

As the question is saying specifications will be in a dedicated configuration file, I see that only in Deployment manager we create configuration file. Are we creating any such file in MIG?

upvoted 2 times

  **drizzydroo** 4 months, 2 weeks ago

I thought his was A. How do you use managed instance groups with a configuration file?

upvoted 3 times

  **neerajgoyal** 4 months, 2 weeks ago

A is the best way

upvoted 2 times



You have a Dockerfile that you need to deploy on Kubernetes Engine. What should you do?


- A. Use kubectl app deploy <dockerfilename>.
- B. Use gcloud app deploy <dockerfilename>.
- C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.
- D. Create a docker image from the Dockerfile and upload it to Cloud Storage. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

**Correct Answer:** C

Reference -



<https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>


-   **Agents89**

Highly Voted 

 1 year, 3 months ago



C is correct

upvoted 22 times
-   **associatecloudexamuser**



Most Recent 

 3 weeks, 1 day ago



C is correct

upvoted 2 times
-   **Finger41** 2 months, 2 weeks ago



A and B are App Engine which doesnt satisfy the use of GKE. You use Container Registry in the e2e process of deploying Docker on GKE.

upvoted 1 times
-   **mcaromit** 2 months, 2 weeks ago



C is correct

upvoted 1 times
-   **yvinisiupacuando** 2 months, 3 weeks ago

It's clearly C, you need to use Container Registry.



upvoted 1 times
-   **[Removed]** 4 months, 1 week ago

C is correct. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.



upvoted 2 times
-   **pondai** 4 months, 2 weeks ago

I vote C

A: kubectl can't deploy "DuckerFile"

upvoted 2 times
-   **GCP\_Student1** 5 months, 1 week ago



C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.

upvoted 1 times
-   **Gohs** 5 months, 4 weeks ago



C.

Container registry is where docker files are stored.

From there you can deploy to GKE, Cloud Run or App Engine flexible.

upvoted 1 times
-   **nherrerab** 6 months, 2 weeks ago

C is correct.

upvoted 1 times
-   **mohdafiuddin** 7 months ago

Answer is C-

A. Not feasible because there is no container image (We only have a Dockerfile mentioned).

B. Not feasible because We don't need gcloud to do this.

C. This is the correct answer and has all the steps. please refer - <https://cloud.google.com/kubernetes-engine/docs/tutorials/hello-app>

D. There is no need to use Cloud Storage. Please refer to above option (step 2 in the link I provided)

upvoted 1 times

- 🗨️ 👤 **Damini51** 7 months, 1 week ago  
c is correct, I was also confused but c is correct  
upvoted 1 times
- 🗨️ 👤 **gh999l** 8 months, 3 weeks ago  
C is the correct , Have checked this in Official guide for ACE EXAM by Dan Sullivan  
upvoted 3 times
- 🗨️ 👤 **hems4all** 9 months ago  
C is OK  
upvoted 1 times
- 🗨️ 👤 **swatititame** 9 months, 1 week ago  
C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.  
upvoted 1 times
- 🗨️ 👤 **glam** 9 months, 3 weeks ago  
C. Create a docker image from the Dockerfile and upload it to Container Registry. Create a Deployment YAML file to point to that image. Use kubectl to create the deployment with that file.  
upvoted 3 times
- 🗨️ 👤 **[Removed]** 10 months, 2 weeks ago  
A - command does not exist.  
B - is to deploy on App Engine  
B and C are good as they talk about creating docker image from the docker file (the purpose of this file is to create image only) we need to store docker image in docker registry so C is correct  
upvoted 1 times
- 🗨️ 👤 **[Removed]** 10 months, 2 weeks ago  
Small correct talking about C and D and not B and C in the third statement.  
upvoted 1 times

Your development team needs a new Jenkins server for their project. You need to deploy the server using the fewest steps possible. What should you do?

- A. Download and deploy the Jenkins Java WAR to App Engine Standard.
- B. Create a new Compute Engine instance and install Jenkins through the command line interface.
- C. Create a Kubernetes cluster on Compute Engine and create a deployment with the Jenkins Docker image.
- D. Use GCP Marketplace to launch the Jenkins solution.

Correct Answer: D

Reference:

<https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine>

- Agents89

Highly Voted

1 year, 3 months ago

D is correct  
upvoted 19 times
- Student7

Most Recent

2 weeks, 1 day ago

Correct Answer is D  
upvoted 2 times
- kashi123

2 weeks, 6 days ago

D is correct  
upvoted 1 times
- vmart

1 month, 3 weeks ago

D is the right answer  
upvoted 1 times
- vmart

1 month, 3 weeks ago

D is the correct answer  
upvoted 1 times
- Agraved

2 months, 2 weeks ago

If Market place is an option it has to be that. So answer is "D".  
upvoted 1 times
- Finger41

2 months, 2 weeks ago

Its faster to get it from the Marketplace. D is the answer  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

D is correct  
upvoted 1 times
- bisla

3 months ago

Yes, D is correct!  
upvoted 1 times
- meh009

3 months, 1 week ago

D is correct  
upvoted 1 times
- Gopi\_GCP

3 months, 1 week ago

D is correct  
upvoted 1 times
- kopper2019

3 months, 2 weeks ago



marketplace de way to go, D  
upvoted 1 times
- [Removed]



4 months, 1 week ago



D is correct. Use GCP Marketplace to launch the Jenkins solution.  
upvoted 1 times
- imranmani

4 months, 1 week ago

D is a correct answer  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago  
Answer is D. Use GCP Marketplace to launch the Jenkins solution.  
upvoted 1 times

  **Hi2ALL** 4 months, 4 weeks ago  
Yes D is the right answer  
upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago  
D. Use GCP Marketplace to launch the Jenkins solution  
upvoted 1 times

You need to update a deployment in Deployment Manager without any resource downtime in the deployment. Which command should you use?

- A. gcloud deployment-manager deployments create --config <deployment-config-path>
- B. gcloud deployment-manager deployments update --config <deployment-config-path>
- C. gcloud deployment-manager resources create --config <deployment-config-path>
- D. gcloud deployment-manager resources update --config <deployment-config-path>

**Correct Answer:** B  
Reference:  
<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/deployments/update>

- Agents89**

Highly Voted

1 year, 3 months ago

B is correct  
upvoted 26 times
- [Removed]**

Highly Voted

10 months, 3 weeks ago

B is correct Additional tip, update and create resource is not even a command under deployment management service.  
upvoted 12 times
- associatecloudexamuser**

Most Recent

3 weeks ago

B is correct  
upvoted 1 times
- vmart**

1 month, 1 week ago

B is correct  
upvoted 1 times
- vmart**

1 month, 3 weeks ago

b is the right choice  
upvoted 1 times
- Finger41**

2 months, 2 weeks ago

B is correct, as per :<https://cloud.google.com/sdk/gcloud/reference/deployment-manager/deployments/update>  
upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- [Removed]**

4 months, 1 week ago

B is correct. gcloud deployment-manager deployments update --config <deployment-config-path>  
upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

B is correct. gcloud deployment-manager deployments update --config <deployment-config-path>  
upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

Ans is B  
upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

B. gcloud deployment-manager deployments update --config <deployment-config-path>  
upvoted 1 times
- victory108**

5 months, 3 weeks ago

B - gcloud deployment-manager deployments update --config <deployment-config-path>  
upvoted 1 times
- INASR**

6 months ago

YEs B is correct  
upvoted 1 times
- raj117**

6 months, 1 week ago

It's correct

upvoted 1 times

  **JKRowlings** 7 months, 2 weeks ago

Correct answer is B. gcloud deployment-manager deployments update - update a deployment based on a provided config file  
upvoted 1 times

  **hems4all** 9 months ago

B is OK  
upvoted 1 times

  **swatititame** 9 months, 1 week ago

B. gcloud deployment-manager deployments update --config <deployment-config-path>  
upvoted 1 times

You need to run an important query in BigQuery but expect it to return a lot of records. You want to find out how much it will cost to run the query. You are using on-demand pricing. What should you do?

- A. Arrange to switch to Flat-Rate pricing for this query, then move back to on-demand.
- B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.
- C. Use the command line to run a dry run query to estimate the number of bytes returned. Then convert that bytes estimate to dollars using the Pricing Calculator.
- D. Run a select count (\*) to get an idea of how many records your query will look through. Then convert that number of rows to dollars using the Pricing Calculator.

**Correct Answer:** B  
Reference:  
<https://cloud.google.com/bigquery/docs/estimate-costs>

- Agents89

Highly Voted

1 year, 2 months ago

B is Correct  
upvoted 20 times
- ESP\_SAP

Highly Voted

11 months, 1 week ago

Correct Answers is (B):  
  
On-demand pricing  
Under on-demand pricing, BigQuery charges for queries by using one metric: the number of bytes processed (also referred to as bytes read). You are charged for the number of bytes processed whether the data is stored in BigQuery or in an external data source such as Cloud Storage, Drive, or Cloud Bigtable. On-demand pricing is based solely on usage.  
  
[https://cloud.google.com/bigquery/pricing#on\\_demand\\_pricing](https://cloud.google.com/bigquery/pricing#on_demand_pricing)  
upvoted 8 times
- vmart

Most Recent

1 month, 3 weeks ago

B is correct answer  
upvoted 1 times
- Ashii

2 months, 1 week ago

Bytes scanned or bytes read always !  
upvoted 1 times
- Finger41

2 months, 2 weeks ago

B is correct as per : <https://cloud.google.com/bigquery/docs/dry-run-queries>  
When you run a query in the bq command-line tool, you can use the --dry\_run flag to estimate the number of bytes read by the query. You can also use the dryRun parameter when submitting a query job using the API or client libraries.  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- yvinisiupacuando

2 months, 3 weeks ago

B is the Right answer, it returns the bytes read.  
upvoted 1 times
- [Removed]

4 months, 1 week ago

B is correct. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.  
upvoted 1 times
- GCP\_Student1















5 months, 1 week ago

B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.  
upvoted 1 times
- INASR

6 months ago

B is Correct  
upvoted 1 times



-   **EABDAJA** 6 months ago  
Answer is B  
upvoted 1 times
-   **raj117** 6 months, 1 week ago  
B is correct  
upvoted 1 times
-   **victory108** 6 months, 4 weeks ago  
B - Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.  
upvoted 2 times
-   **hems4all** 9 months ago  
B is OK  
upvoted 1 times
-   **swatititame** 9 months ago  
B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.  
upvoted 1 times
-   **glam** 9 months, 3 weeks ago  
B. Use the command line to run a dry run query to estimate the number of bytes read. Then convert that bytes estimate to dollars using the Pricing Calculator.  
upvoted 2 times
-   **professor** 1 year, 1 month ago  
Ans is B  
  
Bytes Read  
upvoted 4 times

## Question #16

## Topic 1

You have a single binary application that you want to run on Google Cloud Platform. You decided to automatically scale the application based on underlying infrastructure CPU usage. Your organizational policies require you to use virtual machines directly. You need to ensure that the application scaling is operationally efficient and completed as quickly as possible. What should you do?

- A. Create a Google Kubernetes Engine cluster, and use horizontal pod autoscaling to scale the application.
- B. Create an instance template, and use the template in a managed instance group with autoscaling configured.**
- C. Create an instance template, and use the template in a managed instance group that scales up and down based on the time of day.
- D. Use a set of third-party tools to build automation around scaling the application up and down, based on Stackdriver CPU usage monitoring.

**Correct Answer: B**

  **coldpar** Highly Voted 1 year, 4 months ago

correct is B as you have to use VM instances directly.  
upvoted 33 times

  **mohdafiuddin** Highly Voted 7 months ago

I'll take a simple and logical approach for answering this.  
Let's first break down the question into key requirements -  
1. automatically scale the application based on underlying infrastructure CPU usage.  
2. use virtual machines directly.

A. Not feasible because VMs are not used directly here.  
B. This is the correct answer.  
C. Time of Day... Easy elimination because this does not scale on CPU usage and time of day is mentioned NOWHERE.  
D. Third Party Tools.... Nobody would use GCP if they needed third party tools to do something as simple as scaling based on CPU usage. all popular cloud providers have native solutions for this including GCP.  
upvoted 6 times

  **kopper2019** 3 months, 2 weeks ago

and also D is out because why would I use a third party tool when is a GCP exam  
upvoted 2 times

  **AB15** Most Recent 1 month ago

B is correct.  
upvoted 1 times

  **shankyomre01** 2 months ago

yes B is correct  
upvoted 1 times

  **Ashii** 2 months, 1 week ago


Asked to use VM's directly, otherwise kubernetes would be the choice  
upvoted 1 times

  **Hasaaaaan** 1 month, 3 weeks ago

Kubernetes with Vertical Pod Autoscaling will be right choice as HPA scales pods only.  
upvoted 1 times

  **Finger41** 2 months, 2 weeks ago

Correct answer is B  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago



B is correct  
upvoted 1 times

  **yvinisiupacuando** 2 months, 3 weeks ago

I vote for B.  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

B is correct. Create an instance template, and use the template in a managed instance group with autoscaling configured.  
upvoted 2 times

  **Hi2ALL** 4 months, 4 weeks ago

Tricky but I think answer should be B , but A can be used as option as if there was no managed instances option available

upvoted 1 times

  **sidharthwader** 1 month ago

No. Its mentioned in the question "Your organizational policies require you to use virtual machines directly." so here its not directly used as vm its used as nodes. So its not A

upvoted 1 times



  **cloud\_hunk** 5 months ago

question itself is not complete or not right. here are the catches:

1. Auto scale application that you decide
2. you require to use vm directly as an approved go forward

my thought: if you are the decision maker where you wanted to deploy the app for autos calling GKE HPA is a good option in that case answer is A. but if have been told or mandate to deploy on Vm, the only thing you can do is to go ahead and deploy in VM in that case answer is B.

upvoted 1 times

  **Hasaaaaan** 1 month, 3 weeks ago



Why HPA? Why not VPA?

upvoted 1 times

  **Vinod87** 5 months ago

B is correct

upvoted 1 times

  **Sarin** 5 months, 1 week ago



Correct answer B. It cannot be A because were the questions says we need to use autoscaling on VM. Although GKE uses virtual machines under the hood for its cluster, the autoscaling is different. In this option, although horizontal pod autoscaling is enabled, Cluster Autoscaling is not enabled

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago



B. Create an instance template, and use the template in a managed instance group with autoscaling configured.

upvoted 2 times

  **roy\_123** 5 months, 1 week ago

B. Create an instance template, and use the template in a managed instance group with autoscaling configured.

upvoted 1 times

  **prakashrao008** 5 months, 4 weeks ago

B is the correct answer

upvoted 1 times

  **INASR** 6 months ago

YES ,correct is B as you have to use VM instances directly.

upvoted 1 times

You are analyzing Google Cloud Platform service costs from three separate projects. You want to use this information to create service cost estimates by service type, daily and monthly, for the next six months using standard query syntax. What should you do?

- A. Export your bill to a Cloud Storage bucket, and then import into Cloud Bigtable for analysis.
- B. Export your bill to a Cloud Storage bucket, and then import into Google Sheets for analysis.
- C. Export your transactions to a local file, and perform analysis with a desktop tool.
- D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

Correct Answer: D

- mohdafiuddin

Highly Voted

7 months ago

Solving this by first eliminating the options that don't suit us. By breaking down the question into the key requirements-

1. Analyzing Google Cloud Platform service costs from three separate projects.

2. Using standard query syntax. -> (Relational data and SQL)

A. 'Cloud Storage bucket'.....'Cloud Bigtable'. Not feasible, mainly because cloud BigTable is not good for Structured Data (or Relational Data on which we can run SQL queries as per the question's requirements). BigTable is better suited for Semi Structured data and NoSQL data.

B. 'Cloud Storage bucket'.....'Google Sheets'. Not Feasible because there is no use of SQL in this option, which is one of the requirements.

C. Local file, external tools... this is automatically eliminated because the operation we need is simple, and there has to be a GCP native solution for this. We shouldn't need to rely on going out of the cloud for such a simple thing.

D. 'BigQuery'.....'SQL queries' -> This is the right answer.

upvoted 18 times
- cesar7816

Highly Voted

1 year, 4 months ago

Agreed, BigQuery

upvoted 13 times
- yurstev

7 months, 1 week ago

the key is standard query syntax

upvoted 2 times
- HogwartsTrue

Most Recent

2 months ago

Ans:D BigQuery is for Analytical Purpose

upvoted 1 times
- Finger41

2 months, 2 weeks ago

D as its talking about analysis, which can be performed by BigQuery.

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

D is correct

upvoted 1 times
- yvinisiupacuando

2 months, 3 weeks ago

"D" is the right answer. Talking about analysis and SQL queries it clearly points to D.

upvoted 1 times
- [Removed]

4 months, 1 week ago

D is correct. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

D.Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis

upvoted 1 times
- Hi2ALL

4 months, 4 weeks ago

The given answer is acceptable

upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.

upvoted 2 times
- iPablo26

5 months, 2 weeks ago

D is the correct answer

upvoted 1 times

  **INASR** 6 months ago

D , Agreed, BigQuery  
upvoted 1 times

  **nherrerab** 6 months, 2 weeks ago



D is correct.  
upvoted 1 times

  **hems4all** 9 months ago



D BigQuery  
upvoted 1 times

  **glam** 9 months, 3 weeks ago

D. Export your bill to a BigQuery dataset, and then write time window-based SQL queries for analysis.  
upvoted 2 times

  **prasanu** 9 months, 3 weeks ago

Agreed D  
upvoted 1 times

  **JJ\_ME** 9 months, 3 weeks ago

D.  
BigQuery is the only option that allows SQL.  
upvoted 3 times

## Question #18

## Topic 1

You need to set up a policy so that videos stored in a specific Cloud Storage Regional bucket are moved to Coldline after 90 days, and then deleted after one year from their creation. How should you set up the policy?

- A. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 275 days (365-90).
- B. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 365 days.**
- C. Use gsutil rewrite and set the Delete action to 275 days (365-90).
- D. Use gsutil rewrite and set the Delete action to 365 days.

**Correct Answer: A**

  **Sammigbo** Highly Voted 1 year, 1 month ago

Answer is B. There should be no reason to recalculate the time needed to delete after a year.  
upvoted 23 times

  **JKRowlings** 8 months ago

The correct ans is A.  
upvoted 1 times

  **yvinisiupacundo** 2 months, 3 weeks ago

Right answer is clearly B, "A" does not make any sense.  
upvoted 3 times

  **cloudehthu01** Highly Voted 1 year, 1 month ago

Correct is B.  
You only re-calculate expiry date when objects are re-written using re-write option to another storage class in which case creation date is reset. But in this case objects is moved to Coldline class after 90 days and then we want to delete the object after 365 days.  
upvoted 13 times

  **T\_T\_M** 11 months ago

You can change the storage class of an existing object either by rewriting the object or by using Object Lifecycle Management...Since Object Lifecycle management was used there was no need to recalculate the expiration date and delete action still remains 365 days.

<https://cloud.google.com/storage/docs/storage-classes>  
upvoted 4 times

  **YAS007** Most Recent 1 week ago

B is correct  
from gcp : "The Age condition is satisfied when an object reaches the specified age (in days). Age is measured from the object's creation time. For example, if an object's creation time is 2019/01/10 10:00 UTC and the Age condition is 10 days, then the condition is satisfied for the object on and after 2019/01/20 10:00 UTC. This is true even if the object becomes noncurrent through Object Versioning sometime after its creation."  
<https://cloud.google.com/storage/docs/lifecycle>  
upvoted 1 times

  **YAS007** 1 week ago

B is correct  
from gcp : "The Age condition is satisfied when an object reaches the specified age (in days). Age is measured from the object's creation time. For example, if an object's creation time is 2019/01/10 10:00 UTC and the Age condition is 10 days, then the condition is satisfied for the object on and after 2019/01/20 10:00 UTC. This is true even if the object becomes noncurrent through Object Versioning sometime after its creation."  
<https://cloud.google.com/storage/docs/lifecycle>  
upvoted 1 times

  **akshayonline** 2 weeks, 3 days ago

Correct answer is B ( days calculated from initial date )  
upvoted 1 times

  **sidharthwader** 1 month ago

Answer must be A.  
It says it should be removed 365 days after the creation so 275+90 =365 days.  
Please correct me if am wrong !  
upvoted 1 times

  **sidharthwader** 1 month ago

sorry it should be B. As per age its from the day its created so B is the answer  
upvoted 1 times

- 🗲️ 👤 **tanito83** 1 month, 3 weeks ago  
Gentlemen, the answer is B. Please, modify it.  
upvoted 1 times
- 🗲️ 👤 **vmart** 1 month, 3 weeks ago  
B is the right choice  
upvoted 1 times
- 🗲️ 👤 **Umesh09** 2 months, 1 week ago  
Answer is B you calculated from day the data is initially written on the bucket not when it moves from one tier to other  
upvoted 1 times
- 🗲️ 👤 **Finger41** 2 months, 2 weeks ago  
The keyword is creation, which can be handled by a lifecycle policy.It will be deleted after 365 days. Answer B is correct.  
upvoted 1 times
- 🗲️ 👤 **mcaromit** 2 months, 2 weeks ago  
B is correct  
upvoted 1 times
- 🗲️ 👤 **lxgywil** 2 months, 3 weeks ago  
"A" is correct since modifying the Storage Class creates a new version for the object. The "creation time" reflects the time when that last version was created. Therefore, you do need to subtract those 90 days.  
  
This is pretty easy to test in the GCP Console:  
1. Upload an object to a GCS bucket and see its creation time.  
2. Wait for a few minutes and run `gsutil rewrite -s coldline gs://PATH\_TO\_OBJECT`.  
3. Note that the object's creation time changed.  
4. Run `gsutil ls -ahl gs://BUCKET\_NAME` and you will note that modifying the Storage Class has created a new version for your object (which creation time is, obviously, different).  
upvoted 2 times
- 🗲️ 👤 **lxgywil** 2 months, 3 weeks ago  
I was wrong - the answer is indeed "B".  
  
Changing Storage Class via Lifecycle Management doesn't rewrite the object, thus its creation time stays the same.  
upvoted 6 times
- 🗲️ 👤 **ready2rock** 1 month, 2 weeks ago  
Thank you for digging into it and reposting/correcting your answer. I was trying to determine if changing storage class cause a rewrite (since this is the way to decide whether B or A is the right answer). Looks like you found it, making B the right answer  
upvoted 1 times
- 🗲️ 👤 **mistryminded** 3 months ago  
Answer should be B.  
as per this doc:<https://cloud.google.com/storage/docs/lifecycle#age>  
Age is measured from the object's creation time. Also,This is true even if the object becomes noncurrent through Object Versioning sometime after its creation. Hence, answer should be B not A  
upvoted 2 times
- 🗲️ 👤 **ApaMokus** 3 months, 1 week ago  
The correct answer is B.  
The Age condition is measured from the object's creation time. Unlike changing an object's storage class manually, using SetStorageClass does not rewrite an object.  
<https://cloud.google.com/storage/docs/lifecycle>  
upvoted 1 times
- 🗲️ 👤 **maddy5835** 3 months, 3 weeks ago  
A is the correct answer, Unlike changing an object's storage class manually, using SetStorageClass does not rewrite an object.  
upvoted 1 times
- 🗲️ 👤 **ozeross** 4 months ago  
B is the correct answer no need to recalcuate  
upvoted 3 times
- 🗲️ 👤 **kopper2019** 3 months, 2 weeks ago  
B is the way to go  
upvoted 1 times
- 🗲️ 👤 **[Removed]** 4 months, 1 week ago  
B is correct. Use Cloud Storage Object Lifecycle Management using Age conditions with SetStorageClass and Delete actions. Set the SetStorageClass action to 90 days and the Delete action to 365 days.  
upvoted 2 times



## Question #19

## Topic 1

You have a Linux VM that must connect to Cloud SQL. You created a service account with the appropriate access rights. You want to make sure that the VM uses this service account instead of the default Compute Engine service account. What should you do?

- A. When creating the VM via the web console, specify the service account under the "Identity and API Access" section.
- B. Download a JSON Private Key for the service account. On the Project Metadata, add that JSON as the value for the key compute-engine-service-account.
- C. Download a JSON Private Key for the service account. On the Custom Metadata of the VM, add that JSON as the value for the key compute-engine-service-account.
- D. Download a JSON Private Key for the service account. After creating the VM, ssh into the VM and save the JSON under ~/.gcloud/compute-engine-service-account.json.

**Correct Answer: A**

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

  **Agents89** Highly Voted 1 year, 2 months ago

A is correct  
upvoted 19 times

  **ready2rock** 1 month, 2 weeks ago

How can this be? It says you HAVE a VM, meaning it's already created. A cannot be the solution.  
upvoted 1 times

  **jackdbd** Most Recent 6 days, 15 hours ago

I think it's C.

Not A: the question says "You have", so we cannot create a new VM and specify a service account.

Not B: it says "project metadata". I don't think there is such a thing

C: we can specify a directory of service accounts associated with the VM. We need to stop and restart the VM for the service accounts to take effect though. [https://cloud.google.com/compute/docs/metadata/default-metadata-values#vm\\_instance\\_metadata](https://cloud.google.com/compute/docs/metadata/default-metadata-values#vm_instance_metadata)

Not D: we can SSH into the VM and store the service account JSON key in ~/.gcloud/compute-engine-service-account.json, but I doubt it would take effect without calling gcloud compute instances set-service-account [INSTANCE\_NAME] (and stop and restart the VM)  
<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#changeserviceaccountandscopes>  
upvoted 1 times

  **sidharthwader** 1 month ago

A is the answer. If you have doubt create an instance then u can see a box where we can select the service account it comes under IAM and API access  
upvoted 1 times

  **lomari** 1 month ago

you can't change the service account of a running instance, it has to be stopped first, so A is correct.  
upvoted 1 times

  **Larysmith** 1 month ago

You have a Linux VM that must connect to Cloud SQL.  
i guess the above statment is telling us the vm has been created already. so c should work.  
upvoted 1 times

  **skm444** 2 months ago

can someone please clear which is ans A or C ?  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago

A Correct  
upvoted 1 times

  **Finger41** 2 months, 2 weeks ago


Answer A is the path of least resistance and logically more practical.  
upvoted 1 times

  **jcols** 2 months, 2 weeks ago



It's A. It doesn't say the VM is running. <https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#changeserviceaccountandscopes>

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

A is the only correct option...for an existing VM if you want to change to use a SA, need to use either gcloud compute instances set-service-account or from the console simply update the Service Account section after stopping the VM

upvoted 1 times

  **sandeepmann9** 3 months ago

The correct answer is C as the VM is already created. Therefore, option A and D are gone. In the VM edit option, we have Custom Metadata and not Project Metadata. So option C fits well in this case.

upvoted 4 times

  **kopper2019** 3 months, 2 weeks ago

A, when you create the VM

upvoted 1 times

  **[Removed]** 4 months, 1 week ago


A is correct. When creating the VM via the web console, specify the service account under the "Identity and API Access" section.

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct

upvoted 1 times

  **GCP\_Student1** 4 months, 3 weeks ago

A. When creating the VM via the web console, specify the service account under the "Identity and API Access" section.

upvoted 1 times

  **JackGlemins** 5 months ago

I think C is correct: <https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 3 times

  **neerajgoyal** 5 months, 1 week ago

Should we go with A Or C?

upvoted 1 times

## Question #20

## Topic 1

You created an instance of SQL Server 2017 on Compute Engine to test features in the new version. You want to connect to this instance using the fewest number of steps. What should you do?

- A. Install a RDP client on your desktop. Verify that a firewall rule for port 3389 exists.
- B. Install a RDP client in your desktop. Set a Windows username and password in the GCP Console. Use the credentials to log in to the instance.
- C. Set a Windows password in the GCP Console. Verify that a firewall rule for port 22 exists. Click the RDP button in the GCP Console and supply the credentials to log in.
- D. Set a Windows username and password in the GCP Console. Verify that a firewall rule for port 3389 exists. Click the RDP button in the GCP Console, and supply the credentials to log in.**

**Correct Answer: B**

Reference:

<https://medium.com/falafel-software/sql-server-in-the-google-cloud-a17e8a1f11ce>

  **vnxt** Highly Voted 1 year, 2 months ago

I would say B is correct. RDP is enabled by default when you crate a Windows instance (no need to chek on it). Just make sure you install an RDP client ( chrome ext or RDP) and set windows password.

upvoted 19 times

  **pas77** 3 weeks, 2 days ago

Obviously, B is not the answer because you have to install an RDP client which is an extra step. D is the answer because you can connect directly using the RDP button in the GCP console.

upvoted 1 times

  **ankit89** Highly Voted 1 year, 2 months ago

D seems more correct

upvoted 13 times

  **jcloud965** Most Recent 2 weeks, 1 day ago

D is correct because even if the default firewall is open on 3389, it may have been changed.

B : Similar to D but installing a RDP client is an unnecessary step when you can connect from Chrome

C : RDP works with port 3389. SSH is on port 22.

A : generating a Windows password is mandatory to connect

upvoted 1 times

  **hasib125** 2 weeks, 3 days ago

Crosschecked the issue. For me answer is D.

By default after create a SQL Server 2017 on Compute Engine the firewall rule for port 3389 did not exists. I had to put 3389 on firewall. Also got an alert to RDP to put 3389 on firewall. Hence the right answer is : D

upvoted 2 times

  **hasib125** 2 weeks, 3 days ago

Sorry ! I was wrong cause at GCP after connect using RRP it took time to connect on the instance. also got 3389 on firewall by default. Answer must be : B.

upvoted 1 times

  **BalSan** 2 weeks, 5 days ago

B is not right... Nothing is mentioned about if On prem network (desktop) is connected to Google Cloud... So Answer is D

upvoted 1 times

  **ready2rock** 1 month, 2 weeks ago

B is the fewest number of steps....D ensures that you will actually get a successful connection, where B does not. WTF is the answer here?

upvoted 1 times

  **pas77** 3 weeks, 2 days ago

Obviously, B is not the answer because you have to install an RDP client which is an extra step. D is the answer because you can connect directly using the RDP button in the GCP console.

upvoted 1 times

  **AD\_0525** 1 month, 3 weeks ago

As there is no statement to confirm whether the VM using default vpc or custom vpc, so it should be a best practice to check firewall rules, so D should be the answer.

upvoted 2 times

- 🗨️ 👤 **Ashii** 2 months, 1 week ago  
D is a better bit  
<https://cloud.google.com/compute/docs/instances/connecting-to-windows#remote-desktop-connection-app>  
upvoted 2 times
- 🗨️ 👤 **alperenhasan** 2 months, 1 week ago  
i checked gcp console too, on VPC Firewall has got default-allow-rdp rule so it must be B.  
upvoted 1 times
- 🗨️ 👤 **pas77** 3 weeks, 2 days ago  
But you checked. LOL  
D is fewer steps as you don't need to install an RDP client  
upvoted 1 times
- 🗨️ 👤 **ri\_unhou119** 2 months, 2 weeks ago  
I would say D is correct.  
  
[https://cloud.google.com/architecture/chrome-desktop-remote-windows-compute-engine#connect\\_with\\_rdp](https://cloud.google.com/architecture/chrome-desktop-remote-windows-compute-engine#connect_with_rdp)  
upvoted 2 times
- 🗨️ 👤 **Agraved** 2 months, 2 weeks ago  
B is correct, as check is not required.  
upvoted 1 times
- 🗨️ 👤 **pas77** 3 weeks, 2 days ago  
Obviously, B is not the answer because you have to install an RDP client which is an extra step. D is the answer because you can connect directly using the RDP button in the GCP console.  
upvoted 1 times
- 🗨️ 👤 **mcaromit** 2 months, 2 weeks ago  
D is the obvious correct answer  
upvoted 2 times
- 🗨️ 👤 **lxgywil** 2 months, 3 weeks ago  
I think the answer is "B":  
  
#1 - Port 3389 is enabled for SQL Server VM by default (checking it explicitly is yet another step).  
#2 - You need to have an RDP client (see what happens when you click the "RDP button" without it being installed).  
#3 - The username & password have to be created in GCP Console (under instance details) in order to access the instance.  
  
Now, option A is not applicable because of #3, option C is a mess, and option D - because of #1 and #2.  
upvoted 3 times
- 🗨️ 👤 **FunkyTechnician** 2 months, 3 weeks ago  
Message when creating "Note: You must configure the network firewall to open TCP port 3389 to enable RDP access."  
  
See message here ---> [https://miro.medium.com/max/1400/0\\*D96XBChdmlqTmNM.png](https://miro.medium.com/max/1400/0*D96XBChdmlqTmNM.png)  
upvoted 1 times
- 🗨️ 👤 **Rightsaidfred** 2 months, 4 weeks ago  
B is correct. Remember, it is the fewest amount of steps, and RDP Port 3389 is enabled by default.  
upvoted 2 times
- 🗨️ 👤 **ArizonaClassics** 3 months, 3 weeks ago  
D is the most correct answer ( with a deeper look, you will see it comprises both answers from options A & B to give you full troubleshooting tips should your connection fail)  
upvoted 3 times
- 🗨️ 👤 **passnow** 4 months ago  
I have never had the need to go set up any firewall when using rdp so all the noise with it being D is wrong. The answer is clearly the given option  
upvoted 1 times
- 🗨️ 👤 **maddy5835** 3 months, 3 weeks ago  
This is because the default network has predefined firewall-rule set up for RDP connections on port 3389  
upvoted 1 times
- 🗨️ 👤 **kopper2019** 3 months, 2 weeks ago  
that is correct default VPC has ssh and rdp open but what I created a custom VPC as maddy says you need yo to do both steps, even when D has more steps is the way to go to connect to Widnows VMs  
upvoted 4 times



## Question #21



## Topic 1



You have one GCP account running in your default region and zone and another account running in a non-default region and zone. You want to start a new Compute Engine instance in these two Google Cloud Platform accounts using the command line interface. What should you do?



- A. Create two configurations using `gcloud config configurations create [NAME]`. Run `gcloud config configurations activate [NAME]` to switch between accounts when running the commands to start the Compute Engine instances.
- B. Create two configurations using `gcloud config configurations create [NAME]`. Run `gcloud configurations list` to start the Compute Engine instances.
- C. Activate two configurations using `gcloud configurations activate [NAME]`. Run `gcloud config list` to start the Compute Engine instances.
- D. Activate two configurations using `gcloud configurations activate [NAME]`. Run `gcloud configurations list` to start the Compute Engine instances.



Correct Answer: A



  **coldpar** Highly Voted 1 year, 4 months ago  
A is the correct option  
upvoted 21 times



  **leba** Highly Voted 1 year, 3 months ago  
Correct answer is A as you can create different configurations for each account and create compute instances in each account by activating the respective account. Refer GCP documentation - Configurations Create & Activate Options B, C & D are wrong as `gcloud config configurations list` does not help create instances. It would only lists existing named configurations.  
upvoted 14 times



  **YAS007** Most Recent 1 week ago  
A is correct  
<https://medium.com/google-cloud/how-to-use-multiple-accounts-with-gcloud-848fdb53a39a>  
upvoted 1 times



  **yvinisiupacuando** 2 months, 3 weeks ago  
I vote for A  
upvoted 1 times

  **gspb** 3 months, 1 week ago  
Obviously A is the correct option here.  
Even if u dont know this, u can eliminate the remaining options since by common sense u can say that the "list" command doesnt start but list the required things.  
upvoted 1 times

  **gspb** 3 months, 1 week ago  
Obviously A is the correct option here.  
Even if u dont now this u can eliminate he remaining options since by common sense u can say that the "list" command doesnt start but list the required things.  
upvoted 1 times



  **ApaMokus** 3 months, 1 week ago  
A is the correct anwer.  
B is inappropriate as list commend does not start VMs. C&D are incorrect as you cannot have two active configurations at a time.  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago  
A is correct. Create two configurations using `gcloud config configurations create [NAME]`. Run `gcloud config configurations activate [NAME]` to switch between accounts when running the commands to start the Compute Engine instances.  
upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago  
Vote for 'C' .

A & B - out because they mention 'CREATING' the configurations. Why we need to re-create it?

option D is out - command is wrong  
upvoted 1 times



  **Hi2ALL** 4 months, 4 weeks ago  
C is correct answer we need to focus "activate"  
EXAMPLES

To activate an existing named configuration, run:

```
gcloud config configurations activate my_config
```

To list all properties in the activated configuration, run:

```
gcloud config list --all  
upvoted 1 times
```

  **Hi2ALL** 4 months, 4 weeks ago


C is acceptable  
upvoted 1 times

  **JackGlemins** 5 months ago

A is right:[https://cloud.google.com/sdk/docs/configurations#activating\\_a\\_configuration](https://cloud.google.com/sdk/docs/configurations#activating_a_configuration)  
upvoted 2 times

  **Vinod87** 5 months ago

A is the correct option  
upvoted 1 times

  **Sarin** 5 months, 1 week ago

A is the correct option  
upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

A. Create two configurations using `gcloud config configurations create [NAME]`. Run `gcloud config configurations activate [NAME]` to switch between accounts when running the commands to start the Compute Engine instances.  
upvoted 1 times

  **DucSiu** 6 months ago

Correct answer is A  
upvoted 1 times

  **d18mph23** 6 months ago

A is correct option  
`gcloud configurations *` command doesn't exist  
upvoted 1 times

## Question #22

## Topic 1

You significantly changed a complex Deployment Manager template and want to confirm that the dependencies of all defined resources are properly met before committing it to the project. You want the most rapid feedback on your changes. What should you do?

- A. Use granular logging statements within a Deployment Manager template authored in Python.
- B. Monitor activity of the Deployment Manager execution on the Stackdriver Logging page of the GCP Console.
- C. Execute the Deployment Manager template against a separate project with the same configuration, and monitor for failures.
- D. Execute the** Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.

**Correct Answer: D**

Reference:

<https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments>

 **YashBindlish** Highly Voted 1 year, 2 months ago


Correct answer is D as Deployment Manager provides the preview feature to check on what resources would be created  
upvoted 25 times

 **vmart** Most Recent 1 month, 2 weeks ago

yes D is the correct answer  
upvoted 2 times

 **mcaromit** 2 months, 2 weeks ago

D correct option  
upvoted 1 times

 **[Removed]** 4 months, 1 week ago

D is correct. Execute the Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.  
upvoted 2 times

 **EABDAJA** 4 months, 3 weeks ago

D is correct  
upvoted 1 times

 **GCP\_Student1** 5 months, 1 week ago

D. Execute the Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.  
upvoted 1 times

 **DucSiu** 6 months ago

ANSWER IS D  
upvoted 1 times

 **INASR** 6 months ago

D  
(Optional) Preview an updated configuration  
You can preview the update you want to make before committing any changes, with the `gcloud` command-line tool or the API. The Deployment Manager service previews the configuration by expanding the full configuration and creating "shell" resources.  
upvoted 2 times

 **victory108** 6 months, 1 week ago

D - Execute the Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.  
upvoted 1 times

 **swatitame** 8 months, 2 weeks ago

- D. Execute the Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.

upvoted 2 times



 **glam** 9 months, 3 weeks ago

D. Execute the Deployment Manager template using the `--preview` option in the same project, and observe the state of interdependent resources.  
upvoted 3 times



 **prasanu** 9 months, 3 weeks ago

ANSWER IS D

upvoted 1 times

  **JJ\_ME** 9 months, 3 weeks ago

D.  
[https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments#optional\\_preview\\_an\\_updated\\_configuration](https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments#optional_preview_an_updated_configuration)  
upvoted 1 times

  **SSPC** 11 months, 1 week ago

D is correct  
upvoted 2 times

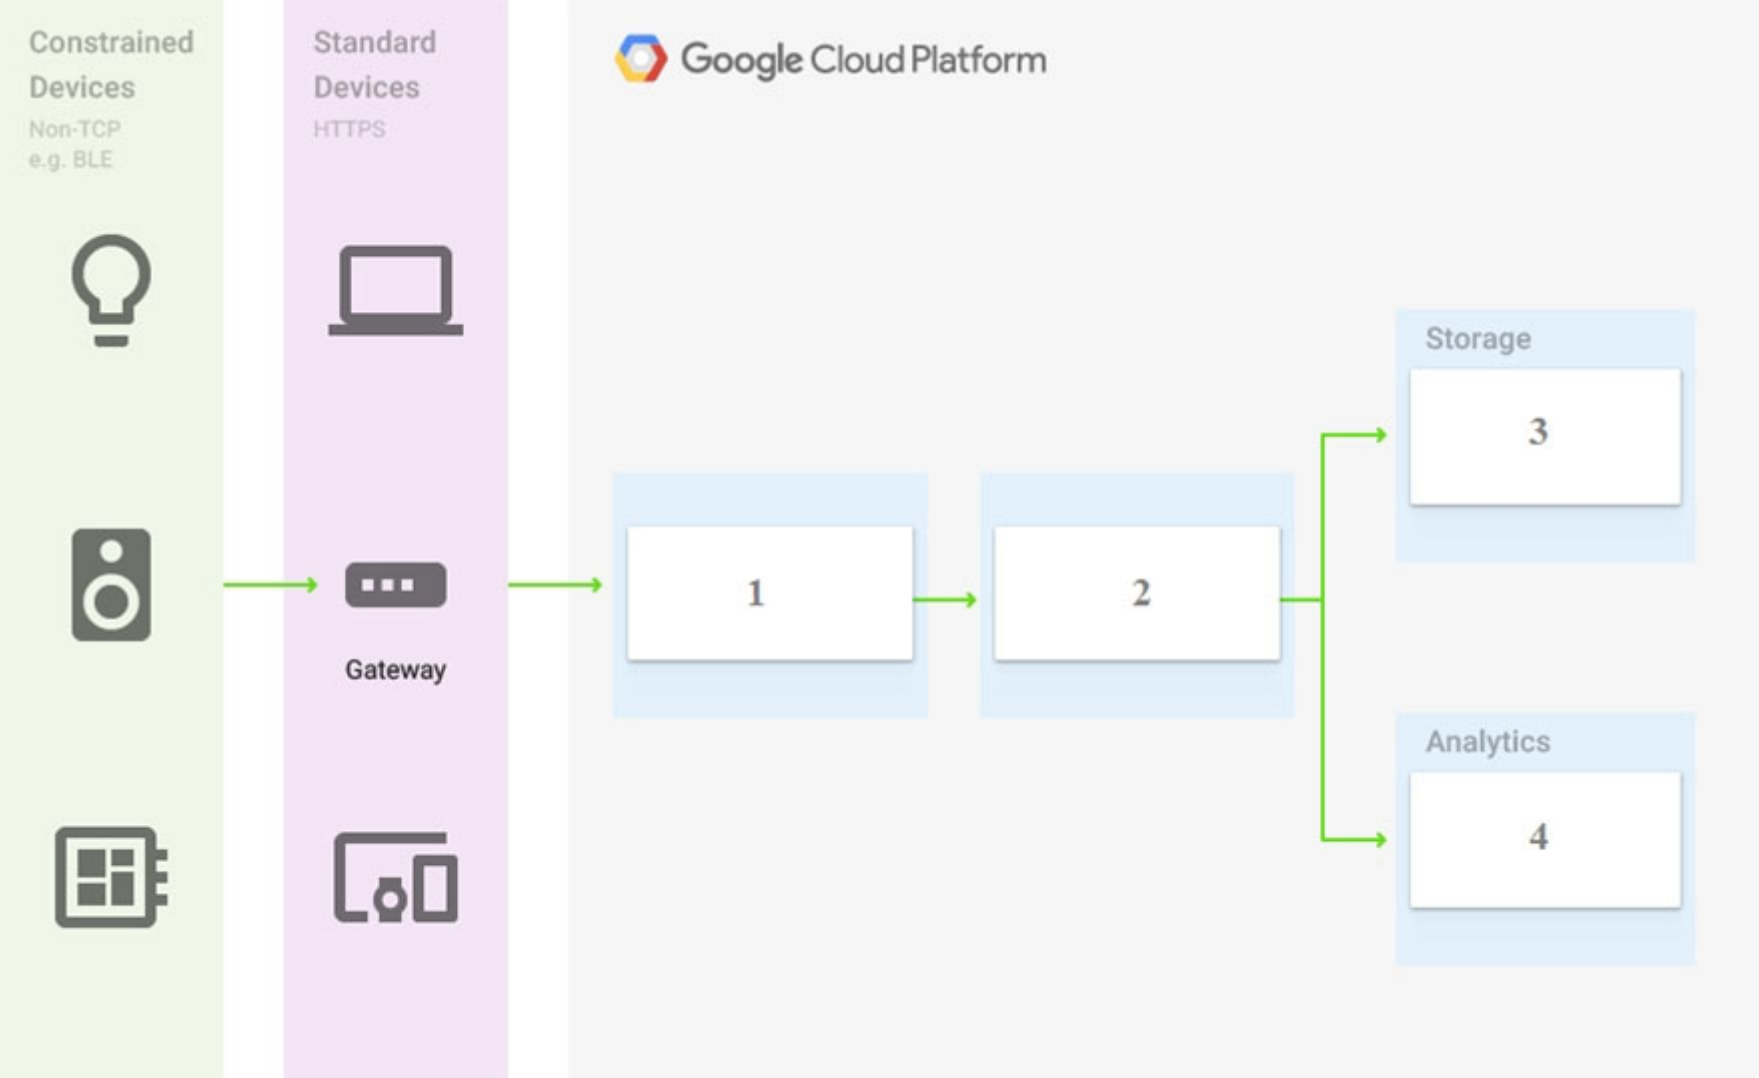
  **professor** 1 year, 1 month ago

<https://cloud.google.com/deployment-manager/docs/configuration/preview-configuration-file>  
upvoted 1 times

  **professor** 1 year, 1 month ago

Ans is D  
upvoted 2 times

You are building a pipeline to process time-series data. Which Google Cloud Platform services should you put in boxes 1,2,3, and 4?



- A. Cloud Pub/Sub, Cloud Dataflow, Cloud Datastore, BigQuery
- B. Firebase Messages, Cloud Pub/Sub, Cloud Spanner, BigQuery
- C. Cloud Pub/Sub, Cloud Storage, BigQuery, Cloud Bigtable
- D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery

**Correct Answer: D**

Reference:  
<https://cloud.google.com/solutions/correlating-time-series-dataflow>

- cloudu01

Highly Voted

1 year, 1 month ago

Without a doubt D.  
Whenever we want to process timeseries data look for BigTable.  
Also you want to perform analytics in Box 4 ..look for BigQuery

Only D provides this option.  
upvoted 28 times
- adedj99

7 months, 4 weeks ago

are we considering bigtable as storage in here , since they expecting some storage

upvoted 1 times
- vlodia

1 year ago

Speaker also looks like an IoT device so D not A

upvoted 2 times
- shuvo\_123

Highly Voted

11 months, 1 week ago

Option D:  
See this reference : <https://cloud.google.com/blog/products/databases/getting-started-with-time-series-trend-predictions-using-gcp>

upvoted 6 times
- Agraved

Most Recent

2 months, 2 weeks ago

awesome question with graphical representation, clearing a lot concepts in one Diagram. Answer has to be D

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

D looks most apt among the given options & considering streaming data

upvoted 1 times



- 🗨️ 👤 **yvinisiupacuando** 2 months, 3 weeks ago  
D is the right choice here.  
upvoted 1 times
- 🗨️ 👤 **kopper2019** 3 months, 2 weeks ago  
just by process of elimination #3 can only be something that ingest data like DynamoDB/BigTable, that way I did the association since I know AWS, so D  
upvoted 1 times
- 🗨️ 👤 **[Removed]** 4 months, 1 week ago  
D is correct. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery  
upvoted 1 times
- 🗨️ 👤 **sesteban** 4 months, 2 weeks ago  
Option A: <https://thecloudgirl.dev/images/analytics.png>  
upvoted 1 times
- 🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago  
D is correct  
upvoted 1 times
- 🗨️ 👤 **cloud\_hunk** 5 months ago  
D is correct:  
Based on the link provided and question ask: we are dealing with time series data,  
"You can store the values by branching with Dataflow pushing the computed values into storage, for example BigQuery or Cloud Bigtable, whenever you need a complete set of values based on the perfect rectangles" - from GCP documentation  
upvoted 1 times
- 🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago  
D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery  
upvoted 1 times
- 🗨️ 👤 **INASR** 6 months ago  
D for sure. DataFlow is required as batch process for realtime data.  
upvoted 3 times
- 🗨️ 👤 **swatititame** 8 months, 2 weeks ago  
• D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery  
upvoted 1 times
- 🗨️ 👤 **adedj99** 7 months, 4 weeks ago  
are we considering bigtable as storage in here , since they expecting some storage  
upvoted 1 times
- 🗨️ 👤 **alpharomeo9** 7 months, 2 weeks ago  
The question is about the time-series data. So BigTable suits the best for time-series database storage. So option D is the right choice  
upvoted 2 times
- 🗨️ 👤 **glam** 9 months, 3 weeks ago  
D. Cloud Pub/Sub, Cloud Dataflow, Cloud Bigtable, BigQuery  
upvoted 2 times

## Question #24

## Topic 1

You have a project for your App Engine application that serves a development environment. The required testing has succeeded and you want to create a new project to serve as your production environment. What should you do?

- A. Use gcloud to create the new project, and then deploy your application to the new project.
- B. Use gcloud to create the new project and to copy the deployed application to the new project.
- C. Create a Deployment Manager configuration file that copies the current App Engine deployment into a new project.
- D. Deploy your application again using gcloud and specify the project parameter with the new project name to create the new project.

Correct Answer: A

  **coldpar** Highly Voted 1 year, 4 months ago

Correct is A.

Option B is wrong as the option to use gcloud app cp does not exist.

Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage

upvoted 26 times

  **manu202020** 1 year ago

you're missing one thing. D isn't about using deployment manager to copy the configuration, instead, using the configuration file to copy the configuration from test project.

upvoted 3 times

  **leba** Highly Voted 1 year, 3 months ago

Correct answer is A as gcloud can be used to create a new project and the gcloud app deploy can point to the new project. Refer GCP documentation - GCloud App Deploy.

Option B is wrong as the option to use gcloud app cp does not exist

.Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yaml

Option D is wrong as gcloud app deploy would not create a new project. The project should be created before usage.

upvoted 6 times

  **vmart** Most Recent 1 month, 2 weeks ago

A is correct answer

upvoted 1 times

  **Umesh09** 1 month, 3 weeks ago

I would have gone with D if the project was an existing project then option D is the best choice since it would have been one step process. Now since the new project has to be created Option A seems to be better choice with the two step process. First create a new project and then deploy your application hence my choice Option A

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

A is correct option

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

A. Use gcloud to create the new project, and then deploy your application to the new project.

The only option that meet the requirements of the questions.

upvoted 1 times

  **DucSiu** 6 months ago

Correct is A.

upvoted 1 times

  **INASR** 6 months ago

I will go with A

upvoted 1 times

  **EABDAJA** 6 months ago

A is Correct

upvoted 2 times

  **Vineetkumar** 7 months, 1 week ago

Does Google team provides answer or should we go with majority votes as in option A for this

upvoted 1 times

  **Bhagirathi** 7 months, 2 weeks ago


A is simple and best !

upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

very confusing - you guys have different opinions and understanding...not sure, whether A ? C ? D ?

upvoted 1 times

  **alpharomeo9** 7 months, 2 weeks ago

most of the people say A and some of them even provided GCP documentation links. So I would go with A.

upvoted 2 times

  **cecocenov** 8 months, 1 week ago

Correct answer is A as gcloud can be used to create a new project and the gcloud app deploy can point to the new project.Refer GCP documentation - GCloud App Deploy--project=&lt;var>PROJECT\_ID&lt;/var>The Google Cloud Platform project name to use for this invocation. If omitted, then the current project is assumed; the current project can be listed using gcloud config list --format='text(core.project)' and can be set using gcloud config set project PROJECTID.--project and its fallback core/project property play two roles in the invocation. It specifies the project of the resource to operate on. It also specifies the project for API enablement check, quota, and billing. To specify a different project for quota and billing, use --billing-project or billing/quota\_project property.Option B is wrong as the option to use gcloud app cp does not exist.Option C is wrong as Deployment Manager does not copy the application, but allows you to specify all the resources needed for your application in a declarative format using yamlOption D is wrong as gcloud app deploy would not create a new project. The project should be created before usage.

upvoted 2 times

  **satish2010a** 8 months, 2 weeks ago

Agreed with A.. Just need to create the new project and in console make sure set the config to new project.

cloud projects create <new-project-name>

gcloud deployment-manager deployments create <deployment-name>

upvoted 3 times

  **swatititame** 8 months, 2 weeks ago

- A. Use gcloud to create the new project, and then deploy your application to the new project.

upvoted 2 times



You need to configure IAM access audit logging in BigQuery for external auditors. You want to follow Google-recommended practices. What should you do?


- A. Add the auditors group to the `logging.viewer` and `bigQuery.dataViewer` predefined IAM roles.
- B. Add the auditors group to two new custom IAM roles.
- C. Add the auditor user accounts to the `logging.viewer` and `bigQuery.dataViewer` predefined IAM roles.
- D. Add the auditor user accounts to two new custom IAM roles.

**Correct Answer:** C



Reference:



<https://cloud.google.com/iam/docs/roles-audit-logging>



-   **coldpar**



Highly Voted 


 1 year, 4 months ago

Correct is A.  
As per google best practices it is recommended to use predefined roles and create groups to control access to multiple users with same responsibility  
upvoted 37 times
-   **droogie** 1 year ago



You assume Auditors Group = External Auditors only. Auditors Group may contain both Internal and External Auditors.  
upvoted 3 times
-   **adeice** 4 months, 1 week ago


I can create External group and Internal group Auditors  
upvoted 2 times
-   **robor97** 8 months ago

The question literally says - External Auditors  
upvoted 6 times
-   **JavierCorrea**



Highly Voted 



 1 year ago



Correct answer is A as per:  
[https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors)  
upvoted 16 times
-   **Student7**



Most Recent 



 2 weeks, 1 day ago

A. Add the auditors group to the `'logging.viewer'` and `'bigQuery.dataViewer'` predefined IAM roles.  
upvoted 1 times
-   **Agraved** 2 months, 2 weeks ago

Answer has to be A, as per Google Best Practice and practically make sense. Would be useless to add individuals to roles one by one  
upvoted 1 times
-   **mcaromit** 2 months, 2 weeks ago


A seems correct, though the question is not entirely clear to me  
upvoted 1 times
-   **manish\_109** 3 months ago

is every 2 of 5 qstn in this forum is wrong .i have to check for every qstn that it is correct or not :(   
upvoted 3 times
-   **here2help** 3 months ago

It's crazy to me how who ever answered these, answers them incorrectly and then puts a link up that directly contradicts themselves.  
The correct answer is A. Google always recommends adding users to a group and then giving the group access. Additionally, Google recommends to use predefined roles since they have been well thought out when created, and are there to save companies the hassle of having to know what granular access every single person needs.  
upvoted 6 times
-   **max94** 3 months ago

Correct is A. And here is explanation why its A and not B:  
Predefined roles are created and maintained by Google. Google automatically updates their permissions as necessary, such as when Google Cloud adds new features or services.  
  
Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.  
<https://cloud.google.com/iam/docs/understanding-custom-roles>

upvoted 1 times

  **r1ck** 3 months, 2 weeks ago



A : always use groups as possible

upvoted 2 times

  **kopper2019** 3 months, 2 weeks ago

A, [https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors)

upvoted 1 times

  **pca2b** 4 months, 1 week ago


A: Use groups when possible, and  
Use predefined roles when possible

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. Add the auditors group to the `logging.viewer` and `bigQuery.dataViewer` predefined IAM roles.

upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

Vote for C



upvoted 2 times

  **tavva\_prudhvi** 4 months, 1 week ago

Refer to this article, it clearly mentioned option A, and the best-recommended practices are to give permissions to groups.

[https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors)

upvoted 1 times

  **tarang3** 4 months, 2 weeks ago

Answer is C

upvoted 1 times

  **tarang3** 4 months ago

@tavva\_prudhvi i agree upon checking the link provided by you that the answer is "A"


upvoted 2 times

  **tavva\_prudhvi** 4 months, 1 week ago

Refer to this article, it clearly mentioned option A, and the best-recommended practices are to give permissions to groups.

[https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors)

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct

upvoted 1 times

  **Jacky\_YO** 5 months ago

My Answer : C

[https://cloud.google.com/iam/docs/job-functions/auditing#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/job-functions/auditing#scenario_external_auditors)

Scenario: External auditors

Role: Logging.viewer , Resource : Organization , Member : Dashboard Service Account , Description :The logging.viewer role permits the service account to read the Admin Activity logs in Cloud Logging.

ex:

```
{
  "bindings": [{
    "role": "roles/logging.viewer",
    "members": [
      "serviceAccount:prod-project-dashboard@admin-resources.iam.gserviceaccount.com"
    ]
  }]
}
```

upvoted 1 times

  **pas77** 3 weeks, 2 days ago

You really didn't read the link you've provided.

It clearly shows to add users to a group. Coinciding with option A.


upvoted 1 times

  **ShakthiGCP** 5 months ago

Ans: A .


upvoted 1 times

- A. Create a service account with an access scope. Use the access scope `https://www.googleapis.com/auth/devstorage.write_only`.
- B. Create a service account with an access scope. Use the access scope `https://www.googleapis.com/auth/cloud-platform`.
- C. Create a service account and add it to the IAM role `storage.objectCreator` for that bucket.**
- D. Create a service account and add it to the IAM role `storage.objectAdmin` for that bucket.

 **coldpar** Highly Voted 1 year, 4 months ago

As per as the least privilege recommended by google, C is the correct Option, A is incorrect because the scope doesnt exist. B incorrect because it will give him full of control



upvoted 26 times


 **robor97** 7 months, 4 weeks ago  
The scope does exist - <https://download.huihoo.com/google/gdgdevkit/DVD1/developers.google.com/compute/docs/api/how-tos/authorization.html>  
upvoted 1 times

Ans B and C


Role from GCP Console:  
ID = roles/storage.objectCreator  
Role launch stage = General Availability  
Description = Access to create objects in GCS.

3 assigned permissions  
resourceManager.projects.get  
resourceManager.projects.list  
storage.objects.create  
upvoted 10 times

  **YAS007** Most Recent 1 week ago  
following google's recommendation : "least privilege"  
<https://cloud.google.com/storage/docs/access-control/iam-roles>  
upvoted 1 times

 **ChrisBelt5** 1 week ago

Answer is C. D also works but doesn't meet "least privilege" Google's recommendation.


 upvoted 1 times

Should be a combination of B+C, but individually should be C. Because scopes provide least privilege, but are not mandatory it is a google best practice.

  **mcaramit** 2 months, 2 weeks ago

C is correct, access scopes are associated with a VM instance & not the service account attached with the instance

upvoted 2 times

 **ri\_unhou119** 2 months, 3 weeks ago

Option C is correct.

- A. [devstorage.write\_only] option does not exist, only devstorage.read\_only.
- B. The authority is too large.
- D. Excessive authority. In addition to write permissions, it also has permissions such as delete.

<https://cloud.google.com/compute/docs/access/service-accounts?hl=ja#accessscopesiam>

upvoted 5 times

<https://www.examtopics.com/exams/google/associate-cloud-engineer/custom-view/>

The correct answer is C. D gives the user too much access and violates the least privilege principle. The benefit to using access scopes is in the case where you already created a service account for a VM and it is running and need to add access. This gives you the benefit of not needing to rebuild the VM with a newly defined service account. In the case of the question you have not created the VM yet, so no need to mess with access scopes. Just create the SA with the right access and assign it to the VM you'll create.

upvoted 1 times

🗲️ 👤 **SaPa2021** 3 months ago

Answer is B

Grant the instance the <https://www.googleapis.com/auth/cloud-platform> scope to allow full access to all Google Cloud APIs, so that the IAM permissions of the instance are completely determined by the IAM roles that you granted to the instance's service account. The service account can only execute API methods that are allowed by both the access scope and the service account's specific IAM roles.

upvoted 1 times

🗲️ 👤 **r1ck** 3 months, 2 weeks ago

Answer.. A

upvoted 2 times

🗲️ 👤 **maddy5835** 3 months, 3 weeks ago

C is the correct answer as Google Cloud now uses IAM, not access scopes, to specify permissions for Compute Engine instances. Moreover .admin role does not follow the google best practice hence D is incorrect

upvoted 2 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago

C is correct. Create a service account and add it to the IAM role `roles/storage.objectCreator` for that bucket.

upvoted 2 times

🗲️ 👤 **HSKR** 4 months, 1 week ago

correct answer is C

It asks just to provide permission to write on to the object, so roles/storage.objectCreator meet the requirement

Refer - <https://cloud.google.com/storage/docs/access-control/iam-role>

upvoted 2 times

🗲️ 👤 **sumanshu** 4 months, 2 weeks ago

Vote for C

upvoted 2 times

🗲️ 👤 **drizzydroo** 4 months, 2 weeks ago

After some testing...When you CREATE a service account you DO NOT have the ability to set individual access scopes. You must give your service account the proper IAM permissions to perform operations. I would choose (C) as admin is outside the range of permissions necessary in this instance. Please tell me if you disagree with this.

upvoted 3 times

🗲️ 👤 **drizzydroo** 4 months, 2 weeks ago

Can someone please enlighten me on where you find these addresses for the Service account Access scopes? Example:

[https://www.googleapis.com/auth/devstorage.write\\_only](https://www.googleapis.com/auth/devstorage.write_only).

I understand that you can set access scopes for different services for the service account such as write only to cloud storage but I don't know how to identify these scopes based on the addresses and I would like to understand that.

upvoted 1 times

🗲️ 👤 **Hi2ALL** 4 months, 4 weeks ago

C should be the correct answer

<https://cloud.google.com/storage/docs/access-control/iam-roles>

Storage Object Creator (roles/storage.objectCreator) Allows users to create objects. Does not give permission to view, delete, or replace objects.

upvoted 1 times



You have sensitive data stored in three Cloud Storage buckets and have enabled data access logging. You want to verify activities for a particular user for these buckets, using the fewest possible steps. You need to verify the addition of metadata labels and which files have been viewed from those buckets. What should you do?

- A. Using the GCP Console, filter the Activity log to view the information.
- B. Using the GCP Console, filter the Stackdriver log to view the information.
- C. View the bucket in the Storage section of the GCP Console.
- D. Create a trace in Stackdriver to view the information.

Correct Answer: A

- iamgcp

Highly Voted

1 year, 2 months ago

A is correct. As mentioned in the question, data access logging is enabled. I tried to download a file from a bucket and was able to view this information in Activity tab in console

upvoted 14 times
- vito9630

1 year, 2 months ago

data access logging don't provide information about addition of metada, so B is correct

upvoted 12 times
- eliteone11

Highly Voted

8 months ago

Answer is A. Activity log does indeed show information about metadata.  
I agree with Eshkrkrkr based on <https://cloud.google.com/storage/docs/audit-logs> Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.

upvoted 7 times
- saikatdatta

Most Recent

1 month, 1 week ago

A is correct

upvoted 1 times
- csdev

1 month, 2 weeks ago

A is correct

upvoted 1 times
- Agravaded

2 months, 2 weeks ago

A is correct, as logging is enabled and Data Access logs provides the ask as below:  
ADMIN\_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.  
DATA\_READ: Entries for operations that read an object.  
DATA\_WRITE: Entries for operations that create or modify an object.  
Ref: <https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times
- Agravaded

2 months, 2 weeks ago

A is correct, as logging is enabled and Data Access logs provides the ask as below:  
ADMIN\_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.  
DATA\_READ: Entries for operations that read an object.  
DATA\_WRITE: Entries for operations that create or modify an object.  
Ref: <https://cloud.google.com/storage/docs/audit-logging>

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is the correct option Logs Explorer lets you search activities for a particular user account

upvoted 2 times
- xkeg9a2oi4ay

3 months, 2 weeks ago

Okay so is it A or B now?  
Looks like A will not contain all of the wanted info ?

upvoted 1 times
- pca2b

4 months, 1 week ago

B: [https://cloud.google.com/logging/docs/audit#viewing\\_audit\\_logs](https://cloud.google.com/logging/docs/audit#viewing_audit_logs)  
  
!A: <https://cloud.google.com/logging/docs/audit#view-activity>  
The actual audit log entries might contain more information than appears on the Activity page.  
  
In the Activity page, where the identity performing logged actions is redacted from the audit log entry, User (anonymized) is displayed.

upvoted 1 times



🗨️ 👤 **[Removed]** 4 months, 1 week ago

B is correct. Using the GCP Console, filter the Stackdriver log to view the information.

upvoted 1 times

🗨️ 👤 **user843983409** 4 months, 1 week ago

vote for A :

sample logs show labels - [https://cloud.google.com/compute/docs/logging/activity-logs#sample\\_log\\_entry](https://cloud.google.com/compute/docs/logging/activity-logs#sample_log_entry)

upvoted 1 times

🗨️ 👤 **ready2rock** 1 month, 2 weeks ago

You've linked a compute engine document. This is regarding storage bucket and audit logs. This link is incorrect

upvoted 1 times

🗨️ 👤 **drizzydroo** 4 months, 2 weeks ago

I was able to see in the activity log which files had been retrieved by which user... i am not able to see the label that I attached to one of the buckets.... can anyone counter this?

upvoted 1 times

🗨️ 👤 **tavva\_prudhvi** 4 months ago

Yes, even I faced the same issue, I couldn't see the label. So, we can choose B.

upvoted 2 times

🗨️ 👤 **neerajgoyal** 4 months, 2 weeks ago

A is the correct answer

<https://cloud.google.com/logging/docs/audit#:~:text=Cloud%20Audit%20Logs%20provides%20the,System%20Event%20audit%20logs>

upvoted 1 times

🗨️ 👤 **MuraliKrishna** 5 months ago

A seems correct

upvoted 1 times

🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago

A. Using the GCP Console, filter the Activity log to view the information.

upvoted 1 times

🗨️ 👤 **Joh3o8an** 6 months ago

I think the answer is A because activity logging refers to app user's request. Activity log and Audit log are both stackdriver logs but while Activity log make reference to user's requests, Audit log makes reference to Admin's requests.

And the question says you want to verify the activity of a particular user on the buckets. So A is a good bet.

upvoted 1 times

🗨️ 👤 **INASR** 6 months ago

it is A, quick access to activities

upvoted 1 times

You are the project owner of a GCP project and want to delegate control to colleagues to manage buckets and files in Cloud Storage. You want to follow Google- recommended practices. Which IAM roles should you grant your colleagues?

- A. Project Editor
- B. Storage Admin
- C. Storage Object Admin
- D. Storage Object Creator

Correct Answer: B

- ESP\_SAP

Highly Voted

11 months, 1 week ago

Correct Answer is (B):

Storage Admin (roles/storage.admin) Grants full control of buckets and objects.  
When applied to an individual bucket, control applies only to the specified bucket and objects within the bucket.

firebase.projects.get  
resourcemanager.projects.get  
resourcemanager.projects.list  
storage.buckets.\*  
storage.objects.\*  
upvoted 20 times
- Agents89

Highly Voted

1 year, 3 months ago

B is correct  
upvoted 11 times
- Student7

Most Recent

2 weeks, 1 day ago

B. Storage Admin  
upvoted 1 times
- Furqon

1 month, 2 weeks ago

B is the correct one  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- [Removed]

4 months, 1 week ago

B is correct. Storage Admin  
upvoted 2 times
- EABDAJA

4 months, 3 weeks ago

Ans is B  
upvoted 1 times
- JackGlemins

5 months ago

B: <https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

B. Storage Admin  
upvoted 1 times
- neerajgoyal

5 months, 1 week ago



B should be promising  
upvoted 1 times
- victory108



5 months, 3 weeks ago



B - Storage Admin  
upvoted 1 times
- INASR



6 months ago

B is correct  
upvoted 1 times

  **nherrerab** 6 months, 2 weeks ago  
B is correct.  
upvoted 1 times

  **swatititame** 8 months, 2 weeks ago  
B. Storage Admin  
upvoted 1 times

  **glam** 9 months, 3 weeks ago  
B. Storage Admin  
upvoted 2 times

  **JJ\_ME** 9 months, 3 weeks ago  
B.  
Storage Admin (roles/storage.admin) - Grants full control of buckets and objects.  
<https://cloud.google.com/storage/docs/access-control/iam-roles>  
upvoted 1 times

You have an object in a Cloud Storage bucket that you want to share with an external company. The object contains sensitive data. You want access to the content to be removed after four hours. The external company does not have a Google account to which you can grant specific user-based access privileges. You want to use the most secure method that requires the fewest steps. What should you do?

- A. Create a signed URL with a four-hour expiration and share the URL with the company.
- B. Set object access to `public` and use object lifecycle management to remove the object after four hours.
- C. Configure the storage bucket as a static website and furnish the object's URL to the company. Delete the object from the storage bucket after four hours.
- D. Create a new Cloud Storage bucket specifically for the external company to access. Copy the object to that bucket. Delete the bucket after four hours have passed.

Correct Answer: A

- JJ\_ME

Highly Voted

9 months, 3 weeks ago

A.  
Signed URLs are used to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account.  
<https://cloud.google.com/storage/docs/access-control/signed-urls>  
upvoted 16 times
- Agents89

Highly Voted

1 year, 3 months ago

A is correct  
upvoted 8 times
- mcaromit

Most Recent

2 months, 2 weeks ago

A correct answer  
upvoted 1 times
- gspb

3 months, 1 week ago

option A is correct  
upvoted 1 times
- [Removed]

4 months, 1 week ago

A is correct. Create a signed URL with a four-hour expiration and share the URL with the company.  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

A. Create a signed URL with a four-hour expiration and share the URL with the company.  
upvoted 1 times
- nliaustemac

5 months ago

A. Is the correct answer:  
Create a signed URL with a four-hour expiration and share the URL with the company  
upvoted 1 times
- nliaustemac

5 months ago

A. Is the correct answer:  
Create a signed URL with a four-hour expiration and share the URL with the company.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

A. Create a signed URL with a four-hour expiration and share the URL with the company.  
upvoted 1 times
- Gohs

5 months, 4 weeks ago


Ans: A  
upvoted 1 times
- DucSiu

6 months ago


A is correct  
upvoted 1 times
- INASR

6 months ago


A for sure  
upvoted 1 times

-  **swatititame** 8 months, 2 weeks ago

- A. Create a signed URL with a four-hour expiration and share the URL with the company.

upvoted 1 times
-  **glam** 9 months, 3 weeks ago

A. Create a signed URL with a four-hour expiration and share the URL with the company.

upvoted 2 times
-  **SSPC** 11 months, 1 week ago

A is correct

upvoted 1 times

You are creating a Google Kubernetes Engine (GKE) cluster with a cluster autoscaler feature enabled. You need to make sure that each node of the cluster will run a monitoring pod that sends container metrics to a third-party monitoring solution. What should you do?

- A. Deploy the monitoring pod in a StatefulSet object.
- B. Deploy the monitoring pod in a DaemonSet object.
- C. Reference the monitoring pod in a Deployment object.
- D. Reference the monitoring pod in a cluster initializer at the GKE cluster creation time.

Correct Answer: B

- Agents89

Highly Voted

1 year, 3 months ago

B is correct  
upvoted 16 times
- JackGlemins

Highly Voted

5 months ago

B is right: <https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>  
Some typical uses of a DaemonSet are:  
  
running a cluster storage daemon on every node  
running a logs collection daemon on every node  
running a node monitoring daemon on every node  
upvoted 8 times
- mcaromit

Most Recent

2 months, 2 weeks ago

B is correct answer  
upvoted 1 times
- max94

3 months, 1 week ago

B is correct [https://cloud.google.com/kubernetes-engine/docs/concepts/daemonset#usage\\_patterns](https://cloud.google.com/kubernetes-engine/docs/concepts/daemonset#usage_patterns)  
upvoted 1 times
- [Removed]

4 months, 1 week ago

B is correct. Deploy the monitoring pod in a DaemonSet object.  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

B is correct  
upvoted 1 times
- nliaustemac

5 months ago

B is the correct answer:  
Deploy the monitoring pod in a DaemonSet object.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

B. Deploy the monitoring pod in a DaemonSet object.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

<https://kubernetes.io/docs/concepts/workloads/controllers/daemonset/>  
upvoted 1 times
- DucSiu

6 months ago

B is correct  
upvoted 1 times
- INASR

6 months ago

B is correct  
upvoted 1 times
- Dineshvon

7 months, 3 weeks ago

B is correct  
upvoted 1 times
- robor97

7 months, 4 weeks ago

Use Daemonset to export all container metrics - <https://cloud.google.com/stackdriver/docs/solutions/gke/legacy-stackdriver/monitoring>

upvoted 4 times

 **glam** 9 months, 3 weeks ago

B. Deploy the monitoring pod in a DaemonSet object.

upvoted 3 times

 **professor** 1 year, 1 month ago

Ans is B

Daemonset runs on every node

upvoted 6 times

You want to send and consume Cloud Pub/Sub messages from your App Engine application. The Cloud Pub/Sub API is currently disabled. You will use a service account to authenticate your application to the API. You want to make sure your application can use Cloud Pub/Sub. What should you do?

- A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.
- B. Rely on the automatic enablement of the Cloud Pub/Sub API when the Service Account accesses it.
- C. Use Deployment Manager to deploy your application. Rely on the automatic enablement of all APIs used by the application being deployed.
- D. Grant the App Engine Default service account the role of Cloud Pub/Sub Admin. Have your application enable the API on the first connection to Cloud Pub/ Sub.

Correct Answer: A

- Bharathy**

Highly Voted

1 year, 2 months ago

We need to enable the pub/sub API, if we are going to use it in your project... then APP engine can able to access it with required ServiceAccount

upvoted 13 times
- ESP\_SAP**

Highly Voted

11 months, 1 week ago

Correct Answer is (A)

Quickstart: using the Google Cloud Console  
This page shows you how to perform basic tasks in Pub/Sub using the Google Cloud Console.

Note: If you are new to Pub/Sub, we recommend that you start with the interactive tutorial.  
Before you begin  
Set up a Cloud Console project.  
Set up a project

Click to:

Create or select a project.  
Enable the Pub/Sub API for that project.  
You can view and manage these resources at any time in the Cloud Console.

Install and initialize the Cloud SDK.  
Note: You can run the gcloud tool in the Cloud Console without installing the Cloud SDK. To run the gcloud tool in the Cloud Console, use Cloud Shell .

<https://cloud.google.com/pubsub/docs/quickstart-console>

upvoted 12 times
- Kaushik\_Mallick**

Most Recent

1 week, 4 days ago

Does anyone has latest questions & answers. Also if you guide me the reading material/blog/portal to pass this exam.

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

As a first step A is the correct answer

upvoted 2 times
- gspb**

3 months, 1 week ago

In the question - "You want to make sure your application can use Cloud Pub/Sub"

Since we want to make sure that our application uses pub/sub, enabling it manually is the best option.

Hence, option A is correct

upvoted 1 times
- ggalvez**

3 months, 2 weeks ago

Es A. Nunca esta habilitada la API PUB/SUB

upvoted 2 times
- [Removed]**

4 months, 1 week ago

A is correct. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

upvoted 1 times
- nliaustemac**

5 months ago

A. is the correct answer:

Enable the Cloud Pub/Sub API in the API Library on the GCP Console.

upvoted 1 times



- 🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago  
A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.  
upvoted 2 times
- 🗲️ 👤 **iPablo26** 5 months, 1 week ago  
Enable the Cloud Pub/Sub Api  
upvoted 1 times
- 🗲️ 👤 **Gohs** 5 months, 4 weeks ago  
Not all API are enabled by default so you have to manually enable the API you'll need.  
upvoted 2 times
- 🗲️ 👤 **INASR** 6 months ago  
Sure A , you need to enable API for any service you want to use in google  
upvoted 1 times
- 🗲️ 👤 **Dineshvon** 7 months, 3 weeks ago  
A is correct  
upvoted 1 times
- 🗲️ 👤 **swatititame** 8 months, 2 weeks ago  
• A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.  
upvoted 1 times
- 🗲️ 👤 **glam** 9 months, 3 weeks ago  
A. Enable the Cloud Pub/Sub API in the API Library on the GCP Console.  
upvoted 2 times
- 🗲️ 👤 **SSPC** 11 months, 1 week ago  
A is correct  
upvoted 1 times
- 🗲️ 👤 **[Removed]** 10 months, 1 week ago  
Only Compute Engine API enable when you access, so no other will enable when you access.. Other enable when you use console.  
upvoted 2 times
- 🗲️ 👤 **[Removed]** 10 months, 1 week ago  
Other enable when you use console or you have to manually enable them.  
upvoted 1 times
- 🗲️ 👤 **[Removed]** 10 months, 1 week ago  
So only A remains.  
upvoted 1 times

## Question #32

## Topic 1

You need to monitor resources that are distributed over different projects in Google Cloud Platform. You want to consolidate reporting under the same Stackdriver Monitoring dashboard. What should you do?

- A. Use Shared VPC to connect all projects, and link Stackdriver to one of the projects.
- B. For each project, create a Stackdriver account. In each project, create a service account for that project and grant it the role of Stackdriver Account Editor in all other projects.
- C. **Configure** a single Stackdriver account, and link all projects to the same account.
- D. Configure a single Stackdriver account for one of the projects. In Stackdriver, create a Group and add the other project names as criteria for that Group.

**Correct Answer: D**

🗲️ 👤 **coldpar** Highly Voted 1 year, 4 months ago  
C is correct not D  
upvoted 22 times

🗲️ 👤 **sahedge** Highly Voted 7 months, 1 week ago  
First of all D is incorrect, Groups are used to define alerts on set of resources(such as VM instances, databases, and load balancers). FYI tried adding Two projects into a group it did not allowed me as the "AND"/"OR" criteria for the group failed with this combination of resources.  
  
C is correct because,  
When you initially click on Monitoring(Stackdriver Monitoring) it creates a workspace(a stackdriver account) linked to the ACTIVE(CURRENT) Project from which it was clicked.  
  
Now if you change the project and again click onto Monitoring it would create an another workspace(a stackdriver account) linked to the changed ACTIVE(CURRENT) Project, we don't want this as this would not consolidate our result into a single dashboard(workspace/stackdriver account).  
  
If you have accidentally created two diff workspaces merge them under Monitoring > Settings > Merge Workspaces > MERGE.  
  
If we have only one workspace and two projects we can simply add other GCP Project under Monitoring > Settings > GCP Projects > Add GCP Projects.  
  
In both of these cases we did not create a GROUP, we just linked GCP Project to the workspace(stackdriver account).  
upvoted 13 times

🗲️ 👤 **vmart** Most Recent 1 month, 2 weeks ago  
C is correct  
upvoted 2 times

🗲️ 👤 **contaexamtopics** 2 months, 1 week ago  
C or D?  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago  
C...create a Workspace for a host project & add the other GCP projects or AWS accounts to monitor  
upvoted 1 times

🗲️ 👤 **r1ck** 3 months, 2 weeks ago  
Isn't the question objective to give a simple, plausible answer?  
C  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago  
C is correct. Configure a single Stackdriver account, and link all projects to the same account.  
upvoted 1 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago  
answer is C  
upvoted 1 times

🗲️ 👤 **shayanahmed** 4 months, 3 weeks ago  
C is correct  
upvoted 1 times



🗲️ 👤 **kipsnak** 4 months, 3 weeks ago

I fear that no answer is OK nowadays for this question, since stackdriver solution does not exist anymore and has been replaced by Google Cloud monitoring.

So, to aggregate more than one project monitoring to the same view, you need to use multi-projects workspace (or merge the automatically created or single project workspace into one workspace).

Anyway, this question may no longer be accurate, otherwise it gonna be quite confusing.

upvoted 3 times

  **Hi2ALL** 4 months, 4 weeks ago

C is correct imo

GCP Console> Monitoring>Settings>at the bottom Merge Workspaces to consolidate all projects

upvoted 3 times

  **Vinod87** 5 months ago

C is correct not D



upvoted 1 times

  **nliaustemac** 5 months ago

C. is the correct answer:

Configure a single Stackdriver account, and link all projects to the same account.

upvoted 1 times

  **nitinz** 5 months, 1 week ago

Correct Answer is C, this is from official GCP ACE Study guide, chapter 18.

If you want to monitor multiple projects in a Workspace, you can optionally select other projects

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

C. Configure a single Stackdriver account, and link all projects to the same account.

upvoted 1 times

  **DucSiu** 6 months ago

C. Configure a single Stackdriver account, and link all projects to the same account.

upvoted 1 times

  **INASR** 6 months ago

C is sure correct , groups on Stackdriver has different purpose

upvoted 1 times

## Question #33

## Topic 1

You are deploying an application to a Compute Engine VM in a managed instance group. The application must be running at all times, but only a single instance of the VM should run per GCP project. How should you configure the instance group?

- A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- B. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 1.
- C. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 2.
- D. Set autoscaling to Off, set the minimum number of instances to 1, and then set the maximum number of instances to 2.

**Correct Answer: A**

  **XRiddlerX** Highly Voted 1 year ago

In my GCP console, I created a managed instance group for each answer. For each answer I deleted the instance that was created as a simple test to prove or disprove each answer.

In answer A, another instance was created after I deleted the instance  
 In answer B, no other instance was created after I deleted the instance  
 In answer C, another instance was created after I deleted the instance  
 In answer D, no other instance was created after I deleted the instance

My observation is A is the correct Answer.

A - Correct - It correctly solves the problem with only a single instance at one time  
 B - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created  
 C - Incorrect - It creates another instance after the delete HOWEVER it 2 VM's could be created even if the target is exceeded  
 D - Incorrect - Does not fit the requirement because AFTER the deletion of the instance, no other instance was created  
 upvoted 69 times

  **Eshkrkrkr** 8 months, 3 weeks ago

Wrong! <https://cloud.google.com/compute/docs/instance-groups> MIGs offer the following advantages:

High availability.  
 Keeping VM instances running. If a VM in the group stops, crashes, or is deleted by an action other than an instance group management command (for example, an intentional scale in), the MIG automatically recreates that VM in accordance with the original instance's specification (same VM name, same template) so that the VM can resume its work.  
 upvoted 3 times



  **here2help** 3 months ago

Good point but this questions says nothing about auto-healing so we need to treat this question within context. The reason that auto-healing is not the correct way to implement this is because you can fall under your min instance number with auto-healing and not with auto-scaling. See below:  
 Autoscaling works independently from autohealing. If you configure autohealing for your group and an instance fails the health check, the autohealer attempts to recreate the instance. Recreating an instance can cause the number of instances in the group to fall below the autoscaling threshold (minNumReplicas) that you specify.

If you autoscale a regional MIG, an instance can be added then immediately deleted from one of the zones. This happens when the utilization in the zone triggers a scale out but the overall utilization in the regional MIG does not require the additional instance or the additional instance is required in a different zone.

Source: <https://cloud.google.com/compute/docs/autoscaler>

In conclusion, the answer is A.  
 upvoted 4 times

  **alpharomeo9** 7 months, 2 weeks ago

Why are you talking about MIG when the options are about AutoScaling on or off?  
 upvoted 1 times

  **Ozymandiox** 7 months, 1 week ago

Because the question states MIG: "you are deploying an application to a Compute Engine VM in a managed instance group "

I am wondering if XRiddlerX got a MIG with the autohealing configured.  
<https://cloud.google.com/compute/docs/autoscaler>

Autoscaling works independently from autohealing.  
 "If you configure autohealing for your group and an instance fails the health check, the autohealer attempts to recreate the instance. Recreating an instance can cause the number of instances in the group to fall below the autoscaling threshold (minNumReplicas) that you specify."  
 upvoted 3 times

  **alpharomeo9** 7 months, 2 weeks ago

I mean why don't you elaborate on Autoscaling part. If it is set to off will the MIG keep instance health?

upvoted 1 times

  **soukumar369** 8 months, 3 weeks ago

Thank you very much for this detail testing and explanation.

upvoted 1 times

  **prasanu** 9 months, 3 weeks ago

thanks for confirming

upvoted 1 times

  **ESP\_SAP** 11 months, 2 weeks ago

thanks so much for your time to test every option.

This kind of collaboration really clarify the doubts!


upvoted 6 times

  **khaliqn**  1 year, 1 month ago

We want the application running at all times. If the VM crashes due to any underlying hardware failure, we want another instance to be added so we need autoscaling ON

Correct answer is A

upvoted 21 times

  **[Removed]** 4 months, 1 week ago

Read this warning message: "The minimum number of instances is equal to maximum number of instances. This means the autoscaler cannot add or remove instances from the instance group. Make sure this is the correct setting."

So when minimum is equal to maximum, it does not matter whether autoscaling is on or off. So now the question is who takes care of running the MINIMUM instances: MIG itself.

upvoted 2 times

  **Ozymandiox** 7 months, 1 week ago



As said in other coment, you can fix this with Autohealing, autoscaling means more machines, autohealing means re creating of VM's

upvoted 1 times

  **GCP\_Student1** 5 months, 3 weeks ago



Even if I agree with your response, it still does not meet the requirements asked in the question, which is "the application should be running all the time", because in your case when the VM instance for whatever reason stops, that mean the application will experience a downtime.

upvoted 1 times

  **Gurnoor** 1 year, 1 month ago

A is wrong - What you are talking about is Autohealing. Autoscale will not rebuild the VM on the crash.

upvoted 4 times

  **Dips\_** 7 months, 1 week ago



won't auto healing take time to recreate that VM? and there should be one VM running all time

upvoted 1 times

  **lomari**  2 weeks, 3 days ago

A is correct, autoscaling must be on. then observe the clause that ONLY one instance must be running.

upvoted 1 times

  **vmart** 1 month, 2 weeks ago

A is correct

upvoted 1 times

  **djgodzilla** 1 month, 3 weeks ago

B is correct . I just tried it with Autoscalling OFF and checked the activity logs here's the proof :

Instance Group Manager 'projects/388783089867/zones/us-east1-b/instanceGroupManagers/instance-group-1' initiated recreateInstance on instance 'projects/388783089867/zones/us-east1-b/instances/instance-group-1-3lrj'.

==>>>Reason: expected instance status RUNNING but STOPPING found.

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago


A is correct

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

upvoted 1 times

  **yuvi69** 4 months, 1 week ago

yes answer is A



upvoted 2 times

  **[Removed]** 4 months, 1 week ago

For those you are arguing with whether to set option A or B, here is the warning message: "The minimum number of instances is equal to maximum number of instances. This means the autoscaler cannot add or remove instances from the instance group. Make sure this is the correct setting."

This warning means that it does not matter whether you set autoscaling or not off, it isn't going to work either way. This means it's better to turn off autoscaling as it has nothing to scale. MIG will take care of keeping minimum instances running.

upvoted 3 times

  **Fidget\_** 4 months, 2 weeks ago

I believe it's B for a couple of reasons:

- Autoscaling uses CPU, HTTP traffic, or Monitoring metrics to scale the MIG up and down.
- I tested this - Created a MIG with autoscaling off, set min and max instances to 1. Deleted the instance via gcloud. After some loading, the instance was deleted and re-created (The instance count read "0 --> 1" during this time). I checked the creation timestamp and it was later than the previous instance. This also worked with a min and max of 2 instances I tried out of curiosity.

upvoted 1 times

  **Fidget\_** 3 months ago



Okay I did this again and it didn't work with Autoscaling Off. It looks like you still need Autoscaling ON in order for instances to be re-created. Testing was done without health checks too.

upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1

upvoted 1 times

  **nitinz** 5 months, 1 week ago

Correct Answer is B.

When you enable Autosclaing you need to give some conditions for it to auto scale. AutoScale is based on CPU, HTTP load or Stack Driver monitoring. What you are thinking about is autohealing here not autoscaling.

upvoted 1 times

  **nitinz** 5 months, 2 weeks ago

A is trap, correct answer is B. Just try it for yourself. B just works.

upvoted 1 times

  **victory108** 5 months, 3 weeks ago

A - Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

upvoted 1 times

  **Gohs** 5 months, 4 weeks ago

If you remove autoscaling, you must enable autohealing.

This will ensure that your instance remains functional at all times.

Now the point here is the length of downtime.

If ONLY autohealing is on and your VM crashes, it'll will recreate the VM and this takes like a minute — more downtime.

Assuming a maximum of one VM.

If you enable autohealing and your VM crashes, it will scale to 1 VM in seconds — less downtime.

upvoted 1 times

  **DucSiu** 6 months ago

A. Set autoscaling to On, set the minimum number of instances to 1, and then set the maximum number of instances to 1.

upvoted 1 times

  **tsoetan** 6 months ago



A is correct, because must run at all times. Hence, Autoscaling need to be turned on.


upvoted 1 times

You want to verify the IAM users and roles assigned within a GCP project named my-project. What should you do?

- A. Run gcloud iam roles list. Review the output section.
- B. Run gcloud iam service-accounts list. Review the output section.
- C. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.
- D. Navigate to the project and then to the Roles section in the GCP Console. Review the roles and status.



Correct Answer: C


-   **coldpar**

Highly Voted 

 1 year, 4 months ago



Correct answer is C as IAM section provides the list of both Members and Roles.Option A is wrong as it would provide information about the roles only.Option B is wrong as it would provide only the service accounts.Option D is wrong as it would provide information about the roles only.


upvoted 40 times
-   **Agents89**

Highly Voted 

 1 year, 3 months ago

C is the correct answer



upvoted 7 times
-   **sidharthwader**

Most Recent 



 3 weeks, 6 days ago

A lists all the roles in gcp  
B lists all the service account



C is right answer

upvoted 2 times
-   **HogwartsTrue** 2 months, 1 week ago



Ans is C

upvoted 2 times
-   **mcaromit** 2 months, 2 weeks ago



C is correct

upvoted 1 times
-   **passnow** 3 months ago



guys running A in gcloud terminal gave a list of each person with their roles for all available

upvoted 1 times
-   **lxgywil** 2 months, 3 weeks ago



That's impossible - `gcloud iam roles list` only lists the roles defined at a parent organization / a project. In particular, it displays DESCRIPTION, ETAG, NAME, STAGE, and TITLE. You wouldn't see the actual members with this command. The correct answer is C.

upvoted 1 times
-   **[Removed]** 4 months, 1 week ago



C is correct. Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 1 times
-   **yuvi69** 4 months, 1 week ago



correct answer is C

upvoted 1 times
-   **EABDAJA** 4 months, 3 weeks ago



C.Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles

upvoted 1 times
-   **nliaustemac** 5 months ago

C. is the correct answer:  
Navigate to the project and then to the IAM section in the GCP Console. Review the members and roles.

upvoted 1 times
-   **ShakthiGCP** 5 months, 1 week ago

Ans is C

upvoted 1 times
-   **iPablo26** 5 months, 1 week ago

I think C is the correct answer

upvoted 1 times

  **INASR** 6 months ago

Correct is 100% C . you need tp pick the project and then click IAm and see users and roles  
upvoted 1 times

  **GCP\_Student1** 6 months ago



Also C is correct because there is no such tab as status in the console. D is talking about status as well.  
upvoted 2 times

  **tardao** 6 months ago

Excuse me, why can't the answer be D ?  
upvoted 1 times

  **abominion** 6 months ago

Ans is D  
upvoted 1 times

  **pas77** 3 weeks, 2 days ago

D is not.  
Is asking for Roles and Users not just Roles.  
upvoted 1 times

  **EABDAJA** 6 months, 1 week ago

C is the Ans.  
upvoted 1 times



## Question #35

## Topic 1



You need to create a new billing account and then link it with an existing Google Cloud Platform project. What should you do?

- A. Verify that you are Project Billing Manager for the GCP project. Update the existing project to link it to the existing billing account.
- B. Verify that you are Project Billing Manager for the GCP project. Create a new billing account and link the new billing account to the existing project.
- C. Verify that you are Billing Administrator for the billing account. Create a new project and link the new project to the existing billing account.
- D. Verify that you are Billing Administrator for the billing account. Update the existing project to link it to the existing billing account.**

Correct Answer: B

  **coldpar** Highly Voted 1 year, 4 months ago

D is correct as the project is already created  
upvoted 30 times

  **lxgywil** 2 months, 3 weeks ago

How come? Option D says "update the existing project to link it to the EXISTING billing account", whereas the task is clearly saying "you need to create a NEW billing account".

Only B meets that criteria.  
upvoted 7 times

  **Agraved** 2 months, 2 weeks ago



check all options have mentioned is as "EXISTING billing account", because they are creating it and so it will exist.  
upvoted 1 times

  **bhavik7899** 1 month ago

billing administration can not create billing account  
so A  
upvoted 1 times

  **spudleymcdudley** 1 year ago

Listen to this guy. I spent 10 minutes and came to the same. <https://cloud.google.com/billing/docs/how-to/billing-access>  
upvoted 6 times

  **lxgywil** 2 months, 3 weeks ago

This documentation is literally saying:

"Billing Account Administrator (roles/billing.admin) - Manage billing accounts (but not create them). "

This role has nothing to do with what is required, while Project Billing Manager allows to link/unlink the project to/from a billing account.  
upvoted 6 times

  **Agents89** Highly Voted 1 year, 3 months ago

Project Billing Manager cannot create a billing account, there is nothing like Billing Administrator it is Billing Account Administrator. Both Project Billing Manager and Billing Account Administrator cannot create a billing account. A is the only answer that make sense. We have to assume the the billing account is already created  
upvoted 15 times

  **jcloud965** 2 weeks, 1 day ago

Wrong because "You need to create a new billing account". You can't assume the billing account is already created.  
upvoted 1 times

  **nightflyer** 7 months, 1 week ago

The answer is A, yes  
upvoted 3 times

  **spudleymcdudley** 1 year ago

'Verify that you are Billing Administrator for the billing account' - aka Billing Account Administrator. A and B do not have privilege to make a new billing account (<https://cloud.google.com/billing/docs/how-to/billing-access>). C is correct, this has been incorrectly up voted  
upvoted 3 times


  **jcloud965** Most Recent 2 weeks, 1 day ago

Why B is correct :

- A. The billing account has to be new (you need to CREATE it), you can't link the project to an EXISTING billing account
- C. The project already exists, you don't want to CREATE a new project
- D. Even if you are a Billing Administrator for the billing account you can't update a project billing account without having access to it + the billing account has to be new so you don't want to link the project to an existing billing account

B. Project Billing Manager is required to manage the GCP project billing settings. We assume you can create a new billing account because you are an organization billing admin or because you are not part of an organization.

upvoted 2 times

  **meriadoc71** 2 weeks, 2 days ago

Billing Account Administrator

(roles/billing.admin) -- Manage billing accounts (but not create them).

Organization or billing account. -- This role is an owner role for a billing account. Use it to manage payment instruments, configure billing exports, view cost information, link and unlink projects and manage other user roles on the billing account.

upvoted 1 times

  **dc2016bte0006** 1 month, 2 weeks ago

Option B is correct

as the question states we need to create a new billing account and link it with an existing project which is done by Project Billing Manager role

Project Billing Manager is the role in charge of link/unlink the project to/from a billing account

upvoted 1 times

  **Aiman\_1808** 1 month, 1 week ago

Manager cannot create billing account.

upvoted 1 times

  **AD\_0525** 1 month, 2 weeks ago

Both the roles can't create a billing account, but it is the project billing manager who can link the projects to billing account. So Ans should be B

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago



None of the Role options can create a Billing Account, so we have to presume the billing account exists...considering this both A & D can link/assign billing account to a project as both have resource manager.projects.createBillingAssignment. But considering GCP Best Practice of Principle of Least Privilege A is the correct option

upvoted 3 times

  **yvinisiupacuando** 2 months, 3 weeks ago

Answer is clearly "B", it reflects the right role and it is the only option that indicates that you create a new billing account as requested on the question.

upvoted 3 times

  **gspb** 3 months, 1 week ago

Form google docs,

Billing Account Administrator(roles/billing.admin) - Manage billing accounts (but not create them).

So, we can eliminate the C & D options.

Now, in question we were asked to create a new billing account and link it to existing project. So, if u check the options A & B, option B must be the correct answer.



Refer : <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 5 times

  **djgodzilla** 1 month, 3 weeks ago

correct , only "Billing Account Creator" Provides access to create billing accounts.

upvoted 1 times

  **sspsp** 3 months, 1 week ago

B is Correct

Project Billing Manager

(roles/billing.projectManager) Link/unlink the project to/from a billing account. Organization, folder, or project. This role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access. <https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

upvoted 6 times

  **yite** 3 months, 3 weeks ago

B - only answer that says create new billing account, like the question asks for.


All other answers say existing billing account, READ the question, easy answer.

upvoted 2 times

  **maddy5835** 3 months, 3 weeks ago

A is the correct answer because both the IAM roles does not give billing account creator power, they can perform link and unlink operations. Following the google best practices of google least privilege access. Administration rights is not a good option

upvoted 1 times

  **Jane111** 3 months, 3 weeks ago

Billing Account Administrator

- Manage billing accounts (but not create them).

upvoted 3 times

  **Jane111** 3 months, 3 weeks ago

(roles/billing.projectManager)

Link/unlink the project to/from a billing account.

This role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access.

upvoted 1 times

  **IamPapulalji** 4 months ago

B. If you are not part of an ORG, you don't need any specific role to create Billing Accounts. Hence we must assume this is the case here.


<https://stackoverflow.com/questions/59559885/project-billing-manager-vs-billing-account-user-gcp#:~:text=If%20you%20are%20not%20a,Billing%20Account%2C%20as%20stated%20here>.

And

[https://cloud.google.com/billing/docs/how-to/manage-billing-account#create\\_a\\_new\\_billing\\_account](https://cloud.google.com/billing/docs/how-to/manage-billing-account#create_a_new_billing_account)

If you are not a member of a Google Cloud Organization but instead are managing your Google Cloud resources or Google Maps Platform APIs using projects, you do not need any specific role or permission to create a Cloud Billing account.

upvoted 1 times

  **khi** 4 months ago

B

[https://cloud.google.com/billing/docs/how-to/modify-project#enable\\_billing\\_for\\_an\\_existing\\_project](https://cloud.google.com/billing/docs/how-to/modify-project#enable_billing_for_an_existing_project)

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

B is correct. Verify that you are Project Billing Manager for the GCP project. Create a new billing account and link the new billing account to the existing project.

upvoted 3 times

You have one project called proj-sa where you manage all your service accounts. You want to be able to use a service account from this project to take snapshots of VMs running in another project called proj-vm. What should you do?

- A. Download the private key from the service account, and add it to each VMs custom metadata.
- B. Download the private key from the service account, and add the private key to each VM’s SSH keys.
- C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.
- D. When creating the VMs, set the service account’s API scope for Compute Engine to read/write.

Correct Answer: C

- kishoredeena

Highly Voted

1 year, 1 month ago

Option C is the right one  
upvoted 15 times
- INASR

Highly Voted

6 months ago

not sure why all is saying C. It says using service account from another project to create snapshots on another project , so not sure how adding IAM role to a service account on the target project will help.  
upvoted 7 times
- jackdbd

Most Recent

6 days, 11 hours ago

C is the correct answer.  
It took me a while to figure it out because I didn't understand how service accounts work across project. This article made it clear for me.  
<https://gtseres.medium.com/using-service-accounts-across-projects-in-gcp-cf9473fef8f0>  
  
You create the service account in proj-sa and take note of the service account email, then you go to proj-vm in IAM > ADD and add the service account's email as new member and give it the Compute Storage Admin role.  
upvoted 3 times
- cuongnd

3 weeks, 1 day ago

C. Compute Storage Admin role has this: compute.snapshots.\*  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

C is the most apt from the given options, though the question is unclear  
upvoted 1 times
- [Removed]

4 months, 1 week ago

C is correct. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.  
upvoted 2 times
- pondai

4 months, 1 week ago

This quest has problem  
upvoted 1 times
- pondai

4 months ago

I'll take it back. I try this solution.It success.C is correct.  
upvoted 1 times
- yuvi69

4 months, 1 week ago

C is correct  
upvoted 1 times
- sumanshu

4 months, 2 weeks ago

How Compute Storage Admin role can take a snapshot of VM ? To take a snapshot - it needs atleast READ access for the Data resides in VM  
upvoted 2 times
- pas77

3 weeks, 2 days ago

If you take a look at the Compute Storage Admin, you will see that one of the permissions is compute.snapshots.\*  
upvoted 1 times
- lxgywil

2 months, 3 weeks ago

<https://cloud.google.com/compute/docs/access/iam#compute.storageAdmin>  
upvoted 2 times
- EABDAJA

4 months, 3 weeks ago

C.Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm

upvoted 1 times

  **nliaustemac** 5 months ago

C. is the correct answer:

Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 2 times

  **victory108** 6 months, 1 week ago

C - Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

  **swatititame** 8 months, 2 weeks ago



C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 1 times

  **glam** 9 months, 3 weeks ago

C. Grant the service account the IAM Role of Compute Storage Admin in the project called proj-vm.

upvoted 2 times



  **JJ\_ME** 9 months, 3 weeks ago

C.

roles/compute.storageAdmin - Permissions to create, modify, and delete disks, images, and snapshots.

<https://cloud.google.com/compute/docs/access/iam#compute.storageAdmin>

upvoted 2 times

  **SSPC** 11 months, 1 week ago

C is correct

upvoted 2 times

You created a Google Cloud Platform project with an App Engine application inside the project. You initially configured the application to be served from the us- central region. Now you want the application to be served from the asia-northeast1 region. What should you do?

- A. Change the default region property setting in the existing GCP project to asia-northeast1.
- B. Change the region property setting in the existing App Engine application from us-central to asia-northeast1.
- C. Create a second App Engine application in the existing GCP project and specify asia-northeast1 as the region to serve your application.
- D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

**Correct Answer: C**

- Bharathy**

Highly Voted

1 year, 4 months ago

Option D is correct, as there can be only one App Engine application inside a project . C is incorrect, as GCP can't have two app engine applications..

upvoted 32 times
- jcloud965**

2 weeks, 1 day ago

Yes, and you can't change an App Engine application region once created

upvoted 1 times
- saaurabh1805**

Highly Voted

1 year, 1 month ago

Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

<https://cloud.google.com/appengine/docs/flexible/nodejs/managing-projects-apps-billing#create>

upvoted 19 times
- dc2016bte0006**

Most Recent

1 month, 2 weeks ago

D is correct since each project can have only one App engine application

upvoted 3 times
- Asita**

2 months, 1 week ago

D is the answer

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

D is correct

upvoted 3 times
- jahnu**

3 months ago

My ans: D NOTE:Important: Each Cloud project can contain only a single App Engine application, and once created you cannot change the location of your App Engine application.

upvoted 2 times
- [Removed]**

4 months, 1 week ago

D is correct. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 1 times
- yuvi69**

4 months, 1 week ago

correct answer is D

upvoted 1 times
- Mini1**

4 months, 2 weeks ago

why not B ?

upvoted 1 times
- pas77**

3 weeks, 2 days ago

because of this

<https://cloud.google.com/appengine/docs/locations#:~:text=you%20cannot%20change%20an%20app's%20region%20after%20you%20set%20i>

t.

upvoted 1 times
- Ashii**

2 months, 1 week ago

1 app engine / project, cannot alter the region once created.

upvoted 1 times

🗨️ 👤 **tavva\_prudhvi** 4 months, 1 week ago

Because once you set the region property you cannot change it again. Hence, you have to create a new project again with the updated region, so option D suits here.

upvoted 1 times

🗨️ 👤 **nliaustemac** 5 months ago

D. is the correct answer:

Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 1 times

🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago

D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 2 times

🗨️ 👤 **DucSiu** 6 months ago

D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 1 times

🗨️ 👤 **INASR** 6 months ago

D makes sense.

upvoted 1 times

🗨️ 👤 **GCP\_Student1** 6 months ago

D is the correct option.

upvoted 1 times

🗨️ 👤 **victory108** 6 months, 1 week ago

D - Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 1 times

🗨️ 👤 **swatititame** 8 months, 2 weeks ago

- D. Create a new GCP project and create an App Engine application inside this new project. Specify asia-northeast1 as the region to serve your application.

upvoted 1 times

🗨️ 👤 **soukumar369** 8 months, 3 weeks ago

With 100% surety answer is D.

<https://cloud.google.com/appengine/docs/standard/python/console#:~:text=Important%3A%20Each%20Cloud%20project%20can,of%20your%20App%20Engine%20application.&text=Enable%20billing%20in%20the%20Billing,see%20Enabling%20billing%20for%20details>.

upvoted 5 times



## Question #38

## Topic 1

You need to grant access for three users so that they can view and edit table data on a Cloud Spanner instance. What should you do?

- A. Run `gcloud iam roles describe roles/spanner.databaseUser`. Add the users to the role.
- B. Run `gcloud iam roles describe roles/spanner.databaseUser`. Add the users to a new group. Add the group to the role.
- C. Run `gcloud iam roles describe roles/spanner.viewer - -project my-project`. Add the users to the role.
- D. Run `gcloud iam roles describe roles/spanner.viewer - -project my-project`. Add the users to a new group. Add the group to the role.

**Correct Answer: A**

🗲️ 👤 **yasu** Highly Voted 1 year, 4 months ago

I think it should be B, setup a group first are suggested way from Google.  
upvoted 39 times

🗲️ 👤 **Agents89** Highly Voted 1 year, 3 months ago

B is the correct option  
upvoted 14 times

🗲️ 👤 **dc2016bte0006** Most Recent 1 month, 2 weeks ago

Option B is correct since we need to apply common roles for each user, so creating a group and assigning role to that group is the best practice.  
upvoted 2 times

🗲️ 👤 **tanito83** 1 month, 3 weeks ago

The correct answer is B. Please, modify it.  
upvoted 1 times

🗲️ 👤 **AD\_0525** 2 months ago

Creating a group is the google suggested best practice, it should be B  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

B is correct  
upvoted 1 times

🗲️ 👤 **gspb** 3 months, 1 week ago

Since we have to follow the google recommended practices, adding the users to a group first and then assigning a role to the group is the right option.  
So, option B is the correct answer.  
upvoted 1 times

🗲️ 👤 **r1ck** 3 months, 2 weeks ago

A is wrong answer  
upvoted 2 times

🗲️ 👤 **YAS007** 6 days, 19 hours ago

why???  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago

B is correct. Run `gcloud iam roles describe roles/spanner.databaseUser`. Add the users to a new group. Add the group to the role.  
upvoted 3 times

🗲️ 👤 **pentol2** 4 months ago

itspanjain is awesome, you always present in every question :)  
upvoted 1 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago

answer is B according to 'google recommended best practises'  
upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

B. Run `gcloud iam roles describe roles/spanner.databaseUser`. Add the users to a new group. Add the group to the role.  
upvoted 1 times

🗲️ 👤 **Jacky\_YO** 4 months, 3 weeks ago

Answer : B  
Run `gcloud iam roles describe roles/spanner.databaseUser`. Add the users to a new group. Add the group to the role.



upvoted 1 times

  **nliaustemac** 5 months ago

B. is the correct answer:

Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to a new group. Add the group to the role.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

D. Run gcloud iam roles describe roles/spanner.viewer - -project my-project. Add the users to a new group. Add the group to the role.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

CORRECTION it is B

B. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to a new group. Add the group to the role.

upvoted 1 times

  **DucSiu** 6 months ago

B. Run gcloud iam roles describe roles/spanner.databaseUser. Add the users to a new group. Add the group to the role.

upvoted 1 times

  **INASR** 6 months ago

it is 100% B , only spanner.databaseUser can read and write to Database. C& D are sure wrong since it says viewer . A is half correct except it is always better to add users to groups

upvoted 1 times

  **nherrerab** 6 months, 2 weeks ago

B is correct.

upvoted 1 times

You create a new Google Kubernetes Engine (GKE) cluster and want to make sure that it always runs a supported and stable version of Kubernetes. What should you do?

- A. Enable the Node Auto-Repair feature for your GKE cluster.
- B. Enable the Node Auto-Upgrades feature for your GKE cluster.
- C. Select the latest available cluster version for your GKE cluster.
- D. Select "Container-Optimized OS (cos)" as a node image for your GKE cluster.

Correct Answer: B

- Lush

Highly Voted

1 year, 2 months ago

The answer is B  
<https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>  
upvoted 23 times
- 4bsolut

Highly Voted

1 year, 1 month ago

"Creating or upgrading a cluster by specifying the version as <latest> does not provide automatic upgrades. Enable automatic node upgrades to ensure that the nodes in your cluster up to date with the latest stable version." --source: <https://cloud.google.com/kubernetes-engine/versioning-and-upgrades>  
  
-Correct answer: B  
upvoted 15 times
- viswanand

Most Recent

2 months, 2 weeks ago

I was trying to create GKE but I dont see that option in Console  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- [Removed]

4 months, 1 week ago

B is correct. Enable the Node Auto-Upgrades feature for your GKE cluster.  
upvoted 1 times
- yuvi69

4 months, 1 week ago

answer is B . <https://cloud.google.com/kubernetes-engine/versioning> in 'specifying cluster version' section  
upvoted 1 times
- sumanshu

4 months, 2 weeks ago

Vote for B  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

B is correct  
upvoted 1 times
- JackGlemins

5 months ago

B is right: <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-upgrades>  
upvoted 1 times
- nliaustemac

5 months ago

B. is the correct answer:  
Enable the Node Auto-Upgrades feature for your GKE cluster.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.  
upvoted 2 times
- ShakthiGCP

5 months, 1 week ago

B is Correct. You want to have a stable version and not a latest version. This can be achieved by enabling Auto-upgrade which matches the Master cluster version configuration.  
upvoted 1 times
- victory108

5 months, 3 weeks ago

B - Enable the Node Auto-Upgrades feature for your GKE cluster.

upvoted 1 times

  **DucSiu** 6 months ago

B. Enable the Node Auto-Upgrades feature for your GKE cluster.

upvoted 1 times

  **INASR** 6 months ago

C is correct .

Creating or upgrading a cluster by specifying the version as latest does not provide automatic upgrades. Enable node auto-upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version.(google documents)

upvoted 1 times

  **Conjoinix** 6 months, 3 weeks ago



Correct Answer is B

[https://cloud.google.com/kubernetes-engine/versioning#specifying\\_cluster\\_version](https://cloud.google.com/kubernetes-engine/versioning#specifying_cluster_version)

Above link clearly mentions, so Option A, C, D are incorrect.

"Creating or upgrading a cluster by specifying the version as latest does not provide automatic upgrades. Enable node auto-upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version."

upvoted 1 times

  **robor97** 7 months ago

Creating or upgrading a cluster by specifying the version as latest does not provide automatic upgrades. Enable node auto-upgrades to ensure that the nodes in your cluster are up-to-date with the latest stable version. Ans = B. <https://cloud.google.com/kubernetes-engine/versioning>

upvoted 1 times

You have an instance group that you want to load balance. You want the load balancer to terminate the client SSL session. The instance group is used to serve a public web application over HTTPS. You want to follow Google-recommended practices. What should you do?

- A. Configure an HTTP(S) load balancer.
- B. Configure an internal TCP load balancer.
- C. Configure an external SSL proxy load balancer.
- D. Configure an external TCP proxy load balancer.

**Correct Answer: A**

Reference:

<https://cloud.google.com/load-balancing/docs/https/>

- Gini** Highly Voted

1 year, 3 months ago

According to the documentation of SSL Proxy Load Balacing on Google, "SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing." in my opinion A should be the most suitable choice.

upvoted 31 times
- yvinisiupacuando**

2 months, 3 weeks ago

Agree with you but, A is not the most suitable choice, it is the only choice, as the other Load Balancers cannot route HTTP(S) traffic.

upvoted 2 times
- Agents89** Highly Voted

1 year, 3 months ago

For HTTP(s) Load balancer, the client SSL session terminates at the load balancer. A is the correct option.

upvoted 14 times
- sidharthwader** Most Recent

3 weeks, 6 days ago

No doubt its answer is (A) in Azure its application gateway. These both are layer 7 traffic load balancer

upvoted 1 times
- contaexamtopics**

2 months, 1 week ago

A or c?

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

A is the correct option. The lb can have target-http-proxy in the forwarding-rule

upvoted 1 times
- meh009**

3 months, 1 week ago

GCP Best practices  
"SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing."  
<https://cloud.google.com/load-balancing/docs/ssl>

Answer is A in my opinion

upvoted 1 times
- phony2**

4 months ago

The answer should be A  
SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Configure an HTTP(S) load balancer.

upvoted 3 times
- yuvi69**

4 months, 1 week ago

SSL proxy load balancer cannot be used for http traffic . so answer is A

upvoted 1 times
- sumanshu**

4 months, 2 weeks ago

Vote for A

upvoted 1 times
- ShakthiGCP**

5 months ago

SSL proxy Load balancing is for non HTTPS . so Answer: A . HTTP(S) Load balancing

upvoted 1 times

  **nliaustemac** 5 months ago

A. is the correct answer:  
Configure an HTTP(S) load balancer.

upvoted 2 times

  **GCP\_Student1** 5 months, 1 week ago

A. Configure an HTTP(S) load balancer.

upvoted 2 times

  **nitinz** 5 months, 2 weeks ago

The correct answer is A, you can use SSL proxy to terminate on SSL Load Balancer BUT best practices say if using HTTPS then use HTTPS Load Balancer. Read the last line: -  
For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

  **victory108** 5 months, 3 weeks ago

A - Configure an HTTP(S) load balancer.

upvoted 1 times

  **lutoa** 5 months, 3 weeks ago

After reading the discussion here I was still confused if A or C is the right answer.  
I read the docs a bit.

As others have said, for HTTPS traffic, google always recommends HTTPS load balancer.

TCP traffic is for other load balancer types.

So I pick A.

[https://cloud.google.com/load-balancing/docs/load-balancing-overview#traffic\\_type](https://cloud.google.com/load-balancing/docs/load-balancing-overview#traffic_type)

upvoted 1 times

  **DucSiu** 6 months ago

C. Configure an external SSL proxy load balancer.

upvoted 1 times

Question #41

Topic 1

You have 32 GB of data in a single file that you need to upload to a Nearline Storage bucket. The WAN connection you are using is rated at 1 Gbps, and you are the only one on the connection. You want to use as much of the rated 1 Gbps as possible to transfer the file rapidly. How should you upload the file?

- A. Use the GCP Console to transfer the file instead of gsutil.
- B. Enable parallel composite uploads using gsutil on the file transfer.
- C. Decrease the TCP window size on the machine initiating the transfer.
- D. Change the storage class of the bucket from Nearline to Multi-Regional.

Correct Answer: B

- leba

Highly Voted

1 year, 3 months ago

Correct answer is B as the bandwidth is good and its a single file, gsutil parallel composite uploads can be used to split the large file and upload in parallel.Refer GCP documentation - Transferring Data to GCP & upvoted 20 times
- berezinsn

Highly Voted

1 year, 2 months ago

Truly B is absolutely correct upvoted 10 times
- NizarIrfani

Most Recent

1 month, 3 weeks ago

Play online games on <http://159.89.187.75/link-alternatif-m88/> upvoted 2 times
- NamiKaze

1 week ago

Yess You Smart upvoted 1 times
- djgodzilla

1 month, 3 weeks ago

B is correct :  
i.e :  
gsutil -o GSUtil:parallel\_composite\_upload\_threshold=200M,parallel\_composite\_upload\_component\_size=100M cp ./localbigfile gs://your-bucket upvoted 2 times
- mcaromit

2 months, 2 weeks ago

B is the obvious choice upvoted 1 times
- ApaMokus

3 months ago

Although parallel composite transfer is not optimal for Nearline storage, option B is correct as the question focuses on quicker upload and not costs. upvoted 1 times
- kopper2019

3 months, 2 weeks ago

it's similar to AWS in order to upload faster using a parallel way so B upvoted 1 times
- [Removed]

4 months, 1 week ago

B is correct. Enable parallel composite uploads using gsutil on the file transfer. upvoted 2 times
- yuvi69

4 months, 1 week ago

answer is B.  
<https://cloud.google.com/storage/docs/uploads-downloads#parallel-composite-uploads> upvoted 2 times
- EABDAJA

4 months, 3 weeks ago

Answer is B. upvoted 1 times
- EABDAJA

4 months, 3 weeks ago


B is corect upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

One strategy for uploading large files is called parallel composite uploads. In such an upload, a file is divided into up to 32 chunks, the chunks are uploaded in parallel to temporary objects, the final object is recreated using the temporary objects, and the temporary objects are deleted.

Parallel composite uploads can be significantly faster if network and disk speed are not limiting factors; however, the final object stored in your bucket is a composite object, which only has a crc32c hash and not an MD5 hash. As a result, you must use crcmod to perform integrity checks when downloading the object with gsutil or other Python applications. You should only perform parallel composite uploads if the following apply:

upvoted 2 times

  **nitinz** 5 months, 2 weeks ago

Tricky, had no clue about it... Given scenario B is the best answer.

upvoted 1 times

  **JackGlemins** 6 months ago

B is correct. <https://cloud.google.com/storage/docs/uploads-downloads#parallel-composite-uploads>

upvoted 2 times

  **DucSiu** 6 months ago



B. Enable parallel composite uploads using gsutil on the file transfer.

upvoted 1 times

  **INASR** 6 months ago

B is correct. cost is not a a factor

upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

- B. Enable parallel composite uploads using gsutil on the file transfer.

upvoted 1 times

You’ve deployed a microservice called myapp1 to a Google Kubernetes Engine cluster using the YAML file specified below:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp1-deployment
spec:
  selector:
    matchLabels:
      app: myapp1
  replicas: 2
  template:
    metadata:
      labels:
        app: myapp1
    spec:
      containers:
        - name: main-container
          image: gcr.io/my-company-repo/myapp1:1.4
          env:
            - name: DB_PASSWORD
              value: "t0ugh2guess!"
          ports:
            - containerPort: 8080
```

You need to refactor this configuration so that the database password is not stored in plain text. You want to follow Google-recommended practices. What should you do?

- A. Store the database password inside the Docker image of the container, not in the YAML file.
- B. Store the database password inside a Secret object. Modify the YAML file to populate the DB\_PASSWORD environment variable from the Secret.
- C. Store the database password inside a ConfigMap object. Modify the YAML file to populate the DB\_PASSWORD environment variable from the ConfigMap.
- D. Store the database password in a file inside a Kubernetes persistent volume, and use a persistent volume claim to mount the volume to the container.

**Correct Answer: C**

- rramani7

Highly Voted

1 year, 1 month ago

it is good practice to use Secrets for confidential data (like API keys) and ConfigMaps for non-confidential data (like port numbers). B is correct.

upvoted 32 times
- saurabh1805

Highly Voted

1 year, 1 month ago

B is correct answer

<https://cloud.google.com/kubernetes-engine/docs/concepts/secret>

upvoted 18 times
- hjyhf

2 days, 17 hours ago

"Storing sensitive data in Secrets is more secure than in plaintext ConfigMaps or in Pod specifications"

upvoted 1 times
- associatecloudexamuser

Most Recent

2 weeks, 6 days ago

Answer is B

upvoted 2 times
- sunilw

1 month, 1 week ago

B is correct answer. Option C is not correct, because ConfigMaps are useful for storing and sharing non-sensitive, unencrypted configuration information. To use sensitive information in your clusters, you must use Secrets.

<https://cloud.google.com/kubernetes-engine/docs/concepts/secret>

upvoted 1 times
- AD\_0525

2 months ago

B should be the answer, as config map is suggested to use for storing the environment configurations, db names etc.

upvoted 1 times
- shankyomre01

2 months ago

please change an answer to B. it's correct

upvoted 1 times



- 🗲️ 👤 **mcaromit** 2 months, 2 weeks ago  
B is correct  
upvoted 1 times
- 🗲️ 👤 **Srinivasd** 3 months, 1 week ago  
Ans is B  
upvoted 1 times
- 🗲️ 👤 **jahnu** 3 months, 1 week ago  
my answer is B because for sensitive information like passwords we use secrets.  
upvoted 1 times
- 🗲️ 👤 **Chetna13** 3 months, 2 weeks ago  
Correct answer is B  
upvoted 1 times
- 🗲️ 👤 **kopper2019** 3 months, 2 weeks ago  
B, Secrets are secure objects which store sensitive data, such as passwords, OAuth tokens, and SSH keys, in your clusters. Storing sensitive data in Secrets is more secure than plaintext ConfigMaps or in Pod specifications. Using Secrets gives you control over how sensitive data is used, and reduces the risk of exposing the data to unauthorized users.  
upvoted 2 times
- 🗲️ 👤 **rkk8s** 3 months, 3 weeks ago  
B is correct to use secrets which is base64 encoded and mount as a volume to the pod, instead of configmaps which is in plain text.  
upvoted 3 times
- 🗲️ 👤 **[Removed]** 4 months, 1 week ago  
B is correct. Store the database password inside a Secret object. Modify the YAML file to populate the DB\_PASSWORD environment variable from the Secret.  
upvoted 2 times
- 🗲️ 👤 **yuvi69** 4 months, 1 week ago  
answer is B  
upvoted 1 times
- 🗲️ 👤 **Durgesh1997** 4 months, 1 week ago  
Correct answer is b  
upvoted 1 times
- 🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago  
B is correct  
upvoted 1 times
- 🗲️ 👤 **nliaustemac** 5 months ago  
B. is the correct answer:  
Store the database password inside a Secret object. Modify the YAML file to populate the DB\_PASSWORD environment variable from the Secret.  
upvoted 1 times

## Question #43



## Topic 1

You are running an application on multiple virtual machines within a managed instance group and have autoscaling enabled. The autoscaling policy is configured so that additional instances are added to the group if the CPU utilization of instances goes above 80%. VMs are added until the instance group reaches its maximum limit of five VMs or until CPU utilization of instances lowers to 80%. The initial delay for HTTP health checks against the instances is set to 30 seconds.

The virtual machine instances take around three minutes to become available for users. You observe that when the instance group autoscales, it adds more instances than necessary to support the levels of end-user traffic. You want to properly maintain instance group sizes when autoscaling. What should you do?



- A. Set the maximum number of instances to 1.
- B. Decrease the maximum number of instances to 3.
- C. Use a TCP health check instead of an HTTP health check.
- D. Increase the initial delay of the HTTP health check to 200 seconds.

**Correct Answer: D**



  **berezinsn** Highly Voted 1 year, 2 months ago  
D is correct answer.  
upvoted 16 times



  **professor** Highly Voted 1 year, 1 month ago  
Ans is D



The virtual machine instances take around three minutes to become available for users.  
upvoted 9 times



  **Abalagu** 1 month, 2 weeks ago  
D is correct. The reason is that when you do health check, you want the VM to be working. Do the first check after initial setup time of 3 mins = 180 s < 200 s is reasonable.  
upvoted 3 times



  **mcaromit** Most Recent 2 months, 2 weeks ago  
D is correct  
upvoted 1 times



  **[Removed]** 4 months, 1 week ago  
D is correct. Increase the initial delay of the HTTP health check to 200 seconds.  
upvoted 2 times



  **yuvi69** 4 months, 1 week ago  
D is correct  
upvoted 1 times



  **EABDAJA** 4 months, 3 weeks ago  
D is Correct.  
upvoted 1 times

  **nliaustemac** 5 months ago  
D. is the correct answer:  
Increase the initial delay of the HTTP health check to 200 seconds  
upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago  
D. Increase the initial delay of the HTTP health check to 200 seconds.  
upvoted 1 times

  **Gohs** 5 months, 3 weeks ago  
By increasing the initial delay to 200s, VMs will boot properly before being checked; instead of judging them "unhealthy" after only 20s  
upvoted 4 times

  **DucSiu** 6 months ago  
D. Increase the initial delay of the HTTP health check to 200 seconds.  
upvoted 1 times



  **INASR** 6 months ago  
D for sure. you need to increase the cool down time.

upvoted 1 times

  **Dineshvon** 7 months, 3 weeks ago

D is correct answer

upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

D. Increase the initial delay of the HTTP health check to 200 seconds.

upvoted 2 times

  **soukumar369** 8 months, 3 weeks ago

This is the easiest one :)

upvoted 3 times

  **glam** 9 months, 3 weeks ago

D. Increase the initial delay of the HTTP health check to 200 seconds.

upvoted 3 times

Question #44

## Topic 1

You need to select and configure compute resources for a set of batch processing jobs. These jobs take around 2 hours to complete and are run nightly. You want to minimize service costs. What should you do?

- A. Select Google Kubernetes Engine. Use a single-node cluster with a small instance type.
- B. Select Google Kubernetes Engine. Use a three-node cluster with micro instance types.
- C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.
- D. Select Compute Engine. Use VM instance types that support micro bursting.

**Correct Answer: C**

 **vnxt** Highly Voted 1 year, 2 months ago

I would say C is the correct answer  
upvoted 21 times

  **gcper** Highly Voted  9 months, 2 weeks ago

As everyone has said the answer is C but here is the source for the information. "For example, batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances."

source: <https://cloud.google.com/compute/docs/instances/preemptible>  
upvoted 15 times

  **dc2016bte0006** Most Recent 1 month, 2 weeks ago

Option C is the correct answer  
upvoted 1 times

 **mcaromit** 2 months, 2 weeks ago

C is correct...For cost-saving & not immediate fault-tolerant workloads like batch jobs use Preemptible VM instances

upvoted 1 times

  **Faizan2991** 2 months, 4 weeks ago

Might be D , because preemptible VM can be terminate at any time with short notice  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

C is correct. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

upvoted 2 times

  **yuvi69** 4 months, 1 week ago

option C is correct  
upvoted 1 times

  **pondai** 4 months, 1 week ago

C is the answer  
upvoted 1 times


  **chithra1990** 5 months ago

Yes C is the right answer. When you see compute engine->cost-effective-> you can strongly go ahead with preemptible  
upvoted 1 times

 **nliaustemac** 5 months ago

C. is the correct answer:  
Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

 **nitinz** 5 months, 2 weeks ago

Reduce cost is the keyword so pre-emptible is the best option... C

upvoted 1 times

  **Gohs** 5 months, 3 weeks ago

Preemptible VMs are the cheapest VMs

upvoted 1 times

  **DucSiu** 6 months ago

C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

upvoted 1 times

  **INASR** 6 months ago



C is correct .

upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

Finally - can someone suggest, if not preemptible - what would be best option to go for?

upvoted 1 times

  **pas77** 2 weeks, 6 days ago



If not preemptible, I would go with D, but in this case says "You want to minimize service costs", so the right answer is definitely C

upvoted 1 times

  **alpharomeo9** 7 months, 2 weeks ago

Definitely not for A & B. And C is the only available best solution. I would choose C

upvoted 1 times

  **swatititame** 8 months, 2 weeks ago

• C. Select Compute Engine. Use preemptible VM instances of the appropriate standard machine type.

upvoted 1 times

You recently deployed a new version of an application to App Engine and then discovered a bug in the release. You need to immediately revert to the prior version of the application. What should you do?

- A. Run gcloud app restore.
- B. On the App Engine page of the GCP Console, select the application that needs to be reverted and click Revert.
- C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.
- D. Deploy the original version as a separate application. Then go to App Engine settings and split traffic between applications so that the original version serves 100% of the requests.

**Correct Answer:** D

Reference:  
<https://medium.com/google-cloud/app-engine-project-cleanup-9647296e796a>

- coldpar

Highly Voted

1 year, 4 months ago

correct is C NOT D.  
Option A is wrong as gcloud app restore was used for backup and restore and has been deprecated.Option B is wrong as there is no application revert functionality available.Option D is wrong as App Engine maintains version and need not be redeployed.  
upvoted 36 times
- Bharathy

Highly Voted

1 year, 4 months ago

App engine maintains versions and to revert back to previous version, traffic can be set to 100% for the prior version.. hence correct answer is C  
upvoted 14 times
- sunilw

Most Recent

1 month, 1 week ago

C is correct.  
  
<https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration#rollback>  
upvoted 1 times
- gcp\_aspirant1

1 month, 1 week ago

C is the only option that will help resolve the issue with no `downtime`.  
upvoted 1 times
- shankyomre01

2 months ago

C is correct  
upvoted 2 times
- mcaromit

2 months, 2 weeks ago

C is correct, presuming while deploying the new version, older version was not stopped  
upvoted 2 times
- [Removed]

4 months, 1 week ago

C is correct. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.  
upvoted 4 times
- yuvi69

4 months, 1 week ago

correct option is C  
upvoted 2 times
- nliaustemac

5 months ago

C. is the correct answer:  
On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.  
upvoted 2 times
- GCP\_Student1

5 months, 1 week ago

C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.  
upvoted 2 times
- nitinz

5 months, 2 weeks ago

You need to immediately revert = reroute traffic. C  
upvoted 2 times
- JackGlemins

6 months ago

C is correct. <https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration#rollback>  
upvoted 2 times

  **DucSiu** 6 months ago

C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.  
upvoted 1 times

  **INASR** 6 months ago



C is correct  
upvoted 1 times

  **victory108** 6 months, 1 week ago

C - On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.  
upvoted 1 times

  **cecocenov** 8 months, 1 week ago

Correct answer is C as you can migrate all the traffic back to the previous version. Refer GCP documentation - App Engine Overview Having multiple versions of your app within each service allows you to quickly switch between different versions of that app for rollbacks, testing, or other temporary events. You can route traffic to one or more specific versions of your app by migrating or splitting traffic. Option A is wrong as gcloud app restore was used for backup and restore and has been deprecated. Option B is wrong as there is no application revert functionality available. Option D is wrong as App Engine maintains version and need not be redeployed.  
upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

- C. On the App Engine Versions page of the GCP Console, route 100% of the traffic to the previous version.

upvoted 1 times

## Question #46

## Topic 1

You deployed an App Engine application using gcloud app deploy, but it did not deploy to the intended project. You want to find out why this happened and where the application deployed. What should you do?

- A. Check the app.yaml file for your application and check project settings.
- B. Check the web-application.xml file for your application and check project settings.
- C. Go to Deployment Manager and review settings for deployment of applications.
- D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.

**Correct Answer: A**

Reference:

<https://cloud.google.com/endpoints/docs/openapi/troubleshoot-aeflex-deployment>

  **Bharathy** Highly Voted 1 year, 4 months ago

I would opt option D : as it would help to check the config details and Option A is not correct, as app.yaml would have only the runtime and script to run parameters and not the Project details  
upvoted 29 times

  **ahmed812** Highly Voted 1 year, 3 months ago



Option D - The config list will give the name of the project  
C:\GCP\appeng>gcloud config list  
[core]  
account = xxx@gmail.com  
disable\_usage\_reporting = False  
project = my-first-demo-xxxx  
upvoted 20 times

  **AD\_0525** Most Recent 2 months ago


In the app.yaml file we just need to provide the runtime, not sure of project name. So option D seems to be the correct one.  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

D is correct  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

D is correct. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.  
upvoted 2 times

  **yuvi69** 4 months, 1 week ago



correct option is D  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

D is correct.  
upvoted 1 times

  **GCP\_Student1** 4 months, 3 weeks ago

D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.  
upvoted 1 times

  **Rimjith** 4 months, 4 weeks ago

There are two parts in the question , we have to find out why and where the app is deployed.

Why : Because App ID was wrongly mentioned in the yaml file  
<https://cloud.google.com/appengine/docs/flexible/python/reference/app-yaml#app-id>  
upvoted 1 times





















  **cloud\_hunk** 5 months ago

from deployment wise i would options are correct but what question is asking is it is deployed but not on the project where you are expecting as well as it is asking you to find out what happened and where it deploy. I would go with option D  
upvoted 1 times

  **ShakthiGCP** 5 months, 1 week ago

Ans: D : Yaml file have application details but not the Project details. To know about the project detail. we need to do gcloud config list  
upvoted 1 times



-   **nitinz** 5 months, 2 weeks ago  
D as question says: -  
You deployed an App Engine application using gcloud app deploy  
  
what an go wrong? Your region in gcloud config on your CLI. Where will you check it? Option D.  
upvoted 1 times
-   **DucSiu** 6 months ago  
D. Go to Cloud Shell and run gcloud config list to review the Google Cloud configuration used for deployment.  
upvoted 1 times
-   **INASR** 6 months ago  
A is correct . Only changes are the YAML config file  
<https://cloud.google.com/sdk/gcloud/reference/app/deploy>  
upvoted 4 times
-   **hasnaa** 3 months, 4 weeks ago  
YAML file don't contain the project info  
upvoted 1 times
-   **guid1984** 6 months, 1 week ago  
Correct answer is D, From google docs:  
By default, the service is deployed the current project configured via:  
gcloud config set core/project PROJECT  
To override this value for a single deployment, use the --project flag:  
gcloud app deploy ~/my\_app/app.yaml --project=PROJECT  
  
<https://cloud.google.com/sdk/gcloud/reference/app/deploy>  
So as per question if the app does not appear in intended project check which is active configuration using option D  
upvoted 1 times
-   **FenixRa73** 6 months, 2 weeks ago  
I think at the current time the correct answer is C.  
But if use deprecated features "application id" in the YAML file.. the correct answer is A.  
application  
Warning: The application element in addition to using the appcfg tooling are now deprecated. See the following instructions to specify your application using the Cloud SDK tooling such as the gcloud command line.  
The recommended approach is to remove the application element from your app.yaml file and instead, use a command-line flag to specify your application ID:  
To use the gcloud app deploy command, you must specify the --project flag:  
gcloud app deploy --project [YOUR\_PROJECT\_ID]  
The application ID is the Cloud Console project ID that you specified when you created the application in the Google Cloud Console.  
<https://cloud.google.com/appengine/docs/standard/python/config/appref>  
upvoted 1 times
-   **mselway** 7 months ago  
A  
Specify all the runtime configuration for your Python app, including versions and URLs, in the app.yaml file. This file acts as a deployment descriptor of a specific service version.  
The question specifies the use of "gcloud app deploy" which uses the app.yaml file as an argument. While there is some implicit config, I think A is the better answer.  
<https://cloud.google.com/appengine/docs/flexible/python/configuring-your-app-with-app-yaml>  
upvoted 2 times
-   **JackGlemins** 6 months ago  
I think that you dont specify de project on the .yaml file.  
<https://cloud.google.com/appengine/docs/flexible/nodejs/configuring-your-app-with-app-yaml>  
A Node.js app in App Engine is configured using an app.yaml file, that contains CPU, memory, network and disk resources, scaling, and other general settings including environment variables.  
  
I was looking for and I dont find nothing about the config project.  
upvoted 1 times
-   **magistrum** 6 months, 2 weeks ago  
looks like you can define your project within the app.yaml file <https://cloud.google.com/appengine/docs/flexible/nodejs/reference/app-yaml>  
upvoted 2 times
-   **pas77** 2 weeks, 6 days ago  
Where in the link you've provided, says that you can specify the project name or ID in the app.yaml file?  
upvoted 2 times

You want to configure 10 Compute Engine instances for availability when maintenance occurs. Your requirements state that these instances should attempt to automatically restart if they crash. Also, the instances should be highly available including during system maintenance. What should you do?

- A. Create an instance template for the instances. Set the `Automatic Restart` to on. Set the `On-host maintenance` to Migrate VM instance. Add the instance template to an instance group.
- B. Create an instance template for the instances. Set `Automatic Restart` to off. Set `On-host maintenance` to Terminate VM instances. Add the instance template to an instance group.
- C. Create an instance group for the instances. Set the `Autohealing` health check to healthy (HTTP).
- D. Create an instance group for the instance. Verify that the `Advanced creation options` setting for `do not retry machine creation` is set to off.

Correct Answer: B

- lio123

Highly Voted

4 months, 2 weeks ago

A

<https://cloud.google.com/compute/docs/instances/setting-instance-scheduling-options>

onHostMaintenance: Determines the behavior when a maintenance event occurs that might cause your instance to reboot.

[Default] MIGRATE, which causes Compute Engine to live migrate an instance when there is a maintenance event.  
TERMINATE, which stops an instance instead of migrating it.  
automaticRestart: Determines the behavior when an instance crashes or is stopped by the system.

[Default] true, so Compute Engine restarts an instance if the instance crashes or is stopped.  
false, so Compute Engine does not restart an instance if the instance crashes or is stopped.

upvoted 9 times
- kashi123

Most Recent

2 weeks, 5 days ago

A is correct

upvoted 2 times
- ranjith3105

3 weeks, 2 days ago

A is the correct answer...

upvoted 1 times
- sunilw

1 month, 1 week ago

A is correct

upvoted 1 times
- vmart

1 month, 2 weeks ago

I vouch for A

upvoted 2 times
- pb11236

2 months, 1 week ago

why there is lot of incorrect answers in this??

upvoted 2 times
- AD\_0525

2 months ago

They want us to explore discussion forum for each questions. :D

upvoted 3 times
- mcaromit

2 months, 2 weeks ago

A is the obvious correct choice

upvoted 3 times
- pavans78

4 months ago

A is correct

upvoted 4 times
- [Removed]











4 months, 1 week ago

A is correct. Create an instance template for the instances. Set the `Automatic Restart` to on. Set the `On-host maintenance` to Migrate VM instance. Add the instance template to an instance group.

upvoted 2 times
- yuvi69

4 months, 1 week ago

correct option is A  
upvoted 2 times

-   **GCP\_Student1** 4 months, 1 week ago  
A. Create an instance template for the instances. Set the "Automatic Restart" to on. Set the "On-host maintenance" to Migrate VM instance. Add the instance template to an instance group.  
upvoted 1 times
-   **Hi2ALL** 4 months, 1 week ago  
A is correct  
upvoted 1 times
-   **sumanshu** 4 months, 2 weeks ago  
Vote for A  
upvoted 1 times
-   **neerajgoyal** 4 months, 2 weeks ago  
A is the correct answer  
upvoted 2 times
-   **NicolaeS** 4 months, 2 weeks ago  
A  
<https://cloud.google.com/compute/docs/instances/setting-instance-scheduling-options#autorestart>  
<https://cloud.google.com/compute/docs/instances/setting-instance-scheduling-options#maintenanceevents>  
upvoted 3 times

You host a static website on Cloud Storage. Recently, you began to include links to PDF files on this site. Currently, when users click on the links to these PDF files, their browsers prompt them to save the file onto their local system. Instead, you want the clicked PDF files to be displayed within the browser window directly, without prompting the user to save the file locally. What should you do?

- A. Enable Cloud CDN on the website frontend.
- B. Enable "Share publicly" on the PDF file objects.
- C. Set Content-Type metadata to application/pdf on the PDF file objects.
- D. Add a label to the storage bucket with a key of Content-Type and value of application/pdf.

Correct Answer: C

- berezinsn** Highly Voted

1 year, 2 months ago

C is correct

upvoted 18 times
- yas\_cloud** Highly Voted

11 months, 1 week ago

C is correct here. tested

upvoted 8 times
- sunilw** Most Recent

1 month, 1 week ago

C is correct.

<https://cloud.google.com/storage/docs/metadata#content-type>

upvoted 1 times
- vmart**

1 month, 2 weeks ago

c is correct

upvoted 1 times
- [Removed]**

4 months, 1 week ago

C is correct. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times
- yuvi69**

4 months, 1 week ago

correct option is C

upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times
- nliaustemac**

5 months ago

C. is the correct answer:

Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times
- victory108**

5 months, 3 weeks ago

C - Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 1 times
- d18mph23**

5 months, 3 weeks ago

C is correct.

<https://cloud.google.com/storage/docs/metadata#content-type>

upvoted 1 times
- DucSiu**

6 months ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 2 times
- INASR**



6 months ago

C is correct

upvoted 1 times

  **EABDAJA** 6 months ago

Answer is C  
upvoted 1 times



  **swatititame** 8 months, 2 weeks ago

- C. Set Content-Type metadata to application/pdf on the PDF file objects.

upvoted 4 times

  **glam** 9 months, 3 weeks ago

C. Set Content-Type metadata to application/pdf on the PDF file objects.  
upvoted 4 times

  **SSPC** 11 months, 1 week ago

C is correct  
upvoted 3 times

## Question #49

## Topic 1

You have a virtual machine that is currently configured with 2 vCPUs and 4 GB of memory. It is running out of memory. You want to upgrade the virtual machine to have 8 GB of memory. What should you do?

- A. Rely on live migration to move the workload to a machine with more memory.
- B. Use gcloud to add metadata to the VM. Set the key to required-memory-size and the value to 8 GB.
- C. Stop the VM, change the machine type to n1-standard-8, and start the VM.
- D. Stop the VM, increase the memory to 8 GB, and start the VM.

**Correct Answer: D**

  **cesar7816** Highly Voted 1 year, 4 months ago

coldpar, why are you getting the people confused? you need to stop teh VM and modify the RAM, that's all  
upvoted 42 times

  **Carls** Highly Voted 1 year, 3 months ago

D is correct. If you pay attention to the question, option C mentions n1-standard-8. That instance type has 8vCPUs and 30 GB RAM, and we only need 8GB. On top of that, it is possible to use custom machine type to adjust current VM RAM to the value we need. Got the answer from this course I did to prepare the exam: <https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?couponCode=21CDE6A4C2B95F79BD97>  
good luck!  
upvoted 20 times

  **Umesh09** Most Recent 1 month, 3 weeks ago

why do we want to stop the VM when live migration is an option I would go for Option A  
upvoted 1 times

  **YAS007** 6 days ago

I don't find how to manually initiate a live migration, its apparently a google managed process  
upvoted 1 times

  **Umesh09** 1 month, 3 weeks ago



Why not A GCP does give a live migration option.  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago


D is correct  
upvoted 1 times

  **ri\_unhou119** 2 months, 3 weeks ago

Option A [Live migration] Why is it not applicable?  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

D is correct. Stop the VM, increase the memory to 8 GB, and start the VM.  
upvoted 2 times

  **yuvi69** 4 months, 1 week ago

correct option is D  
upvoted 2 times

  **cloud\_hunk** 5 months ago















ques is tricky , it did not provided enough information of VM type.  
if it is a custom VM then yes, you can stop and just update the memory to 8 GB. If it is standard VM then you need to change the config to n1-standard-8 which I believe is not the case here.  
upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

n1-standard-8 is not 8GB, its 32GB and 8 CPU  
  
<https://cloud.google.com/compute/docs/machine-types>  
upvoted 1 times

  **nliaustemac** 5 months ago

D. is the correct answer:  
Stop the VM, increase the memory to 8 GB, and start the VM.  
upvoted 1 times

-   **GCP\_Student1** 5 months, 1 week ago  
D. Stop the VM, increase the memory to 8 GB, and start the VM.  
upvoted 1 times
-   **nitinz** 5 months, 2 weeks ago  
Its D, question says make 8 GB RAM. C is n1-standard-8 is wrong. Check the specs for n1-standard-8 it starts at 30 GB RAM.  
upvoted 1 times
-   **DucSiu** 6 months ago  
D. Stop the VM, increase the memory to 8 GB, and start the VM.  
upvoted 1 times
-   **INASR** 6 months ago  
D is the right answer  
upvoted 2 times
-   **Conjoinix** 6 months, 3 weeks ago  
Correct Answer is D. Tried and tested  
n1-standard-8 has 30 GB memory and question specifically asks for 8 GB memory which can be upgraded choosing correct image e.g. e2-standard-2(2 vCPU, 8 GB)  
upvoted 2 times
-   **Bhagirathi** 7 months, 3 weeks ago  
ah it should be D.  
upvoted 1 times
-   **Dineshvon** 7 months, 3 weeks ago  
correct answer is D  
upvoted 1 times

You have production and test workloads that you want to deploy on Compute Engine. Production VMs need to be in a different subnet than the test VMs. All the VMs must be able to reach each other over Internal IP without creating additional routes. You need to set up VPC and the 2 subnets. Which configuration meets these requirements?

- A. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.
- B. Create a single custom VPC with 2 subnets. Create each subnet in the same region and with the same CIDR range.
- C. Create 2 custom VPCs, each with a single subnet. Create each subnet in a different region and with a different CIDR range.
- D. Create 2 custom VPCs, each with a single subnet. Create each subnet in the same region and with the same CIDR range.

Correct Answer: A

- JamesBond

Highly Voted

1 year, 4 months ago

A is correct  
upvoted 23 times
- nwk

Highly Voted

9 months ago

Vote A  
<https://cloud.google.com/vpc/docs/using-vpc#subnet-rules>  
Primary and secondary ranges for subnets cannot overlap with any allocated range, any primary or secondary range of another subnet in the same network, or any IP ranges of subnets in peered networks.  
upvoted 9 times
- sunilw

Most Recent

1 month, 1 week ago

Vote A  
<https://cloud.google.com/vpc/docs/using-vpc#subnet-rules>  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct...no additional routing needed as same VPC  
upvoted 1 times
- [Removed]

4 months, 1 week ago

A is correct. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.  
upvoted 1 times
- yuvi69

4 months, 1 week ago

correct option is A  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

A is correct  
upvoted 1 times
- cloud\_hunk

5 months ago

A is the correct as que it self answer many silos like "Production VMs need to be in a different subnet than the test VMs"  
upvoted 1 times
- nliaustemac

5 months ago

A. is the correct answer:  
Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

A. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.  
upvoted 1 times
- DucSiu

6 months ago

A. Create a single custom VPC with 2 subnets. Create each subnet in a different region and with a different CIDR range.  
upvoted 1 times
- GCP\_Student1

6 months ago

I vote for A.  
upvoted 2 times
- INASR

6 months ago



A is correct  
upvoted 2 times

🗨️ 👤 **Conjoinix** 6 months, 3 weeks ago

Correct Answer is A.  
VPC is global in nature. Subnets are regional but are connected globally.  
So creating 1 VPC and different CIDR for different subnets should suffice the purpose  
upvoted 4 times

🗨️ 👤 **mselway** 7 months ago

B  
Question specifies the need for 2 subnets, so C & D cannot be correct.  
The question also says "All the VMs must be able to reach each other over Internal IP without creating additional routes", so they have to be in the same CIDR block. The CIDR block can be carved up into different subnets.  
upvoted 3 times

🗨️ 👤 **sadas159** 7 months, 2 weeks ago

Hello every on in the question referents CIRD Range no IP, is importar read correctly , you need to deploy two subnets with the same CIRD Range and this is the correct answer B  
upvoted 3 times

🗨️ 👤 **magistrum** 7 months ago

So, once you created 10.0.0.0/9 in us-central1, using the "same CIDR" means creating another subnet as 10.0.0.0/9 in us-central1  
upvoted 2 times

🗨️ 👤 **MEHULSP** 2 weeks, 6 days ago

So, acc to this should it not be A?  
upvoted 1 times

🗨️ 👤 **MEHULSP** 2 weeks, 6 days ago

since you told we can't have same cidr range in same region  
upvoted 1 times

🗨️ 👤 **magistrum** 7 months ago

which is not allowed  
upvoted 2 times

🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago

finally - what to choose?  
A<  
upvoted 1 times

You need to create an autoscaling managed instance group for an HTTPS web application. You want to make sure that unhealthy VMs are recreated. What should you do?

- A. Create a health check on port 443 and use that when creating the Managed Instance Group.
- B. Select Multi-Zone instead of Single-Zone when creating the Managed Instance Group.
- C. In the Instance Template, add the label `health-check`.
- D. In the Instance Template, add a startup script that sends a heartbeat to the metadata server.

**Correct Answer:** C

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>

- cesar7816**

Highly Voted

1 year, 4 months ago

I'll go with A, MIGs support autohealing, load balancing, autoscaling, and auto-updating. no the Images templates, this is set up in the MIG

upvoted 37 times
- baudrual**

Highly Voted

12 months ago

Hi Evangelist, I would say that is to let us discuss on ;-)

upvoted 5 times
- tanito83**

Most Recent

1 month, 3 weeks ago

The correct answer is A. Please, modify it.

upvoted 3 times
- Rocket\_science**

2 months ago

Why not c?

upvoted 1 times
- Finger41**

1 month, 4 weeks ago

You are just adding a label, ie Key value pair for description purposes, wont do anything from a healthcheck perspective eg Environment = Developement. Answer is A, as this gives you loadbalancing, and autohealing on the HTTPS port.

upvoted 1 times
- AD\_0525**

2 months ago

It should be A, adding a label will not recreate the VM.

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

A is correct

upvoted 1 times
- r1ck**

3 months, 2 weeks ago

Isn't health check done in port 80 - Http?!

upvoted 1 times
- Amara\_Michy**

3 months, 1 week ago

If the health check's protocol is TCP or HTTP , it uses --port=80 . If the health check's protocol is SSL , HTTPS , or HTTP2 , it uses --port=443

upvoted 4 times
- [Removed]**

4 months, 1 week ago

A is correct. Create a health check on port 443 and use that when creating the Managed Instance Group.

upvoted 1 times
- yuvi69**

4 months, 1 week ago

correct option is A

upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

A is Correct.

upvoted 1 times
- nliaustemac**

5 months ago

A. is the correct answer:  
Create a health check on port 443 and use that when creating the Managed Instance Group.

upvoted 1 times

- 🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago  
A - Create a health check on port 443 and use that when creating the Managed Instance Group.  
upvoted 1 times
- 🗨️ 👤 **JackGlemins** 6 months ago  
I think C is righth.  
[https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs#example\\_health\\_check\\_set\\_up](https://cloud.google.com/compute/docs/instance-groups/autohealing-instances-in-migs#example_health_check_set_up)  
  
A dont talk about recreate de VM, just check  
upvoted 1 times
- 🗨️ 👤 **aurora21** 1 month, 3 weeks ago  
Managed instance group - will recreate the VM automatically with right health check  
upvoted 1 times
- 🗨️ 👤 **arsav** 5 months, 2 weeks ago  
Jack, C just adds label. You first need to create a health check and then label it. Creation of health check is missing in option C. So only right choice is option A  
upvoted 3 times
- 🗨️ 👤 **DucSiu** 6 months ago  
A. Create a health check on port 443 and use that when creating the Managed Instance Group.  
upvoted 1 times
- 🗨️ 👤 **INASR** 6 months ago  
it is A for sure. while you are creating the MIG, set the Autohealing and add the port 443 .  
upvoted 2 times
- 🗨️ 👤 **victory108** 6 months, 1 week ago  
A - Create a health check on port 443 and use that when creating the Managed Instance Group.  
upvoted 2 times
- 🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
why not C ?  
upvoted 1 times
- 🗨️ 👤 **magistrum** 7 months ago  
labels by themselves don't do anything  
upvoted 5 times

Your company has a Google Cloud Platform project that uses BigQuery for data warehousing. Your data science team changes frequently and has few members.

You need to allow members of this team to perform queries. You want to follow Google-recommended practices. What should you do?

- A. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery jobUser role to the group.
- B. 1. Create an IAM entry for each data scientist's user account. 2. Assign the BigQuery dataViewer user role to the group.
- C. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery jobUser role to the group.
- D. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery dataViewer user role to the group.

Correct Answer: D

Reference:

<https://cloud.google.com/bigquery/docs/cloud-sql-federated-queries>

- Gini

Highly Voted

1 year, 2 months ago

C is correct because dataViewer does not allow user to perform queries. jobUser can.

upvoted 38 times
- rbrto

1 year, 2 months ago

dataviewer can perform queries as well. D is correct

upvoted 3 times
- FunkyTechnician

2 months, 3 weeks ago

<https://cloud.google.com/bigquery/docs/access-control>

roles/bigquery.jobUser

BigQuery Job User

Provides permissions to run jobs, including queries, within the project.

upvoted 1 times
- Dips\_

7 months, 1 week ago

no it does not.

BigQuery Job User Provides permissions to run jobs, including queries, within the project.

upvoted 1 times
- nightflyer

7 months, 1 week ago

incorrect

upvoted 2 times
- zukko78

Highly Voted

1 year, 2 months ago

C is correct, doc's said: When applied to a dataset, dataViewer provides permissions to:

Read the dataset's metadata and to list tables in the dataset.

Read data and metadata from the dataset's tables.

When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 21 times
- spudleymcdudley

1 year ago

listen to this guy, he's right

upvoted 2 times
- spudleymcdudley

1 year ago

Ref 'D' data viewer = When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

upvoted 3 times
- MONICA291012

Most Recent

1 week, 5 days ago

I think it should be C

BigQuery Data Viewer: When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.

This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.  
Read data and metadata from the dataset's tables.

BigQuery Job User: Provides permissions to run jobs, including queries, within the project.  
reference: <https://cloud.google.com/bigquery/docs/access-control>  
upvoted 1 times

🗲️ 👤 **arsh1916** 2 months, 1 week ago  
C is correct  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago  
C is correct  
upvoted 1 times

🗲️ 👤 **aa\_desh** 3 months ago  
D is correct. Data science teams does not need to change existing data state, They can perform all analysis with dataViewer role and export result.  
upvoted 2 times

🗲️ 👤 **kopper2019** 3 months, 2 weeks ago  
C, roles/bigquery.jobUser BigQuery Job User Provides permissions to run jobs, including queries, within the project.  
upvoted 1 times

🗲️ 👤 **Pawar\_Shah** 3 months, 4 weeks ago  
<https://www.udemy.com/course/google-cloud-associate-engineer-exam-practice-tests/?couponCode=21CDE6A4C2B95F79BD97>  
upvoted 2 times

🗲️ 👤 **muhasinem** 4 months ago  
C is Correct .  
tested on GCP , with data viewer role it shows error " User does not have bigquery.jobs.create permission in project" and working with job query user .  
upvoted 1 times

🗲️ 👤 **Rajusrinivasa** 4 months ago  
Ans) C  
roles/bigquery.jobUser BigQuery Job User - Provides permissions to run jobs, including queries, within the project.

BigQuery dataViewer user - When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.

<https://cloud.google.com/bigquery/docs/access-control>  
upvoted 2 times

🗲️ 👤 **squishy\_fishy** 2 months ago  
This is most detailed answer. C is correct.  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago  
C is correct. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery jobUser role to the group.  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago  
Taking above one back. Correct answer is D. D. 1. Create a dedicated Google group in Cloud Identity. 2. Add each data scientist's user account to the group. 3. Assign the BigQuery dataViewer user role to the group.  
upvoted 1 times

🗲️ 👤 **tavva\_prudhvi** 4 months ago  
roles/bigquery.jobUser provides permissions to run jobs, including queries, within the project.  
bigquery.jobs.create  
resourceManager.projects.get  
resourceManager.projects.list  
  
Still you wanna reside with D?  
upvoted 1 times

🗲️ 👤 **j1c4b** 4 months, 1 week ago  
You want to follow Google-recommended practices - minimum permission.  
BigQuery Data Viewer - allow you to access to view datasets and all of their contents.  
BigQuery Job User - allow you much more. run jobs, including queries, within the project. This role can check the existence of all jobs, enumerate their own jobs, and cancel their own jobs.  
Answer D  
upvoted 1 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago  
correct answer is C. because 'BigQuery Job User' Provides permissions to run jobs, including queries, within the project.

upvoted 2 times

  **sumanshu** 4 months, 2 weeks ago

Vote for 'C'

upvoted 1 times



  **neerajgoyal** 4 months, 2 weeks ago

C should be the answer

As per doc

BigQuery Job User Provides permissions to run jobs, including queries, within the project.

upvoted 2 times

  **lio123** 4 months, 2 weeks ago

D



jobUser has permission to run jobs. It is additional access as it did not specifically mention in the question.

upvoted 1 times

  **tavva\_prudhvi** 4 months, 1 week ago

In the question, they have mentioned the phrase ' run the queries' which is actually the jobUser but the dataViewer cannot run the queries!!!

upvoted 1 times

  **Jacky\_YO** 4 months, 3 weeks ago

Answer : D

Predefined roles and permissions

<https://cloud.google.com/bigquery/docs/access-control>

1. BigQuery permissions and predefined IAM roles

To grant access to a BigQuery resource, assign one or more roles to a user, group, or service account. You can grant access at the following

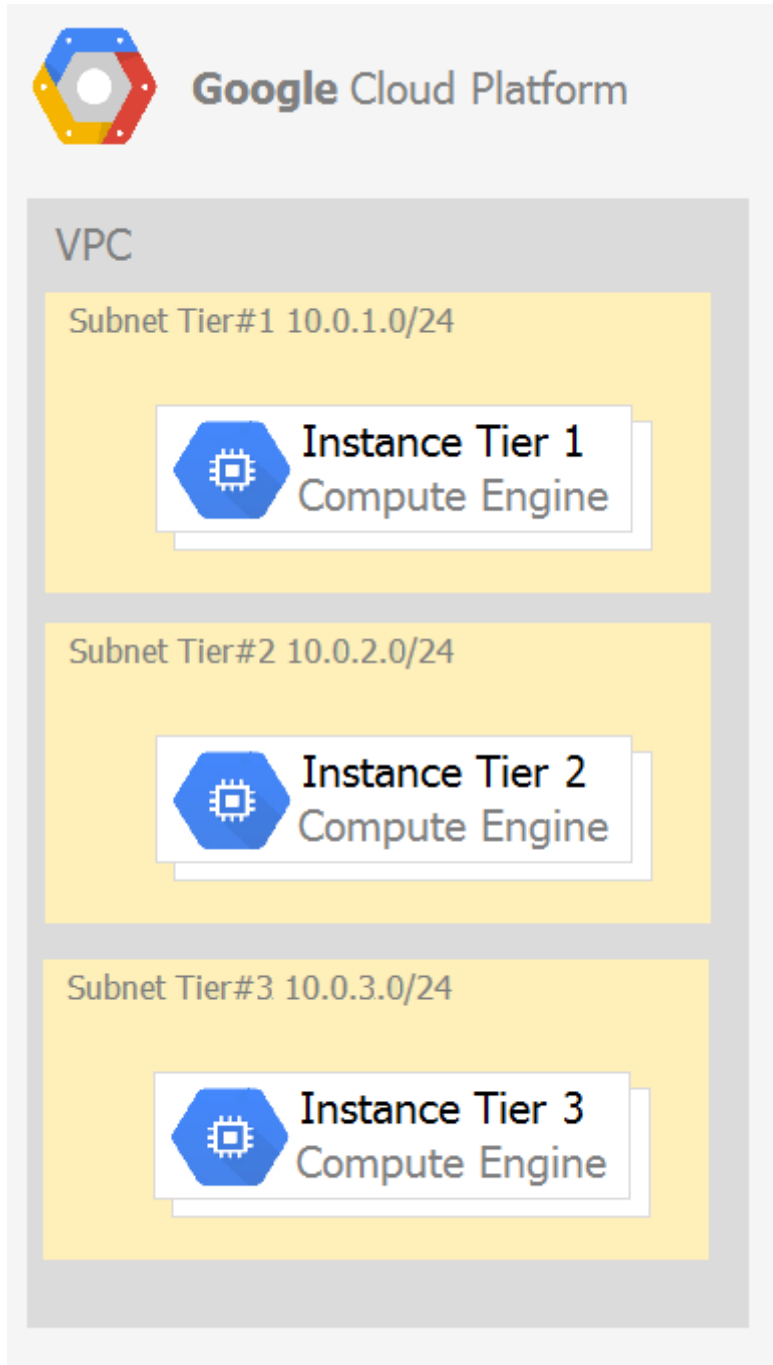
BigQuery resource levels:

- organization or Google Cloud project level
- dataset level
- table or view level

2. BigQuery predefined IAM roles

upvoted 1 times

Your company has a 3-tier solution running on Compute Engine. The configuration of the current infrastructure is shown below.





Each tier has a service account that is associated with all instances within it. You need to enable communication on TCP port 8080 between tiers as follows:



- Instances in tier #1 must communicate with tier #2.
- Instances in tier #2 must communicate with tier #3.

What should you do?

- A. 1. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow all
- B. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow TCP: 8080
- C. 1. Create an ingress firewall rule with the following settings: Targets: all instances with tier #2 service account Source filter: all instances with tier #1 service account Protocols: allow all 2. Create an ingress firewall rule with the following settings: Targets: all instances with tier #3 service account Source filter: all instances with tier #2 service account Protocols: allow all
- D. 1. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.2.0/24) Protocols: allow TCP: 8080 2. Create an egress firewall rule with the following settings: Targets: all instances Source filter: IP ranges (with the range set to 10.0.1.0/24) Protocols: allow TCP: 8080

**Correct Answer:** B

  **sunilw** 1 month, 1 week ago  
B is correct  
upvoted 3 times



  **vmart** 1 month, 2 weeks ago  
B is correct

upvoted 2 times

  **shankyomre01** 2 months ago



B is correct

upvoted 4 times

  **mcaromit** 2 months, 2 weeks ago


B is correct







upvoted 3 times

  **kopper2019** 3 months, 2 weeks ago



if you see closely, port 8080 and service account is required so B is the answer without reading all answers

upvoted 3 times

  **[Removed]** 4 months, 1 week ago


B is correct. 1. Create an ingress firewall rule with the following settings:  Targets: all instances with tier #2 service account  Source filter: all instances with tier #1 service account  Protocols: allow TCP:8080 2. Create an ingress firewall rule with the following settings:  Targets: all instances with tier #3 service account  Source filter: all instances with tier #2 service account  Protocols: allow TCP: 8080

upvoted 2 times

  **yuvi69** 4 months, 1 week ago



Correct answer is B

upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago



Vote for 'B'

upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

Option D is egress - So eliminated

upvoted 1 times

  **The\_Bro** 4 months, 2 weeks ago

B is correct.

upvoted 4 times



## Question #54

## Topic 1

You are given a project with a single Virtual Private Cloud (VPC) and a single subnetwork in the us-central1 region. There is a Compute Engine instance hosting an application in this subnetwork. You need to deploy a new instance in the same project in the europe-west1 region. This new instance needs access to the application. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- B. 1. Create a VPC and a subnetwork in europe-west1. 2. Expose the application with an internal load balancer. 3. Create the new instance in the new subnetwork and use the load balancer's address as the endpoint.
- C. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Use Cloud VPN to connect the two subnetworks. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.
- D. 1. Create a VPC and a subnetwork in europe-west1. 2. Peer the 2 VPCs. 3. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

**Correct Answer: A**

  **Agents89** Highly Voted 1 year, 3 months ago

A is correct  
upvoted 23 times

  **vmart** Most Recent 1 month, 2 weeks ago

A is the correct option  
upvoted 2 times

  **AD\_0525** 2 months ago

Could someone please explain if C can be an option as well.  
upvoted 1 times

  **Finger41** 1 month, 4 weeks ago

Cloud VPN is used when connecting on-premise infrastrucutre to the cloud. In this case creating a new subnet in a new region, and using the internal IP to communicate is sufficient enough to fulfillthe requirements.  
upvoted 2 times

  **sanhoo** 1 month, 4 weeks ago

cloud VPN is not to connect two subnets in same VPN. It is usually to connect on-premises computer to cloud. we just need a separate subnet because of a separate region . The subnet can stay in same VPC that we they can talk to each other with private IP. Hence A is correct.  
upvoted 2 times

  **arsh1916** 2 months, 1 week ago

A is correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago


A is the correct option  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct  
upvoted 1 times

  **padma29** 4 months, 3 weeks ago

why not option D? Someone explain please  
upvoted 2 times

  **kopper2019** 3 months, 2 weeks ago

no need to create another VPC, just another subnet in another region in the same VPC they will communicate transparently  
upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

Why you want to create a separate NETWORK (VPC) , when we can do connection by creating a separate subnet inside the same VPC..And when in same VPC, resources can access each other using internal IP adress

upvoted 2 times

  **nliaustemac** 5 months ago

A. is the correct answer:

1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

  **Gohs** 5 months, 3 weeks ago

Ans: A

Adding an internal load balancer will increase total cost which isn't even necessary at all.

upvoted 1 times

  **DucSiu** 6 months ago

A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 2 times

  **victory108** 6 months, 1 week ago



A - 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

can you pls explain - why you choose B & C ?

upvoted 1 times



  **Wachy** 7 months, 1 week ago

@Bhagirathi

C is incorrect because:

- Cloud VPN requires that addresses be reserved for connection thus, you will pay for reserved IP and also the resource been used. VPN is best when performing data transfer between cloud and on prem with 3GBps max. Nothing is happening with on Prem and this, it is not needed.

upvoted 1 times

  **Wachy** 7 months, 1 week ago

@Bhagirathi

B is incorrect because:

- VPC is a global resource and thus, there is no need to create a new VPC. Her subnetworks are regional too.  
- Exposing the application to an internal load balancer will cost more and it is not needed at all.

upvoted 2 times

  **Bhagirathi** 7 months, 3 weeks ago

A ? why not B/C ..

upvoted 1 times

  **Dineshvon** 7 months, 3 weeks ago

A is correct

upvoted 1 times

  **swatititame** 8 months, 1 week ago

A. 1. Create a subnetwork in the same VPC, in europe-west1. 2. Create the new instance in the new subnetwork and use the first instance's private address as the endpoint.

upvoted 1 times

Your projects incurred more costs than you expected last month. Your research reveals that a development GKE container emitted a huge number of logs, which resulted in higher costs. You want to disable the logs quickly using the minimum number of steps. What should you do?

- A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.
- B. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE Cluster Operations resource.
- C. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Logging.
- D. 1. Go to the GKE console, and delete existing clusters. 2. Recreate a new cluster. 3. Clear the option to enable legacy Stackdriver Monitoring.

**Correct Answer: A**

- Gini

Highly Voted

1 year, 2 months ago

The question mentioned that "GKE container emitted a huge number of logs", in my opinion A is correct.  
upvoted 30 times
- berezinsn

Highly Voted

1 year, 2 months ago

A correct  
upvoted 9 times
- vmart

Most Recent

1 month, 2 weeks ago

A is right  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- [Removed]

4 months, 1 week ago

A is correct. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.  
upvoted 1 times
- yuvi69

4 months, 1 week ago

Correct answer is A  
upvoted 1 times
- sumanshu

4 months, 2 weeks ago

Vote for A  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

A is the Answer.  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

A is the Answer.  
upvoted 1 times
- kartikjena31

4 months, 4 weeks ago

Ans is A  
upvoted 1 times
- nliaustemac

5 months ago

A. is the correct answer:  
1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.  
upvoted 1 times
- GCP\_Student1

5 months, 1 week ago

A. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.  
upvoted 1 times
- JackGlemins

5 months, 4 weeks ago

I think A is right.  
<https://cloud.google.com/logging/docs/api/v2/resource-list>  
  
GKE Containers have more log than GKE Cluster Operations:

.-GKE Containe:

cluster\_name: An immutable name for the cluster the container is running in.

namespace\_id: Immutable ID of the cluster namespace the container is running in.

instance\_id: Immutable ID of the GCE instance the container is running in.

pod\_id: Immutable ID of the pod the container is running in.

container\_name: Immutable name of the container.

zone: The GCE zone in which the instance is running.

VS

.-GKE Cluster Operations

project\_id: The identifier of the GCP project associated with this resource, such as "my-project".

cluster\_name: The name of the GKE Cluster.


location: The location in which the GKE Cluster is running.

upvoted 5 times

  **DucSiu** 6 months ago

B. 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE Cluster Operations resource.

upvoted 1 times

  **pas77** 2 weeks, 6 days ago

Why disabling logging for the cluster? The question is not about a cluster but about a container. A is the correct answer.

upvoted 1 times

  **victory108** 6 months, 1 week ago

A - 1. Go to the Logs ingestion window in Stackdriver Logging, and disable the log source for the GKE container resource.

upvoted 1 times

  **msselway** 7 months ago

Answer: B

Cloud Logging monitored resource type is "gke\_cluster", "GKE Cluster Operations" - <https://cloud.google.com/logging/docs/api/v2/resource-list#resource-types>

Though a container is the source of the logs, the smallest unit of deployment in GKE is the Pod. Therefore I don't think you can make the change at the container level, ruling out answer A.

C & D are incorrect because they delete the clusters.

upvoted 2 times

  **pca2b** 4 months ago

containers are ephemeral, you can disable logging for the pod containing the specific offending container, but that could be recreated anytime, resulting in a new pod/container to chase. You could get offending log surge again, even before figuring out the reason in the first place.

upvoted 1 times

  **pca2b** 3 months, 3 weeks ago

re-reading <https://cloud.google.com/logging/docs/api/v2/resource-list>.

i change vote to A, as it is possible to go to StackDriver Logging and disable the source for a specific container. This will stop logging on that container asap, while you troubleshoot

A

upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

You guys again confused every one...A or B ? what is the best option to choose?


upvoted 1 times

  **pas77** 2 weeks, 6 days ago

If you read the question you'll notice the correct answer is A.

Container vs Cluster

upvoted 1 times

  **alpharomeo9** 7 months, 2 weeks ago

I would go with A.

<https://cloud.google.com/logging/docs/api/v2/resource-list>

As mentioned by Fidget\_ in the above comments, GKE Container has more labels and the question is container specific.

upvoted 2 times

## Question #56

## Topic 1

You have a website hosted on App Engine standard environment. You want 1% of your users to see a new test version of the website. You want to minimize complexity. What should you do?

- A. Deploy the new version in the same application and use the --migrate option.
- B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.
- C. Create a new App Engine application in the same project. Deploy the new version in that application. Use the App Engine library to proxy 1% of the requests to the new version.
- D. Create a new App Engine application in the same project. Deploy the new version in that application. Configure your network load balancer to send 1% of the traffic to that new application.

**Correct Answer: B**

🗲️ 👤 **yasu** Highly Voted 1 year, 4 months ago

I will prefer B as the answer.. why we need create new application?  
upvoted 34 times

🗲️ 👤 **YAS007** 5 days, 22 hours ago

more over, in app engine we cannot create "new application", we have to create a new Project to do that, an app engine projet has 1 application (which can have multiple versions and services)  
upvoted 1 times

🗲️ 👤 **sanhoo** 1 month, 4 weeks ago

Agree B is correct. creating a new application in the same project for app engine is anyways not possible.  
upvoted 2 times

🗲️ 👤 **Gini** Highly Voted 1 year, 3 months ago

I agree with yasu. And only one app engine can exist in one project. B is the best choice, simple and easy.  
upvoted 12 times

🗲️ 👤 **vmart** Most Recent 1 month, 2 weeks ago

B is correct  
upvoted 1 times

🗲️ 👤 **arsh1916** 2 months, 1 week ago

B is correct  
upvoted 1 times

🗲️ 👤 **Agraved** 2 months, 2 weeks ago

Answer is B, new application is not possible at all.  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

B is the only correct option  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago

B is correct. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.  
upvoted 1 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago

Correct answer is B  
upvoted 1 times

🗲️ 👤 **nliaustemac** 5 months ago

B. is the correct answer:  
Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.  
upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.  
upvoted 1 times

🗨️ 👤 **iPablo26** 5 months, 1 week ago

I think B is the answer  
upvoted 1 times

🗨️ 👤 **d18mph23** 5 months, 3 weeks ago

B is correct.  
--splits=SPLITS,[SPLITS,...]  
Key-value pairs describing what proportion of traffic should go to each version. The split values are added together and used as weights. The exact values do not matter, only their relation to each other. For example, v1=2,v2=2 is equivalent to v1=.5,v2=.5  
upvoted 3 times

🗨️ 👤 **DucSiu** 6 months ago

B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.  
upvoted 4 times

🗨️ 👤 **INASR** 6 months ago

B is right  
upvoted 2 times

🗨️ 👤 **swatititame** 8 months, 2 weeks ago

• B. Deploy the new version in the same application and use the --splits option to give a weight of 99 to the current version and a weight of 1 to the new version.  
upvoted 3 times

🗨️ 👤 **Kuchaarun** 9 months ago

B should be the right answer  
upvoted 1 times

🗨️ 👤 **nwk** 9 months ago

(B) [https://cloud.google.com/appengine/docs/standard/python3/splitting-traffic#splitting\\_traffic\\_across\\_multiple\\_versions](https://cloud.google.com/appengine/docs/standard/python3/splitting-traffic#splitting_traffic_across_multiple_versions)  
upvoted 2 times

## Question #57

## Topic 1

You have a web application deployed as a managed instance group. You have a new version of the application to gradually deploy. Your web application is currently receiving live web traffic. You want to ensure that the available capacity does not decrease during the deployment. What should you do?

- A. Perform a rolling-action start-update with maxSurge set to 0 and maxUnavailable set to 1.
- B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.
- C. Create a new managed instance group with an updated instance template. Add the group to the backend service for the load balancer. When all instances in the new managed instance group are healthy, delete the old managed instance group.
- D. Create a new instance template with the new application version. Update the existing managed instance group with the new instance template. Delete the instances in the managed instance group to allow the managed instance group to recreate the instance using the new instance template.

**Correct Answer: C**

  **Carls** Highly Voted 1 year, 3 months ago

Correct option is B. We need to ensure the global capacity remains intact, for that reason we need to establish maxUnavailable to 0. On the other hand, we need to ensure new instances can be created. We do that by establishing the maxSurge to 1. Option C is more expensive and more difficult to set up and option D won't meet requirements since it won't keep global capacity intact.

upvoted 48 times

  **JavierCorrea** Highly Voted 10 months, 3 weeks ago

I take my own previous comment back. It's definitely B.

upvoted 12 times

  **YAS007** Most Recent 5 days, 21 hours ago

B is correct

[https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups#max\\_unavailable](https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups#max_unavailable)

upvoted 2 times

  **shankyomre01** 2 months ago


B is correct. GCP provides max surge and max unavailable options for this kind of update only.

upvoted 1 times

  **arsh1916** 2 months, 1 week ago

B is Correct

upvoted 1 times

  **Agraved** 2 months, 2 weeks ago

Answer is b, explanation: If you do not want any unavailable machines during an update, set the maxUnavailable value to 0 and the maxSurge value to greater than 0. With these settings, Compute Engine removes each old machine only after its replacement new machine is created and running.

Ref: <https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups>

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago


B is correct

upvoted 1 times

  **kopper2019** 3 months, 2 weeks ago

B, <https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups#options>

upvoted 1 times

  **[Removed]** 4 months, 1 week ago



B is correct. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.

upvoted 1 times

  **yuvi69** 4 months, 1 week ago

Correct answer is B

upvoted 1 times

  **pondai** 4 months, 1 week ago

B is correct

<https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups>

If you do not want any unavailable machines during an update, set the maxUnavailable value to 0 and the maxSurge value to greater than 0. With these settings, Compute Engine removes each old machine only after its replacement new machine is created and running.



upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

B is correct

upvoted 1 times

  **ShakthiGCP** 4 months, 4 weeks ago

Ans: B : "If you do not want any unavailable machines during an update, set the maxUnavailable value to 0 and the maxSurge value to greater than 0. With these settings, Compute Engine removes each old machine only after its replacement new machine is created and running." and "Note: If you set both maxSurge and maxUnavailable properties and both properties resolve to 0, the Updater automatically sets maxUnavailable=1, to ensure that the automated update can always proceed." [https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups?\\_ga=2.33763493.-527346271.1610729770#max\\_surge](https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups?_ga=2.33763493.-527346271.1610729770#max_surge)

upvoted 3 times

  **nliaustemac** 5 months ago

B. is the correct answer:

Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.

upvoted 1 times

  **Gohs** 5 months, 3 weeks ago

I choose C.

This is because managed instance groups are usually attached to a load balancer — so that all of them will be under one IP address.

So there's no significant increase in cost if you initially get your update running on new VMs, then later deleting the old VMs.

C is more explanatory than B.

In B, setting maxUnavailable to zero means don't delete any VM.

Setting maxSurge to one means create one new VM for each old VM.

You have essentially duplicated your VMs.

B didn't say anything about deleting old VMs.

This will generate more cost in the long run.

upvoted 2 times

  **arsav** 5 months, 2 weeks ago

Gohs, question says your web application is receiving live web traffic. If you choose option C, if your old managed instance group was serving traffic to the web application. During its deletion it could impact the application. Also the question says ensure the available capacity does not decrease during deployment. Setting the Max surge to 1 here will ensure even one node is taken down, MIG will ensure one active instance is serving the live traffic.

upvoted 3 times

  **DucSiu** 6 months ago

B. Perform a rolling-action start-update with maxSurge set to 1 and maxUnavailable set to 0.

upvoted 1 times



Question #58

Topic 1

You are building an application that stores relational data from users. Users across the globe will use this application. Your CTO is concerned about the scaling requirements because the size of the user base is unknown. You need to implement a database solution that can scale with your user growth with minimum configuration changes. Which storage solution should you use?

- A. Cloud SQL
- B. Cloud Spanner
- C. Cloud Firestore
- D. Cloud Datastore

Correct Answer: B

- Fidget\_**

Highly Voted

12 months ago

B

Cloud SQL for small relational data, scaled manually

Cloud Spanner for relational data, scaled automatically

Cloud Firestore for app-based data(?)

Cloud Datastore for non-relational data

Correct me if i'm wrong

upvoted 37 times
- karol\_wu**

Highly Voted

1 year, 4 months ago

in my opinion correct is B

upvoted 25 times
- Bala13**

Most Recent

2 months ago

for globalization and relational data...Cloud spanner is correct Option B

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

B is correct

upvoted 1 times
- jahnu**

3 months, 1 week ago

B is correct

upvoted 1 times
- [Removed]**

4 months, 1 week ago

B is correct. Cloud Spanner

upvoted 1 times
- yuvi69**

4 months, 1 week ago

Correct answer is B. because cloud spanner is global

upvoted 1 times
- pondai**

4 months, 1 week ago

When you see "relational data" ,"globe"the answer always cloud spanner

upvoted 3 times
- pas77**

2 weeks, 6 days ago

Not really. Cloud SQL is also relational. The answer is B, but not for that specific reason.

upvoted 1 times
- Devgela**

4 months, 2 weeks ago

Keywords global. Is B

upvoted 1 times
- chithra1990**

5 months ago

When you see Global and Relational - and requires Scaling keywords - Option B is correct Cloud Spanner

upvoted 2 times
- nliaustemac**

5 months ago

B. is the correct answer:

Cloud Spanner

upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

B. Cloud Spanner  
upvoted 2 times

🗨️ 👤 **iPablo26** 5 months, 1 week ago

The answer is B. Cloud Spanner is a relational data model, a globally distributed database service and storage solution.  
upvoted 1 times

🗨️ 👤 **nitinz** 5 months, 2 weeks ago

Cloud Spanner is RDBMS, Cloud Datastore is No-SQL. So B because option B can scale as well.  
upvoted 1 times

🗨️ 👤 **DucSiu** 6 months ago

B. Cloud Spanner  
upvoted 1 times

🗨️ 👤 **INASR** 6 months ago

B is correct for scalability . I don't know why someone will say D and confuse everyone. Relational data are cloud SQL and Spanner only.  
upvoted 2 times



🗨️ 👤 **EABDAJA** 6 months, 2 weeks ago

B is the Answer.Cloud Spanner is best for Large scale data Application. Uses:Global consistency, Financial trading & E-commerce.  
upvoted 1 times



You are the organization and billing administrator for your company. The engineering team has the Project Creator role on the organization. You do not want the engineering team to be able to link projects to the billing account. Only the finance team should be able to link a project to a billing account, but they should not be able to make any other changes to projects. What should you do?

- A. Assign the finance team only the Billing Account User role on the billing account.
- B. Assign the engineering team only the Billing Account User role on the billing account.
- C. Assign the finance team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.
- D. Assign the engineering team the Billing Account User role on the billing account and the Project Billing Manager role on the organization.

Correct Answer: D



-   **Bharathy** Highly Voted 1 year, 4 months ago

Option A is correct, as we don't want the engineering team to link projects to billing account and want only the Finance team. Billing Account User role will help to link projects to the billing account...



upvoted 36 times
-   **naveedpk00** 11 months, 2 weeks ago

Option A is incorrect: Reason-  
This role has very restricted permissions, so you can grant it broadly, typically in combination with Project Creator. These two roles allow a user to create new projects linked to the billing account on which the role is granted.  
Reference: <https://cloud.google.com/billing/docs/how-to/billing-access>



I will go with option C.

upvoted 11 times
-   **Hasaaaaan** 1 month, 3 weeks ago


Billing Account User also enables the user to make changes in resources.

upvoted 1 times
-   **djgodzilla** 1 month, 3 weeks ago



wrong : you can't link project to billing accounts without Project billing manager.  
C is Correct

upvoted 1 times
-   **djgodzilla** 1 month, 3 weeks ago



my bad , it seems A,C are both technically correct

upvoted 1 times
-   **mwwoodm** 10 months, 3 weeks ago


Option A makes the most sense since Billing Account User can link projects to the billing account and the question reinforces principle of least privilege. Source: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 6 times
-   **measme** Highly Voted 1 year, 1 month ago


for me is C:  
[https://cloud.google.com/billing/docs/how-to/modify-project#permissions\\_required\\_for\\_this\\_task\\_2](https://cloud.google.com/billing/docs/how-to/modify-project#permissions_required_for_this_task_2)  
"Roles with adequate permissions to perform this task:  
\* Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account."

upvoted 16 times
-   **PB101** Most Recent 2 weeks, 4 days ago



Option A is correct. Question states that -"Only the finance team should be able to link a project to a billing account, but they should not be able to make any other changes to projects.". So finance team needs a restrictive access on billing account and should limits to link project to billing account. Refer: <https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 1 times
-   **tibi1111** 2 weeks, 4 days ago

Is it option A or C?

upvoted 1 times
-   **associatecloudexamuser** 2 weeks, 5 days ago

Correct answer is C. Need to have project billing manager role each different project

upvoted 2 times
-   **IamPapulalji** 1 month, 1 week ago

I created an org and tested. This question is wrong as stated. A & D does not work as you still cannot see the Project with the stated options.

Project Billing Manager role stated in option D also does not show the project... it only shows the org

What works is :

1. Assign the finance team only the Billing Account User role on the billing account
2. Assign the finance team Project Browser role. Without the Browser role finance cannot see the project to set billing.

Hence option D probably should be Billing Account User role + Project Browser role

upvoted 1 times

  **jcloud965** 2 weeks, 1 day ago

You don't need to have Project Browser role to attach a billing account.  
Only Project Billing Manager on the project is required.

upvoted 1 times

  **PeppaPig** 1 month, 3 weeks ago

C is the correct answer.

A is not enough. By granting only billing account user role, the finance team would not even be able to see the projects, not to mention linking the projects to billing

upvoted 1 times

  **happyBoo** 1 month, 3 weeks ago

Ans is C as que is talking about organization and don't wanna allow other changes to PROJECTS it should be a broad access

upvoted 2 times

  **shankyomre01** 2 months ago

Option A is right.

<https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

check this

upvoted 1 times

  **viswanand** 2 months, 1 week ago

Option A has to be correct because the Billing Account is separate account leaves outside of Project and it does not belong to Organization as well

upvoted 1 times

  **viswanand** 2 months, 1 week ago

Also the question states only to link projects to Billing Account so User Role is enough

upvoted 1 times

  **Enzo** 2 months, 1 week ago

A is correct : Project Billing Manager



(roles/billing.projectManager) Link/unlink the project to/from a billing account. Organization, folder, or project. This role allows a user to attach the project to the billing account, but does not grant any rights over resources. Project Owners can use this role to allow someone else to manage the billing for the project without granting them resource access. And this is not required

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago

A is correct...Project Billing Manager would additionally grant the permission to disable billing for the project which is redundant in this scenario

upvoted 1 times

  **lxgywil** 2 months, 3 weeks ago

The answer is "A".

The question says that the finance team should (ONLY) be able to link a project to a billing account, but not be able to make any other changes to projects.

The purpose of the Billing Account User role is to link projects to billing accounts. Also, from its description:

"This role has very restricted permissions, so you can grant it broadly, typically in combination with Project Creator. These two roles allow a user to create new projects linked to the billing account on which the role is granted."

Reference: <https://cloud.google.com/billing/docs/how-to/billing-access#overview-of-cloud-billing-roles-in-cloud-iam>

upvoted 1 times

  **[Removed]** 4 months, 1 week ago


A is correct. Assign the finance team only the Billing Account User role on the billing account.

upvoted 2 times

  **[Removed]** 4 months ago



Taking it back. Correct answer is C.



upvoted 2 times

  **yuvi69** 4 months, 1 week ago

Correct answer is C

upvoted 3 times

  **sumanshu** 4 months, 2 weeks ago  
vote for 'C'  
upvoted 3 times

  **CloudMom** 4 months, 2 weeks ago  
sounds like a pick-2 answers type of question, no?  
upvoted 1 times

## Question #60

## Topic 1

You have an application running in Google Kubernetes Engine (GKE) with cluster autoscaling enabled. The application exposes a TCP endpoint. There are several replicas of this application. You have a Compute Engine instance in the same region, but in another Virtual Private Cloud (VPC), called gce-network, that has no overlapping IP ranges with the first VPC. This instance needs to connect to the application on GKE. You want to minimize effort. What should you do?

- A. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Set the service's externalTrafficPolicy to Cluster. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- B. 1. In GKE, create a Service of type NodePort that uses the application's Pods as backend. 2. Create a Compute Engine instance called proxy with 2 network interfaces, one in each VPC. 3. Use iptables on this instance to forward traffic from gce-network to the GKE nodes. 4. Configure the Compute Engine instance to use the address of proxy in gce-network as endpoint.
- C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.
- D. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add a Cloud Armor Security Policy to the load balancer that whitelists the internal IPs of the MIG's instances. 3. Configure the Compute Engine instance to use the address of the load balancer that has been created.

**Correct Answer: A**

  **someoneinthecloud** Highly Voted 12 months ago

I believe it's A. It's never mentioned in the question that traffic cannot go through the Internet but it's mentioned that effort should be minimized. A requires a lot less effort than C to accomplish the same (no VPC peering, per example).  
upvoted 17 times

  **ShakthiGCP** 4 months, 3 weeks ago

Ans: A . This sounds correct and avoids unnecessary steps in C. C is also correct but compared to it, A is much easier to achieve. Go over Kubernetes Loadbalancer concepts to get more details. Initially i was thinking C is the Answer. but after putting some time on K8's Network - changed my mind to A.  
upvoted 4 times

  **AmitKM** 11 months, 1 week ago

Yeah, I feel the same. Nowhere does it say that the traffic has to be internal. But it does say "minimal effort" which I feel is option A.  
upvoted 4 times

  **juancambb** Highly Voted 1 year, 2 months ago

i think C is better solution, the solution A pass traffic through public internet, also C by internal network and the "no overlap ips" in the statement suggest that.  
upvoted 16 times

  **WindDriver** Most Recent 2 weeks, 2 days ago

I go with A, because C is "wrong" :)  
<https://cloud.google.com/kubernetes-engine/docs/how-to/internal-load-balancing>  
Using an internal TCP/UDP load balancer  
This page explains how to create a Compute Engine internal TCP/UDP load balancer on Google Kubernetes Engine.  
Internal TCP/UDP Load Balancing makes your cluster's services accessible to applications outside of your cluster that use the same VPC network and are located in the same Google Cloud region.  
SAME VPC NETWORK IS NEEDED, and "another Virtual Private Cloud (VPC)" is mentioned.  
upvoted 1 times

  **contaexamtopics** 2 months, 1 week ago

Very confusing - finally what to choose...

A or C ??

upvoted 1 times

  **Agraved** 2 months, 2 weeks ago

Answer should be C, Opting A changes the exposure of GKE endpoints, thus even if it less work makes it incorrect. Also the question indicates to explore VPC sharing, as this is one of the most suitable case.  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

Both A & C would work, but if effort is to be minimized then A is the best choice  
upvoted 2 times

- 🗳️ 👤 **meh009** 3 months, 1 week ago  
Another tricky one between A and C however again, approaching this from hints within this specific Q (No Conflicting IPs) I believe the answer is C.  
upvoted 2 times
- 🗳️ 👤 **Furqon** 1 month, 2 weeks ago  
Indeed. I think previously is A, but there is a "no conflicting IPs" statement.  
upvoted 1 times
- 🗳️ 👤 **[Removed]** 4 months, 1 week ago  
C is correct. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.  
upvoted 2 times
- 🗳️ 👤 **yuvi69** 4 months, 1 week ago  
Correct answer is A . option C is wrong because internal loadbalancers are used within same VPC network. but here in the question it is mentioned that compute engine instance is present in another VPC .  
<https://cloud.google.com/kubernetes-engine/docs/how-to/internal-load-balancing>  
upvoted 3 times
- 🗳️ 👤 **sanhoo** 1 month, 4 weeks ago  
In the link you provided it is mentioned as "In addition, clients in a VPC network connected to the LoadBalancer network using VPC Network Peering can also access the Service."  
In option C they are talking about VPC peering. So technically it is possible. I tend to believe C is correct answer  
upvoted 2 times
- 🗳️ 👤 **pondai** 4 months, 1 week ago  
If use vpc peer. Use internal loadbalancer is a good plan(Need to configure firewall).  
<https://cloud.google.com/load-balancing/docs/l7-internal/internal-https-lb-and-other-networks>  
But we need min effort so I think A is right.  
upvoted 1 times
- 🗳️ 👤 **Pravin3c** 4 months, 1 week ago  
NOT C - Global access is an optional parameter for internal LoadBalancer Services that allows clients from any region in your VPC network to access the internal TCP/UDP load balancer.  
upvoted 1 times
- 🗳️ 👤 **sumanshu** 4 months, 2 weeks ago  
Vote for 'A'  
upvoted 2 times
- 🗳️ 👤 **GCP\_Student1** 4 months, 3 weeks ago  
C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.  
upvoted 1 times
- 🗳️ 👤 **tavva\_prudhvi** 3 months, 3 weeks ago  
See, both A,C are right but in the question it says to minimize effort, then we can go with A, i guess!  
upvoted 1 times
- 🗳️ 👤 **viswanand** 2 months, 1 week ago  
But Google always recommends security as the best practise and VPC peering is always better  
upvoted 1 times
- 🗳️ 👤 **nliaustemac** 5 months ago  
C. is the correct answer:  
1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.  
upvoted 1 times
- 🗳️ 👤 **nitinz** 5 months, 1 week ago  
common sense A, C will do the same. Requirement is less work, requirements never say you can not go to internet. A and C does the same. A via routing traffic via Internet and C using internally. But less effort in A. I will stick to A.  
upvoted 3 times
- 🗳️ 👤 **tavva\_prudhvi** 3 months, 2 weeks ago  
As you said, A is doing via routing traffic via internet and C is using it internally, so Acc to your statement, it's C,right?  
upvoted 1 times
- 🗳️ 👤 **nitinz** 5 months, 1 week ago  
A is simplest to achieve. I do not see a reason of doing extra work with option C.  
upvoted 2 times
- 🗳️ 👤 **DucSiu** 6 months ago

C. 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.

upvoted 1 times

  **victory108** 6 months, 1 week ago

C - 1. In GKE, create a Service of type LoadBalancer that uses the application's Pods as backend. 2. Add an annotation to this service: cloud.google.com/load-balancer-type: Internal 3. Peer the two VPCs together. 4. Configure the Compute Engine instance to use the address of the load balancer that has been created.

upvoted 1 times



Your organization is a financial company that needs to store audit log files for 3 years. Your organization has hundreds of Google Cloud projects. You need to implement a cost-effective approach for log file retention. What should you do?

- A. Create an export to the sink that saves logs from Cloud Audit to BigQuery.
- B. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.
- C. Write a custom script that uses logging API to copy the logs from Stackdriver logs to BigQuery.
- D. Export these logs to Cloud Pub/Sub and write a Cloud Dataflow pipeline to store logs to Cloud SQL.

**Correct Answer: A**

Reference:

<https://cloud.google.com/logging/docs/audit/>

- yasu

Highly Voted

1 year, 4 months ago

Why not B? cost effective

upvoted 33 times
- Gini

Highly Voted

1 year, 3 months ago

if it is all about cost, B is the best. However, speaking of "audit" you probably need to access the data once in a while, which Coldline storage might not be ideal for this case I guess? I would go for A in the exam though.

upvoted 16 times
- Ale1973

10 months, 4 weeks ago

Be strong!!! If B is the best, go for B!!!

upvoted 7 times
- Nelson2080

Most Recent

3 weeks, 1 day ago

I Agree with B I found this snippet:

When you load data into BigQuery from Cloud Storage, you are not charged for the load operation, but you do incur charges for storing the data in Cloud Storage. After the data is loaded into BigQuery, the data is subject to BigQuery's storage pricing.

<https://cloud.google.com/bigquery/docs/best-practices-storage>

upvoted 1 times
- sunilw

1 month, 1 week ago

Option B:  
Supported destinations for logs exports: Cloud Storage, BigQuery and Pub/Sub out of which Cloud Storage more suitable.  
Cloud Storage: JSON files stored in Cloud Storage buckets; provides inexpensive, long-term storage.  
<https://cloud.google.com/logging/docs/export#supported-destinations>

upvoted 1 times
- Umesh09

2 months, 1 week ago

Option B: Coldline is the most cost effective option for data that is not required for analysis

upvoted 1 times
- arsh1916

2 months, 1 week ago

B is correct

upvoted 1 times
- katos

2 months, 1 week ago

cloud storage cold line is 5 times cheaper than bigquery storage

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct & cheapest option

upvoted 1 times
- jahnu

3 months ago

My Ans: A because coldline storage store only one yaear.so I selected BigQuery.

upvoted 2 times
- Umesh09

2 months, 1 week ago

no it is recommended to store data that is not required for next 1 year to be stored in coldline for cost effectiveness

upvoted 1 times
- mistryminded

3 months ago

no documentation explains that it stores for only 1 year sorry to correct you but answer is B.

upvoted 1 times

🗨️ 👤 **ayush\_1995** 3 months, 2 weeks ago

A

Admin Activity audit logs and System Event audit logs are free.

Data Access audit logs and Policy Denied audit logs are chargeable.

<https://cloud.google.com/logging/docs/audit/>

upvoted 2 times

🗨️ 👤 **r1ck** 3 months, 2 weeks ago

Logs are required for Audit, not regular analysis - cold line is a better option

upvoted 1 times

🗨️ 👤 **r1ck** 3 months, 2 weeks ago

Also requirement here is retention, would this be different because it's a Financial organization?

upvoted 1 times

🗨️ 👤 **[Removed]** 4 months, 1 week ago

B is correct. Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.

upvoted 2 times

🗨️ 👤 **yuvi69** 4 months, 1 week ago

correct option is B. because in the only talked about storing at cheap-cost, so coldline has low storage cost

upvoted 1 times

🗨️ 👤 **pondai** 4 months, 1 week ago

I vote B. It didn't say the file need to audit .It just say store file 3 year & cost-effective .

upvoted 1 times

🗨️ 👤 **sumanshu** 4 months, 2 weeks ago

vote 4 'B'

upvoted 2 times

🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago

B is correct

upvoted 1 times

🗨️ 👤 **nliaustemac** 5 months ago

B. is the correct answer:

Create an export to the sink that saves logs from Cloud Audit to a Coldline Storage bucket.

upvoted 1 times

🗨️ 👤 **nitinz** 5 months, 1 week ago

B it is, its written on the wall - "You need to implement a cost-effective approach for log file retention"

upvoted 1 times

## Question #62

## Topic 1

You want to run a single caching HTTP reverse proxy on GCP for a latency-sensitive website. This specific reverse proxy consumes almost no CPU. You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes. You want to minimize cost. How should you run this reverse proxy?

- A. Create a Cloud Memorystore for Redis instance with 32-GB capacity.
- B. Run it on Compute Engine, and choose a custom instance type with 6 vCPUs and 32 GB of memory.
- C. Package it in a container image, and run it on Kubernetes Engine, using n1-standard-32 instances as nodes.
- D. Run it on Compute Engine, choose the instance type n1-standard-1, and add an SSD persistent disk of 32 GB.

**Correct Answer: B**

  **jzh** Highly Voted 11 months, 2 weeks ago

Go to cloud console and create instance

select Memorystore with Basic tier, select us-central1 and us-central1-a, and capacity 32GB, the cost estimate is \$0.023/GB/hr

select VM instance with custom machine type with 6 vCPUs and 32 GB memory, the same region and zone as Memorystore setting, the cost estimate is \$0.239/hr

Option B will definitely cost more as it adds on CPU usage cost even it uses little in this scenario, but still charge you. So answer is A from real practice example.

upvoted 25 times

  **mexblood1** 10 months, 3 weeks ago

Using pricing calculator matching 730 hrs per month for both.. Memorystore is 537.28 per month and vm (6 cpus 32 gb memory) is 174.41 per month. So vm is still cheaper even with 6 cpus.

upvoted 3 times

  **[Removed]** 10 months, 3 weeks ago

I agree its cheaper, but 2 drawbacks, 1 hit latency, 2 you need to install cache plain won't help even if check. So still think A.

upvoted 2 times

  **[Removed]** 10 months, 1 week ago

Typo correct my statements. There are 2 issues If you create a VM how every cheap there will be a hit with latency when communicating with it. Second, you will need to install/implement a caching system on that.

Whereas cloud memorystore for redis is designed for the sole purpose of HTTP caching which has very low latency compared to any other solution we are thinking of doing.



upvoted 2 times

  **FenixRa73** 6 months, 2 weeks ago

$\$0.023 * 32 = \$0.736$

is it cheaper?

upvoted 1 times

  **SSPC** 11 months, 1 week ago

I agree with you

upvoted 1 times

  **ESP\_SAP** Highly Voted 11 months, 2 weeks ago

Correct Answer should be A:

The question mention "You want to have a 30-GB in-memory cache, and need an additional 2 GB of memory for the rest of the processes"

What is Google Cloud Memorystore?

Overview. Cloud Memorystore for Redis is a fully managed Redis service for Google Cloud Platform. Applications running on Google Cloud Platform can achieve extreme performance by leveraging the highly scalable, highly available, and secure Redis service without the burden of managing complex Redis deployments.

upvoted 17 times

  **ESP\_SAP** 11 months, 1 week ago

Just to complement the answer:

We are looking for "latency-sensitive website"

What it's good for

Memorystore for Redis provides a fast, in-memory store for use cases that require fast, real-time processing of data. From simple caching use cases to real time analytics, Memorystore for Redis provides the performance you need.

Caching: Cache is an integral part of modern application architectures. Memorystore for Redis provides low latency access and high throughput

for heavily accessed data, compared to accessing the data from a disk based backend store. Session management, frequently accessed queries, scripts, and pages are common examples of caching.

[https://cloud.google.com/memorystore/docs/redis/redis-overview#what\\_its\\_good\\_for](https://cloud.google.com/memorystore/docs/redis/redis-overview#what_its_good_for)  
upvoted 12 times

  **[Removed]** 10 months, 3 weeks ago

I agree with your reasoning. Given that the question stresses that this is for a 'latency sensitive website', that's a clue that Redis is part of the answer. Even if spinning up a similarly sized VM were more cost effective, I can't find any documentation that this would provide sufficiently low latency as a memory cache. Yes, you want to keep costs low, but not if it causes your latency-sensitive website problems. Thus I agree that option A is the answer.  
upvoted 1 times

  **magistrum** 7 months ago

Agree, don't think anything you create with the Compute Engine will meet the "You want to have a 30-GB in-memory cache" requirement...that's a very different technology  
upvoted 2 times

  **ready2rock** Most Recent 1 month, 2 weeks ago


I'm sorry - is Redis a reverse proxy server? I don't think it is so I go with B. If it is a reverse cache proxy then A, but I don't think it is  
upvoted 2 times

  **accuracy23** 1 month, 2 weeks ago

Correct. Redis is an in memory key/value store. It won't proxy requests.  
upvoted 1 times

  **Umesh09** 2 months, 1 week ago

Option B : Caching/latency is the key work  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

A is correct...microsecond latency  
upvoted 1 times

  **tavva\_prudhvi** 4 months ago

Here we have 2 main points to think if price -> Vm is cheaper than memstore but when it comes to latency-sensitive website -> memory store is the best.



No, what should i choose?  
upvoted 1 times

  **tavva\_prudhvi** 4 months ago

It's simple, even if you create a VM how cheap it is there will be a hit with latency when communicating with it. Second, you will need to install/implement a caching system on that.  
upvoted 2 times

  **[Removed]** 4 months ago



A is correct. Create a Cloud Memorystore for Redis instance with 32-GB capacity.  
upvoted 2 times

  **yuvi69** 4 months, 1 week ago



correct option is A  
upvoted 2 times

  **sumanshu** 4 months, 2 weeks ago

vote for 'A'  
upvoted 1 times

  **Hi2ALL** 4 months, 3 weeks ago

A is the correct answer as questions clearly mentioned it doesn't use CPU  
upvoted 1 times

  **Hi2ALL** 4 months, 3 weeks ago

Correction: doesn't "consume CPU"  
upvoted 1 times

  **kartikjena31** 4 months, 4 weeks ago

A is correct  
upvoted 1 times



  **nliaustemac** 5 months ago

A. is the correct answer:  
Create a Cloud Memorystore for Redis instance with 32-GB capacity.  
upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

A. Create a Cloud Memorystore for Redis instance with 32-GB capacity.

upvoted 1 times

  **adedj99** 5 months, 2 weeks ago

i am with A, Redish bit chipper and a fully manage in memory data store . Any how application doesnt require CPU , then it waste some money for cost and operation cost for non fully managed services

upvoted 1 times

  **JackGlemins** 5 months, 4 weeks ago

I think A is right. The key word is "for a latency-sensitive website."

[https://cloud.google.com/memorystore/docs/redis/redis-overview#what\\_its\\_good\\_for](https://cloud.google.com/memorystore/docs/redis/redis-overview#what_its_good_for)

Memorystore for Redis provides a fast, in-memory store for use cases that require fast, real-time processing of data. From simple caching use cases to real time analytics, Memorystore for Redis provides the performance you need.

Caching: Cache is an integral part of modern application architectures. Memorystore for Redis provides low latency access and high throughput for heavily accessed data, compared to accessing the data from a disk based backend store. Session management, frequently accessed queries, scripts, and pages are common examples of caching

upvoted 3 times

  **victory108** 6 months, 1 week ago

A - Create a Cloud Memorystore for Redis instance with 32-GB capacity.

upvoted 1 times

  **FenixRa73** 6 months, 2 weeks ago

how about the key phrase " need an additional 2 GB of memory for the rest of the processes"?

A - not suitable, Memorystore for Redis = without background processes.

B - suitable,

price:

per month Memorystore for Redis 32 Gb(basic): \$537.28 (standard):\$1,074.56

per month GCE 6vCPU 32 Ram: \$199.78

upvoted 8 times

## Question #63

## Topic 1

You are hosting an application on bare-metal servers in your own data center. The application needs access to Cloud Storage. However, security policies prevent the servers hosting the application from having public IP addresses or access to the internet. You want to follow Google-recommended practices to provide the application with access to Cloud Storage. What should you do?

- A. 1. Use nslookup to get the IP address for storage.googleapis.com. 2. Negotiate with the security team to be able to give a public IP address to the servers. 3. Only allow egress traffic from those servers to the IP addresses for storage.googleapis.com.
- B. 1. Using Cloud VPN, create a VPN tunnel to a Virtual Private Cloud (VPC) in Google Cloud. 2. In this VPC, create a Compute Engine instance and install the Squid proxy server on this instance. 3. Configure your servers to use that instance as a proxy to access Cloud Storage.
- C. 1. Use Migrate for Compute Engine (formerly known as Velostrata) to migrate those servers to Compute Engine. 2. Create an internal load balancer (ILB) that uses storage.googleapis.com as backend. 3. Configure your new instances to use this ILB as proxy.
- D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in Google Cloud. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.

**Correct Answer: C**

  **poogcp** Highly Voted 1 year, 1 month ago


D is the correct one as per Ref: <https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>  
upvoted 28 times

  **sarva** Highly Voted 1 year, 1 month ago

D VPN/interconnect without exposing to public  
upvoted 8 times

  **jitu12113** Most Recent 2 weeks, 4 days ago

Can someone pls explain why D?  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago



D is Correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

D is correct  
upvoted 1 times

  **[Removed]** 4 months ago


D is correct. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in GCP. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.  
upvoted 4 times

  **[Removed]** 4 months, 1 week ago

D is correct. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in GCP. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.  
upvoted 2 times

  **ACE\_ASPIRE** 4 months, 1 week ago

Option D  
upvoted 1 times

  **yuvi69** 4 months, 1 week ago

correct option is D  
upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

vote for D'  
upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

D is correct  
upvoted 1 times

  **nliaustemac** 5 months ago

D. is the correct answer:



1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in GCP. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago



D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in GCP. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.

upvoted 1 times

  **nitinz** 5 months, 2 weeks ago

D, anytime you see 199.36.153.4/30 thats is the right answer. thats the only way for Private into on-perm form cloud.

upvoted 3 times

  **DucSiu** 5 months, 4 weeks ago

D. 1. Using Cloud VPN or Interconnect, create a tunnel to a VPC in GCP. 2. Use Cloud Router to create a custom route advertisement for 199.36.153.4/30. Announce that network to your on-premises network through the VPN tunnel. 3. In your on-premises network, configure your DNS server to resolve \*.googleapis.com as a CNAME to restricted.googleapis.com.

upvoted 2 times

  **INASR** 6 months ago

it is D

upvoted 2 times

  **Bhagirathi** 7 months, 3 weeks ago

can you pls explain - why you choose D ?

Why not C ?

upvoted 1 times

  **magistrum** 7 months ago

read poogcp comment

upvoted 1 times



## Question #64

## Topic 1

You want to deploy an application on Cloud Run that processes messages from a Cloud Pub/Sub topic. You want to follow Google-recommended practices. What should you do?

- A. 1. Create a Cloud Function that uses a Cloud Pub/Sub trigger on that topic. 2. Call your application on Cloud Run from the Cloud Function for every message.
- B. 1. Grant the Pub/Sub Subscriber role to the service account used by Cloud Run. 2. Create a Cloud Pub/Sub subscription for that topic. 3. Make your application pull messages from that subscription.
- C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.
- D. 1. Deploy your application on Cloud Run on GKE with the connectivity set to Internal. 2. Create a Cloud Pub/Sub subscription for that topic. 3. In the same Google Kubernetes Engine cluster as your application, deploy a container that takes the messages and sends them to your application.

**Correct Answer: D**

  **Meix** Highly Voted 1 year, 1 month ago

C looks right for me as per <https://cloud.google.com/run/docs/tutorials/pubsub#integrating-pubsub>  
upvoted 21 times

  **ChrisBelt5** 4 days, 7 hours ago

great doc, its' C  
upvoted 1 times

  **Bhagirathi** Highly Voted 7 months, 1 week ago

why c ?  
explained>>  
You can use Pub/Sub to push messages to the endpoint of your Cloud Run service, where the messages are subsequently delivered to containers as HTTP requests. You cannot use Pub/Sub pull subscriptions because Cloud Run only allocates CPU during the processing of a request.  
upvoted 12 times

  **jcloud965** Most Recent 2 weeks, 1 day ago



C is correct  
Using Pub/Sub with Cloud Run tutorial : <https://cloud.google.com/run/docs/tutorials/pubsub>  
upvoted 2 times

  **sunilw** 4 weeks, 1 day ago

D is correct  
  
<https://cloud.google.com/vpc/docs/configure-private-google-access-hybrid>  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

B is correct  
upvoted 2 times

  **[Removed]** 4 months, 1 week ago

C is correct. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.  
upvoted 2 times

  **GCP\_Student1** 4 months, 3 weeks ago

C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.  
upvoted 1 times

  **cloud\_hunk** 5 months ago

it's C, you cannot use Pub/Sub pull subscriptions because Cloud Run only allocates CPU during the processing of a request.  
upvoted 1 times

  **JackGlemins** 5 months ago

C is right: <https://cloud.google.com/run/docs/tutorials/pubsub#integrating-pubsub>  
upvoted 2 times

  **nliaustemac** 5 months ago





C. is the correct answer:

Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint  
upvoted 2 times

  **victory108** 6 months, 1 week ago

Tricky one C is correct - 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.  
upvoted 2 times

  **DucSiu** 6 months, 1 week ago

I think its B  
upvoted 1 times

  **huyhoang8344** 7 months, 3 weeks ago

C is correct as explained here <https://cloud.google.com/run/docs/triggering/pubsub-push>  
upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

I think - it should be B.  
who can help explain - why you chose C ?  
upvoted 1 times

  **devscorpio2001** 8 months, 1 week ago

I think its B as it mentions "PULL"  
upvoted 2 times

  **Bhagirathi** 8 months, 1 week ago

can someone elaborate - why you think it is " C " ?  
upvoted 4 times

  **swatititame** 8 months, 2 weeks ago

• C. 1. Create a service account. 2. Give the Cloud Run Invoker role to that service account for your Cloud Run application. 3. Create a Cloud Pub/Sub subscription that uses that service account and uses your Cloud Run application as the push endpoint.  
upvoted 2 times

You need to deploy an application, which is packaged in a container image, in a new project. The application exposes an HTTP endpoint and receives very few requests per day. You want to minimize costs. What should you do?

- A. Deploy the container on Cloud Run.
- B. Deploy the container on Cloud Run on GKE.
- C. Deploy the container on App Engine Flexible.
- D. Deploy the container on GKE with cluster autoscaling and horizontal pod autoscaling enabled.

Correct Answer: B

- Gurnoor**

Highly Voted

1 year, 1 month ago

A should be cheapest as no infra needed.  
upvoted 40 times
- spudleymcdudley**

1 year ago

Listen to this guy. Google says "Cloud Run abstracts away all infrastructure management by automatically scaling up and down from zero almost instantaneously—depending on traffic. Cloud Run only charges you for the exact resources you use."  
upvoted 16 times
- ESP\_SAP**

Highly Voted

11 months, 2 weeks ago

Correct Answer should be A:  
  
Cloud Run takes any container images and pairs great with the container ecosystem: Cloud Build, Artifact Registry, Docker. ... No infrastructure to manage: once deployed, Cloud Run manages your services so you can sleep well. Fast autoscaling. Cloud Run automatically scales up or down from zero to N depending on traffic.  
  
<https://cloud.google.com/run>  
upvoted 15 times
- arsh1916**

Most Recent

2 months, 1 week ago

A is correct  
upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- Finger41**

2 months, 2 weeks ago

Answer is A. No need for infrastructure, and run when you need it.  
upvoted 1 times
- SaPa2021**

3 months ago

Answer B is correct since it is nothing but Cloud Run for Anthos - The Cloud Run for Anthos provides a flexible serverless development platform on Google Kubernetes Engine (GKE). Cloud Run for Anthos is powered by Knative, an open source project that supports serverless workloads on Kubernetes. Cloud Run is also available as a fully managed serverless platform, without Kubernetes. So this clearly means you cannot deploy container image directly on Cloud Run.  
upvoted 1 times
- FunkyTechnician**

3 months ago

<https://medium.com/google-cloud/deploying-containers-to-cloud-run-in-5mins-b03f1d8d4a64>  
  
"gcloud beta run deploy demo-app --image gcr.io/<PROJECT\_ID>/demo-image --region us-central1 --platform managed --allow-unauthenticated --quiet"  
upvoted 1 times
- Amu89**

2 months, 3 weeks ago

Deploy the container on Cloud Run is the correct one  
upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Deploy the container on Cloud Run.  
upvoted 2 times
- yuvi69**

4 months, 1 week ago

correct option is A  
upvoted 1 times

🗲️ 👤 **sumanshu** 4 months, 2 weeks ago

vote for 'A'

<https://www.youtube.com/watch?v=gx8VTa1c8DA>

upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

A. Deploy the container on Cloud Run.

upvoted 1 times

🗲️ 👤 **victory108** 4 months, 3 weeks ago

A - Deploy the container on Cloud Run.

upvoted 1 times

🗲️ 👤 **Hi2ALL** 4 months, 3 weeks ago

A is the probably correct answer since question asking minimize cost and low traffic per day so GKE is not cheap wont comes in this picture  
Cloud Run is better choice

A service exposes a unique endpoint and automatically scales the underlying infrastructure to handle incoming requests. Deployment platform and service name cannot be changed.

upvoted 1 times

🗲️ 👤 **cloud\_hunk** 5 months ago

It's A, GCP Cloud Run : Develop and deploy highly scalable containerized applications on a fully managed serverless platform.

upvoted 1 times

🗲️ 👤 **nliaustemac** 5 months ago

A. is the correct answer:

Deploy the container on Cloud Run.

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

A. Deploy the container on Cloud Run.

upvoted 2 times

🗲️ 👤 **nitinz** 5 months, 2 weeks ago

Correct answer is A. GKE is out of scope because of cost, left with App engine and Cloud Run. Cloud run you pay only for requests that comes in, App engine you pay for full month. Question is about minimum cost and Cloud Run is the only viable option in it.

upvoted 1 times

🗲️ 👤 **INASR** 6 months ago

A is correct

upvoted 2 times

## Question #66

## Topic 1

Your company has an existing GCP organization with hundreds of projects and a billing account. Your company recently acquired another company that also has hundreds of projects and its own billing account. You would like to consolidate all GCP costs of both GCP organizations onto a single invoice. You would like to consolidate all costs as of tomorrow. What should you do?

- A. Link the acquired company's projects to your company's billing account.
- B. Configure the acquired company's billing account and your company's billing account to export the billing data into the same BigQuery dataset.
- C. Migrate the acquired company's projects into your company's GCP organization. Link the migrated projects to your company's billing account.
- D. Create a new GCP organization and a new billing account. Migrate the acquired company's projects and your company's projects into the new GCP organization and link the projects to the new billing account.

**Correct Answer: D**

Reference:

<https://cloud.google.com/resource-manager/docs/migrating-projects-billing>


  **XRiddlerX** Highly Voted 11 months, 3 weeks ago

I could be missing something but where does it say in the question that the two orgs want to migrate projects? I believe the question and key points are "consolidate all GCP costs" and "consolidate all costs as of tomorrow". With that said, C and D would not be a 24 hour task and seems a bit cumbersome to perform for something simple as "creating a single invoice" AND that's a migration and not a consolidation of cost. With A, I can't find anywhere in GCP docs that this is a best practice, only a medium.com blog. IMHO, I won't go down this route because "Just because you can do something, doesn't mean you should." and I would consult GCP support for best practices on A before I do something like that.

That leaves B which is to export both detailed billing to BigQuery and create a invoice/report. This would be a temporary solution until you migrate Organizations. IMHO

I go with B.

upvoted 30 times

  **TAvenger** 5 months, 1 week ago

I am not sure that exporting some statistical data to BigQuery means anything for Google who creates the invoice. With "A" you are right, that is not the best practice, but the key word "for tomorrow" allows this custom approach. So the answer is "A"

upvoted 4 times

  **GunjGupta** Highly Voted 1 year, 1 month ago

To me, A looks correct. projects are linked to another organization as well in the acquired company so migrating would need google cloud support. we can not do ourselves. however, we can link other company projects to an existing billing account to generate total cost.



<https://medium.com/google-cloud/google-cloud-platform-cross-org-billing-41c5db8fefa6>

upvoted 20 times

  **spudleymcdudley** 1 year ago

Listen to this guy. It's 'A' as moving projects can take some time from Google. There's no need to create a new organisation and other options don't make any sense

upvoted 5 times

  **lxgywil** 2 months, 3 weeks ago

You're saying it as if "moving projects" was a viable option. What about B?

upvoted 1 times

  **PriyaGCP** Most Recent 1 month ago

Answer C

<https://medium.com/google-cloud/migrating-a-project-from-one-organization-to-another-gcp-4b37a86dd9e6>

upvoted 2 times

  **JenishD** 1 month ago

This is absolutely correct. I have performed this in the past & did the same procedure mentioned in the blog.

upvoted 1 times

  **Kaazim** 2 months ago

A Cloud Billing account:

Is a cloud-level resource managed in the Cloud Console.

Tracks all of the costs (charges and usage credits) incurred by your Google Cloud usage

A Cloud Billing account can be linked to one or more projects.

Project usage is charged to the linked Cloud Billing account.

Results in a single invoice per Cloud Billing account

Operates in a single currency  
 Defines who pays for a given set of resources  
 Is connected to a Google Payments Profile, which includes a payment instrument, defining how you pay for your charges  
 Has billing-specific roles and permissions to control accessing and modifying billing-related functions (established by IAM roles)

upvoted 1 times

🗲️ 👤 **ArizonaClassics** 2 months, 1 week ago

C- i will go with C  
 upvoted 2 times

🗲️ 👤 **learnazureportal** 2 months, 2 weeks ago

The correct answer is C. We need to apply migration.  
 upvoted 2 times

🗲️ 👤 **jcloud965** 2 weeks, 1 day ago

Migrating projects from one org to an other is time consuming.  
 better and faster to link projects to existing billing account  
 upvoted 1 times

🗲️ 👤 **Agraved** 2 months, 2 weeks ago

Answer has to be A as we need common invoice and not just determine total cost, so B would not suffice.  
 upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

A is correct....B will assist in analysis but not consolidating cost into 1 invoice  
 upvoted 1 times

🗲️ 👤 **SaPa2021** 3 months ago

Answer is A :- Change the billing account for a project  
 Cloud Billing accounts can be used across organizations. Moving a project from one organization to another won't impact billing, and charges will continue against the old billing account. However, organization moves often also include a requirement to move to a new billing account.

To change the billing account for an existing project, you must have the roles/owner role on the project, and the roles/billing.admin role on the destination billing account. To change the billing account:

Go to the Billing page in the Cloud Console.

Go to the Billing page

Click the name of the billing account you want to change.

Under Projects linked to this billing account, find the name of the Project to move and then click the menu button to the right.

Click Change billing, and then select the new billing account.

Click Set account.

Charges already incurred that have not yet been reported in the transaction history will be billed to the former billing account. This can include charges from up to two days prior to when the project was moved.

upvoted 2 times

🗲️ 👤 **[Removed]** 4 months ago

A is correct. Link the acquired company's projects to your company's billing account.  
 upvoted 3 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago

correct option is A  
 upvoted 1 times

🗲️ 👤 **ShakthiGCP** 4 months, 4 weeks ago

Ans: A . Spent more time in this entire discussion. Finally going with - <https://medium.com/google-cloud/google-cloud-platform-cross-org-billing-41c5db8fefa6> . This makes sense.

upvoted 2 times

🗲️ 👤 **nliaustemac** 5 months ago

A. is the correct answer:

Link the acquired company's projects to your company's billing account.

upvoted 2 times

🗲️ 👤 **rmout** 5 months, 1 week ago

A seems right.

For B,

"Timing is important. To have access to a more comprehensive set of Google Cloud billing data for your analysis needs, we recommend that you enable Cloud Billing data export to BigQuery at the same time that you create a Cloud Billing account. Be aware that your BigQuery dataset only reflects Google Cloud billing data incurred from the date you set up Cloud Billing export, and after. That is, Google Cloud billing data is not added retroactively, so you won't see Cloud Billing data from before you enable export."

upvoted 3 times

🗲️ 👤 **lxgywil** 2 months, 3 weeks ago

The best explanation on why B is not appropriate for this situation

upvoted 1 times

🗲️ 👤 **nitinz** 5 months, 1 week ago

This a fake question which makes no sense at all. Here is google documentation on Billing: -

If a project is associated with an organization, you can't change it back to No organization or migrate it to another organization without involving Google Cloud support

upvoted 2 times

  **GCP\_Student1** 5 months, 1 week ago

Answer is A, because: <https://cloud.google.com/billing/docs/concepts#relationships-between-resources> Note: Payment linkage of a project linked to a Cloud Billing account is not limited by organization ownership. It is possible for a Cloud Billing account to pay for projects that belong to an organization that is different than the organization that owns the Cloud Billing account.

upvoted 5 times

  **TAvenger** 5 months, 1 week ago

Guys, who chooses "B" - this is incorrect!! please read the article how you can get invoice

<https://cloud.google.com/billing/docs/how-to/get-invoice>

If you export your data into your local BigQuery dataset it does not change anything!

you only export logs! your old projects are still linked to the only billing account. And google will create invoice based on their information. Google will NOT rely on the data that is stored in your local BigQuery table. "B" is incorrect!!

Answer is "A"

upvoted 6 times

You built an application on Google Cloud that uses Cloud Spanner. Your support team needs to monitor the environment but should not have access to table data.

You need a streamlined solution to grant the correct permissions to your support team, and you want to follow Google-recommended practices. What should you do?

- A. Add the support team group to the roles/monitoring.viewer role
- B. Add the support team group to the roles/spanner.databaseUser role.
- C. Add the support team group to the roles/spanner.databaseReader role.
- D. Add the support team group to the roles/stackdriver.accounts.viewer role.

Correct Answer: B

- poogcp

Highly Voted

1 year, 1 month ago

its A, As you need to monitor only

upvoted 29 times
- WindDriver

2 weeks, 2 days ago

A, right, correct answer.

B and C are incorrect because allow to read data.

D also incorrect: Not for monitoring. roles/stackdriver.accounts.viewer Stackdriver Accounts Viewer: Read-only access to get and list information about Stackdriver account structure (resourcemanager.projects.get, resourcemanager.projects.list and stackdriver.projects.get)

upvoted 2 times
- WindDriver

2 weeks, 2 days ago

<https://cloud.google.com/iam/docs/understanding-roles>

upvoted 1 times
- Gurnoor

Highly Voted

1 year, 1 month ago

A is correct as user should not have any access to data, so B and C cant be used in this scenario.

upvoted 14 times
- sunilw

Most Recent

4 weeks ago

A is correct

[https://cloud.google.com/monitoring/access-control#mon\\_roles\\_desc](https://cloud.google.com/monitoring/access-control#mon_roles_desc)

Predefined roles

The following IAM roles are predefined by Cloud Monitoring. They grant permissions only for Monitoring.

Monitoring

The following roles grant general permissions for Monitoring:

Name	Title Description
roles/monitoring.viewer	Monitoring Viewer Gives you read-only access to Monitoring in the Google Cloud Console and API.
roles/monitoring.editor	Monitoring Editor Gives you read-write access to Monitoring in the Google Cloud Console and API, and lets you modify a metrics scope.
roles/monitoring.admin	Monitoring Admin Gives you full access to Monitoring in the Google Cloud Console and lets you modify a metrics scope.

The following role is used by service accounts for write-only access:

Name	Title Description
roles/monitoring.metricWriter	Monitoring Metric Writer Permits writing monitoring data to a metrics scope; doesn't permit access to Monitoring in the Google Cloud Console. For service accounts.

upvoted 3 times
- arsh1916

2 months, 1 week ago

A is correct

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct



upvoted 1 times

🗨️ 👤 **Finger41** 2 months, 2 weeks ago

A, as they need to monitor

upvoted 1 times

🗨️ 👤 **passnow** 4 months ago

roles/monitoring.viewer†

Monitoring Viewert Gives you read-only access to Monitoring in the Google Cloud Console and API. A is super leg

upvoted 2 times

🗨️ 👤 **[Removed]** 4 months, 1 week ago

A is correct. Add the support team group to the roles/monitoring.viewer role

upvoted 2 times

🗨️ 👤 **yuvi69** 4 months, 1 week ago

correct option is A , as we need to monitor environment like CPU usage etc.

upvoted 1 times

🗨️ 👤 **sumanshu** 4 months, 2 weeks ago

vote for 'A'

B & C - gives read access to table data - but question says - support team should NOT have access to table data

upvoted 1 times

🗨️ 👤 **Hi2ALL** 4 months, 3 weeks ago

C is the correct answer

upvoted 1 times

🗨️ 👤 **Finger41** 1 month, 4 weeks ago

Unfortunately, this role gives access to the data. You need the monitoring role to satisfy the requirements.

upvoted 1 times

🗨️ 👤 **cloud\_hunk** 5 months ago

A should be the answer. Here is a catch in Que: Your support team needs to monitor the environment (not the spanner service only )

upvoted 1 times

🗨️ 👤 **aji15** 5 months ago

A is correct. For C, roles/spanner.databaseReader- members with this role can read data from CloudSpanner, execute SQL queries and view database schema. So C is incorrect. Though according to Google's best practices, least privileges should be given, but A is the only correct one where the user won't be able to access the table which is the basic requirement in this case. So A is the correct one.

upvoted 1 times

🗨️ 👤 **nliaustemac** 5 months ago

A. is the correct answer:

Add the support team group to the roles/monitoring.viewer role

upvoted 1 times

🗨️ 👤 **lutoa** 5 months, 3 weeks ago

It's got to be A - specifically needs to allow support access to "environment" and NOT data - therefore monitoring the instance should be enough.

Stackdriver could potentially allow access to data which question specifically says not to allow.

upvoted 2 times

🗨️ 👤 **victory108** 6 months, 1 week ago

Tricky one A is correct - Add the support team group to the roles/monitoring.viewer role

upvoted 2 times

🗨️ 👤 **nherrerab** 6 months, 2 weeks ago

A is correct.

upvoted 3 times





## Question #68

## Topic 1

For analysis purposes, you need to send all the logs from all of your Compute Engine instances to a BigQuery dataset called platform-logs. You have already installed the Cloud Logging agent on all the instances. You want to minimize cost. What should you do?

- A. 1. Give the BigQuery Data Editor role on the platform-logs dataset to the service accounts used by your instances. 2. Update your instances' metadata to add the following value: logs-destination: bq://platform-logs.
- B. 1. In Cloud Logging, create a logs export with a Cloud Pub/Sub topic called logs as a sink. 2. Create a Cloud Function that is triggered by messages in the logs topic. 3. Configure that Cloud Function to drop logs that are not from Compute Engine and to insert Compute Engine logs in the platform-logs dataset.
- C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.
- D. 1. Create a Cloud Function that has the BigQuery User role on the platform-logs dataset. 2. Configure this Cloud Function to create a BigQuery Job that executes this query: `INSERT INTO dataset.platform-logs (timestamp, log) SELECT timestamp, log FROM compute.logs WHERE timestamp > DATE_SUB(CURRENT_DATE(), INTERVAL 1 DAY)` 3. Use Cloud Scheduler to trigger this Cloud Function once a day.

**Correct Answer: C**

  **vmart** 1 month, 2 weeks ago

I vote for C  
upvoted 3 times

  **Finger41** 1 month, 4 weeks ago

Its C.  
upvoted 3 times

  **arsh1916** 2 months, 1 week ago

C should be correct  
upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago



C is correct as it would restrict the volume of data to be scanned by bq, thus reducing cost  
upvoted 2 times

  **[Removed]** 4 months, 1 week ago

C is correct. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.  
upvoted 4 times

  **GCP\_Student1** 4 months, 1 week ago


C. 1. In Cloud Logging, create a filter to view only Compute Engine logs. 2. Click Create Export. 3. Choose BigQuery as Sink Service, and the platform-logs dataset as Sink Destination.  
upvoted 3 times

  **Hi2ALL** 4 months, 1 week ago

C is correct  
upvoted 1 times

  **drizzydroo** 4 months, 2 weeks ago

C right? No need to establish pub/sub topic?  
upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

vote for "C"

[https://cloud.google.com/logging/docs/export/configure\\_export\\_v2](https://cloud.google.com/logging/docs/export/configure_export_v2)  
upvoted 3 times



You are using Deployment Manager to create a Google Kubernetes Engine cluster. Using the same Deployment Manager deployment, you also want to create a DaemonSet in the kube-system namespace of the cluster. You want a solution that uses the fewest possible services. What should you do?


- A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.
- B. Use the Deployment Manager Runtime Configurator to create a new Config resource that contains the DaemonSet definition.
- C. With Deployment Manager, create a Compute Engine instance with a startup script that uses kubectl to create the DaemonSet.
- D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.

**Correct Answer:** C

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/cluster-access-for-kubectl>

-   **ESP\_SAP**

Highly Voted 



 11 months, 1 week ago



Correct Answer is (A):


Adding an API as a type provider  
This page describes how to add an API to Google Cloud Deployment Manager as a type provider. To learn more about types and type providers, read the Types overview documentation.

A type provider exposes all of the resources of a third-party API to Deployment Manager as base types that you can use in your configurations. These types must be directly served by a RESTful API that supports Create, Read, Update, and Delete (CRUD).



If you want to use an API that is not automatically provided by Google with Deployment Manager, you must add the API as a type provider.


<https://cloud.google.com/deployment-manager/docs/configuration/type-providers/creating-type-provider>  
upvoted 27 times
-   **magistrum** 7 months ago

very good find, sounds like you hit the nail in the head  
upvoted 4 times
-   **kishoredeena**



Highly Voted 



 1 year, 1 month ago

Option A is the right answer  
upvoted 9 times
-   **mcaromit**



Most Recent 



 2 months, 2 weeks ago



A is correct  
upvoted 1 times
-   **Linus11** 4 months ago



A is correct.  
upvoted 1 times
-   **pca2b** 4 months ago

A: in addition to ESP\_SAP's comment, the following link explains exactly how to...

<https://github.com/GoogleCloudPlatform/deploymentmanager-samples/blob/master/examples/v2/gke/daemonsets/daemonset.jinja>  
upvoted 2 times
-   **[Removed]** 4 months ago

A is correct. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.  
upvoted 2 times
-   **sumanshu** 4 months, 2 weeks ago

vote for 'D'  
upvoted 1 times
-   **GCP\_Student1** 4 months, 3 weeks ago

A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.  
upvoted 2 times
-   **JackGlemins** 5 months ago

<https://cloud.google.com/deployment-manager/docs/deployments/updating-deployments>  
upvoted 1 times

- 🗲️ 👤 **nliaustemac** 5 months ago  
A. is the correct answer:  
Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet  
upvoted 2 times
- 🗲️ 👤 **nitinz** 5 months, 1 week ago  
Answer is D  
add this to your yml file: -  
kind: DeamonSet  
and you are done!  
upvoted 2 times
- 🗲️ 👤 **JackGlemins** 5 months, 4 weeks ago  
Why not B?  
upvoted 1 times
- 🗲️ 👤 **victory108** 6 months, 1 week ago  
A - Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.  
upvoted 1 times
- 🗲️ 👤 **nherrerab** 6 months, 2 weeks ago  
D is correct.  
upvoted 1 times
- 🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
it is like no one sure of answer...all of you confusing.  
  
I can chose A.  
upvoted 2 times
- 🗲️ 👤 **Bhagirathi** 8 months, 1 week ago  
all of you hoose A, B C & D ...what is final take away?  
more than finding what is correct - why A/B/C/D you choose ?  
upvoted 2 times
- 🗲️ 👤 **swatititame** 8 months, 2 weeks ago  
• A. Add the cluster's API as a new Type Provider in Deployment Manager, and use the new type to create the DaemonSet.  
upvoted 1 times

## Question #70

## Topic 1

You are building an application that will run in your data center. The application will use Google Cloud Platform (GCP) services like AutoML. You created a service account that has appropriate access to AutoML. You need to enable authentication to the APIs from your on-premises environment. What should you do?

- A. Use service account credentials in your on-premises application.
- B. Use gcloud to create a key file for the service account that has appropriate permissions.
- C. Set up direct interconnect between your data center and Google Cloud Platform to enable authentication for your on-premises applications.
- D. Go to the IAM & admin console, grant a user account permissions similar to the service account permissions, and use this user account for authentication from your data center.

**Correct Answer: B**

Reference:

<https://cloud.google.com/vision/automl/docs/before-you-begin>

  **ESP\_SAP** Highly Voted 11 months, 2 weeks ago

Correct answer should be (B):

To use a service account outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal.

<https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 21 times

  **sunilw** Most Recent 4 weeks ago

B is correct.

Creating service account keys


To use a service account from outside of Google Cloud, such as on other platforms or on-premises, you must first establish the identity of the service account. Public/private key pairs provide a secure way of accomplishing this goal. When you create a service account key, the public portion is stored on Google Cloud, while the private portion is available only to you. For more information about public/private key pairs, see Service account keys.

upvoted 2 times

  **arsh1916** 2 months, 1 week ago

B is correct

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

B is correct as a 1st step...the key file is to be referenced in the env variable GOOGLE\_APPLICATION\_CREDENTIALS which would then provide access to on-prem application using ADC library

upvoted 1 times

  **[Removed]** 4 months, 1 week ago

B is correct. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 2 times

  **nliaustemac** 5 months ago

B. is the correct answer:

Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 1 times

  **GCP\_Student1** 5 months, 1 week ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.



upvoted 2 times

  **lutoa** 5 months, 3 weeks ago

I think it is B. I have seen other variations on this question where the scenario specifies that the on-premise has no internet, in which case interconnect is required.

This question does NOT specify, and since it is using API connection (so HTTP) it must have internet so the JSON key should be fine in answer B.

upvoted 1 times

  **DucSiu** 5 months, 4 weeks ago

B. Use gcloud to create a key file for the service account that has appropriate permissions.

upvoted 1 times

- 🗲️ 👤 **DucSiu** 5 months, 4 weeks ago  
D. In the cluster's definition in Deployment Manager, add a metadata that has kube-system as key and the DaemonSet manifest as value.  
upvoted 1 times
- 🗲️ 👤 **victory108** 6 months ago  
B - Use gcloud to create a key file for the service account that has appropriate permissions.  
upvoted 1 times
- 🗲️ 👤 **DucSiu** 6 months, 1 week ago  
B = 100%  
upvoted 1 times
- 🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
B = 100%  
upvoted 1 times
- 🗲️ 👤 **swatititame** 8 months, 2 weeks ago  
• B. Use gcloud to create a key file for the service account that has appropriate permissions.  
upvoted 1 times
- 🗲️ 👤 **nwk** 9 months ago  
B, should use key file  
upvoted 1 times
- 🗲️ 👤 **glam** 9 months, 3 weeks ago  
B. Use gcloud to create a key file for the service account that has appropriate permissions.  
upvoted 3 times
- 🗲️ 👤 **francisco\_guerra** 11 months, 2 weeks ago  
The answer is B:  
Accessing private data on behalf of a service account outside Google Cloud environments Service account key  
You need to create a service account, and download its private key as a JSON file. You need to pass the file to Google Cloud Client Libraries, so they can generate the service account credentials at runtime.  
upvoted 1 times

You are using Container Registry to centrally store your company's container images in a separate project. In another project, you want to create a Google Kubernetes Engine (GKE) cluster. You want to ensure that Kubernetes can download images from Container Registry. What should you do?

A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.

B. When you create the GKE cluster, choose the Allow full access to all Cloud APIs option under "Access scopes".

C. Create a service account, and give it access to Cloud Storage. Create a P12 key for this service account and use it as an imagePullSecrets in Kubernetes.

D. Configure the ACLs on each image in Cloud Storage to give read-only access to the default Compute Engine service account.

Correct Answer: C

- ESP\_SAP

Highly Voted

11 months, 2 weeks ago

Correct Answer (A):  
IAM permissions  
IAM permissions determine who can access resources. All users, service accounts, and other identities that interact with Container Registry must have the appropriate Cloud Storage permissions.  
  
By default, Google Cloud use default service accounts to interact with resources within the same project. For example, the Cloud Build service account can both push and pull images when Container Registry is in the same project.  
  
You must configure or modify permissions yourself if:  
  
You are using a service account in one project to access Container Registry in a different project  
You are using a default service account with read-only access to storage, but you want to both pull and push images  
You are using a custom service account to interact with Container Registry  
  
<https://cloud.google.com/container-registry/docs/access-control>  
upvoted 23 times
- [Removed]

10 months, 1 week ago

A is correct, practical implementation in video <https://www.youtube.com/watch?v=R16z7Sjrkxs>  
upvoted 3 times
- XRiddlerX

Highly Voted

11 months, 3 weeks ago

A is correct...  
Container Registry uses Cloud Storage buckets as the underlying storage for container images. You control access to your images by granting appropriate Cloud Storage permissions to a user, group, service account, or other identity.  
  
If the service account needs to access Container Registry in another project, you must grant the required permissions in the project with Container Registry.  
  
Reference:  
<https://cloud.google.com/container-registry/docs/access-control#permissions>  
upvoted 11 times
- viswanand

Most Recent

2 months, 1 week ago

A looks more sensible option  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- ssgcp

2 months, 4 weeks ago

C should be the answer. It is the correct procedure to pull images from gcr in another project.  
upvoted 1 times
- [Removed]

4 months, 1 week ago

A is correct. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.  
upvoted 2 times
- sumanshu

4 months, 2 weeks ago

vote for 'A'  
  
C - does not mention which access provided?

upvoted 1 times

🗨️ 👤 **jackycc** 4 months, 2 weeks ago  
<https://kubernetes.io/docs/tasks/configure-pod-container/pull-image-private-registry/>  
I guess its C  
upvoted 1 times

🗨️ 👤 **GCP\_Student1** 4 months, 3 weeks ago  
A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.  
upvoted 2 times

🗨️ 👤 **AGHPE** 5 months ago  
Storage Object Viewer to pull images only  
Storage Object Admin to push and pull images  
  
A is correct!  
upvoted 1 times

🗨️ 👤 **nliaustemac** 5 months ago  
A. is the correct answer:  
In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.  
upvoted 1 times

🗨️ 👤 **DucSiu** 5 months, 4 weeks ago  
A. In the project where the images are stored, grant the Storage Object Viewer IAM role to the service account used by the Kubernetes nodes.  
upvoted 1 times

🗨️ 👤 **guid1984** 6 months, 1 week ago  
Answer is A, please read it here <https://cloud.google.com/container-registry/docs/using-with-google-cloud-platform>  
upvoted 1 times

🗨️ 👤 **nherrerab** 6 months, 2 weeks ago  
A is correct.  
upvoted 2 times

🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
A to choose<  
upvoted 2 times

🗨️ 👤 **Bhagirathi** 8 months, 1 week ago  
why this page carries some answer and you guys find a new one?  
  
what shall we understand finally ? for me it only confusions  
upvoted 3 times

🗨️ 👤 **ayj** 8 months, 1 week ago  
A  
If the VM and Container Registry are in separate projects  
You must grant the service account with IAM permissions to access the storage bucket used by Container Registry.  
  
<https://cloud.google.com/container-registry/docs/access-control#integration>  
upvoted 2 times



You deployed a new application inside your Google Kubernetes Engine cluster using the YAML file specified below.

```
apiVersion: apps/v1      apiVersion: v1
kind: Deployment         kind: Service
metadata:               metadata:
  name: myapp-deployment  name: myapp-service
spec:                  spec:
  selector:            ports:
    matchLabels:        - port: 8000
      app: myapp         targetPort: 80
  replicas: 2           protocol: TCP
  template:             selector:
    metadata:            app: myapp
    labels:
      app: myapp
    spec:
      containers:
        - name: myapp
          image: myapp:1.1
          ports:
            - containerPort: 80
```

You check the status of the deployed pods and notice that one of them is still in PENDING status:

```
kubectl get pods -l app=myapp
NAME                                READY   STATUS    RESTART  AGE
myapp-deployment-58ddbbb995-lp86m  0/1     Pending   0        9m
myapp-deployment-58ddbbb995-qjpkg  1/1     Running   0        9m
```

You want to find out why the pod is stuck in pending status. What should you do?

- A. Review details of the myapp-service Service object and check for error messages.
- B. Review details of the myapp-deployment Deployment object and check for error messages.
- C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.
- D. View logs of the container in myapp-deployment-58ddbbb995-lp86m pod and check for warning messages.

**Correct Answer: C**

Reference:  
<https://cloud.google.com/run/docs/gke/troubleshooting>

- spudleymcdudley** Highly Voted 1 year ago

It's C - <https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#debugging-pods>  
upvoted 20 times
- someoneinthecloud** Highly Voted 1 year ago

Answer is C - You can't view logs of a pod that isn't deployed, so D is incorrect.  
C allows you to check the pod deployment messages and look for errors  
upvoted 13 times
- sidharthwader** 3 weeks, 4 days ago

What u said is incorrect you can view pod's log even in pending state.  
kubectl logs <pon-name> -n <namespace>  
upvoted 1 times
- latislal** Most Recent 3 weeks ago

Why not D.  
upvoted 1 times
- associatecloudexamuser** 2 weeks, 4 days ago

You can't view logs of a pod that isn't deployed  
upvoted 1 times
- mcaromit** 2 months, 2 weeks ago

C is correct



upvoted 1 times

🗳️ 👤 **[Removed]** 4 months, 1 week ago

C is correct. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 2 times

🗳️ 👤 **sumanshu** 4 months, 2 weeks ago

vote for 'C'

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/>

upvoted 1 times

🗳️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 2 times

🗳️ 👤 **nliaustemac** 5 months ago

C. is the correct answer:

Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 1 times

🗳️ 👤 **DucSiu** 5 months, 4 weeks ago

D. View logs of the container in myapp-deployment-58ddbbb995-lp86m pod and check for warning messages.

upvoted 1 times

🗳️ 👤 **Bhagirathi** 8 months, 1 week ago

who is going with D ? and why ?

upvoted 1 times

🗳️ 👤 **swatititame** 8 months, 2 weeks ago

• C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 1 times

🗳️ 👤 **glam** 9 months, 3 weeks ago

C. Review details of myapp-deployment-58ddbbb995-lp86m Pod and check for warning messages.

upvoted 5 times

🗳️ 👤 **[Removed]** 10 months ago

C is the answer.

When a pod goes in pending, waiting you need to describe the pod and view the information

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#my-pod-stays-pending>

If the pod crashes or is running but not doing anything or what it is supposed to do then the way to proceed is view the container logs.

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#debugging-pods>

upvoted 4 times

🗳️ 👤 **raj231** 10 months, 1 week ago

Option C is correct as option D just says View.. but Option C has Review the logs...

upvoted 1 times

🗳️ 👤 **mwwoodm** 10 months, 3 weeks ago

The reference given in the answer (<https://cloud.google.com/run/docs/gke/troubleshooting>) does indicate that C is the correct troubleshooting action: "Choose one and use the following command to see detailed information for its status. Some useful fields are conditions and containerStatuses: kubectl get pod POD-NAME -o yaml"

upvoted 1 times

🗳️ 👤 **Eshkrkrkr** 8 months, 3 weeks ago

Are you even reading the links you provide? Its about Anthos on Cloud Run, it doesn't metter its back-ended with Kubernetes.

upvoted 4 times

🗳️ 👤 **SSPC** 11 months, 1 week ago

I would go with the answer "C" If a Pod is stuck in Pending it means that it can not be scheduled onto a node. Generally this is because there are insufficient resources of one type or another that prevent scheduling. Look at the output of the kubectl describe ... command above.

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#my-pod-stays-pending>

upvoted 4 times

🗳️ 👤 **XRiddlerX** 11 months, 3 weeks ago

The answer is B

A (incorrect) - The service object has nothing to do with the pod in a pending status



B (correct) - If a Pod is stuck in Pending it means that it can not be scheduled onto a node. Generally this is because there are insufficient resources of one type or another that prevent scheduling. Look at the output of the kubectl describe deployment. There should be messages from the scheduler about why it can not schedule your pod.

C (incorrect) - Since the pod has not been scheduled there will be nothing in the logs for the pod/container.

D (incorrect) - Since the pods has not been schedule there will be nothing in the container log.

<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#my-pod-stays-pending>

upvoted 6 times

  **Eshkrkrkr** 8 months, 3 weeks ago

The qustion don't ask us to check the logs! You want to find out why the pod is stuck, that's all. Following the link you've gave: "kubectl describe ... command. There should be messages from the scheduler about why it can not schedule your pod." - Answer is C.  
upvoted 3 times

  **glam** 9 months, 3 weeks ago

You have described 'C' as 'B'.  
upvoted 3 times

You are setting up a Windows VM on Compute Engine and want to make sure you can log in to the VM via RDP. What should you do?

- A. After the VM has been created, use your Google Account credentials to log in into the VM.
- B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.
- C. When creating the VM, add metadata to the instance using `~windows-password` as the key and a password as the value.
- D. After the VM has been created, download the JSON private key for the default Compute Engine service account. Use the credentials in the JSON file to log in to the VM.

Correct Answer: D

- John\_Iam

Highly Voted

1 year, 1 month ago

Correct Answer is B.

B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.

<https://cloud.google.com/sdk/gcloud/reference/beta/compute/reset-windows-password>

upvoted 38 times
- ragu123

11 months, 3 weeks ago

Correct answer is B.

gcloud beta compute reset-windows-password allows a user to reset and retrieve a password for a Windows virtual machine instance. If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned.

upvoted 3 times
- voler

1 year, 1 month ago

Yes! "If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

upvoted 2 times
- dan80

1 year, 1 month ago

did you even look at the link you provide ? it clearly say gcloud beta compute reset-windows-password my-instance and not gcloud compute reset-windows-password. D is correct - <https://cloud.google.com/iam/docs/creating-managing-service-account-keys>

upvoted 2 times
- lxgywil

2 months, 3 weeks ago

Oh yes? Then what about this link (for non-beta command)?  
<https://cloud.google.com/sdk/gcloud/reference/compute/reset-windows-password>

"If the Windows account does not exist, this command will cause the account to be created and the password for that new account will be returned."

The answer is obviously B. Just test it and it'll become very clear

upvoted 2 times
- dan80

1 year, 1 month ago

nobody talk on reset the password but how to access the Windows - best way - Service Account

upvoted 1 times
- PriyaGCP

Most Recent

1 month ago

Answer B

<https://cloud.google.com/compute/docs/quickstart-windows>

upvoted 2 times
- alperenhasan

2 months, 1 week ago

folks B is correct. you can try on console google cloud.

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct

upvoted 1 times
- jahnu

3 months, 1 week ago

B is correct in GCP we use only login username and password

upvoted 1 times
- Linus11

4 months ago

B only.

upvoted 1 times

- 🗨️ 👤 **j1c4b** 4 months ago  
Correct Answer is B.  
Windows Server instances use password authentication instead of SSH authentication. To prevent unauthorized access to new Windows instances, Compute Engine requires that you generate a new Windows password for that instance before you connect to it.  
upvoted 1 times
- 🗨️ 👤 **[Removed]** 4 months, 1 week ago  
B is correct. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.  
upvoted 2 times
- 🗨️ 👤 **Rabbitfur** 4 months, 1 week ago  
B is correct answer <https://cloud.google.com/compute/docs/instances/windows/creating-passwords-for-windows-instances>  
upvoted 1 times
- 🗨️ 👤 **X100** 4 months, 2 weeks ago  
B as windows password needs to be set even if done from console  
<https://cloud.google.com/compute/docs/quickstart-windows>  
upvoted 1 times
- 🗨️ 👤 **pentium2000** 4 months, 2 weeks ago  
The answer is B. A, C and D doesn't make any sense.  
upvoted 1 times
- 🗨️ 👤 **GCP\_Student1** 4 months, 3 weeks ago  
B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.  
upvoted 3 times
- 🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago  
B is correct.  
upvoted 1 times
- 🗨️ 👤 **victory108** 4 months, 3 weeks ago  
B - After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.  
upvoted 1 times
- 🗨️ 👤 **Rimjith** 4 months, 4 weeks ago  
For me the correct answer is D  
<https://cloud.google.com/compute/docs/instances/windows/automate-pw-generation#python>  
  
Option B says to use "reset-windows-password" to retrieve the login credentials for the VM.  
We wont be able to "retrieve" credentials we will be able to reset and create new credentials.  
But from the JASON private key we can decrypt the correct password  
upvoted 1 times
- 🗨️ 👤 **badetoye** 5 months ago  
@John\_lam is absolutely correct. gcloud compute reset-windows-password "INSTANCE\_NAME". this will generate a new credential details. you can as well change the zone of the VM.  
upvoted 1 times
- 🗨️ 👤 **DucSiu** 5 months, 4 weeks ago  
B. After the VM has been created, use gcloud compute reset-windows-password to retrieve the login credentials for the VM.  
upvoted 1 times

## Question #74

## Topic 1

You want to configure an SSH connection to a single Compute Engine instance for users in the dev1 group. This instance is the only resource in this particular

Google Cloud Platform project that the dev1 users should be able to connect to. What should you do?

- A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.
- B. Set metadata to enable-oslogin=true for the instance. Set the service account to no service account for that instance. Direct them to use the Cloud Shell to ssh to that instance.
- C. Enable block project wide keys for the instance. Generate an SSH key for each user in the dev1 group. Distribute the keys to dev1 users and direct them to use their third-party tools to connect.
- D. Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

**Correct Answer: D**

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

  **poogcp** Highly Voted 1 year, 1 month ago

A correct one  
upvoted 28 times

  **spudleymcdudley** 1 year ago

For further evidence... <https://cloud.google.com/compute/docs/instances/managing-instance-access>  
upvoted 3 times

  **nithinpb180** 1 year, 1 month ago

Agree with that  
upvoted 2 times

  **student002** Highly Voted 9 months, 1 week ago

Pure from logic thinking: A can't be right. If the group get access to that instance with enable-oslogin=true, then they could have access to every instance that has enable-oslogin=true. Or do I miss something?  
upvoted 6 times

  **magistrum** 7 months ago


I'm convinced with this logic  
upvoted 1 times

  **squishy\_fishy** Most Recent 1 month, 4 weeks ago



A is the correct answer. It is recommended by Google at here <https://cloud.google.com/compute/docs/instances/connecting-advanced#linux-and-macos>  
upvoted 3 times

  **mcaromit** 2 months, 2 weeks ago

A is correct  
upvoted 1 times

  **Linus11** 4 months ago

A is the most apt one.  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago

A is correct. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.  
upvoted 2 times

  **drizzydroo** 4 months, 2 weeks ago

Has anyone actually tested A? I gave one of my test accounts the same role as in the answer and enabled oslogin. Yet when i try to access through cloud shell it doesnt let me ssh in. Error says I am missing 'compute.projects.get' even when i run the correct command with proejct id and compute zone  
upvoted 1 times

  **X100** 4 months, 2 weeks ago

A because oslogin is the preferred method for access because it allows IAM policies to be used.

upvoted 1 times

🗨️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.

upvoted 2 times

🗨️ 👤 **lutoa** 5 months, 3 weeks ago

A is correct.

NEVER EVER DISTRIBUTE private keys or generate SSH KEYS for others.

This automatically excludes two answers (C + D), it has to be A because B talks about service accounts which has nothing to do with dev group needing to SSH to the instance.

upvoted 5 times

🗨️ 👤 **hicham** 5 months ago

totally agree with you and your logical reasoning

upvoted 1 times

🗨️ 👤 **DucSiu** 5 months, 4 weeks ago

A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.

upvoted 2 times

🗨️ 👤 **victory108** 6 months, 1 week ago

D - Enable block project wide keys for the instance. Generate an SSH key and associate the key with that instance. Distribute the key to dev1 users and direct them to use their third-party tools to connect.

upvoted 1 times

🗨️ 👤 **nitinz** 5 months, 1 week ago

os.login is preferred method and recommended by Google. Hence A.

upvoted 1 times

🗨️ 👤 **swayam** 7 months, 1 week ago

it should be "D" - Allow or block project-wide public SSH keys: Use this metadata value to block or allow project-wide public SSH keys on a specific instance. By default, an instance accepts all project-wide public SSH keys. Block project-wide public SSH keys if you want to restrict the instance to only users with instance-level public SSH keys. Refer - <https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

upvoted 3 times

🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago

A should be best one ??

upvoted 3 times

🗨️ 👤 **Bhagirathi** 8 months, 1 week ago

what shall I choose ?

upvoted 1 times

🗨️ 👤 **swatititame** 8 months, 2 weeks ago

• A. Set metadata to enable-oslogin=true for the instance. Grant the dev1 group the compute.osLogin role. Direct them to use the Cloud Shell to ssh to that instance.

upvoted 2 times

🗨️ 👤 **gh999l** 8 months, 3 weeks ago

third party tool options are usually incorrect

upvoted 1 times

You need to produce a list of the enabled Google Cloud Platform APIs for a GCP project using the gcloud command line in the Cloud Shell. The project name is my-project. What should you do?

- A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.
- B. Run gcloud init to set the current project to my-project, and then run gcloud services list --available.
- C. Run gcloud info to view the account value, and then run gcloud services list --account <Account>.
- D. Run gcloud projects describe <project ID> to verify the project value, and then run gcloud services list --available.

Correct Answer: A

- dan80** Highly Voted

1 year, 1 month ago

A is the correct answer, log to gcloud and run the commands, doesnt make sense to run cloud init and gcloud services list --available gives you the full services that are available.

upvoted 32 times
- raffiq**

1 year, 1 month ago

Yes, Answer A correct. it shows only enabled services of API

upvoted 3 times
- mcaromit** Most Recent

2 months, 2 weeks ago

A is correct

upvoted 2 times
- lxgywil**

2 months, 3 weeks ago

"A" is correct.

For those, who have doubts:

`gcloud services list --available` returns not only the enabled services in the project but also services that CAN be enabled. Therefore, option B is incorrect.

<https://cloud.google.com/sdk/gcloud/reference/services/list#--available>

upvoted 3 times
- squishy\_fishy**

1 month, 3 weeks ago

Best answer!

upvoted 1 times
- Rabbitfur**

3 months, 1 week ago

A is correct answer --project option available in gcloud command

upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.

upvoted 2 times
- Vic1043**

4 months, 3 weeks ago

B is the right answer

Run the following command to list the APIs and services available to you in your current project:

gcloud services list --available

<https://cloud.google.com/service-usage/docs/list-services#gcloud>

upvoted 1 times
- Abalagu**

1 month, 2 weeks ago

B is wrong here. The question asks about enabled APIs, not available APIs. The `gcloud list` command returns the enabled APIs by default.

upvoted 1 times
- sumanshu**

4 months, 2 weeks ago

B wrong as there is no need to run init command

upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

A is correct

upvoted 1 times



- 🗨️ 👤 **GCP\_Student1** 4 months, 3 weeks ago  
A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.  
upvoted 2 times
- 🗨️ 👤 **ShakthiGCP** 4 months, 3 weeks ago  
I would say both A and C will work. Ans : A and C . Any denial that C wont work ?  
upvoted 1 times
- 🗨️ 👤 **ShakthiGCP** 4 months, 3 weeks ago  
Sorry. Ans : A only. C will give the account details of the current configuration only. but if you want to check the services list of a different project..it wont work unless you modify he configuration. So will go with Ans: A alone.  
upvoted 1 times
- 🗨️ 👤 **GCP\_Student1** 5 months, 1 week ago  
A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.  
upvoted 1 times
- 🗨️ 👤 **DucSiu** 5 months, 4 weeks ago  
A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.  
upvoted 1 times
- 🗨️ 👤 **darren123** 7 months ago  
A and B both works you can set the default project and also get the activated services. You can try and see.  
upvoted 1 times
- 🗨️ 👤 **adedj99** 5 months, 1 week ago  
B for available service not enable service yet  
upvoted 1 times
- 🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
A to select?  
upvoted 1 times
- 🗨️ 👤 **swatititame** 8 months, 2 weeks ago  
• A. Run gcloud projects list to get the project ID, and then run gcloud services list --project <project ID>.  
upvoted 1 times
- 🗨️ 👤 **gcpace** 8 months, 2 weeks ago  
A is the correct answer...executed the given command and below is the output :  
@cloudshell:~ (test-practice-project)\$ gcloud services list --project test-practice-project  
NAME TITLE  
compute.googleapis.com Compute Engine API  
oslogin.googleapis.com Cloud OS Login API  
upvoted 4 times
- 🗨️ 👤 **Anand2608** 8 months, 3 weeks ago  
The answer is B as the rest of the options do not fit the requirements of the question.  
upvoted 1 times
- 🗨️ 👤 **nwk** 9 months ago  
Vote A  
<https://cloud.google.com/sdk/gcloud/reference/--project>  
upvoted 3 times
- 🗨️ 👤 **Khoka** 8 months, 2 weeks ago  
A for sure: <https://cloud.google.com/sdk/gcloud/reference/services/list>  
It's a clear GCP documentation. Available will list enabled + available. end enabled is the default, so doesn't need for explicit enabled flag for option A  
upvoted 1 times



## Question #76

## Topic 1

You are building a new version of an application hosted in an App Engine environment. You want to test the new version with 1% of users before you completely switch your application over to the new version. What should you do?

- A. Deploy a new version of your application in Google Kubernetes Engine instead of App Engine and then use GCP Console to split traffic.
- B. Deploy a new version of your application in a Compute Engine instance instead of App Engine and then use GCP Console to split traffic.
- C. Deploy a new version as a separate app in App Engine. Then configure App Engine using GCP Console to split traffic between the two apps.
- D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

**Correct Answer: D**

🗲️ 👤 **SIX** Highly Voted 👍 1 year, 1 month ago

Correct answer is D  
upvoted 39 times

🗲️ 👤 **mohdafiuddin** Highly Voted 👍 7 months ago

Splitting the question to the key requirements

1. new version of an application hosted in an App Engine environment.
2. test the new version with 1% of users

App engine supports versioning and traffic splitting so no need to involve anything else  
(source - <https://cloud.google.com/appengine#all-features>)

- A. ....'Google Kubernetes Engine'.... - No need to involve GKE. Not the right option
- B. ....'Compute Engine instance'.... - No need to involve Compute Engine.
- C. ....'Separate app in App Engine'....- No need to deploy as a separate app. versioning is supported already. Not the right option.
- D. This is the right answer.

upvoted 12 times

🗲️ 👤 **arsh1916** Most Recent ⌚ 2 months, 1 week ago

Ans: D  
upvoted 2 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

D is the only correct option  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months, 1 week ago

D is correct. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 2 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 2 times

🗲️ 👤 **lutoa** 5 months, 3 weeks ago

D. Remember, you can only have a single App Engine app per project.  
And the question has nothing to do with GKE. Traffic splitting over versions is the obvious answer (D)

upvoted 2 times

🗲️ 👤 **DucSiu** 5 months, 4 weeks ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 2 times

🗲️ 👤 **victory108** 6 months, 1 week ago

D - Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 1 times

🗲️ 👤 **nightflyer** 7 months, 2 weeks ago

Who posts these answers? this is very misleading, they have the guts to choose A. The answer is 100% D

upvoted 3 times

🗲️ 👤 **alpharomeo9** 7 months, 3 weeks ago

Correct Answer is D. The question is about App Engine. In the App Engine splitting the traffic will do the job. Why would someone go to Kubernetes in this scenario?

upvoted 2 times

🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago

not sure - how it can be A ?

Is it not D ?

upvoted 1 times

🗲️ 👤 **luerangel** 7 months, 4 weeks ago

D because you can use split by cookie to do this. With cookie you have more control about user traffic.

upvoted 1 times

🗲️ 👤 **swatititame** 8 months, 2 weeks ago

- D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 3 times

🗲️ 👤 **gcpace** 8 months, 2 weeks ago

I think option D is correct. Split traffic is common between option A & D, but what is the need to deploy the new version in Kubernetes engine, when we have version management in APP Engine?

upvoted 1 times

🗲️ 👤 **Anand2608** 8 months, 3 weeks ago

D is the correct answer as that is the easiest way to accomplish this task.

upvoted 2 times

🗲️ 👤 **gh999l** 8 months, 3 weeks ago

No doubt D is the answer

upvoted 2 times

You need to provide a cost estimate for a Kubernetes cluster using the GCP pricing calculator for Kubernetes. Your workload requires high IOPs, and you will also be using disk snapshots. You start by entering the number of nodes, average hours, and average days. What should you do next?

- A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.
- B. Fill in local SSD. Add estimated cost for cluster management.
- C. Select Add GPUs. Fill in persistent disk storage and snapshot storage.
- D. Select Add GPUs. Add estimated cost for cluster management.

**Correct Answer: C**

Reference:

<https://cloud.google.com/products/calculator#tab=container>

- dan80**

Highly Voted

1 year, 1 month ago

This one is Tricky, local SSD is require for High IOPS - <https://cloud.google.com/compute/docs/disks/local-ssd> , but it say using disk snapshots. A is correct.

upvoted 35 times
- poogcp**

Highly Voted

1 year, 1 month ago

A is correct .

upvoted 10 times
- PB101**

Most Recent

5 days, 21 hours ago

Option A is the correct ans. 100% sure.

Try below steps to understand why-

1) Click - <https://cloud.google.com/products/calculator/>

2) Select GKE Standard

3) Check the help '?' icon for "Local SSD" field- "Local solid state disks, providing very high IOPS and very low latency block storage."

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

only correct option is A

upvoted 1 times
- FunkyTechnician**

2 months, 3 weeks ago

Reference:

<https://cloud.google.com/products/calculator/>

High IOP's = needing local SSD.

"you will also be using disk snapshots" = PDS + Snapshot Storage

upvoted 2 times
- [Removed]**

4 months, 1 week ago

A is correct. Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 2 times
- yuvi69**

4 months, 1 week ago

correct option is C . because local SSD option isn't available in kubernetes pricing calculator section.

<https://cloud.google.com/products/calculator/>

upvoted 1 times
- pondai**

4 months, 1 week ago

Wrong check again. GKE can use Local SSD

upvoted 1 times
- CloudMom**

4 months, 2 weeks ago

C is correct. SSD for snapshots

upvoted 1 times
- M\_DD8**

5 months ago

Only C is valid as snapshots can only be created from persistent disks. Local SSD is not persistent - ruling out A and B.

"Create snapshots to periodically back up data from your zonal persistent disks or regional persistent disks." -->

<https://cloud.google.com/compute/docs/disks/create-snapshots>

upvoted 1 times
- pondai**

4 months, 1 week ago

Local SSD is a boot disk. So it can help to high IOPs job .And A also choose persistent disk "snapshot storage" for snapshot .So A is correct.

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 3 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

D. Deploy a new version of your application in App Engine. Then go to App Engine settings in GCP Console and split traffic between the current version and newly deployed versions accordingly.

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

This answer belongs to the previous question.

upvoted 1 times

🗲️ 👤 **DucSiu** 5 months, 4 weeks ago

A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 2 times

🗲️ 👤 **ehizren** 6 months ago

C. Is the answer.

Local SSD ensure high IOPS but do not allow for "Snapshots".

"GPU allow for very high IOPS, then Persistent disks allow for Snapshots"

Hence, answer is C

upvoted 1 times

🗲️ 👤 **victory108** 6 months, 1 week ago

A - Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 2 times

🗲️ 👤 **alpharomeo9** 7 months, 3 weeks ago

Option A is correct.

Local solid state disks, providing very high IOPS and very low latency block storage. As said by dan80 and krakozia bl option A is right. You can check by going to <https://cloud.google.com/products/calculator> and click on question mark next to local ssd option. And you can learn more by going to <https://cloud.google.com/compute/docs/disks/local-ssd>

Since the question states workload requires high IOPs and disk snapshots.

upvoted 4 times

🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago

My take away is " confusion ". every one of you differ with others.

upvoted 3 times

🗲️ 👤 **swatititame** 8 months, 2 weeks ago

• A. Fill in local SSD. Fill in persistent disk storage and snapshot storage.

upvoted 2 times

You are using Google Kubernetes Engine with autoscaling enabled to host a new application. You want to expose this new application to the public, using HTTPS on a public IP address. What should you do?

- A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.
- B. Create a Kubernetes Service of type ClusterIP for your application. Configure the public DNS name of your application using the IP of this Service.
- C. Create a Kubernetes Service of type NodePort to expose the application on port 443 of each node of the Kubernetes cluster. Configure the public DNS name of your application with the IP of every node of the cluster to achieve load-balancing.
- D. Create a HAProxy pod in the cluster to load-balance the traffic to all the pods of the application. Forward the public traffic to HAProxy with an iptable rule. Configure the DNS name of your application using the public IP of the node HAProxy is running on.

**Correct Answer: A**  
Reference:  
<https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>

- dan80** Highly Voted

1 year, 1 month ago

A is correct.  
upvoted 21 times
- magistrum**

7 months ago

Saw this which provides good context <https://medium.com/google-cloud/kubernetes-nodeport-vs-loadbalancer-vs-ingress-when-should-i-use-what-922f010849e0>  
upvoted 7 times
- nitinz**

5 months, 1 week ago

you nailed it.  
upvoted 1 times
- arsav** Highly Voted

5 months, 1 week ago

HAProxy is HTTP only, doesnt support HTTPS, so you can reject option D  
<https://www.haproxy.org/#desc>

Cluster IP - is an internal IP, you cannot expose public externally. reject option B

out of option A and C

C, port 443 is https but public DNS is not going to give you a load balancing  
A is the right choice,  
kubernets ingress exposes HTTPS  
<https://kubernetes.io/docs/concepts/services-networking/ingress/>

and cloud load balancer is the right choice which will help to expose the app to public  
upvoted 10 times
- mcaromit** Most Recent

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- [Removed]**

4 months, 1 week ago

A is correct. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.  
upvoted 4 times
- yuvi69**

4 months, 1 week ago

correct option is A.  
upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago

A is correct  
upvoted 1 times
- JackGlemins**

5 months ago

A is right: <https://kubernetes.io/docs/concepts/services-networking/service/#publishing-services-service-types>

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months, 1 week ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 2 times

🗲️ 👤 **DucSiu** 5 months, 4 weeks ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 2 times

🗲️ 👤 **victory108** 6 months ago

A - Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 2 times

🗲️ 👤 **guid1984** 6 months, 1 week ago

Both A and D can be solutions but A is preferred as setting up an ingress object which will work with single Cloud Load Balancer will be much simpler. CORRECT CHOICE IS A

upvoted 1 times

🗲️ 👤 **T\_K** 6 months, 3 weeks ago

A is correct

upvoted 1 times

🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago

kindly re-check and confirm the right one, else it adds confusions for everyone else.

upvoted 1 times

🗲️ 👤 **ladh** 8 months ago

why not D?

upvoted 1 times

🗲️ 👤 **swatititame** 8 months, 2 weeks ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 1 times

🗲️ 👤 **gh999I** 8 months, 3 weeks ago

expose option works with Cloud LB so A is correct

upvoted 1 times

🗲️ 👤 **glam** 9 months, 3 weeks ago

A. Create a Kubernetes Service of type NodePort for your application, and a Kubernetes Ingress to expose this Service via a Cloud Load Balancer.

upvoted 3 times

## Question #79

## Topic 1

You need to enable traffic between multiple groups of Compute Engine instances that are currently running two different GCP projects. Each group of Compute Engine instances is running in its own VPC. What should you do?

- A. Verify that both projects are in a GCP Organization. Create a new VPC and add all instances.
- B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.
- C. Verify that you are the Project Administrator of both projects. Create two new VPCs and add all instances.
- D. Verify that you are the Project Administrator of both projects. Create a new VPC and add all instances.

**Correct Answer: B**

  **dan80** Highly Voted 1 year, 1 month ago

B - <https://cloud.google.com/vpc/docs/shared-vpc>  
upvoted 21 times

  **glam** Highly Voted 9 months, 3 weeks ago


B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 8 times

  **vmart** Most Recent 1 month, 2 weeks ago

b is correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

B is correct  
upvoted 1 times

  **[Removed]** 4 months, 1 week ago



B is correct. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 2 times

  **chithra1990** 5 months ago

So when the question talks about communication between 2 VPC's in same organization - go for Shared VPC If diff organization, we can consider VPC Peering  
upvoted 3 times

  **GCP\_Student1** 5 months, 1 week ago



B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 3 times

  **DucSiu** 5 months, 4 weeks ago


B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 2 times

  **victory108** 6 months ago

B - Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 1 times

  **don\_v** 7 months, 3 weeks ago

I think the authors of this question meant B as the correct answer, - Shared VPC.  
However, it says "Each group of Compute Engine instances is running in its own VPC", so there are multiple VPCs. Shared VPC is just a single VPC created in the host project and shared with all service projects. RTFM  
upvoted 4 times

  **swatititame** 8 months, 2 weeks ago


• B. Verify that both projects are in a GCP Organization. Share the VPC from one project and request that the Compute Engine instances in the other project use this shared VPC.  
upvoted 1 times

  **nwk** 9 months ago

(B) - <https://cloud.google.com/vpc/docs/provisioning-shared-vpc>

Shared VPC allows you to export subnets from a VPC network in a host project to other service projects in the same organization. Instances in the service projects can have network connections in the shared subnets of the host project

upvoted 2 times

  **stepkurniawan** 11 months, 1 week ago

Its either B or D, but B requires less "creating a VPC", so B is a better way.

upvoted 2 times

  **raj231** 10 months, 1 week ago

creating VPC is cost effective use the current VPC so B is correct.

upvoted 1 times

  **Ciumela** 1 year, 1 month ago

yes B is correct

upvoted 6 times



You want to add a new auditor to a Google Cloud Platform project. The auditor should be allowed to read, but not modify, all project items. How should you configure the auditor's permissions?

- A. Create a custom role with view-only project permissions. Add the user's account to the custom role.
- B. Create a custom role with view-only service permissions. Add the user's account to the custom role.
- C. Select the built-in IAM project Viewer role. Add the user's account to this role.
- D. Select the built-in IAM service Viewer role. Add the user's account to this role.

**Correct Answer:** C

Reference:

<https://cloud.google.com/resource-manager/docs/access-control-proj>

- cloudenth01**

Highly Voted

1 year, 1 month ago

C is correct  
roles/Viewer role provides access to all resources under the projects but do not alter the state of these resources  
upvoted 27 times
- glam**

Highly Voted

9 months, 3 weeks ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.  
upvoted 12 times
- A\_ksh\_a\_Y**

Most Recent

2 months, 1 week ago

there is nothing like a project Viewer role ? its only a projectBroswer role.  
upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

C is correct  
upvoted 1 times
- [Removed]**

4 months, 1 week ago

C is correct. Select the built-in IAM project Viewer role. Add the user's account to this role.  
upvoted 2 times
- EABDAJA**

4 months, 3 weeks ago

C is correct  
upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.  
upvoted 2 times
- DucSiu**

5 months, 4 weeks ago

C. Select the built-in IAM project Viewer role. Add the user's account to this role.  
upvoted 2 times
- wanmad**

6 months ago

C - for sure  
upvoted 1 times
- swatititame**

8 months, 2 weeks ago

- C. Select the built-in IAM project Viewer role. Add the user's account to this role.upvoted 3 times
- nmohapatra87**

8 months, 4 weeks ago

C is correct  
Select the built-in IAM project Viewer role. Add the user's account to this role.  
upvoted 3 times
- GopinathM**

10 months, 1 week ago

C is Correct  
upvoted 2 times
- SSPC**

11 months, 1 week ago

C is correct  
upvoted 2 times



You are operating a Google Kubernetes Engine (GKE) cluster for your company where different teams can run non-production workloads. Your Machine Learning (ML) team needs access to Nvidia Tesla P100 GPUs to train their models. You want to minimize effort and cost. What should you do?

- A. Ask your ML team to add the `accelerator: gpu` annotation to their pod specification.
- B. Recreate all the nodes of the GKE cluster to enable GPUs on all of them.
- C. Create your own Kubernetes cluster on top of Compute Engine with nodes that have GPUs. Dedicate this cluster to your ML team.
- D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

Correct Answer: B

- John\_Iam**

Highly Voted

1 year, 1 month ago

D is the correct answer.

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

upvoted 34 times
- tablet444**

11 months, 3 weeks ago

the documentation states "Limitations  
Before using GPUs on GKE, keep in mind the following limitations:

You cannot add GPUs to existing node pools.  
GPU nodes cannot be live migrated during maintenance events."

upvoted 6 times
- nightflyer**

7 months, 2 weeks ago

In this case it is about adding a GPU enabled node pool not a GPU to an existing node-pool

upvoted 8 times
- glam**

Highly Voted

9 months, 3 weeks ago

D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

upvoted 9 times
- shankyomre01**

Most Recent

2 months ago

D is a very wrong candidate but I would say B. because we also need to care take of cost .  
if we create a new node pool with GPU and don't delete exiting node pool it's increase pricing.  
i would go with B to recreate node pool with GPU as it's test env so I don't need to care about downtime.

upvoted 3 times
- mcaromit**

2 months, 2 weeks ago

D is correct

upvoted 1 times
- [Removed]**

4 months ago

D is correct. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

upvoted 2 times
- pondai**

4 months, 1 week ago

D & B can add gpu to this cluster ,but "minimize effort" so I pick D

upvoted 1 times
- ARDY**

5 months ago

D for sure

upvoted 1 times
- GCP\_Student1**

5 months, 1 week ago

D - Add a new, GPU-enabled node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.

upvoted 2 times
- DucSiu**

5 months, 4 weeks ago

Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the `cloud.google.com/gke-accelerator: nvidia-tesla-p100` nodeSelector to their pod specification.


upvoted 1 times

  **rvgcp** 6 months ago


D is the correct answer  
upvoted 1 times

  **EABDAJA** 6 months ago

D is Correct  
upvoted 1 times

  **guid1984** 6 months, 1 week ago



D is correct as nodepools are way to logically separate your workloads and machinetypes based on your organization team/project structure and requirements.  
upvoted 1 times

  **victory108** 6 months, 1 week ago


D - Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.  
upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

can you pls verify once..to confirm whether D is correct.  
upvoted 3 times

  **swatititame** 8 months, 2 weeks ago

• D. Add a new, GPU-enabled, node pool to the GKE cluster. Ask your ML team to add the cloud.google.com/gke -accelerator: nvidia-tesla-p100 nodeSelector to their pod specification.  
upvoted 1 times

  **ayj** 8 months, 2 weeks ago

I think D:  
With GKE, you can create node pools equipped with NVIDIA Tesla® K80, P100, P4, V100, and T4 GPUs.  
You cannot add GPUs to existing node pools. So creating a new one seems like the right answer. B is a bit meh as an answer I feel  
upvoted 2 times

  **ayj** 8 months, 2 weeks ago

I think D:  
With GKE, you can create node pools equipped with NVIDIA Tesla® K80, P100, P4, V100, and T4 GPUs.  
You cannot add GPUs to existing node pools. So creating a new one seems like the right answer. B is a bit meh as an answer I feel  
upvoted 2 times

## Question #82

## Topic 1

Your VMs are running in a subnet that has a subnet mask of 255.255.255.240. The current subnet has no more free IP addresses and you require an additional 10 IP addresses for new VMs. The existing and new VMs should all be able to reach each other without additional routes. What should you do?

- A. Use gcloud to expand the IP range of the current subnet.
- B. Delete the subnet, and recreate it using a wider range of IP addresses.
- C. Create a new project. Use Shared VPC to share the current network with the new project.
- D. Create a new subnet with the same starting IP but a wider range to overwrite the current subnet.

**Correct Answer: C**

🗲️ 👤 **JustLearning** Highly Voted 👍 1 year, 1 month ago

A: Expand the existing subnet.

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 35 times

🗲️ 👤 **glam** Highly Voted 👍 9 months, 3 weeks ago

A. Use gcloud to expand the IP range of the current subnet.

upvoted 9 times

🗲️ 👤 **kashi123** Most Recent ⌚ 2 weeks, 5 days ago

A is correct , expand the existing subnet .  
subnets can expand but can not shrink .

upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

A is correct

upvoted 1 times

🗲️ 👤 **NARWAL** 3 months, 3 weeks ago

A: Expand the existing subnet

upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago

A is correct. Use gcloud to expand the IP range of the current subnet.

upvoted 2 times

🗲️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

A. Use gcloud to expand the IP range of the current subnet.

upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

A is correct.

upvoted 1 times

🗲️ 👤 **raheja21** 5 months ago

its A, subnet can be expanded.

upvoted 1 times

🗲️ 👤 **ARDY** 5 months ago

it's A.

Expanding subnet is making it from /28 to /26 or less. Not /30.

upvoted 2 times

🗲️ 👤 **DucSiu** 5 months, 4 weeks ago

A. Use gcloud to expand the IP range of the current subnet.

upvoted 1 times

🗲️ 👤 **LearningGCP** 6 months ago

A: Expand the existing subnet.

upvoted 1 times



🗲️ 👤 **Bhagirathi** 7 months, 1 week ago

200% A.



upvoted 2 times

  **Range2019** 7 months, 3 weeks ago

The correct answer is A. I thought C was the right option but again this is GCP and not AWS.  
<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>  
upvoted 2 times

  **sanhoo** 1 month, 4 weeks ago

I like your comment. It is GCP not AWS  
upvoted 1 times

  **sanhoo** 1 month, 4 weeks ago

and thanks for providing the link.  
upvoted 1 times



  **swatititame** 8 months, 2 weeks ago

- A. Use gcloud to expand the IP range of the current subnet.

upvoted 2 times

  **nwk** 9 months ago

Vote A  
upvoted 2 times

  **rungcpnow** 9 months, 3 weeks ago

You can increase to /24.. Create a vpc and add a subnet 192.168.10.0/28. Save it.. Again go back and expand it /24. It is possible. So Ans is A.  
upvoted 3 times

Your organization uses G Suite for communication and collaboration. All users in your organization have a G Suite account. You want to grant some G Suite users access to your Cloud Platform project. What should you do?

A. Enable Cloud Identity in the GCP Console for your domain.

B. Grant them the required IAM roles using their G Suite email address.

C. Create a CSV sheet with all users' email addresses. Use the gcloud command line tool to convert them into Google Cloud Platform accounts.

D. In the G Suite console, add the users to a special group called cloud-console-users@yourdomain.com. Rely on the default behavior of the Cloud Platform to grant users access if they are members of this group.

**Correct Answer: B**  
Reference:  
<https://cloud.google.com/resource-manager/docs/creating-managing-organization>

- austini**

Highly Voted

1 year, 1 month ago

B is correct  
upvoted 21 times
- Ciumela**

Highly Voted

1 year, 1 month ago

B is correct: To actively adopt the Organization resource, the G Suite or Cloud Identity super admins need to assign the Organization Administrator Cloud IAM role to a user or group  
upvoted 10 times
- mcaromit**

Most Recent

2 months, 2 weeks ago

B is correct  
upvoted 2 times
- Christiank**

4 months ago

B. Grant them the required IAM roles using their G Suite email address.  
upvoted 2 times
- [Removed]**

4 months ago

B is correct. Grant them the required IAM roles using their G Suite email address.  
upvoted 3 times
- GCP\_Student1**

5 months ago

B. Grant them the required IAM roles using their G Suite email address.  
upvoted 1 times
- victory108**

6 months, 1 week ago

B - Grant them the required IAM roles using their G Suite email address.  
upvoted 1 times
- swatititame**

8 months, 2 weeks ago

• A. Use gcloud to expand the IP range of the current subnet.

upvoted 3 times

**TAvenger**

5 months, 1 week ago

this is for previous question :)

upvoted 4 times

**nwk**

9 months ago

Should be B

Default behavior does not grant access to the "your GCP Project"

Default behavior allow only create billing account and project - When the organization is created, all users in your domain are automatically granted Project Creator and Billing Account Creator IAM roles at the organization level. This enables users in your domain to continue creating projects with no disruption.

upvoted 2 times

**glam**

9 months, 3 weeks ago

B. Grant them the required IAM roles using their G Suite email address.  
upvoted 6 times

**[Removed]**

10 months, 1 week ago

D is better with best practice.
- <https://www.examtopycs.com/exams/google/associate-cloud-engineer/custom-view/>

175/369

upvoted 1 times

  **[Removed]** 10 months, 1 week ago

Take back B is correct. There are a lot of uncertain words in the option.

upvoted 1 times



## Question #84

## Topic 1

You have a Google Cloud Platform account with access to both production and development projects. You need to create an automated process to list all compute instances in development and production projects on a daily basis. What should you do?

- A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.
- B. Create two configurations using gsutil config. Write a script that sets configurations as active, individually. For each configuration, use gsutil compute instances list to get a list of compute resources.
- C. Go to Cloud Shell and export this information to Cloud Storage on a daily basis.
- D. Go to GCP Console and export this information to Cloud SQL on a daily basis.

Correct Answer: A

  **cloudehthu01** Highly Voted 1 year, 1 month ago

A is correct  
upvoted 26 times

  **glam** Highly Voted 9 months, 3 weeks ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.  
upvoted 10 times

  **mcaromit** Most Recent 2 months, 2 weeks ago

A is correct  
upvoted 1 times

  **[Removed]** 4 months ago



A is correct. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.  
upvoted 2 times

  **GCP\_Student1** 5 months ago



A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.  
upvoted 2 times

  **GCP\_Student1** 5 months, 3 weeks ago

A is the correct answer  
upvoted 1 times

  **DucSiu** 5 months, 4 weeks ago

A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.  
upvoted 2 times

  **EABDAJA** 6 months, 1 week ago

Answer A  
upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

200% A.  
upvoted 4 times

  **swatititame** 8 months, 2 weeks ago

• A. Create two configurations using gcloud config. Write a script that sets configurations as active, individually. For each configuration, use gcloud compute instances list to get a list of compute resources.  
upvoted 1 times

  **stepkurniawan** 11 months, 1 week ago

since gsutil is for storage, B is definitely wrong.  
upvoted 5 times

  **ESP\_SAP** 11 months, 2 weeks ago

Correct Answer is (A):

gcloud compute instances list - list Google Compute Engine instances  
gcloud compute instances list displays all Google Compute Engine instances in a project.

By default, instances from all zones are listed. The results can be narrowed down using a filter: --filter="zone:( ZONE ... )".

<https://cloud.google.com/sdk/gcloud/reference/compute/instances/list>  
upvoted 6 times

## Question #85

## Topic 1

You have a large 5-TB AVRO file stored in a Cloud Storage bucket. Your analysts are proficient only in SQL and need access to the data stored in this file. You want to find a cost-effective way to complete their request as soon as possible. What should you do?

- A. Load data in Cloud Datastore and run a SQL query against it.
- B. Create a BigQuery table and load data in BigQuery. Run a SQL query on this table and drop this table after you complete your request.
- C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.
- D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries.



**Correct Answer: C**

  **mohdafiuddin** Highly Voted 7 months ago

Breaking down the question into key points -

1. 5-TB AVRO file stored in a Cloud Storage bucket.
2. Analysts are proficient only in SQL
3. cost-effective way to complete their request as soon as possible

- A. ....Load data in Cloud Datastore... (Not Correct because Cloud Datastore is not a good option to run SQL Queries)
  - B. ...Load data in BigQuery.... (Not Cost Effective because loading the data which is already present in the bucket into BigQuery again is expensive)
  - C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request. (This is the right answer as it meets all the requirements from the question)
  - D. Create a Hadoop cluster and copy the AVRO file to NDFS by compressing it. Load the file in a hive table and provide access to your analysts so that they can run SQL queries. (Too roundabout and indirect. Not the right option)
- upvoted 33 times

  **pondai** 4 months, 1 week ago

listen this guy  
upvoted 3 times

  **Ciumela** Highly Voted 1 year, 1 month ago

C is correct: <https://cloud.google.com/bigquery/external-data-sources>  
upvoted 18 times

  **mcaromit** Most Recent 2 months, 2 weeks ago

C is the most cost effective option  
upvoted 3 times

  **mistryminded** 3 months ago

C seems to be correct [https://cloud.google.com/bigquery/docs/loading-data-cloud-storage-avro#avro\\_schemas](https://cloud.google.com/bigquery/docs/loading-data-cloud-storage-avro#avro_schemas)  
upvoted 2 times

  **[Removed]** 4 months ago

C is correct. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.  
upvoted 2 times

  **drizzydroo** 4 months, 2 weeks ago

has anyone actually tried this? I don't know how to have a table point to a bucket without actually loading the data into the table.  
upvoted 1 times

  **GCP\_Student1** 5 months ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.  
upvoted 2 times

  **swatitame** 8 months, 2 weeks ago

• C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.  
upvoted 3 times

  **nwk** 9 months ago

(C) - An external data source (also known as a federated data source) is a data source that you can query directly even though the data is not stored in BigQuery. Instead of loading or streaming the data, you create a table that references the external data source.

upvoted 4 times

  **glam** 9 months, 3 weeks ago

C. Create external tables in BigQuery that point to Cloud Storage buckets and run a SQL query on these external tables to complete your request.

upvoted 6 times

  **stepkurniawan** 11 months, 1 week ago

B will be too expensive to load all the 5 TB data to BigQuery, and then dropping it again directly.

upvoted 3 times

  **[Removed]** 1 year, 1 month ago

C is correct

upvoted 8 times

You need to verify that a Google Cloud Platform service account was created at a particular time. What should you do?

- A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.
- B. Filter the Activity log to view the Configuration category. Filter the Resource type to Google Project.
- C. Filter the Activity log to view the Data Access category. Filter the Resource type to Service Account.
- D. Filter the Activity log to view the Data Access category. Filter the Resource type to Google Project.

Correct Answer: D

- John\_Iam**

Highly Voted

1 year, 1 month ago

Correct Answer is A.  
Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.  
upvoted 39 times
- mlantonis**

1 year, 1 month ago

I agree A  
upvoted 6 times
- shafiqee1**

Highly Voted

1 year ago

A - I reproduced in my project.  
upvoted 15 times
- AD\_0525**

Most Recent

2 months ago

Data access category will only have the details of the service accounts which and when access the data, so option should be A  
upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- viswanand**

3 months, 1 week ago

Tried and tested A is correct  
upvoted 1 times
- [Removed]**

4 months ago

A is correct. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.  
upvoted 2 times
- GCP\_Student1**

5 months, 3 weeks ago

Vote for A  
upvoted 2 times
- rvgcp**

6 months ago

A is the right answer  
upvoted 1 times
- Bhagirathi**

7 months, 1 week ago

200% A.  
upvoted 3 times
- Likhith**

8 months, 1 week ago

A ne correct  
upvoted 2 times
- swatitame**

8 months, 2 weeks ago

• A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.

upvoted 3 times
- glam**

9 months, 3 weeks ago

A. Filter the Activity log to view the Configuration category. Filter the Resource type to Service Account.  
upvoted 10 times
- stepkurniawan**

11 months, 1 week ago

I tested it, A is correct.  
upvoted 5 times



## Question #87

## Topic 1

You deployed an LDAP server on Compute Engine that is reachable via TLS through port 636 using UDP. You want to make sure it is reachable by clients over that port. What should you do?

- A. Add the network tag allow-udp-636 to the VM instance running the LDAP server.
- B. Create a route called allow-udp-636 and set the next hop to be the VM instance running the LDAP server.
- C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.
- D. Add a network tag of your choice to the instance running the LDAP server. Create a firewall rule to allow egress on UDP port 636 for that network tag.

**Correct Answer: C**

🗲️ 👤 **kishoredeena** Highly Voted 👍 1 year, 1 month ago

Option C is the right one  
upvoted 24 times

🗲️ 👤 **cloudenth01** Highly Voted 👍 1 year, 1 month ago

C is correct  
You tag the instances ,then create ingress firewall rules to allow udp on desired port for target-tags name applied to instances  
upvoted 14 times

🗲️ 👤 **devopsbatch** Most Recent ⌚ 1 month, 3 weeks ago

C A tag is simply a character string added to a tags field in a resource, such as Compute Engine virtual machine (VM) instances or instance templates. A tag is not a separate resource, so you cannot create it separately. All resources with that string are considered to have that tag. Tags enable you to make firewall rules and routes applicable to specific VM instances.  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

C is correct  
upvoted 1 times

🗲️ 👤 **kopper2019** 3 months, 2 weeks ago

C, D is for egress so for outgoing traffic  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago

C is correct. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.  
upvoted 2 times

🗲️ 👤 **bart54** 4 months, 2 weeks ago

Defo C as the LDAP server will sit internally in the network protected by a Firewall, so an ingress rule will allow traffic to be routed internally to the LDAP Server.  
upvoted 1 times

🗲️ 👤 **pentium2000** 4 months, 3 weeks ago

C, LDAP can be reached by customers means ingress traffic.  
upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

C is correct  
upvoted 1 times

🗲️ 👤 **victory108** 4 months, 3 weeks ago

C - Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.  
upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months ago











C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.  
upvoted 2 times

🗲️ 👤 **Aravind8898** 5 months, 3 weeks ago

it should be D. because it has to reachable to clients with egress  
upvoted 2 times

🗲️ 👤 **nitinz** 5 months, 1 week ago

its C, ingress is incoming, egress is outgoing. for ldap servers its ingress rule.  
upvoted 1 times

-  **DucSiu** 5 months, 4 weeks ago  
D. Add a network tag of your choice to the instance running the LDAP server. Create a firewall rule to allow egress on UDP port 636 for that network tag.  
upvoted 2 times
-  **nherrerab** 6 months, 2 weeks ago  
C is correct.  
upvoted 1 times
-  **Bhagirathi** 7 months, 1 week ago  
200% C.  
upvoted 2 times
-  **swatititame** 8 months, 2 weeks ago
  - C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.upvoted 3 times
-  **glam** 9 months, 2 weeks ago  
C. Add a network tag of your choice to the instance. Create a firewall rule to allow ingress on UDP port 636 for that network tag.  
upvoted 11 times



You need to set a budget alert for use of Compute Engine services on one of the three Google Cloud Platform projects that you manage. All three projects are linked to a single billing account. What should you do?

- A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.
- B. Verify that you are the project billing administrator. Select the associated billing account and create a budget and a custom alert.
- C. Verify that you are the project administrator. Select the associated billing account and create a budget for the appropriate project.
- D. Verify that you are project administrator. Select the associated billing account and create a budget and a custom alert.

Correct Answer: B

- kishoredeena

Highly Voted

1 year, 1 month ago

I think the answer is A, You can rely on default alert. No need for custom alert

upvoted 23 times
- Eshkrkrkr

8 months, 2 weeks ago

One point - there is no such role as Project Billing Administrator - it should be Project Billing Manager but he can't create budgets, the only one who can - Billing Account Administrator. Nor Project Administrator exists. Very tricky question, maybe the option a wrong, hope smb will catch it on exam and pass some light on real variants.

<https://cloud.google.com/iam/docs/understanding-roles#billing-roles>

upvoted 3 times
- Wachy

7 months, 1 week ago

Eshkrkrkr read the question calmly. The role there is Billing Administrator. Not Project Billing Administrator.

It's more like: "Verify you are the project; billing administrator"

upvoted 5 times
- alan9999

10 months ago

Right its not asking to set custom alert

upvoted 1 times
- Ciumela

Highly Voted

1 year, 1 month ago

A is correct, as you can set a default alert also on a single project: <https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 9 times
- mwwoodm

10 months, 3 weeks ago

Agreed. Per the link included: "To create a budget for your Cloud Billing account, you must be a Billing Account Administrator on the Cloud Billing account." So that eliminates C & D. Then no need for custom alert, eliminating B. The answer is A.

upvoted 2 times
- Bha9886

Most Recent

2 months, 1 week ago

Why admin is not correcting the wrong ones

upvoted 3 times
- mcaromit

2 months, 2 weeks ago

A is correct

upvoted 1 times
- [Removed]

4 months ago

A is correct. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 2 times
- yuvi69

4 months, 1 week ago

correct option is A. because default alert or custom alert only deals with percentages of amount spent. they has nothing to do with notifications

upvoted 1 times
- ShakthiGCP

4 months, 4 weeks ago

Ans: A . Tried it in GCP console. You can selecte the required project and also select the services that you need for alert.

upvoted 1 times
- arsav



5 months, 1 week ago

go for A = 100%

watch more at


<https://www.youtube.com/watch?v=F4omjjMZ54k>

upvoted 1 times

  **DucSiu** 5 months, 4 weeks ago

A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 2 times

  **EABDAJA** 6 months, 1 week ago



A is Correct

upvoted 1 times

  **nherrerab** 6 months, 2 weeks ago

A is correct.

upvoted 1 times

  **cRobert** 7 months, 1 week ago

If you just go into the Google Cloud console, you can easily see that there's no custom alert.  
For Alerts, you first set the project, then the service, you have the option to select All Services if you want.

Correct answer is A

upvoted 4 times

  **anushkasf** 7 months, 2 weeks ago

@Bhagirathi did you got finalized answer?

upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

No. It is not clear.

upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago


confusing dont know A or B ?

upvoted 1 times

  **swatititame** 8 months, 2 weeks ago

• A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 3 times

  **JCGO** 9 months, 2 weeks ago

Supposedly A. billing alerts to different project is not something that is 'custom'. standard feature

upvoted 1 times

  **glam** 9 months, 2 weeks ago

A. Verify that you are the project billing administrator. Select the associated billing account and create a budget and alert for the appropriate project.

upvoted 7 times

You are migrating a production-critical on-premises application that requires 96 vCPUs to perform its task. You want to make sure the application runs in a similar environment on GCP. What should you do?

- A. When creating the VM, use machine type n1-standard-96.
- B. When creating the VM, use Intel Skylake as the CPU platform.
- C. Create the VM using Compute Engine default settings. Use gcloud to modify the running instance to have 96 vCPUs.
- D. Start the VM using Compute Engine default settings, and adjust as you go based on Rightsizing Recommendations.

Correct Answer: B

- dan80**

Highly Voted

1 year, 1 month ago

A is correct - <https://cloud.google.com/compute/docs/machine-types>  
upvoted 26 times
- glam**

Highly Voted

9 months, 2 weeks ago

A. When creating the VM, use machine type n1-standard-96.  
upvoted 8 times
- mcaromit**

Most Recent

2 months, 2 weeks ago

A is correct  
upvoted 1 times
- [Removed]**

4 months ago

A is correct. When creating the VM, use machine type n1-standard-96.  
upvoted 2 times
- emerp**

4 months, 1 week ago

A is correct  
upvoted 1 times
- Hi2ALL**

4 months, 3 weeks ago

A is correct  
upvoted 1 times
- GCP\_Student1**

5 months ago

A. When creating the VM, use machine type n1-standard-96.  
upvoted 2 times
- antomastro**

5 months, 1 week ago

I think A is correct  
upvoted 2 times
- Tez1**

5 months, 2 weeks ago

I think D would be correct if this wasn't a mission critical server, and thus would not allow you to spend time adjusting the VM.  
  
Therefore A is correct.  
upvoted 1 times
- DucSiu**

5 months, 4 weeks ago

A. When creating the VM, use machine type n1-standard-96.  
upvoted 2 times
- desekis**

7 months ago

go for A  
upvoted 1 times
- Harish57**

7 months, 3 weeks ago

A is correct  
upvoted 1 times
- Bhagirathi**


7 months, 3 weeks ago

A or C - what is finally to go for ?  
upvoted 1 times
- alpharomeo9**

7 months, 3 weeks ago


definitely A

upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

- A. When creating the VM, use machine type n1-standard-96.

upvoted 3 times

  **prasanu** 9 months, 3 weeks ago

A is correct , just verified in pricing calculator, we can create up to 96vcpu and 624 gb ram

upvoted 1 times

  **GopinathM** 10 months, 1 week ago

C is correct

upvoted 1 times

  **szakaria** 1 year ago

For me C and A are correct. However, with C you are more flexible on the RAM. Hence, I vote for C.

upvoted 1 times

  **someoneinthecloud** 12 months ago

A is correct. The question never asked about RAM. Plus, you need to stop the instance to change its family and C says "change the running instance" which invalidates this alternative.

upvoted 8 times

  **[Removed]** 10 months, 1 week ago

You cannot do C as you need to select a type of CE while creating that we have to start by selecting a family of CE. Hence A.

upvoted 2 times



You want to configure a solution for archiving data in a Cloud Storage bucket. The solution must be cost-effective. Data with multiple versions should be archived after 30 days. Previous versions are accessed once a month for reporting. This archive data is also occasionally updated at month-end. What should you do?



- A. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Coldline Storage.
- B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.
- C. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Coldline Storage.
- D. Add a bucket lifecycle rule that archives data from regional storage after 30 days to Nearline Storage.



**Correct Answer:** B



Reference:



<https://cloud.google.com/storage/docs/managing-lifecycles>



  **neelesh88** Highly Voted 1 year, 1 month ago  
B is correct  
upvoted 23 times



  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer (B):  
  
NumberOfNewerVersions  
The NumberOfNewerVersions condition is typically only used in conjunction with Object Versioning. If the value of this condition is set to N, an object version satisfies the condition when there are at least N versions (including the live version) newer than it. For a live object version, the number of newer versions is considered to be 0. For the most recent noncurrent version, the number of newer versions is 1 (or 0 if there is no live object version), and so on.  
  
Important: When specifying this condition in a .json configuration file, you must use numNewerVersions instead of NumberOfNewerVersions.  
  
<https://cloud.google.com/storage/docs/lifecycle#numberofnewerversions>  
upvoted 15 times



  **jcols** Most Recent 2 months ago  
B is correct. C and D don't mention newer object versions. Then, it's Coldline vs. Nearline, the best option is Nearline because:  
\* "Previous versions are accessed once a month for reporting"  
\* "Nearline Storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline Storage is a great choice." --  
<https://cloud.google.com/storage/docs/storage-classes#nearline>  
upvoted 2 times



  **mcaromit** 2 months, 2 weeks ago  
B is correct  
upvoted 1 times

  **[Removed]** 4 months ago  
B is correct. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.  
upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago  
B is correct  
upvoted 1 times

  **GCP\_Student1** 5 months ago  
B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.  
upvoted 2 times

  **arsav** 5 months, 1 week ago  
"Previous versions are accessed once a month for reporting"  
  
only option C and D says, that you are storing the data in "regional" storage and then trying to move from either "coldline" or "nearline"  
  
so, i will go with option D. to hold the data which is accessed once a month in "regional" storage and then after 30 days i archive it and move to "nearline"  
  
also note, question says you occasionally update the archive data. so near-line would be right choice.  
upvoted 3 times

  **nitinz** 5 months, 1 week ago



I disagree with all of you, correct answer is D. The key line in question is " The solution must be cost-effective." When you are doing multi-region bucket then you pay for data at rest per location. It will be nearline storage but single region.

upvoted 2 times

  **user843983409** 4 months, 1 week ago

D doesnt refer to the "data with versions" hence it applies to all data which is not what is asked in the question. My vote for B

upvoted 2 times

  **DucSiu** 5 months, 4 weeks ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 2 times

  **mohdafiuddin** 7 months ago

Key Requirements -

1. Data with \*\*\*multiple versions\*\*\* should be archived after 30 days.
2. Previous versions are accessed once a month for reporting. This archive data is also occasionally updated at month-end.

A. ...Coldline Storage... Not suitable because data is accessed once a month

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

This is the right option



Ref - <https://cloud.google.com/storage/docs/storage-classes#nearline> and

<https://cloud.google.com/storage/docs/lifecycle>

C. ...Coldline Storage....Not suitable because data is accessed once a month

D. ...regional storage after 30 days to Nearline Storage... Nothing is mentioned about regional storage in the question. where did we get this from in the options? Not suitable

upvoted 3 times

  **nitinz** 5 months, 1 week ago

D, question says cost effective solution.... you are paying for data at rest per location.

upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

- B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 2 times

  **glam** 9 months, 2 weeks ago

B. Add a bucket lifecycle rule that archives data with newer versions after 30 days to Nearline Storage.

upvoted 5 times

  **Ciumela** 1 year, 1 month ago



B is correct answer

upvoted 2 times

  **nithinpb180** 1 year, 1 month ago

B is correct, mentioned that data with multiple versions need to be archived.

upvoted 2 times

  **raffiq** 1 year, 1 month ago



I think C is correct answer because archival data

upvoted 1 times

  **[Removed]** 10 months, 3 weeks ago

C would be correct if Archival data would not be accessed, as cold line have retention of 90 days.

upvoted 1 times

  **raffiq** 1 year, 1 month ago

sorry, its Answer B Nearline storage is correct answer.

upvoted 5 times

## Question #91

## Topic 1

Your company's infrastructure is on-premises, but all machines are running at maximum capacity. You want to burst to Google Cloud. The workloads on Google Cloud must be able to directly communicate to the workloads on-premises using a private IP range. What should you do?

- A. In Google Cloud, configure the VPC as a host for Shared VPC.
- B. In Google Cloud, configure the VPC for VPC Network Peering.
- C. Create bastion hosts both in your on-premises environment and on Google Cloud. Configure both as proxy servers using their public IP addresses.
- D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.

**Correct Answer: D**

  **SIX** Highly Voted 1 year, 1 month ago

I believe D is the right answer  
upvoted 39 times

  **dan80** 1 year, 1 month ago

B is correct - <https://cloud.google.com/solutions/best-practices-vpc-design> . this answer also on all machines are running at maximum capacity.  
upvoted 1 times

  **JustLearning** 1 year, 1 month ago

vpc network peering does not connect to on-prem. Cloud VPN is the correct solution.  
<https://cloud.google.com/vpn/docs/concepts/overview>  
upvoted 19 times

  **mlantonis** 1 year, 1 month ago

You need VPN, so D is the correct. VPC network peering is between VPCs.  
upvoted 14 times

  **ESP\_SAP** Highly Voted 11 months, 2 weeks ago

Correct Answer is (D):

Access internal IPs directly

Your VPC network's internal (RFC 1918) IP addresses are directly accessible from your on-premises network with peering, no NAT device or VPN tunnel required.

Hybrid made easy

Today's business climate demands flexibility. Connecting your on-premises resources to your cloud resources seamlessly, with minimum latency or interruption, is a business-critical requirement. The speed and reliability of Cloud Interconnect lets you extend your organization's data center network into Google Cloud, simply and easily, while options such as Cloud VPN provide flexibility for all your workloads. This unlocks the potential of hybrid app development and all the benefits the cloud has to offer.

In the graphic below: What GCP Connection is right for you? shows clearly what is the method for extend your on premise network (IP Private communication).



What GCP Connection is right for you?

<https://cloud.google.com/hybrid-connectivity>

upvoted 12 times

  **mcaromit** Most Recent 2 months, 2 weeks ago

D is correct  
upvoted 1 times

  **[Removed]** 3 months, 3 weeks ago

Not A and B, because for network peering, shared VPC you need to either be in the same project or same organization. Here we are in on-premise and cloud.



Maybe C, but why do you need an extra server on both GCP and On-premise for a proxy?. We already have solutions for proxy inside GCP.

D is correct because, we generally connect GCP to on-premise using VPN























upvoted 3 times

  **[Removed]** 4 months ago

D is correct. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 2 times

  **Hi2ALL** 4 months, 3 weeks ago

D is correct since VPC peering only works on GCP cloud network  
upvoted 1 times

-   **GCP\_Student1** 5 months ago  
D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 2 times
-   **DucSiu** 5 months, 4 weeks ago  
D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 2 times
-   **rvgcp** 6 months ago  
D is the right answer  
upvoted 1 times
-   **LearningGCP** 6 months ago  
D correct Answer  
upvoted 1 times
-   **INASR** 6 months ago  
D is correct  
upvoted 2 times
-   **victory108** 6 months ago  
D - Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 3 times
-   **swatititame** 8 months, 2 weeks ago  
• D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 3 times
-   **gcpaace** 8 months, 2 weeks ago  
I will go for D. VPN is required to connect to on-premises network.  
upvoted 2 times
-   **awscloudgeek** 8 months, 2 weeks ago  
D is the right answer  
upvoted 2 times
-   **nwk** 9 months ago  
Network peering is for connection between VPC  
Vote for D  
upvoted 3 times
-   **glam** 9 months, 2 weeks ago  
D. Set up Cloud VPN between the infrastructure on-premises and Google Cloud.  
upvoted 6 times



You want to select and configure a solution for storing and archiving data on Google Cloud Platform. You need to support compliance objectives for data from one geographic location. This data is archived after 30 days and needs to be accessed annually. What should you do?

- A. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.
- B. Select Multi-Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- C. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.
- D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

Correct Answer: D

- SIX

Highly Voted

1 year, 1 month ago

D

Google Cloud Coldline is a new cold-tier storage for archival data with access frequency of less than once per year. Unlike other cold storage options, Nearline has no delays prior to data access, so now it is the leading solution among competitors.

upvoted 33 times
- JustLearning

1 year, 1 month ago

D is correct. Coldline is a better choice.

upvoted 6 times
- dan80

1 year, 1 month ago

C is correct - This data is archived after 30 days - Nearline Storage 30 days , Coldline Storage 90 days

<https://cloud.google.com/storage/docs/storage-classes>

upvoted 8 times
- lxgywil

2 months, 3 weeks ago

The answer is D.

The main thing here is how often the data is retrieved. The question is saying that data needs to be accessed annually - i.e. once a year. Therefore, you should choose Coldline Storage, as it implies less frequent access than Nearline. (Archival Storage would fit even better but there's no such option)

From the link you provided:

"Nearline Storage is ideal for data you plan to read or modify on average once per month or less."

and

"Coldline Storage is ideal for data you plan to read or modify at most once a quarter. "

upvoted 3 times
- gh999l

8 months, 3 weeks ago

you have misunderstood minimum storage period here, nearline storage class minimum you have to plan for 30 days

upvoted 4 times
- y2kniel

9 months, 2 weeks ago

D, It is saying AFTER 30 days. We should use coldline storage

upvoted 5 times
- ESP\_SAP

Highly Voted

11 months, 2 weeks ago

Correct Answer is (D):

<https://cloud.google.com/storage/docs/storage-classes>

Nearline Storage

Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. Nearline Storage is a better choice than Standard Storage in scenarios where slightly lower availability, a 30-day minimum storage duration, and costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Nearline Storage is ideal for data you plan to read or modify on average once per month or less. For example, if you want to continuously add files to Cloud Storage and plan to access those files once a month for analysis, Nearline Storage is a great choice.

Nearline Storage is also appropriate for data backup, long-tail multimedia content, and data archiving. Note, however, that for data accessed less frequently than once a quarter, Coldline Storage or Archive Storage are more cost-effective, as they offer lower storage costs.

<https://cloud.google.com/storage/docs/storage-classes#nearline>

upvoted 9 times
- ESP\_SAP

11 months, 1 week ago

CORRECTION.

Correct Answer is (D):

The Real description is about Coldline storage Class:

Coldline Storage

Coldline Storage is a very-low-cost, highly durable storage service for storing infrequently accessed data. Coldline Storage is a better choice than Standard Storage or Nearline Storage in scenarios where slightly lower availability, a 90-day minimum storage duration, and higher costs for data access are acceptable trade-offs for lowered at-rest storage costs.

Coldline Storage is ideal for data you plan to read or modify at most once a quarter. Note, however, that for data being kept entirely for backup or archiving purposes, Archive Storage is more cost-effective, as it offers the lowest storage costs.

<https://cloud.google.com/storage/docs/storage-classes#coldline>

upvoted 6 times

  **bunnyabi** Most Recent 1 month ago

D is correct

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

D is correct

upvoted 1 times

  **[Removed]** 4 months ago

D is correct. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

upvoted 1 times

  **Devgela** 4 months, 2 weeks ago

D for sure.

upvoted 1 times

  **drizzydroo** 4 months, 3 weeks ago

Can someone verify that Coldline storage is now archive? I beleive the storage classes have chnages to

1.standard



2.nearline 30 days

3.coldline 90 days

4.archive 365 days.

Ans for this is D but think this might be an outdated one. Please correct me if im wrong.

upvoted 1 times

  **Hi2ALL** 4 months, 3 weeks ago

Question is asking Geo Location when data needed so I think A is correct in this scenario

upvoted 1 times

  **neerajgoyal** 4 months, 4 weeks ago


D is the correct answer

upvoted 1 times

  **GCP\_Student1** 5 months ago

D. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Coldline Storage.

upvoted 2 times

  **\_batman\_** 5 months, 1 week ago



People who are choosing option D is forgetting about minimum storage duration for coldline is 90 days. In the question the requirement is to archive the data after 30 days which is eligible for nearline. :)

upvoted 1 times

  **rmout** 5 months ago


A minimum storage duration applies to data stored using one of these storage classes. You can delete the object before it has been stored for this duration, but at the time of deletion you are charged as if the object was stored for the minimum duration

upvoted 2 times

  **DucSiu** 5 months, 4 weeks ago

C. Select Regional Storage. Add a bucket lifecycle rule that archives data after 30 days to Nearline Storage.

upvoted 2 times

  **RKS\_2021** 5 months, 1 week ago

Ans should be D from given choices other wise the right answer is Archive storage.

upvoted 1 times

  **rvgcp** 6 months ago

D - One geography and annually ..so regional and cold line

upvoted 2 times

  **INASR** 6 months ago

D , annually for coldline  
upvoted 1 times

  **desekis** 7 months ago

go for D  
upvoted 1 times

  **DucSiu** 7 months ago

D is best choice  
upvoted 2 times

  **mohdafiuddin** 7 months ago

Breaking the question into key points -

1. Storing and archiving data on Google Cloud Platform.
2. One geographic location.
3. Data is archived after 30 days
4. Needs to be accessed annually.

A. Select Multi-Regional Storage. -Not needed because we only need one geographic location.

B. Select Multi-Regional Storage. -Not needed because we only need one geographic location.

C. Regional...Nearline Storage - Not needed because we don't plan on accessing the data frequently (like once a month)

D. Regional....Coldline Storage - Right answer because we are only accessing data once a year.

upvoted 4 times

  **mohdafiuddin** 7 months ago

Resources -

[https://cloud.google.com/storage/docs/storage-classes#available\\_storage\\_classes](https://cloud.google.com/storage/docs/storage-classes#available_storage_classes)

Note - Don't confuse the minimum storage duration requirement with what's mentioned in the question. The question only speaks about the age at which the data needs to be moved out of our regular storage into one of the other storage options. we are going to keep this data in the other storage option for really long and that's why the question mentions that the data will be accessed once a year AFTER we move it to a different storage.

upvoted 4 times

Your company uses BigQuery for data warehousing. Over time, many different business units in your company have created 1000+ datasets across hundreds of projects. Your CIO wants you to examine all datasets to find tables that contain an employee\_ssn column. You want to minimize effort in performing this task.

What should you do?

- A. Go to Data Catalog and search for employee\_ssn in the search box.
- B. Write a shell script that uses the bq command line tool to loop through all the projects in your organization.
- C. Write a script that loops through all the projects in your organization and runs a query on INFORMATION\_SCHEMA.COLUMNS view to find the employee\_ssn column.
- D. Write a Cloud Dataflow job that loops through all the projects in your organization and runs a query on INFORMATION\_SCHEMA.COLUMNS view to find employee\_ssn column.

Correct Answer: D

- poogcp

Highly Voted

1 year, 1 month ago

Its A.

upvoted 20 times
- filco72

Highly Voted

1 year ago

Correct is A.

I tested on my account following this procedure: <https://cloud.google.com/bigquery/docs/quickstarts/quickstart-web-ui?authuser=4>

I created a data set and through Data Catalog I easily and effortlessly searched for the column name "gender"

upvoted 11 times
- dc2016bte0006

Most Recent

1 month, 2 weeks ago

IDK why the admin has set all the wrong answers..Like really?

upvoted 4 times
- MEHULSP

1 week, 5 days ago

To increase the discussion maybe (to show users more traffic on site)

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

A is correct

upvoted 1 times
- [Removed]

4 months ago

A is correct. Go to Data Catalog and search for employee\_ssn in the search box.

upvoted 2 times
- yuvi69

4 months, 1 week ago

i'll go with A

upvoted 1 times
- GCP\_Student1

4 months, 3 weeks ago

A. Go to Data Catalog and search for employee\_ssn in the search box.

upvoted 1 times
- Hi2ALL

4 months, 3 weeks ago

I think A is acceptable since writing code is time consuming effort

You want to minimize effort in performing this task.

What should you do?

A. Go to Data Catalog and search for employee\_ssn in the search box.

upvoted 2 times
- victory108

5 months, 3 weeks ago

A - Go to Data Catalog and search for employee\_ssn in the search box.

upvoted 1 times
- dhanamurthy

5 months, 4 weeks ago

Answer is C. Question is to search for a column in a table. Datacatalog as mentioned in the url <https://cloud.google.com/data-catalog/docs/how-to/search>

can search for only datasets, tables and views and not for columns, hence A cannot be correct. Since with minimum cost we cannot go for dataflow, hence D can also be rejected.

upvoted 1 times

  **Gowtham789** 5 months, 1 week ago

We can search through column as well, so answer is A.

upvoted 1 times

  **guid1984** 6 months, 1 week ago

Its A refer to this from Data Catalog section from gcp docs  
Overview

Data Catalog allows you to discover, manage, and understand data assets across Google Cloud Platform. Data Catalog API natively indexes Cloud BigQuery, Cloud Storage, and Cloud Pub/Sub data assets.

The Data Catalog API can be used to:

Search for data assets across different projects and GCP resources

Create tags and tag templates to add structured business metadata unique to your organization

Automate the tagging process of data assets as part of a data ingestion pipeline

upvoted 2 times

  **desekis** 7 months ago



A is better choice to scan across projects and it is an easy way to find the list of tables having this column

upvoted 1 times

  **PhilipAWS** 7 months, 1 week ago



D is making sense... Don't know why people confusing others here.

upvoted 1 times

  **katayoonnt** 8 months, 1 week ago



A is correct, as per [https://cloud.google.com/data-catalog/docs/how-to/search#how\\_to\\_search\\_for\\_data\\_assets](https://cloud.google.com/data-catalog/docs/how-to/search#how_to_search_for_data_assets)

upvoted 3 times

  **Bhagirathi** 8 months, 1 week ago



finally A or D ?

upvoted 1 times

  **Harish57** 7 months, 3 weeks ago

option D is correct

upvoted 2 times

  **swatititame** 8 months, 2 weeks ago

• A. Go to Data Catalog and search for employee\_ssn in the search box.

upvoted 2 times

  **ayj** 8 months, 2 weeks ago

I choose D!

A would be a ballache across hundreds of projects!

upvoted 1 times

  **ayj** 8 months, 2 weeks ago

I take it back, you can search across multiple projects in your org! So A is nice

upvoted 2 times

You create a Deployment with 2 replicas in a Google Kubernetes Engine cluster that has a single preemptible node pool. After a few minutes, you use kubectl to examine the status of your Pod and observe that one of them is still in Pending status:

```
$ kubectl get pods -l app=myapp
NAME                                READY   STATUS    RESTART   AGE
myapp-deployment-58ddbbb995-lp86m   0/1     Pending   0         9m
myapp-deployment-58ddbbb995-qjpkg   1/1     Running   0         9m
```

What is the most likely cause?

- A. The pending Pod's resource requests are too large to fit on a single node of the cluster.
- B. Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.
- C. The node pool is configured with a service account that does not have permission to pull the container image used by the pending Pod.
- D. The pending Pod was originally scheduled on a node that has been preempted between the creation of the Deployment and your verification of the Pods' status. It is currently being rescheduled on a new node.

Correct Answer: B

- ESP\_SAP

Highly Voted

11 months, 2 weeks ago

Correct Answer is (B):

Reasons for a Pod Status Pending:  
Troubleshooting Reason #1: Not enough CPU  
Troubleshooting Reason #2: Not enough memory  
Troubleshooting Reason #3: Not enough CPU and memory  
<https://managedk8s.com/kubernetes/k8sbot/troubleshooting/pending/pod/2019/02/22/pending-pod.html>  
upvoted 22 times
- Linus11

3 months ago

The real crux of this question is the mention about "Pre-emptible Node pool". That need to take into consider while determining the answer. If we choose B, then the importance of "Pre-emptible node pool" is not there. Whether the node pool is pre-emptible or not, resource scarcity can lead to pending pods.

When we consider the mention of "Pre-emptible Node Poll" , then the answer is obviously D. if a pre-meptible Node get pre-empted there will be a delay in cluster to sync it.

Answer is D.  
upvoted 3 times
- Finger41

1 month, 4 weeks ago

This is to throw you off, when there is insufficient resources for a Pod to stand up, then the status will equal pending :  
<https://kubernetes.io/docs/tasks/debug-application-cluster/debug-application/#:~:text=If%20a%20Pod%20is%20stuck,be%20scheduled%20onto%20a%20node.&text=You%20don't%20have%20enough,new%20nodes%20to%20your%20cluster.>  
upvoted 1 times
- SSPC

11 months, 1 week ago

I agree with you. The correct answer is B  
upvoted 3 times
- [Removed]

10 months, 1 week ago

D gives you the reason why the resource could not be available a it was preempted  
upvoted 3 times
- Finger41

1 month, 4 weeks ago

Its in a deployment, the pod will be recreated. There is insufficient resources in the node, not because its preemptible but because there is no memory/cpu.....  
upvoted 1 times
- tavva\_prudhvi

3 months, 4 weeks ago































If it was preempted, then it has to be restarted right? then it will show its failing not pending, check the articles mentioned by ESP\_SAP  
upvoted 2 times
- cloudeenthu01

Highly Voted

1 year, 1 month ago

D is correct as the node on which pod was scheduled to run was preempted & now this pod is scheduled to run on different preemptible node from the node-pool  
upvoted 12 times



- 

**gcp\_aspirant1** Most Recent 1 month, 1 week ago  
 Will go with B!  
 upvoted 1 times
- 

**Hasaaaaan** 1 month, 3 weeks ago  
 Why not C?  
 Pod can be failed for the following reasons;  
 1. If there is not enough memory  
 2. Kubernetes couldn't pull the images  
 3. Not appropriate permissions.  
 upvoted 1 times
- 

**viswanand** 2 months, 1 week ago  
 B is the correction option I see some people opting for D .... if the node goes down both pods would have been in pending state  
 upvoted 1 times
- 

**Finger41** 2 months, 2 weeks ago  
 Its B, 100%. The pre-emptible thing was added to trick you. See : [https://cloud.google.com/kubernetes-engine/docs/troubleshooting#node\\_allocatable](https://cloud.google.com/kubernetes-engine/docs/troubleshooting#node_allocatable)  
 upvoted 1 times
- 

**mcaromit** 2 months, 2 weeks ago  
 B is correct...A pod is never rescheduled to a node, instead a new pod with a new UID is assigned to the node.  
 upvoted 1 times
- 

**Konnon** 3 months, 3 weeks ago  
 D is incorrect because a rescheduled pod will be a new pod with younger age. I vote for B.  
 upvoted 2 times
- 

**[Removed]** 4 months ago  
 B is correct. Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.  
 upvoted 2 times
- 

**yuvi69** 4 months, 1 week ago  
 correct answer is B , because in the question it is given 'single preemptible node pool'  
 upvoted 1 times
- 

**GCP\_Student1** 4 months, 3 weeks ago  
 B - Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.  
 upvoted 2 times
- 

**EABDAJA** 4 months, 3 weeks ago  
 B is correct.  
 upvoted 1 times
- 

**drizzydroo** 4 months, 3 weeks ago  
 If there are "Too many Pods" then why aren't those pods showing up when you run the command to get pods? It displays one running and one pending and the replica is set for 2. I said D originally but now I am just confused.  
 upvoted 2 times
- 

**TAvenger** 5 months, 1 week ago  
 It could be "D"  
 Please note, it is said that "node pool". And the pool may have many nodes. In theory it may happen that some nodes were preempted.  
  
 It's really suspicious, that it was mentioned that nodes are preemptible, that we waited for a few minutes and decided to check the status. Also we have the link  
[https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#best\\_practices](https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#best_practices)  
 where it is mentioned that "It may take several minutes for GKE to detect that the node was preempted and that the Pods are no longer running, which will delay the rescheduling of the Pods to a new node."  
 I have a feeling that they talk about answer "D"  
 upvoted 2 times
- 

**nitinz** 5 months, 1 week ago  
 read the question, "You create a Deployment with 2 replicas in a Google Kubernetes Engine cluster that has a single preemptible node pool." It is single NODE. if it was preempted then both pods should be gone not one..... Answer is B.  
 upvoted 1 times
- 

**TAvenger** 4 months, 3 weeks ago  
 It is not "pool with as single preemptible node", this is single preemptible NODE POOL. "Node pool" can be the the only pool (single pool) that is present on cluster.  
 In this case GKE can have multiple nodes. And in this current situation one node stayed alive (and another was preempted)  
 upvoted 2 times
- 

**arsav** 5 months, 1 week ago  
 you are on "single preemptible node pool" - so if that preemptible node pool is evicted, both the pod should not running.  
  
 as you see one pod is already running and only one is in pending status,

which clearly tells the node was not preempted.

so the reason for pending could be due to "not enough resources" is the obvious reason.

upvoted 2 times

🗨️ 👤 **TAvenger** 5 months, 1 week ago

Please note, it is said that "node pool". And the pool may have many nodes. In theory it may happen that all nodes except the last one were preempted.

It's really suspicious, that it was mentioned that nodes are preemptible, that we waited for a few minutes and decided to check the status. Also we have the link

[https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#best\\_practices](https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms#best_practices)

where it is mentioned that "It may take several minutes for GKE to detect that the node was preempted and that the Pods are no longer running, which will delay the rescheduling of the Pods to a new node."

I have a feeling that they talk about answer "D"

upvoted 1 times

🗨️ 👤 **arsav** 5 months, 1 week ago

we are attaching the preemptible VM to the "node pool" hence if that VM is evicted it becomes unavailable to the whole "node pool" all the nodes which tries to perform its work load / batch becomes unavailable.

Also the document says

\*\*\*\*If you want to ensure that your jobs or workloads are processed even if no preemptible VMs are available, you can create both non-preemptible and preemptible node pools in your cluster.\*\*\*\*\*

so you need to have both preemptible and non-preemptible VM in the node pool for uninterrupted pods to be servered.

in this case of this question, one pod is already running which clearly confirms preemption didnt happen yet.

the reason for the "pending" could be different due to other resource constraints.

upvoted 2 times

🗨️ 👤 **victory108** 5 months, 3 weeks ago

B - Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.

upvoted 1 times

🗨️ 👤 **DucSiu** 5 months, 4 weeks ago

B. Too many Pods are already running in the cluster, and there are not enough resources left to schedule the pending Pod.

upvoted 1 times

🗨️ 👤 **PhilipAWS** 7 months, 1 week ago

B is right answer if you want to pass exam.

upvoted 4 times

🗨️ 👤 **gcp\_aspirant1** 1 month, 1 week ago

haha yess

upvoted 1 times



You want to find out when users were added to Cloud Spanner Identity Access Management (IAM) roles on your Google Cloud Platform (GCP) project. What should you do in the GCP Console?

- A. Open the Cloud Spanner console to review configurations.
- B. Open the IAM & admin console to review IAM policies for Cloud Spanner roles.
- C. Go to the Stackdriver Monitoring console and review information for Cloud Spanner.
- D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.

Correct Answer: B

- Meix** Highly Voted

1 year, 1 month ago

I think the answer is D

upvoted 28 times
- Anand2608**

8 months, 3 weeks ago

As per the Cloud Audit logs documentation.

upvoted 2 times
- samvegas** Highly Voted

1 year ago

Answer = D, I have simple rule; if metrics then Monitoring, if Auditing then Logging.

upvoted 22 times
- YAS007** Most Recent

3 days, 23 hours ago

D is correct

<https://stackoverflow.com/questions/58936403/way-to-see-when-a-user-was-added-to-a-gcp-project>

upvoted 1 times
- YAS007**

3 days, 23 hours ago

also

<https://cloud.google.com/logging/docs/audit/#admin-activity>

upvoted 1 times
- janman4pl**

1 month ago

why B and how? The answer should be D..

upvoted 1 times
- gcpengineer**

2 days, 11 hours ago

because in iam console we can view any roles, so B is correct

upvoted 1 times
- gcpengineer**

2 days, 11 hours ago

ok, its clear now the question asked 'when', so D

upvoted 1 times
- learnazureportal**

2 months, 2 weeks ago

The key puzzle refers to "when users were added to Cloud Spanner ". therefore, Answer D is correct.

upvoted 2 times
- Agraved**

2 months, 2 weeks ago

D, as the ask is about "When"

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

D is correct

upvoted 1 times
- bhuvana123**

3 months ago

it's D



















upvoted 1 times
- Christiank**

4 months ago

D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.

The correct is 'D' and not 'B'.

upvoted 1 times

-   **[Removed]** 4 months ago  
D is correct. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.  
upvoted 3 times
-   **nikhil1235** 4 months, 2 weeks ago  
answer is B.verified  
upvoted 1 times
-   **pondai** 4 months, 1 week ago  
Check quest again. It says "find out when",so go to IAM console can't help you.  
vote D  
upvoted 1 times
-   **EABDAJA** 4 months, 3 weeks ago  
D is correct  
upvoted 1 times
-   **hicham** 5 months ago  
the Answer is D. don't confuse people  
upvoted 1 times
-   **GCP\_Student1** 5 months ago  
D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.  
upvoted 2 times
-   **DucSiu** 5 months, 4 weeks ago  
D. Go to the Stackdriver Logging console, review admin activity logs, and filter them for Cloud Spanner IAM roles.  
upvoted 2 times
-   **Bhagirathi** 7 months, 1 week ago  
it is D.  
upvoted 2 times
-   **Bhagirathi** 7 months, 3 weeks ago  
its D - pls do not confuse with some other choices.  
upvoted 3 times

## Question #96

## Topic 1

Your company implemented BigQuery as an enterprise data warehouse. Users from multiple business units run queries on this data warehouse. However, you notice that query costs for BigQuery are very high, and you need to control costs. Which two methods should you use? (Choose two.)

- A. Split the users from business units to multiple projects.
- B. Apply a user- or project-level custom query quota for BigQuery data warehouse.
- C. Create separate copies of your BigQuery data warehouse for each business unit.
- D. Split your BigQuery data warehouse into multiple data warehouses for each business unit.
- E. Change your BigQuery query model from on-demand to flat rate. Apply the appropriate number of slots to each Project.

**Correct Answer:** *BE*

CPBach Highly Voted 1 year, 1 month ago

I'd say B and E. So either you do B or E to reduce costs.  
upvoted 21 times

Jignesh\_Gamdha Highly Voted 8 months, 4 weeks ago

B & E  
Refer below link - first of all you can define quotas on project or user level and 2nd one is you can change from on demand to flat rate model and define the parameters based on your requirement ---

<https://cloud.google.com/bigquery/docs/custom-quotas>  
[https://cloud.google.com/bigquery/pricing#flat\\_rate\\_pricing](https://cloud.google.com/bigquery/pricing#flat_rate_pricing)  
upvoted 12 times

kyo Most Recent 6 days, 4 hours ago

Is there a User who is thinking of answering A and B like me? B is currently unable to function on its own. So why not choose with A? E is certainly a cost-cutting measure, but I think there is insufficient evidence that this option will cost less than on-demand in this pattern.  
upvoted 1 times

BalSan 3 weeks, 2 days ago

Those who suggest option D, while splitting into multiple DW might reduce query costs, we don't know how many job(s) are populating the BigQuery DW and the impact the split will have on all those source jobs... Hence D is not the best option.  
upvoted 1 times

mcaromit 2 months, 2 weeks ago

B & E are correct  
upvoted 1 times

yuvarajmrec 4 months ago

D and E are correct  
upvoted 3 times

Christiank 4 months ago

B. Open the IAM & admin console to review IAM policies for Cloud Spanner roles.  
upvoted 1 times

[Removed] 4 months ago

B, E are correct.  
upvoted 2 times

Fidget\_ 4 months, 1 week ago

Surely splitting the data over projects (D) is pointless as you pay per byte anyway?  
upvoted 1 times

ShakthiGCP 4 months, 4 weeks ago

So we are trying to control cost here. So Do 'B' - <https://cloud.google.com/bigquery/docs/custom-quotas> and 'E' - [https://cloud.google.com/bigquery/pricing#flat\\_rate\\_pricing](https://cloud.google.com/bigquery/pricing#flat_rate_pricing) . D is something out of the box that you want to try when the solution is already thr.  
Go with - B,E  
upvoted 2 times



JackGlemins 5 months, 1 week ago

E: <https://cloud.google.com/blog/products/data-analytics/how-to-manage-bigquery-flat-rate-slots-within-a-project>  
B  
upvoted 1 times

nitinz 5 months, 1 week ago

E for sure, B does not help to reduce cost jut control cost. Big data query are based on reads, if you have a smaller dataset then less price per query. You can split the dataset so that quires are cheaper to run, hence D&E.

upvoted 1 times

  **Linus11** 3 months ago

The question does not state about REDUCING Cost. But about CONTROLLING cost.

B, E correct.

upvoted 2 times

  **gcp\_aspirant1** 1 month, 1 week ago

exactly, people pls think before you write or post it

upvoted 1 times



  **lutoa** 5 months, 2 weeks ago

I think it is D & E.

Lots of people saying B but the adding a quota doesn't stop the requirement for their queries or keep cost down.

Bigquery costs go up the more data that is read so splitting the data out should help, and switching to flat costing should also help.

upvoted 3 times

  **Linus11** 3 months ago

The question does not state about REDUCING Cost. But about CONTROLLING cost.

B, E correct.

upvoted 2 times

  **victory108** 5 months, 3 weeks ago

B - Apply a user- or project-level custom query quota for BigQuery data warehouse.

E - Change your BigQuery query model from on-demand to flat rate. Apply the appropriate number of slots to each Project.

upvoted 1 times

  **rvgcp** 6 months ago

IMO , its D & E ..in line with some of the comments below

upvoted 1 times

  **tavva\_prudhvi** 3 months, 1 week ago

Any article mentioning about D, i haven't seen any google cloud document saying that.

upvoted 1 times

  **Yashv** 7 months ago

DE is right Bigquery is charged based on the read operation

upvoted 2 times

  **Yashv** 7 months ago

DE is right

upvoted 3 times

## Question #97

## Topic 1

You are building a product on top of Google Kubernetes Engine (GKE). You have a single GKE cluster. For each of your customers, a Pod is running in that cluster, and your customers can run arbitrary code inside their Pod. You want to maximize the isolation between your customers' Pods. What should you do?

- A. Use Binary Authorization and whitelist only the container images used by your customers' Pods.
- B. Use the Container Analysis API to detect vulnerabilities in the containers used by your customers' Pods.
- C. Create a GKE node pool with a sandbox type configured to gvisor. Add the parameter `runtimeClassName: gvisor` to the specification of your customers' Pods.
- D. Use the `cos_containerd` image for your GKE nodes. Add a `nodeSelector` with the value `cloud.google.com/gke-os-distribution: cos_containerd` to the specification of your customers' Pods.

**Correct Answer: C**

Reference:

<https://cloud.google.com/kubernetes-engine/sandbox/>

  **[Removed]**  4 months ago

C is correct. Create a GKE node pool with a sandbox type configured to gvisor. Add the parameter `runtimeClassName: gvisor` to the specification of your customers' Pods.

upvoted 12 times

  **lio123**  4 months, 2 weeks ago

I think it is C

GKE Sandbox provides an extra layer of security to prevent untrusted code from affecting the host kernel on your cluster nodes.

GKE Sandbox is an especially good fit for the following types of applications. See Limitations for more information to help you decide which applications to sandbox.

Untrusted or third-party applications using runtimes such as Rust, Java, Python, PHP, Node.js, or Golang  
 Web server front-ends, caches, or proxies  
 Applications processing external media or data using CPUs  
 Machine-learning workloads using CPUs  
 CPU-intensive or memory-intensive applications

<https://cloud.google.com/kubernetes-engine/docs/concepts/sandbox-pods>

upvoted 8 times

  **YAS007**  3 days, 23 hours ago

I'm not sur about C, because, from google : "GKE Sandbox provides an extra layer of security to prevent untrusted code from affecting the host kernel on your cluster nodes." thus it help protecting cluster's kernel, not allow isolation between your customers' Pods

upvoted 1 times

  **djgodzilla** 1 month, 1 week ago

This is so screwed up questions . I've just spent an hour understanding the relationship between `cos_contained` and `Gvisor` .  
 Actually C implies D . You can't run a gvisor sandbox without using a `cos_containerd` node image . So why are they making life so hard for a silly Associate level exam , ludacris.

<https://youtu.be/uDOu6rK4yOk?t=1179>

<https://cloud.google.com/kubernetes-engine/docs/how-to/sandbox-pods>

<https://youtu.be/kxUZ4lVFuVo?t=1755>

```
gcloud container node-pools create node-pool-name \
--cluster=cluster-name \
--node-version=node-version \
--machine-type=machine-type \
--image-type=cos_containerd \
--sandbox type=gvisor \
upvoted 3 times
```

  **mcaromit** 2 months, 2 weeks ago

C is correct

upvoted 1 times



  **guruguru** 4 months ago



D, as it isolate from host level too. Sandbox isolate the container and kernel process. Not between pods.

upvoted 1 times

  **GCP\_Student1** 4 months, 1 week ago

C. Create a GKE node pool with a sandbox type configured to gvisor. Add the parameter runtimeClassName: gvisor to the specification of your customers' Pods.  
upvoted 4 times

  **raman23125** 4 months, 2 weeks ago  
correct: c  
  
<https://hub.packtpub.com/gke-sandbox-a-gvisor-based-feature-to-increase-security-and-isolation-in-containers/>  
upvoted 3 times

  **cloud\_\_guru** 4 months, 2 weeks ago  
On Container isolation, GCP usesGKE Sandbox for a second layer of defense between containerized workloads on GKE for enhanced workload security.  
upvoted 2 times

## Question #98

## Topic 1

Your customer has implemented a solution that uses Cloud Spanner and notices some read latency-related performance issues on one table. This table is accessed only by their users using a primary key. The table schema is shown below.

```
CREATE TABLE Persons (  
  person_id INT64 NOT NULL,    // sequential number based on number of registration  
  account_creation_date DATE,  // system date  
  birthdate DATE,             // customer birthdate  
  firstname STRING (255),      // first name  
  lastname STRING (255),       // last name  
  profile_picture BYTES (255)   // profile picture  
) PRIMARY KEY (person_id)
```

You want to resolve the issue. What should you do?

- A. Remove the profile\_picture field from the table.
- B. Add a secondary index on the person\_id column.
- C. Change the primary key to not have monotonically increasing values.

- D. Create a secondary index using the following Data Definition Language (DDL):

```
CREATE INDEX person_id_ix  
ON Persons (  
  person_id,  
  firstname,  
  lastname  
) STORING (  
  profile_picture  
)
```

Correct Answer: D

 **ri\_unhou119** Highly Voted 2 months, 3 weeks ago

C is the correct answer.

As mentioned in Schema and data model, you should be careful when choosing a primary key to not accidentally create hotspots in your database. One cause of hotspots is having a column whose value monotonically increases as the first key part, because this results in all inserts. This pattern is desirable because Cloud Spanner divides data among servers by key ranges, which means all your inserts will be directed at a single server that will end up doing all the work.

Reference URL:

<https://cloud.google.com/spanner/docs/schema-design>

<https://medium.com/google-cloud/cloud-spanner-choosing-the-right-primary-keys-cd2a47c7b52d>

upvoted 10 times

 **jackdbd** Most Recent 4 days, 16 hours ago

D is the correct answer.

Avoiding monotonically increasing values is a good idea to avoid index hotspotting (answer C), but the question mentions read latency-related performance issues, while hotspotting is a problem for write operations.

My rationale for choosing option D stems from this passage in the docs: "if you need to quickly look up a set of SingerId values for a given range of LastName values, you should create a secondary index on LastName, so Cloud Spanner does not need to scan the entire table." To me, this is clearly a solution to read latency-related performance issues.

<https://cloud.google.com/spanner/docs/secondary-indexes>

upvoted 2 times

 **ayush\_1995** 2 weeks, 2 days ago

My guess is C

stated in the documentation: <https://cloud.google.com/spanner/docs/secondary-indexes>

Cloud Spanner automatically creates an index for each table's primary key. For example, you don't need to do anything to index the primary key of Singers, because it's automatically indexed for you.

You can also create secondary indexes for other columns. Adding a secondary index on a column makes it more efficient to look up data in that column. For example, if you need to quickly look up a set of SingerId values for a given range of LastName values, you should create a secondary index on LastName, so Cloud Spanner does not need to scan the entire table.

upvoted 1 times

 **JH86** 1 month, 2 weeks ago

the answer should be D. C speaks of monotonically increasing primary key given you have a timestamp that causes a hotspot in the database as part of the identifier. <https://cloud.google.com/spanner/docs/schema-design>

however in order to increase performance, it is recommended to have a secondary index to make it more efficient to look up in the database, for instance the below excerpts that has been extracted from google doc:

You can also create secondary indexes for other columns. Adding a secondary index on a column makes it more efficient to look up data in that column. For example, if you need to quickly look up a set of SingerId values for a given range of LastName values, you should create a secondary index on LastName, so Cloud Spanner does not need to scan the entire table.

Although my take is the DDL with the store (profile pic) looks a bit odd, i will still go with secondary indexing.

<https://cloud.google.com/spanner/docs/secondary-indexes>  
upvoted 2 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

C is correct....use a non-sequence generator as an index column  
upvoted 2 times

🗲️ 👤 **lxgywil** 2 months, 3 weeks ago

The latency issue must be caused by hotspotting here as the values of `person\_id` increase sequentially (monotonically) and are used as a primary key.

The answer must be "C" as Cloud Spanner documentation states:

"You should be careful when choosing a primary key to not accidentally create hotspots in your database. One cause of hotspots is having a column whose value monotonically increases as the first key part, because this results in all inserts occurring at the end of your key space. "

Ref: <https://cloud.google.com/spanner/docs/schema-design>  
upvoted 2 times

🗲️ 👤 **bhuvana123** 3 months ago

IT'S D?  
upvoted 1 times

🗲️ 👤 **ApaMokus** 3 months ago

To me, none of the options make sense.  
A. is modifying table structure. It helps only if profile picture is not used regularly.  
B. Secondary index won't help if the table is queried by the PK.  
C. Does not speed up lookup  
D. Secondary index won't help  
upvoted 1 times

🗲️ 👤 **Amu89** 2 months, 3 weeks ago

C might be the Ans  
upvoted 1 times



Your finance team wants to view the billing report for your projects. You want to make sure that the finance team does not get additional permissions to the project. What should you do?

- A. Add the group for the finance team to roles/billing user role.
- B. Add the group for the finance team to roles/billing admin role.
- C. Add the group for the finance team to roles/billing viewer role.
- D. Add the group for the finance team to roles/billing project/Manager role.

Correct Answer: C

- miniso8153

Highly Voted

1 year, 1 month ago

c

"Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account."

<https://cloud.google.com/billing/docs/how-to/billing-access>

upvoted 34 times
- dan80

Highly Voted

1 year, 1 month ago

Answer is C - Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account.

upvoted 16 times
- krsourav

5 months, 3 weeks ago

Hey, look at this .....)

upvoted 2 times
- Katheri7689

Most Recent

2 months, 2 weeks ago

C is the correct answer

upvoted 1 times
- mcaromit

2 months, 2 weeks ago

C is correct

upvoted 1 times
- [Removed]

4 months ago

C is correct. Add the group for the finance team to roles/billing viewer role.

upvoted 1 times
- cloud\_\_guru

4 months, 2 weeks ago

Obvious choice here is C-- viewer to make sure they dont have additional access.  
Source: this is my job

upvoted 1 times
- Vic1043

4 months, 3 weeks ago

ANS - C

upvoted 1 times
- GCP\_Student1

5 months ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times
- swatititame

8 months, 2 weeks ago

• C. Add the group for the finance team to roles/billing viewer role.

upvoted 2 times
- glam

9 months, 2 weeks ago

C. Add the group for the finance team to roles/billing viewer role.

upvoted 4 times
- GopinathM

10 months, 1 week ago

C is correct

upvoted 2 times
- [Removed]

10 months, 3 weeks ago

C is correct.  
B also qualifies to view report but gives additional privileges

[https://cloud.google.com/billing/docs/how-to/reports#required\\_permissions\\_to\\_access\\_reports](https://cloud.google.com/billing/docs/how-to/reports#required_permissions_to_access_reports)

upvoted 2 times

  **ESP\_SAP** 11 months, 2 weeks ago



Correct Answer is (C):

<https://cloud.google.com/billing/docs/how-to/billing-access>

Billing Account Viewer



(roles/billing.viewer) View billing account cost information and transactions. Organization or billing account. Billing Account Viewer access would usually be granted to finance teams, it provides access to spend information, but does not confer the right to link or unlink projects or otherwise manage the properties of the billing account.

upvoted 11 times

  **Kaazim** 1 month, 4 weeks ago

Please tell me your answer for question 98

upvoted 2 times

  **qafro** 11 months, 3 weeks ago

C should be correct

upvoted 2 times

  **poogcp** 1 year, 1 month ago



Its C . only view billing report

upvoted 3 times

  **mlantonis** 1 year, 1 month ago

Yes it is C

upvoted 4 times

  **SIX** 1 year, 1 month ago

The answer is D.

upvoted 1 times

  **stepkurniawan** 11 months, 1 week ago

The correct answer is c

upvoted 4 times

  **afr88** 1 year ago

No answer is C. with project/Manager (answer D) they could still link/unlink account. Moreover: "Billing Account Viewer access would usually be granted to finance teams"

upvoted 3 times

## Question #100

## Topic 1

Your organization has strict requirements to control access to Google Cloud projects. You need to enable your Site Reliability Engineers (SREs) to approve requests from the Google Cloud support team when an SRE opens a support case. You want to follow Google-recommended practices. What should you do?

- A. Add your SREs to roles/iam.roleAdmin role.
- B. Add your SREs to roles/accessapproval.approver role.
- C. Add your SREs to a group and then add this group to roles/iam.roleAdmin.role.
- D. Add your SREs to a group and then add this group to roles/accessapproval.approver role.

**Correct Answer: B**

  **reinocd21** Highly Voted 1 year, 1 month ago



D. Add your SREs to a group and then add this group to roles/accessapproval approver role.  
-Google recommendation.  
upvoted 31 times

  **sidharthwader** Most Recent 3 weeks, 3 days ago

B AND D are right but i think GCP recommends to use groups all the time  
upvoted 3 times

  **arsh1916** 2 months, 1 week ago

D is correct  
upvoted 3 times

  **mcaromit** 2 months, 2 weeks ago

D is correct  
upvoted 1 times

  **Finger41** 2 months, 2 weeks ago



Answer is D  
upvoted 1 times

  **[Removed]** 4 months ago

D is correct. Add your SREs to a group and then add this group to roles/accessapproval approver role.  
upvoted 1 times

  **Pravin3c** 4 months, 1 week ago


D - Grant whoever will be performing approvals for the project (either a service account or a human user) the IAM role Access Approval Approver on the project folder, or organization that you would like the person to have the role for.  
upvoted 1 times

  **cloud\_\_guru** 4 months, 2 weeks ago

Definitely D. Google recommended practices lead us to use groups rather than granting each individual SRE permissions  
upvoted 1 times

  **GCP\_Student1** 5 months ago


D. Add your SREs to a group and then add this group to roles/accessapproval approver role.  
upvoted 2 times

  **victory108** 6 months ago

D - Add your SREs to a group and then add this group to roles/accessapproval approver role.  
upvoted 3 times

  **Bhagirathi** 7 months, 3 weeks ago

D sounds is the option to choose .  
upvoted 1 times

  **swatititame** 8 months, 2 weeks ago

• D. Add your SREs to a group and then add this group to roles/accessapproval approver role.  
upvoted 2 times

  **nwk** 9 months ago

D - <https://cloud.google.com/iam/docs/understanding-roles#access-approval-roles>  
upvoted 1 times

  **glam** 9 months, 2 weeks ago

D. Add your SREs to a group and then add this group to roles/accessapproval approver role.

upvoted 4 times

🗨️ 👤 **JJ\_ME** 9 months, 3 weeks ago

B is correct.

Grant the IAM role Access Approvals Approver on the project, folder, or organization to the principal (either a service account or a human) who will perform approvals.

Therefore, this roles is NOT granted to a group.

[https://cloud.google.com/access-approval/docs/approve-requests#configuring\\_access\\_approvers\\_in\\_your\\_organization](https://cloud.google.com/access-approval/docs/approve-requests#configuring_access_approvers_in_your_organization)

upvoted 2 times

🗨️ 👤 **ARDY** 5 months ago

does that mean that we cannot have Humans in a Group?

upvoted 3 times

🗨️ 👤 **protn1221** 9 months, 3 weeks ago

"Therefore, this roles is NOT granted to a group."

where you found this information !??

upvoted 3 times

🗨️ 👤 **Rimjith** 4 months, 3 weeks ago

<https://cloud.google.com/access-approval/docs/quickstart#email>

upvoted 1 times

🗨️ 👤 **GopinathM** 10 months, 1 week ago

D is recommended

upvoted 1 times

🗨️ 👤 **stepkurniawan** 11 months, 1 week ago

using group is always recommended -> D

upvoted 1 times

You need to host an application on a Compute Engine instance in a project shared with other teams. You want to prevent the other teams from accidentally causing downtime on that application. Which feature should you use?

- A. Use a Shielded VM.
- B. Use a Preemptible VM.
- C. Use a sole-tenant node.
- D. Enable deletion protection on the instance.

**Correct Answer:** D

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (D):

Preventing Accidental VM Deletion  
This document describes how to protect specific VM instances from deletion by setting the deletionProtection property on an Instance resource. To learn more about VM instances, read the Instances documentation.



As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.



By setting the deletionProtection flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the deletionProtection flag, the request fails. Only a user that has been granted a role with compute.instances.create permission can reset the flag to allow the resource to be deleted.

<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>  
upvoted 18 times



  **professor** Highly Voted 1 year, 1 month ago  
Agree with D



You can enabale Termination protection  
upvoted 11 times



  **sunilw** Most Recent 1 month, 1 week ago  
D: Correct  
upvoted 1 times



  **mcaromit** 2 months, 2 weeks ago  
D is correct  
upvoted 1 times

  **Finger41** 2 months, 2 weeks ago  
Answer is D, prevents accidental termination.  
upvoted 1 times

  **EJJ** 4 months ago  
the answer is D.  
  
ref: <https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>  
upvoted 1 times



  **[Removed]** 4 months ago  
D is correct. Enable deletion protection on the instance.  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago  
D is correct.  
upvoted 1 times

  **GCP\_Student1** 5 months ago  
D. Enable deletion protection on the instance.  
  
"Accidently" is keyword for the D to be the correct answer.  
upvoted 3 times



  **Tez1** 5 months, 1 week ago  
I don't think any answer is completely correct, but D seems the closest.

upvoted 1 times

  **EABDAJA** 5 months, 1 week ago

D is correct

upvoted 1 times

  **Sentinel** 5 months, 3 weeks ago

Should be A.

It lets you add additional protections and safeguards against downtime like encrypted disks and access policies:

<https://cloud.google.com/compute/docs/nodes/sole-tenant-nodes>

B - will definitely go down

C - It gives u additional control over hardware nothing else.

D - Users might not be able to delete VM, but they can still shut/restart VM that will cause downtime

upvoted 4 times

  **j1c4b** 4 months ago

Shielded VM is Windows Server 2016 Hyper-V feature and can't be used for non-windows based VM's. That's why - D

upvoted 2 times

  **bosslady** 6 months ago

a sole tenant node is not to be shared in the first place its for the sole tenant i vote D

upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

200 % D.

upvoted 3 times

  **Bhagirathi** 7 months, 3 weeks ago

D is the option to choose.

upvoted 2 times

  **swatititame** 8 months, 1 week ago

- D. Enable deletion protection on the instance.

upvoted 3 times

  **nightflyer** 7 months, 2 weeks ago

But downtime doesn't necessarily mean deletion

upvoted 1 times

  **magistrum** 7 months ago

True, but deletion 100% becomes downtime

upvoted 3 times

  **nmohapatra87** 8 months, 4 weeks ago

Correct Answer is (D)

upvoted 2 times

## Question #102

## Topic 1

Your organization needs to grant users access to query datasets in BigQuery but prevent them from accidentally deleting the datasets. You want a solution that follows Google-recommended practices. What should you do?

- A. Add users to roles/bigquery user role only, instead of roles/bigquery dataOwner.
- B. Add users to roles/bigquery dataEditor role only, instead of roles/bigquery dataOwner.
- C. Create a custom role by removing delete permissions, and add users to that role only.
- D. Create a custom role by removing delete permissions. Add users to the group, and then add the group to the custom role.

**Correct Answer: B**

  **someoneinthecloud** Highly Voted 1 year ago

I believe the key part is the "following Google Best Practices" phrase.



A - Works, but doesn't follow GCP best practices

B - Doesn't work as the role grants permission to delete datasets

C - Works, but is more complicated than A and doesn't follow Google best practices

D - Correct, more complicated than A, but it follows Google Best Practices.

upvoted 34 times

  **kyo** 5 days, 21 hours ago

I don't think A works properly.

roles/bigquery.user has bigquery.datasets.create. And the documentation states:

> Additional, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

If bigquery.user creates a new dataset, it's likely that bigquery.user will get permission to delete that dataset. This means that bigquery.user may have permission to delete data.

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 1 times

  **YuvarajK** 1 month, 2 weeks ago

I think A is the Answer and it follow GCP best practices.

<https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>

We do have the role - BigQuery User which does the below permissions

When applied to a project, this role also provides the ability to run jobs, including queries, within the project.

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

upvoted 2 times

  **TAvenger** 4 months, 3 weeks ago

Read description carefully "prevent from accidentally deleting the datasets". Not tables, datasets! option B does not allow to delete datasets either.

Check dataset permissions in the roles/bigquery.dataEditor:

bigquery.datasets.create

bigquery.datasets.get

bigquery.datasets.getIamPolicy

bigquery.datasets.updateTag

You CANNOT delete dataset with option "B"

upvoted 3 times

  **ESP\_SAP** Highly Voted 11 months, 2 weeks ago

Correct Answer is (D):

The proper answer regarding to bigquery roles is the listed in the options, the proper role that resolve this requirement is: roles/bigquery.dataViewer  
[https://cloud.google.com/bigquery/docs/access-control#custom\\_roles](https://cloud.google.com/bigquery/docs/access-control#custom_roles)

on the other hand, the question explicitly is asking to use the GCP best practices on IAM :

GCP Best Practices explain clearly these rules:

Policy management

☐ Set organization-level IAM policies to grant access to all projects in your organization.

☐ Grant roles to a Google group instead of individual users when possible. It is easier to add members to and remove members from a Google group instead of updating an IAM policy to add or remove users.

☐ If you need to grant multiple roles to allow a particular task, create a Google group, grant the roles to that group, and then add users to that group.

[https://cloud.google.com/iam/docs/using-iam-securely#policy\\_management](https://cloud.google.com/iam/docs/using-iam-securely#policy_management)

upvoted 22 times

  **JackGlemins** 5 months, 1 week ago

Other best practice is use predefined roles over custom roles. Maybe A is correct




upvoted 2 times

  **JackGlemins** 5 months, 1 week ago

I correct myself: <https://cloud.google.com/iam/docs/understanding-custom-roles>

Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.



upvoted 1 times

  **prashuG** 10 months, 4 weeks ago

Answer is A: roles/bigquery.user is a BigQuery User role which when applied to a project provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project.

Ref: <https://cloud.google.com/iam/docs/understanding-roles#bigquery-roles>

upvoted 3 times

  **j1c4b** 3 months, 4 weeks ago

you can create data set with bigquery.user role because it has bigquery.datasets.create permissions. And if a user has bigquery.datasets.create permissions, when that user creates a dataset, they are granted bigquery.dataOwner access to it. So A is NOT a choice

upvoted 1 times

  **Nelson2080** Most Recent 2 weeks, 4 days ago

I agree with D being the correct answer.

The predefined role for answer A grants the user to delete any dataset they create, as they are assigned the owner. That alone seems to remove A from consideration.

B. Would work as well and may be fine if it's a single user, however, it makes sense to create the custom role in case you need to add additional users to the role.

from <https://cloud.google.com/bigquery/docs/access-control>

Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.

upvoted 1 times


  **gcp\_aspirant1** 1 month, 1 week ago

Dear people, please relax and think!

When you are creating a custom role, why will you even add more permissions and then delete them!? Does that even make any sense??

I'm going with A

upvoted 2 times

  **JH86** 1 month, 2 weeks ago

ill go with A, use predefined roles over custom roles:

Creation

Before you decide to create a custom role, check whether the service has a predefined role—or a combination of multiple predefined roles—that meets your needs. For a complete list of predefined roles, as well as the permissions that are included in each predefined role, see the predefined roles reference.

If there is no predefined role—or combination of predefined roles—that meets your needs, you can create a custom role that includes only the permissions you need to grant. For details, see Creating and managing custom roles.

<https://cloud.google.com/iam/docs/understanding-custom-roles>

upvoted 1 times

  **AD\_0525** 2 months ago

Answer should be D, A can be an option but I have not found any such rule as roles/bigquery user in IAM roles.

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

D is correct...BigQuery User implicitly grants DataOwner role as well

upvoted 1 times

  **Finger41** 2 months, 2 weeks ago

As google has highlighted, in <https://cloud.google.com/bigquery/docs/access-control>, both A and B are likely to allow users to delete datasets. D provides the ability to remove permissions of deletion all together..... I vote D.

upvoted 1 times

  **ri\_unhou119** 2 months, 3 weeks ago

Personally think that the correct answer is D.

the reason:

Option A [BigQuery User] has such a description in the description text.

-----  
When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.  
-----

Therefore, the [BigQuery User] role has the ability to delete datasets.

upvoted 2 times



- 🗨️ 👤 **Rightsaidfred** 3 months ago  
Its A. dataEditor role Can delete both Tables and Datasets:  
<https://cloud.google.com/bigquery/docs/access-control#bigquery>  
upvoted 1 times
- 🗨️ 👤 **kopper2019** 3 months, 2 weeks ago  
Definitely A, I know people say Google recommend secrets but if you have a rule that allows read/query only why would you create a new one? So have to maintain it or people get confuse and assign role to both groups the custom one and the google, A  
upvoted 2 times
- 🗨️ 👤 **relaxm** 4 months ago  
<https://cloud.google.com/datastore/docs/tools/datastore-emulator>  
D  
upvoted 2 times
- 🗨️ 👤 **[Removed]** 4 months ago  
A is correct. Add users to roles/bigquery user role only, instead of roles/bigquery dataOwner.  
upvoted 3 times
- 🗨️ 👤 **yuvi69** 4 months, 1 week ago  
answer is D. see the question carefully, it says google recommended practise which is adding users to groups and assigning permission to that group  
upvoted 2 times
- 🗨️ 👤 **drizzydroo** 4 months, 2 weeks ago  
actually just tested this and it looks like User can delete datasets.... I feel like custom roles are never recommended but for this tricky one I would pick D  
upvoted 1 times
- 🗨️ 👤 **ackolandjian** 2 months, 4 weeks ago  
No, because User : "allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets."  
So if the User creates new datasets and tables, they can delete them.  
upvoted 2 times
- 🗨️ 👤 **djgodzilla** 1 month, 3 weeks ago  
you're correct , but the question is tricky . what did they mean by "accidentally deleting the datasets" . the new datasets created by these users , or existing datasets ?  
upvoted 1 times
- 🗨️ 👤 **bart54** 4 months, 2 weeks ago  
Based on least priv i think the answer is B as per below  
  
when the user creates a dataset, they are granted bigquery.dataOwner access to it. bigquery.dataOwner access gives the user the ability to delete datasets and tables they create.  
upvoted 2 times
- 🗨️ 👤 **tarang3** 4 months, 2 weeks ago  
Answer is A  
proof: <https://cloud.google.com/bigquery/docs/access-control>  
upvoted 2 times

## Question #103

## Topic 1

You have a developer laptop with the Cloud SDK installed on Ubuntu. The Cloud SDK was installed from the Google Cloud Ubuntu package repository. You want to test your application locally on your laptop with Cloud Datastore. What should you do?

- A. Export Cloud Datastore data using `gcloud datastore export`.
- B. Create a Cloud Datastore index using `gcloud datastore indexes create`.
- C. Install the `google-cloud-sdk-datastore-emulator` component using the `apt get install` command.
- D. Install the `cloud-datastore-emulator` component using the `gcloud components install` command.

Correct Answer: D

🗨️  **someoneinthecloud** Highly Voted 1 year ago

I believe answer is C  
<https://cloud.google.com/sdk/docs/downloads-apt-get>

The question is not about the datastore command itself but from where we should run the update command on the Ubuntu to install the component.  
upvoted 14 times

🗨️  **XRiddlerX** 12 months ago

I agree with this comment. The answer is C.  
If you installed the SDK from the Ubuntu repo and try to do the following:  
`$ gcloud components install cloud-datastore-emulator`

You will receive this message:  
ERROR: (gcloud.components.install)  
You cannot perform this action because the Cloud SDK component manager is disabled for this installation. You can run the following command to achieve the same result for this installation:

`sudo apt-get install google-cloud-sdk-datastore-emulator`  
upvoted 19 times

🗨️  **Ale1973** 10 months, 3 weeks ago

WOW!!! Today I have learned a new and interesting thing thanks to you...  
upvoted 2 times

🗨️  **stepkurniawan** 11 months, 1 week ago

it says that in your Ubuntu, you have Cloud SDK installed already. So it should be able to run the command in D  
upvoted 2 times

🗨️  **Ale1973** 10 months, 3 weeks ago

Yes, but it says that "The Cloud SDK was installed from the Google Cloud Ubuntu package repository", then to install datastore emulator you should use the command in Option C.  
upvoted 6 times

🗨️  **professor** Highly Voted 1 year, 1 month ago

Ans is D

<https://cloud.google.com/datastore/docs/tools/datastore-emulator>  
upvoted 13 times

🗨️  **Eshkrkrkr** 8 months, 3 weeks ago

Wrong! The answer is C! When you install SDK using apt Cloud SDK Component Manager is disabled and you need to install extra packages again using apt.  
[https://cloud.google.com/sdk/docs/components#managing\\_cloud\\_sdk\\_components](https://cloud.google.com/sdk/docs/components#managing_cloud_sdk_components)  
Note: These instructions will not work if you have installed Cloud SDK using a package manager such as APT or yum because Cloud SDK Component Manager is disabled when using that method of installation.  
upvoted 13 times

🗨️  **SWObaby** 7 months, 3 weeks ago

I believe the answer is C...

It is a tricky question!! The question states, "The Cloud SDK was installed from the Google Cloud Ubuntu package repository." For those, who aren't that familiar with Debian/Ubuntu, D seems like an attractive answer. It works as a way to install Datastore...but it does NOT fit the context of the question.

I recommend looking back to G Cloud SDK installation (Debian/Ubuntu): <https://cloud.google.com/sdk/docs/install#deb>

Read the "Installation Steps" in the documentation. In Step 3, "`sudo apt-get update && sudo apt-get install google-cloud-sdk`". Then, Step 4

is additionally adding other components, such as "sudo apt-get install google-cloud-sdk-datastore-emulator".

Proving C the correct answer.

upvoted 16 times

  **ShakthiGCP** 4 months, 4 weeks ago

Go With 'C' ... just tried creating a ubuntu server and verified these.. Dont worry about any other options.

<https://cloud.google.com/sdk/docs/quickstart#deb> check this link .

upvoted 3 times

  **MONICA291012** Most Recent 1 week, 4 days ago

D

Installing the emulator

The Datastore emulator is a component of the Google Cloud SDK's gcloud tool. Use the gcloud components install command to install the Datastore emulator

gcloud components install cloud-datastore-emulator



upvoted 1 times

  **associatecloudexamuser** 2 weeks, 4 days ago

Answer is D.

gcloud components install cloud-datastore-emulator

upvoted 2 times

  **JH86** 1 month, 2 weeks ago



my final asnwer is C, due to the following:

[https://cloud.google.com/sdk/docs/components#additional\\_components](https://cloud.google.com/sdk/docs/components#additional_components)

For example, the google-cloud-sdk-app-engine-java component can be installed as follows:

sudo apt-get install google-cloud-sdk-app-engine-java

upvoted 1 times

  **JH86** 1 month, 2 weeks ago

confused with C or D. google docs has recommended D.

[https://cloud.google.com/sdk/docs/components#additional\\_components](https://cloud.google.com/sdk/docs/components#additional_components)

Installing components

To install a component at the current version of your Cloud SDK installation, run gcloud components install:

gcloud components install [COMPONENT-ID] where component-id = cloud-datastore-emulator

upvoted 1 times

  **Caribou** 2 months ago

B

B. Create a Cloud Datastore index using gcloud datastore indexes create.

to use datastore, you need configure index first: <https://cloud.google.com/datastore/docs/tools/indexconfig>

upvoted 1 times

  **Bha9886** 2 months ago

Installing the emulator

The Datastore emulator is a component of the Google Cloud SDK's gcloud tool. Use the gcloud components install command to install the Datastore emulator:

gcloud components install cloud-datastore-emulator

upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

C is correct

upvoted 1 times

  **Finger41** 2 months, 2 weeks ago

I belive the answer is C:

I think there is a tricck to this question, as per google "If you installed the Cloud SDK through the apt or yum package managers, you can also install additional Cloud SDK components using those same package managers."

[https://cloud.google.com/sdk/docs/components#external\\_package\\_managers](https://cloud.google.com/sdk/docs/components#external_package_managers)

upvoted 1 times

  **ViMan** 2 months, 2 weeks ago

D.

[https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing\\_the\\_emulator](https://cloud.google.com/datastore/docs/tools/datastore-emulator#installing_the_emulator)

upvoted 1 times

  **jahnu** 3 months ago

Ans C: Update the package list and install the Cloud SDK

sudo apt-get update && sudo apt-get install google-cloud-sdk

upvoted 1 times

🗨️ 👤 **jahnu** 3 months ago

Ans C : Updat e the package list and install the Cloud SDK  
sudo apt-get update && sudo apt-get install google-cloud-sdk  
upvoted 1 times

🗨️ 👤 **Biju1** 3 months, 1 week ago

Answer C  
Last login: Sun Apr 18 14:48:59 2021 from 74.125.45.70  
bijumathewt@instance-1:~\$ sudo apt-get install google-cloud-sdk-datastore-emulator  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
E: Unable to locate package google-cloud-sdk-datastore-emulator  
bijumathewt@instance-1:~\$  
upvoted 2 times

🗨️ 👤 **cbarg** 3 months, 3 weeks ago

To me it's D, cause it says that the SDK is already installed in the Ubuntu. You just need to install de emulator using gcloud.

You can see <https://cloud.google.com/datastore/docs/tools/datastore-emulator>:

To use the Datastore emulator you need:

A Java JRE (version 8 or greater)  
The Google Cloud SDK  
An application built using the Google Cloud Client Libraries

Installing the emulator

The Datastore emulator is a component of the Google Cloud SDK's gcloud tool. Use the gcloud components install command to install the Datastore emulator:

gcloud components install cloud-datastore-emulato  
upvoted 1 times

🗨️ 👤 **Jacky\_YO** 4 months ago

ANS : C  
Google Recommended installation  
Installing Google Cloud SDK  
<https://cloud.google.com/sdk/docs/install#deb>  
upvoted 1 times

🗨️ 👤 **[Removed]** 4 months ago

D is correct. Install the cloud-datastore-emulator component using the gcloud components install command.  
upvoted 1 times

## Question #104

## Topic 1

Your company set up a complex organizational structure on Google Cloud. The structure includes hundreds of folders and projects. Only a few team members should be able to view the hierarchical structure. You need to assign minimum permissions to these team members, and you want to follow Google-recommended practices. What should you do?

- A. Add the users to roles/browser role.
- B. Add the users to roles/iam.roleViewer role.
- C. Add the users to a group, and add this group to roles/browser.
- D. Add the users to a group, and add this group to roles/iam.roleViewer role.

**Correct Answer: C**

🗲️ 👤 **SIX** Highly Voted 👍 1 year, 1 month ago

C is the better answer.  
upvoted 21 times

🗲️ 👤 **ESP\_SAP** Highly Voted 👍 11 months, 1 week ago

Correct Answer is (C):

We need to apply the GCP Best practices.  
roles/browser Browser Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.

<https://cloud.google.com/iam/docs/understanding-roles>  
upvoted 11 times

🗲️ 👤 **satyamgupta11692** Most Recent 🕒 1 month, 1 week ago

C is the better answer.  
upvoted 2 times

🗲️ 👤 **arsh1916** 2 months, 1 week ago

C is correct  
upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

C is the recommended way  
upvoted 1 times

🗲️ 👤 **Finger41** 2 months, 2 weeks ago

C is the answer. As per <https://cloud.google.com/iam/docs/understanding-roles> - "Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project."  
upvoted 1 times

🗲️ 👤 **ViMan** 2 months, 2 weeks ago

C.  
[https://cloud.google.com/resource-manager/docs/access-control-proj#using\\_predefined\\_roles](https://cloud.google.com/resource-manager/docs/access-control-proj#using_predefined_roles)  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago

C is correct. Add the users to a group, and add this group to roles/browser.  
upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

C is correct.  
upvoted 1 times

🗲️ 👤 **cloud\_hunk** 5 months ago

C is the answer: here are the facts  
quick tip from que: Only a few team members should be able to view the hierarchical structure.  
From Doc: roles/browser - Read access to browse the hierarchy for a project, including the folder, organization, and IAM policy. This role doesn't include permission to view resources in the project.  
<https://cloud.google.com/iam/docs/understanding-roles>  
upvoted 2 times

🗲️ 👤 **mohdafiuddin** 7 months ago

Breaking down the question into key requirements -

1. hundreds of folders and projects.
2. Only a few team members should be able to view the hierarchical structure.
3. assign minimum permissions to these team members
4. follow Google- recommended practices.

A. Not the correct answer. Dealing with Users is not recommended. Dealing with Groups is recommended.  
B. Not the correct answer. Dealing with Users is not recommended. Dealing with Groups is recommended.  
C. Add the users to a group, and add this group to roles/browser role. -> This is the right answer. we need to view the resources (point 2 from above) and the users are in a group.  
D. 'roles/iam.roleViewer'. -> Not enough permissions for what we need.  
upvoted 3 times

🗲️ 👤 **JKRowlings** 7 months, 2 weeks ago

Ans. C: Based on Google best practice, it is easier to manage user via group.  
upvoted 1 times

🗲️ 👤 **swatititame** 8 months, 1 week ago

• C. Add the users to a group, and add this group to roles/browser role.  
upvoted 3 times

🗲️ 👤 **glam** 9 months, 2 weeks ago

C. Add the users to a group, and add this group to roles/browser role.  
upvoted 2 times

🗲️ 👤 **RockAJ** 9 months, 4 weeks ago

C also for me!  
upvoted 2 times

🗲️ 👤 **hiteshrup** 10 months, 3 weeks ago

A correct and better answer is C as per Google guidelines  
upvoted 2 times

🗲️ 👤 **XRiddlerX** 12 months ago

C is the best answer:  
According to the Google Docs on "Best Practices for Enterprise Orgs":  
<https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#groups-and-service-accounts>

"We recommend collecting users with the same responsibilities into groups and assigning IAM roles to the groups rather than to individual users."  
upvoted 3 times

Your company has a single sign-on (SSO) identity provider that supports Security Assertion Markup Language (SAML) integration with service providers. Your company has users in Cloud Identity. You would like users to authenticate using your company's SSO provider. What should you do?

- A. In Cloud Identity, set up SSO with Google as an identity provider to access custom SAML apps.
- B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.
- C. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Mobile & Desktop Apps.
- D. Obtain OAuth 2.0 credentials, configure the user consent screen, and set up OAuth 2.0 for Web Server Applications.

Correct Answer: A

- poogcp

Highly Voted

1 year, 1 month ago

For me its B option  
upvoted 19 times
- PhilipAWS

Highly Voted

7 months, 1 week ago

Only option B make sense to me as per - [https://support.google.com/cloudidentity/answer/6262987?hl=en&ref\\_topic=7558767](https://support.google.com/cloudidentity/answer/6262987?hl=en&ref_topic=7558767)  
upvoted 13 times
- nitinz

5 months, 1 week ago

you nailed it. B is correct.  
upvoted 3 times
- Otark

Most Recent

1 month, 4 weeks ago

B is correct:  
<https://support.google.com/a/answer/6349809>  
upvoted 1 times
- tikkitakka

2 months ago

B  
  
<https://cloud.google.com/architecture/identity/single-sign-on>  
  
When you use SSO for Cloud Identity or Google Workspace, your external IdP is the SAML IdP and Google is the SAML service provider.  
upvoted 1 times
- arsh1916

2 months, 1 week ago

B is correct  
upvoted 1 times
- mcaromit

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- chalkys

4 months ago

A. source: <https://cloud.google.com/identity/solutions/enable-ss>  
upvoted 2 times
- [Removed]

4 months ago

B is correct. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.  
upvoted 3 times
- yuvi69

4 months, 1 week ago

answer is A  
<https://cloud.google.com/identity-platform/docs/web/saml>  
upvoted 2 times
- Vic1043

4 months, 3 weeks ago

ANS-A  
  
To provide users with SSO-based access to selected cloud apps, Cloud Identity as your IdP supports the OpenID Connect (OIDC) and Security Assertion Markup Language 2.0 (SAML) protocols.  
  
Cloud Identity has a large catalog of SAML apps. G Suite users can get OIDC apps in the G Suite Marketplace. While most cloud apps only support one of these protocols, a few support both.  
upvoted 2 times



🗨️ 👤 **GCP\_Student1** 4 months, 3 weeks ago

B. In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.  
upvoted 3 times

🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago

Answer is B.  
upvoted 1 times

🗨️ 👤 **cloud\_hunk** 5 months ago

B is the correct answer, based on following facts  
que tip: SSO identity provider that supports Security Assertion Markup Language (SAML) integration ("which is third party working with your company") with service providers ( in this case it is Google )

facts from google support: google acts as the service provider and provides services such as Gmail and Start Pages. Google partners act as identity providers and control usernames, passwords and other information used to identify, authenticate and authorize users for web applications that Google hosts.

Referece: [https://support.google.com/cloudidentity/answer/6262987?hl=en&ref\\_topic=7558767](https://support.google.com/cloudidentity/answer/6262987?hl=en&ref_topic=7558767)

upvoted 3 times

🗨️ 👤 **dirtyjohnny** 4 months ago

Yep, B is the correct answer. Thank you for the reference link.  
upvoted 1 times

🗨️ 👤 **adedj99** 5 months ago

A --  
<https://youtu.be/CK0ZGSh2Lh8> <-- from google  
upvoted 2 times

🗨️ 👤 **victory108** 6 months ago

A - In Cloud Identity, set up SSO with Google as an identity provider to access custom SAML apps.  
upvoted 1 times

🗨️ 👤 **victory108** 3 months, 3 weeks ago

B is correct - In Cloud Identity, set up SSO with a third-party identity provider with Google as a service provider.  
upvoted 1 times

🗨️ 👤 **vara3dk** 6 months, 4 weeks ago

Google acts as the service provider and provides services such as Gmail and Start Pages. Google partners act as identity providers and control usernames, passwords and other information used to identify, authenticate and authorize users for web applications that Google hosts. Answer is surely B  
upvoted 2 times

🗨️ 👤 **Yashv** 7 months ago

We will have Google as service provider and a third party idp like OKTA for authentication  
upvoted 1 times



## Question #106

## Topic 1

Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. You need to assign this person the minimum role for projects. What should you do?

- A. Add the user to roles/iam.roleAdmin role.
- B. Add the user to roles/iam.securityAdmin role.
- C. Add the user to roles/iam.serviceAccountUser role.
- D. Add the user to roles/iam.serviceAccountAdmin role.

**Correct Answer: C**

Reference:

<https://cloud.google.com/iam/docs/creating-managing-service-accounts>

  **SIX** Highly Voted 1 year, 1 month ago

The right answer is D.

upvoted 30 times

  **PhilipAWS** Highly Voted 7 months, 1 week ago

Whoever say C is right answer, please read the question 1000000000 times if not understand - "Your organization has a dedicated person who creates and manages all service accounts for Google Cloud projects. " Dedicated person who creates and manages all service... Now read below; To allow a user to manage service accounts, grant one of the following roles:

Service Account User (roles/iam.serviceAccountUser): Includes permissions to list service accounts, get details about a service account, and impersonate a service account.

Service Account Admin (roles/iam.serviceAccountAdmin): Includes permissions to list service accounts and get details about a service account. Also includes permissions to create, update, and delete service accounts, and to view or change the IAM policy on a service account.

Now look in which role mentioned "CREATE"?

Obviously - roles/iam.serviceAccountAdmin..... So Answer is????

1M% - D only

upvoted 20 times

  **mcaromit** Most Recent 2 months, 2 weeks ago


D is correct

upvoted 2 times

  **[Removed]** 4 months ago

D is correct. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 3 times

  **yuvi69** 4 months, 1 week ago

option D is correct. because it is the only one in the 4 options that can create and manage service accounts

upvoted 1 times

  **GCP\_Student1** 4 months, 3 weeks ago

D. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

D is correct

upvoted 1 times

  **JackGlemins** 5 months, 1 week ago

C is wrong: <https://cloud.google.com/iam/docs/impersonating-service-accounts#allow-impersonation>

Service Account User (roles/iam.serviceAccountUser): Allows members to indirectly access all the resources that the service account can access. For example, if a member has the Service Account User role on a service account, and the service account has the Cloud SQL Admin role (roles/cloudsql.admin) on the project, then the member can impersonate the service account to create a Cloud SQL instance.

I believe that D is better option.

upvoted 1 times


  **mohdafiuddin** 7 months ago

A and B are automatically eliminated because they don't have the right role/permissions.

Out of C and D, since C is just a user level access, that does not include Creation privileges, which are needed for our task. Therefore D is the right option here.



Ref - <https://cloud.google.com/iam/docs/creating-managing-service-accounts#permissions>

upvoted 3 times

  **tuta** 6 months, 4 weeks ago

have you taken exam ? how accurate is this dump ?

upvoted 1 times

  **DR15** 7 months, 1 week ago

C is the correct one.

The question states "You need to assign this person the minimum role for projects", so by providing "serviceAccountAdmin" role wont satisfy it.

upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

Its D to choose.

upvoted 2 times

  **swatititame** 8 months, 1 week ago

• D. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 2 times

  **ayj** 8 months, 3 weeks ago

roles/iam.serviceAccountUser gives:

iam.serviceAccounts.actAs

iam.serviceAccounts.get

iam.serviceAccounts.list

resourceManager.projects.get

resourceManager.projects.list

roles/iam.serviceAccountAdmin gives:

iam.serviceAccounts.create \*\*\*\*

iam.serviceAccounts.delete

iam.serviceAccounts.disable

iam.serviceAccounts.enable

iam.serviceAccounts.get

iam.serviceAccounts.getIamPolicy

iam.serviceAccounts.list

iam.serviceAccounts.setIamPolicy

iam.serviceAccounts.undelete


iam.serviceAccounts.update

resourceManager.projects.get

resourceManager.projects.list

So D, but actually think roles/iam.serviceAccountCreator would be a better role for them if an option?!

upvoted 4 times

  **cuban123** 8 months, 3 weeks ago

Its D

Not C because:

Users granted the Service Account User role on a service account can use it to indirectly access all the resources to which the service account has access. For example, if a service account has been granted the Compute Admin role (roles/compute.admin), a user that has been granted the Service Account Users role (roles/iam.serviceAccountUser) on that service account can act as the service account to start a Compute Engine instance. In this flow, the user impersonates the service account to perform any tasks using its granted roles and permissions.

upvoted 2 times

  **nwk** 9 months ago


D - roles/iam.serviceAccountAdmin

upvoted 2 times

  **glam** 9 months, 2 weeks ago


D. Add the user to roles/iam.serviceAccountAdmin role.

upvoted 2 times

  **Surjit24** 9 months, 3 weeks ago

Service Account User (roles/iam.serviceAccountUser): Allows members to indirectly access all the resources that the service account can access. For example, if a member has the Service Account User role on a service account, and the service account has the Cloud SQL Admin role (roles/cloudsql.admin) on the project, then the member can impersonate the service account to create a Cloud SQL instance. C IS CORRECT.

upvoted 1 times

  **Surjit24** 9 months, 3 weeks ago

Correct it is D

upvoted 3 times

## Question #107

## Topic 1

You are building an archival solution for your data warehouse and have selected Cloud Storage to archive your data. Your users need to be able to access this archived data once a quarter for some regulatory requirements. You want to select a cost-efficient option. Which storage option should you use?

- A. Cold Storage
- B. Nearline Storage
- C. Regional Storage
- D. Multi-Regional Storage

**Correct Answer: A**

Nearline, Coldline, and Archive offer ultra low-cost, highly-durable, highly available archival storage. For data accessed less than once a year, Archive is a cost-effective storage option for long-term preservation of data. Coldline is also ideal for cold storage data your business expects to touch less than once a quarter.

For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year. All storage classes are available across all GCP regions and provide unparalleled sub-second access speeds with a consistent API.

Reference:

<https://cloud.google.com/storage/archival>

  **Teegongkia** Highly Voted 1 year ago

Took ACE last week and the exact question came out. I go with B as i felt A is a trick answer. There is no Cold Storage in GCP.  
upvoted 13 times

  **BenAji** 3 months, 1 week ago

Cold data tiering refers to the storage of less frequently, or sporadically accessed data in low cost media such as HDFS (Hadoop Distributed File System) and cloud storage options including Amazon Web Services (AWS), Google Cloud Platform (GCP), and Azure Data Lake Storage (ADLS) that are managed separately from the SAP HANA database, but still accessible at any time. [blogs.sap.com/2018/12/03/what-is-sap-hana-cold-data-tiering/](https://blogs.sap.com/2018/12/03/what-is-sap-hana-cold-data-tiering/)  
upvoted 1 times

  **ssankar** 11 months, 2 weeks ago

Hello Teegongkia ,  
is the questions are still valid ??  
Thanks  
upvoted 1 times


  **droogie** Highly Voted 1 year, 1 month ago

This one is confusing. First, there's no 'Cold' storage. It's Coldline.  
Nearline Storage is ideal for data you plan to read or modify on average once per month or less. Coldline Storage is ideal for data you plan to read or modify at most once a quarter.  
<https://cloud.google.com/storage/docs/storage-classes>

So with the misspelling of 'Cold' and these guys accessing it every 90 days, I'm leaning towards Nearline  
upvoted 9 times

  **DickDastardly** 4 months, 3 weeks ago

"Cold" is not a typo. I took the exam today and the answers appeared exactly as listed here.  
upvoted 1 times



  **TAvenger** 5 months, 1 week ago

I believe the question is old, when Regional and Multi-Regional were also storage classes of the GCS.  
Before changes: (Multi-Region, Regional, Nearline, Coldline)  
After recent changes we have  
- Storage Classes (Standard, Nearline, Coldline, Archive)  
- Storage Locations (Regional, Dual-region, Multi-Region)  
It's tricky for exam because we don't to answer according to old version or new version.  
For the latest version, costs for 1Gb for storing (3 month) + retrieval  
Nearline:  $0.01 * 3 + 0.01 = 0.04$   
Coldline:  $0.004 * 3 + 0.02 = 0.032$   
Coldline is more cost effective.  
If "Cold" means Coldline (not Archive) the asnwer is A  
If "Cold" means Archive the answer is B

I hope that "Cold" means Coldline. I would try wirh A  
upvoted 1 times


  **ri\_unhou119** 2 months, 2 weeks ago

A:  
Google Cloud doc:  
<https://cloud.google.com/storage/docs/storage-classes#coldline>  
upvoted 1 times



  **sarahf** 7 months, 2 weeks ago

At the page for data archiving (<https://cloud.google.com/storage/archival>) the first paragraph says: "Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter."  
So there is such thing as Cold storage according to Google.


Also at (<https://cloud.google.com/storage/docs/storage-classes#archive>) they talk about Cold storage: "Cold data storage - Archived data, such as data stored for legal or regulatory reasons, can be stored at low cost as Archive Storage, yet still be available if you need it."  
upvoted 4 times

  **sanhoo** 1 month, 3 weeks ago



Thanks for bringing this up. Really helpfull.  
upvoted 1 times

  **sanhoo** 1 month, 3 weeks ago

This line indicates that cold storage term is used for - archival / coldline  
"With low latency and a consistent API across Cloud Storage, Archive and Coldline introduce cold storage you can actually use"  
<https://cloud.google.com/storage/archival>  
upvoted 1 times

  **Eshkrkrkr** 8 months, 3 weeks ago



It's a typo. Google wouldn't force to consume knowledge that is a non-best practice from Google. Asnwer is A.  
upvoted 2 times

  **lxgywil** 2 months, 3 weeks ago


For Google, these exams are just another business.  
upvoted 2 times

  **tanito83** Most Recent 1 month, 2 weeks ago

The correct answer is B. Please, modify it.  
upvoted 2 times

  **rrrr\_** 2 months ago

<https://cloud.google.com/storage/archival>  
Coldline is also ideal for cold storage—data  
A is the correct answer.  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago

A is correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

A is correct  
upvoted 1 times

  **FunkyTechnician** 2 months, 3 weeks ago

Let's stop making a simple question confusing.

Cold - Accessed a few times a year  
Near - Accessed at least once every quarter ( 90 days )  
Regional - Accessed all the time ( without redundancy )  
Multi-Regional - Accessed all the time ( with regional redundancy )

The question asks:  
Your users need to be able to access this archived data ONCE A QUARTER ( aka 90 days) for some regulatory requirements.

The Answer is Clearly B.  
upvoted 2 times


  **EJJ** 4 months ago

"Coldline Storage is ideal for data you plan to read or modify at most once a quarter. Note, however, that for data being kept entirely for backup or archiving purposes, Archive Storage is more cost-effective, as it offers the lowest storage costs."

<https://cloud.google.com/storage/docs/storage-classes>  
upvoted 1 times



  **[Removed]** 4 months ago

A is correct. Cold Storage  
upvoted 1 times

  **yuvi69** 4 months, 1 week ago

answer is B. because coldline provides access once a year. so nearline is best option

upvoted 2 times

  **yuvi69** 4 months, 1 week ago

if we see the options four options are four storage classes so among those four storage classes nearline is best and cost effective

upvoted 1 times

  **sumanshu** 4 months, 2 weeks ago

vote for 'B'

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct

upvoted 1 times

  **arsav** 5 months, 2 weeks ago

you can see this wording in below link.

<https://cloud.google.com/storage/archival>

\*\*\*\*Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter. For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year.\*\*\*\*\*

it clearly says,

Coldline is also ideal for cold storage—data your business expects to touch less than once a quarter.

you should read the next line as well

For warmer storage, choose Nearline: data you expect to access less than once a month, but possibly multiple times throughout the year

as you in these wordings, if you possibly try to access the data multiple times throughout the year you can go with Nearline.

so you can choose option B with 100% confidence.

upvoted 2 times

  **Joh3o8an** 5 months, 2 weeks ago

Cold Storage are Coldline and Archive. So choosing A means choosing both. But Archive has a minimum storage as 365 days.

C&D are standard storage and meant for request access storage.



B. Nearline can be used for archive, cost effective compared to C&D, is ideal for storage accessed once a month but can as well be accessed quartely. Ths is correct answer.

upvoted 1 times

  **AnasAloan84** 7 months ago

Correct is B. Nearline means access less than once a month, which means the access is once every 2 or 3 months. is once quarterly still applied to Nearline.

upvoted 1 times

  **lxgywil** 2 months, 3 weeks ago

Well, Coldline is intended for accessing objects once in 90 days (i.e. every quarter)...

upvoted 1 times

  **PhilipAWS** 7 months, 1 week ago

B is right answer - As per <https://cloud.google.com/storage/archival>

upvoted 1 times

  **Yashv** 7 months, 2 weeks ago

Answer A as access is required every quarter 90 days. If mentioned less than 90 its nearline

upvoted 1 times

## Question #108

## Topic 1

A team of data scientists infrequently needs to use a Google Kubernetes Engine (GKE) cluster that you manage. They require GPUs for some long-running, non- restartable jobs. You want to minimize cost. What should you do?

- A. Enable node auto-provisioning on the GKE cluster.
- B. Create a VerticalPodAutscaler for those workloads.
- C. Create a node pool with preemptible VMs and GPUs attached to those VMs.
- D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

**Correct Answer: C**

Reference:

<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>

  **Polok** Highly Voted 1 year, 1 month ago

If you need something for long-running, non- restartable jobs you dont use preemptible VMs

Think answer is D.

upvoted 35 times

  **mcaromit** Most Recent 2 months, 2 weeks ago

D is correct

upvoted 3 times

  **[Removed]** 3 months, 2 weeks ago

Incorrect options are

- B. VerticalPodAutscaler scales PODS based on the app you deploy.  
For handle infrequently GPU access, you need infrequently GPU nodes  
VerticalAutscaler Pod deployed on a non GPU node it useless,  
[We cant have the node always have GPU for infrequent requests]
- C. Preemptible VMs cant last long
- D. For infrequent access, you don't want to have a permanent homogenous cluster.

The correct option is "A"

auto-provisioning = Attaches and deletes node pools to cluster based on the requirements.  
Hence creating a GPU node pool, and auto-scaling would be better

<https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

upvoted 3 times

  **kyo** 5 days, 21 hours ago

I think using NAP is the correct answer.

→Node Auto Provisioning (NAP a.k.a., Nodepool Auto Provisioning)

There is an introduction of NAP described below on the blog.

>The above recommendations optimize for cost. NAP, for instance, reduces costs by taking down nodes during underutilized periods.



<https://cloud.google.com/blog/products/containers-kubernetes/best-practices-for-creating-a-highly-available-gke-cluster>

upvoted 1 times

  **EJJ** 4 months ago

it says "non-restartable job", ANS is definitely letter D. by enabling auto scaling with max 1 instances, you can ensure that there is always 1 instance running.

upvoted 1 times

  **lxgywil** 2 months, 3 weeks ago

Since the usage is infrequent that doesn't like the most optimal solution price-wise.

upvoted 1 times

  **khi** 4 months ago

C

<https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms>

upvoted 1 times

  **Ashii** 2 months, 2 weeks ago

non-restartable, has to be D

upvoted 1 times



  **khi** 4 months ago



D is correct.

upvoted 1 times

  **[Removed]** 4 months ago


D is correct. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

upvoted 2 times

  **cloud\_\_guru** 4 months, 2 weeks ago

Definitely D is the answer. You shouldn't use preemptible VMs for long running jobs that can't be restarted. Thanks a lot examtopics

upvoted 1 times

  **Vic1043** 4 months, 3 weeks ago

ANS-D

some long-running, non- restartable jobs

not possible with preemptible VMs

upvoted 1 times

  **GCP\_Student1** 4 months, 3 weeks ago

D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

D is correct.

upvoted 1 times

  **ShakthiGCP** 4 months, 4 weeks ago

Answer is D . Spent more time to look into Option A since someone confused .

I choose - D . Check how to get a GPUs and set autoscale enabled in this link [https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#gpu\\_pool](https://cloud.google.com/kubernetes-engine/docs/how-to/gpus#gpu_pool)

upvoted 1 times

  **JackGlemins** 5 months ago

A is the correct. <https://cloud.google.com/kubernetes-engine/docs/how-to/node-auto-provisioning>

Node auto-provisioning creates node pools based on the following information:

CPU, memory and ephemeral storage resource requests.

GPU requests

Pending Pods' node affinities and label selectors.

Pending Pods' node taints and tolerations.

upvoted 3 times

  **JackGlemins** 5 months ago

Overview

Node auto-provisioning automatically manages a set of node pools on the user's behalf. Without node auto-provisioning, GKE considers starting new nodes only from the set of user created node pools. With node auto-provisioning, new node pools can be created and deleted automatically.

upvoted 1 times

  **JackGlemins** 5 months ago

B- Wrong. Do Nothing about the question

C- preemptible instances cannot be automatically restarted. They require GPUs for some long-running, non- restartable jobs.

D- Is right but is more expensive. You want to minimize cost.

upvoted 1 times

  **JackGlemins** 5 months ago

Key word: "infrequently needs to use". You dont need create a pool of instance with GPU. D is wrong

upvoted 1 times

  **prasadjblin** 5 months, 1 week ago



D is the correct answer. Preemptible vms are restartable during maintainece

upvoted 1 times

  **victory108** 6 months ago

D - Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.

upvoted 1 times

  **vara3dk** 6 months, 4 weeks ago

During maintenance events, preemptible instances with GPUs are preempted by default and cannot be automatically restarted. If you want to recreate your instances after they have been preempted, use a managed instance group. Managed instance groups recreate your instances if the vCPU, memory, and GPU resources are available.

If you want a warning before your instance is preempted, or want to configure your instance to automatically restart after a maintenance event, use a non-preemptible instance with a GPU. For non-preemptible instances with GPUs, Google provides one hour advance notice before preemption.

Hence Answer is C

upvoted 1 times

  **PhilipAWS** 7 months, 1 week ago

According to info in below urls, answer is C ONLY... Whoever wants to bet, let's have....  
<https://cloud.google.com/kubernetes-engine/docs/how-to/gpus>  
<https://cloud.google.com/kubernetes-engine/docs/how-to/preemptible-vms>

Sorry, this time GOPINATHM WON :)  
upvoted 1 times

  **swatititame** 8 months, 1 week ago

- D. Create a node pool of instances with GPUs, and enable autoscaling on this node pool with a minimum size of 1.
- upvoted 1 times



Your organization has user identities in Active Directory. Your organization wants to use Active Directory as their source of truth for identities. Your organization wants to have full control over the Google accounts used by employees for all Google services, including your Google Cloud Platform (GCP) organization. What should you do?

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.
- B. Use the cloud Identity APIs and write a script to synchronize users to Cloud Identity.
- C. Export users from Active Directory as a CSV and import them to Cloud Identity via the Admin Console.
- D. Ask each employee to create a Google account using self signup. Require that each employee use their company email address and password.

**Correct Answer: A**  
Reference:  
<https://cloud.google.com/solutions/federating-gcp-with-active-directory-introduction>

- professor**

Highly Voted

1 year, 1 month ago

Ans is A

<https://tools.google.com/dlpage/dirsync/>  
upvoted 18 times
- ESP\_SAP**

Highly Voted

11 months, 1 week ago

Correct Answer (A):

Directory Sync  
Google Cloud Directory Sync enables administrators to synchronize users, groups and other data from an Active Directory/LDAP service to their Google Cloud domain directory

<https://tools.google.com/dlpage/dirsync/>  
upvoted 12 times
- mcaromit**

Most Recent

2 months, 2 weeks ago

A is correct

upvoted 1 times
- [Removed]**

4 months ago

A is correct. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times
- GCP\_Student1**

5 months ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 3 times
- devscorpio2001**

8 months, 1 week ago

This is A , you can use Google Cloud Sync

upvoted 1 times
- swatititame**

8 months, 1 week ago

- A. Use Google Cloud Directory Sync (GCDS) to synchronize users into Cloud Identity.

upvoted 1 times
- RockAJ**

9 months, 4 weeks ago

A also for me!

upvoted 1 times
- hiteshrup**

10 months, 3 weeks ago

Straight Answer: A

Google has Cloud Directory Sync service to link LDAP in GCP

upvoted 2 times
- SSPC**

11 months, 2 weeks ago

A is correct for me

upvoted 2 times
- Verve**

12 months ago

A it is

upvoted 2 times

## Question #110

## Topic 1

You have successfully created a development environment in a project for an application. This application uses Compute Engine and Cloud SQL. Now you need to create a production environment for this application. The security team has forbidden the existence of network routes between these 2 environments and has asked you to follow Google-recommended practices. What should you do?

- A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.
- B. Create a new production subnet in the existing VPC and a new production Cloud SQL instance in your existing project, and deploy your application using those resources.
- C. Create a new project, modify your existing VPC to be a Shared VPC, share that VPC with your new project, and replicate the setup you have in the development environment in that new project in the Shared VPC.
- D. Ask the security team to grant you the Project Editor role in an existing production project used by another division of your company. Once they grant you that role, replicate the setup you have in the development environment in that project.

**Correct Answer: A**

🗲️ 👤 **poogcp** Highly Voted 1 year, 1 month ago

Correct answer is A.  
upvoted 25 times

🗲️ 👤 **pYWORLD** 1 day ago

Correct answer!  
upvoted 1 times

🗲️ 👤 **[Removed]** Highly Voted 10 months, 3 weeks ago

A - correct. Best practice is to create a new project for each environment, such as production and testing. There are no routes between VPCs in these projects by default, so that satisfies the requirement by the security team.  
B. Nope. not best practice and allows communication.  
C. While this is best practice to create a new project for a different environment, it explicitly breaks the security team's rule of having no path between environments by nature of the shared VPC. The shared VPC allows entities in both VPCs to communicate as if they were in the same VPC. That's definitely wrong.  
D. One - not best practice to replicate in the setup in that project. Two - why do they suddenly need the project editor rule? Just a bad answer. Wrong.  
upvoted 12 times

🗲️ 👤 **JieHeng** Most Recent 1 month ago

Should be A  
it's a best practice "to have one project per application per environment." - <https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations#project-structure>  
upvoted 2 times

🗲️ 👤 **arsh1916** 2 months, 1 week ago

A answer  
upvoted 1 times

🗲️ 👤 **mcaramit** 2 months, 2 weeks ago

A is correct  
upvoted 1 times

🗲️ 👤 **EJJ** 4 months ago

i will go with A  
upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago

A is correct. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.  
upvoted 1 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

C is correct  
upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months ago

A. Create a new project, enable the Compute Engine and Cloud SQL APIs in that project, and replicate the setup you have created in the development environment.  
upvoted 2 times

🗲️ 👤 **piipo** 7 months ago

Ans is A.

upvoted 1 times

🗲️ 👤 **nightflyer** 7 months, 2 weeks ago

Guys please be aware of Google recommended practices, yes A is right but it's not recommended in this case . Use a shared vpc with a single project as the host, these are the fine margins for passing the exam

upvoted 1 times

🗲️ 👤 **ShakthiGCP** 4 months, 4 weeks ago

"The security team has forbidden the existence of network routes between these 2 environments " . Forbidden means preventing access across environment and why u want to link those network using shared VPC ???

upvoted 1 times

🗲️ 👤 **DickDastardly** 4 months, 4 weeks ago

From <https://cloud.google.com/solutions/best-practices-vpc-design>

Create a shared services VPC if multiple VPC networks need access to common resources but not each other

A VPC network provides a full mesh of global reachability. For this reason, shared services and continuous integration pipelines residing in the same VPC network don't require special consideration when it comes to connectivity—they are inherently reachable. Shared VPC extends this concept, allowing shared services to reside in an isolated project while providing connectivity to other services or consumers.

This points to C

upvoted 1 times

🗲️ 👤 **cRobert** 7 months, 1 week ago

Why would you use a shared VPC?

upvoted 1 times

🗲️ 👤 **RockAJ** 9 months, 4 weeks ago

Ans is A

upvoted 1 times

🗲️ 👤 **hiteshrup** 10 months, 3 weeks ago

The answer is A only. Google never recommend you to share Dev and Prod resources to be shared. For Production env you should have a new project and a new set of rules. This can be more efficient by using Deployment Manager use, however, that is not a case. So sticking with Answer A. None other will satisfy any condition.

Though Shared VPC can achieve things which are asked that is not the recommended approach and same time, we should have Cloud SQL and Compute Engine to be deployed on prod project eventually to make the boundary between Prod and Dev. So option C should be avoided

upvoted 4 times

🗲️ 👤 **ESP\_SAP** 11 months, 1 week ago

Correct Answer is (C):

Shared VPC overview

Shared VPC allows an organization to connect resources from multiple projects to a common Virtual Private Cloud (VPC) network, so that they can communicate with each other securely and efficiently using internal IPs from that network. When you use Shared VPC, you designate a project as a host project and attach one or more other service projects to it. The VPC networks in the host project are called Shared VPC networks. Eligible resources from service projects can use subnets in the Shared VPC network.

<https://cloud.google.com/vpc/docs/shared-vpc>

upvoted 3 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago

we cannot use shared VPC because it provides network routes between those two projects which violates the question.

upvoted 2 times

🗲️ 👤 **ShakthiGCP** 4 months, 4 weeks ago

"The security team has forbidden the existence of network routes between these 2 environments " . Forbidden means preventing access across environment and why u want to link those network using shared VPC ???

upvoted 1 times

🗲️ 👤 **[Removed]** 10 months, 2 weeks ago

It is recommended the projects are isolated why we are talking about sharing the resources?

upvoted 2 times

🗲️ 👤 **bachor** 11 months ago

But "The security team has forbidden the existence of network routes between these 2 environments" hence there shouldn't be any communication between development and production. I would say the correct answer is A

upvoted 6 times

🗲️ 👤 **raffiq** 1 year, 1 month ago

Seems answer is C: Shared VPC which allows replicate

upvoted 1 times

🗲️ 👤 **jskumar** 1 year ago

According to Google recommendation we should create separate project for each environment . Hence option A looks good.  
upvoted 6 times

  **professor** 1 year, 1 month ago



Ans is A


instances communicate in Shared VPC  
upvoted 2 times

Your management has asked an external auditor to review all the resources in a specific project. The security team has enabled the Organization Policy called Domain Restricted Sharing on the organization node by specifying only your Cloud Identity domain. You want the auditor to only be able to view, but not modify, the resources in that project. What should you do?



- A. Ask the auditor for their Google account, and give them the Viewer role on the project.
- B. Ask the auditor for their Google account, and give them the Security Reviewer role on the project.
- C. Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.
- D. Create a temporary account for the auditor in Cloud Identity, and give that account the Security Reviewer role on the project.



Correct Answer: C


-  **dan80**

Highly Voted 



 1 year, 1 month ago


C - [https://cloud.google.com/iam/docs/roles-audit-logging#scenario\\_external\\_auditors](https://cloud.google.com/iam/docs/roles-audit-logging#scenario_external_auditors)  
upvoted 32 times
-  **spudleymcdudley** 1 year ago

This guy is right!  
upvoted 5 times
-  **ESP\_SAP**



Highly Voted 



 11 months, 1 week ago



Correct Answer is (C):  
  
roles/viewer Read access to all resources. Get and list access for all resources.  
  
Using primitive roles  
The following table lists the primitive roles that you can grant to access a project, the description of what the role does, and the permissions bundled within that role. Avoid using primitive roles except when absolutely necessary. These roles are very powerful, and include a large number of permissions across all Google Cloud services. For more details on when you should use primitive roles, see the Identity and Access Management FAQ.  
  
IAM predefined roles are much more granular, and allow you to carefully manage the set of permissions that your users have access to. See Understanding Roles for a list of roles that can be granted at the project level. Creating custom roles can further increase the control you have over user permissions.  
  
[https://cloud.google.com/resource-manager/docs/access-control-proj#using\\_primitive\\_roles](https://cloud.google.com/resource-manager/docs/access-control-proj#using_primitive_roles)  
upvoted 12 times
-  **Cthakker**



Most Recent 



 2 weeks ago

C is right  
upvoted 1 times
-  **nana1995** 1 month, 3 weeks ago

Why is not D?  
upvoted 1 times
-  **happyBoo** 1 month, 2 weeks ago

roles/iam.securityReviewer will give access to view IAM roles in addition, which is not required.  
upvoted 3 times
-  **mcaromit** 2 months, 2 weeks ago

C is correct  
upvoted 1 times
-  **meh009** 3 months, 1 week ago

Easy - C  
  
Domain Retricted Sharing is on ( restrict the set of identities that are allowed to be used in Identity and Access Management policies) which takes out A/B of the equation.  
  
Easy choice between C/D  
upvoted 1 times
-  **[Removed]** 4 months ago



C is correct. Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.  
upvoted 1 times



- 🗲️ 👤 **yuvi69** 4 months, 1 week ago  
option C is correct. because viewer role grants viewing all the resources in the project whereas securityreviewer role only grants permission to list the resources  
<https://cloud.google.com/iam/docs/understanding-roles#iam-roles>  
upvoted 1 times
- 🗲️ 👤 **SSunny** 4 months, 2 weeks ago  
C- <https://cloud.google.com/iam/docs/understanding-custom-roles>  
  
The IAM Security Reviewer role (roles/iam.securityReviewer) enables the ability to view custom roles but not administer them.  
upvoted 1 times
- 🗲️ 👤 **GCP\_Student1** 4 months, 3 weeks ago  
C - Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.  
upvoted 2 times
- 🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago  
Answer C  
upvoted 1 times
- 🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago  
Answer C  
upvoted 1 times
- 🗲️ 👤 **victory108** 5 months, 3 weeks ago  
C - Create a temporary account for the auditor in Cloud Identity, and give that account the Viewer role on the project.  
upvoted 1 times
- 🗲️ 👤 **nherrerab** 6 months, 2 weeks ago  
C is Correct.  
upvoted 1 times
- 🗲️ 👤 **ayj** 8 months, 2 weeks ago  
Security Reviewer:  
Provides permissions to list all resources and IAM policies on them.  
  
Viewer: Permissions for read-only actions that do not affect state, such as viewing (but not modifying) existing resources or data.  
  
Think C makes sense then  
upvoted 3 times
- 🗲️ 👤 **Yamac** 8 months, 3 weeks ago  
For sure it's C  
upvoted 1 times
- 🗲️ 👤 **[Removed]** 10 months, 3 weeks ago  
A and B are just bad. A temporary account is much safer. What if they forgot to revoke access once the auditor is done? Yikes.  
  
C - seems to grant only the necessary permissions.  
D - close, but with the ability to view Cloud IAM policies as spudleymcdudley pointed out, it's over-permissive when the need is to just view project resources.  
  
answer: C  
upvoted 3 times



You have a workload running on Compute Engine that is critical to your business. You want to ensure that the data on the boot disk of this workload is backed up regularly. You need to be able to restore a backup as quickly as possible in case of disaster. You also want older backups to be cleaned automatically to save on cost. You want to follow Google-recommended practices. What should you do?



- A. Create a Cloud Function to create an instance template.
- B. Create a snapshot schedule for the disk using the desired interval.
- C. Create a cron job to create a new disk from the disk using gcloud.
- D. Create a Cloud Task to create an image and export it to Cloud Storage.

**Correct Answer:** B

  **DarioFama23** Highly Voted 1 year ago  
B is correct for this question  
upvoted 14 times

  **stepkurniawan** 11 months, 1 week ago  
Question: One cannot delete the old disk when using snapshot, right?  
upvoted 1 times

  **Ale1973** 10 months, 3 weeks ago  
Snapshots and disks are independent objects con GCP, you could create a snapshot form disk and then delete the disk, the snapshot will stay in place. Actually, you could use this snapshot to create a new disk, assign to another VM, mount it, and use it (all the information that the original disk had at the time of the snapshot will still be there).  
upvoted 3 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer (B):  
Best practices for persistent disk snapshots  
You can create persistent disk snapshots at any time, but you can create snapshots more quickly and with greater reliability if you use the following best practices.

Creating frequent snapshots efficiently  
Use snapshots to manage your data efficiently.

Create a snapshot of your data on a regular schedule to minimize data loss due to unexpected failure.

Improve performance by eliminating excessive snapshot downloads and by creating an image and reusing it.



Set your snapshot schedule to off-peak hours to reduce snapshot time.



Snapshot frequency limits  
Creating snapshots from persistent disks  
You can snapshot your disks at most once every 10 minutes. If you want to issue a burst of requests to snapshot your disks, you can issue at most 6 requests in 60 minutes.



If the limit is exceeded, the operation fails and returns the following error:



<https://cloud.google.com/compute/docs/disks/snapshot-best-practices>  
upvoted 12 times

  **mcaromit** Most Recent 2 months, 2 weeks ago  
B is correct  
upvoted 1 times

  **[Removed]** 4 months ago  
B is correct. Create a snapshot schedule for the disk using the desired interval.  
upvoted 1 times

  **cloud\_\_guru** 4 months, 2 weeks ago  
B is correct due to snapshots being the perfect solution to the problem question  
upvoted 1 times

  **GCP\_Student1** 5 months ago  
B. Create a snapshot schedule for the disk using the desired interval.  
upvoted 2 times

  **EABDAJA** 5 months, 1 week ago  
B is Correct



upvoted 1 times

  **Vikash211982** 5 months, 4 weeks ago

B is correct

upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

B 200%

upvoted 1 times

  **hiteshrup** 10 months, 3 weeks ago

Answer B is correct. Snapshot can be best option available from listed option

- Option A is just creating instance template and not taking boot disk backup

- Option C is to create a disk, but to run this corn job, you need another gcloud resource which is not cost effective and introducing another single failure point.

- Option D is not a definite not a good cost effective solution.

upvoted 3 times

  **[Removed]** 10 months, 3 weeks ago

Definitely B. That's Google's way of helping you. Why do C or D, which requires a lot of custom coding and configuration when B is pre-built?

A - that just misses the mark. An instnace template have any backup data. Even if it did, it's still a manual way to go, just like C and D.

Definitely B

upvoted 1 times

You need to assign a Cloud Identity and Access Management (Cloud IAM) role to an external auditor. The auditor needs to have permissions to review your Google Cloud Platform (GCP) Audit Logs and also to review your Data Access logs. What should you do?

- A. Assign the auditor the IAM role roles/logging.privateLogViewer. Perform the export of logs to Cloud Storage.
- B. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.
- C. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Perform the export of logs to Cloud Storage.
- D. Assign the auditor's IAM user to a custom role that has logging.privateLogEntries.list permission. Direct the auditor to also review the logs for changes to Cloud IAM policy.

Correct Answer: C

ESP\_SAP

Highly Voted

11 months, 1 week ago

Correct Answer is (B):

Background  
Google Cloud provides Cloud Audit Logs, which is an integral part of Cloud Logging. It consists of two log streams for each project: Admin Activity and Data Access.  
Admin Activity logs contain log entries for API calls or other administrative actions that modify the configuration or metadata of resources. Admin Activity logs are always enabled. There is no charge for your Admin Activity audit logs.  
Data Access logs record API calls that create, modify, or read user-provided data. Data Access audit logs are disabled by default because they can be large.

logging.viewer: The logging.viewer role gives the security admin team the ability to view the Admin Activity logs.  
logging.privateLogViewer : The logging.privateLogViewer role gives the ability to view the Data Access logs.  
upvoted 21 times

ESP\_SAP

11 months, 1 week ago

Correct Answer is (B): (Continuation).  
Scenario: External auditors  
In this scenario, audit logs for an organization are aggregated and exported to a central sink location. A third-party auditor is granted access several times a year to review the organization's audit logs. The auditor is not authorized to view PII data in the Admin Activity logs.

During normal access, the auditors' Google group is only granted access to view the historic logs stored in BigQuery. If any anomalies are discovered, the group is granted permission to view the actual Cloud Logging Admin Activity logs via the dashboard's elevated access mode. At the end of each audit period, the group's access is then revoked.

Data is redacted using Cloud DLP before being made accessible for viewing via the dashboard application.  
upvoted 12 times

ESP\_SAP

11 months, 1 week ago

Correct Answer is (B): (Continuation).  
The table below explains IAM logging roles that an Organization Administrator can grant to the service account used by the dashboard, as well as the resource level at which the role is granted:

logging.viewer Organization Dashboard service account The logging.viewer role permits the service account to read the Admin Activity logs in Cloud Logging.  
bigquery.dataViewer BigQuery dataset Dashboard service account The bigquery.dataViewer role permits the service account used by the dashboard application to read the exported Admin Activity logs.  
upvoted 11 times

DarioFama23

Highly Voted

1 year ago

for me B is the correct answer..  
upvoted 16 times

Eshkrkrkr

8 months, 2 weeks ago

Yes, B is correct because:  
1) Question doesn't ask us to export and store logs for any long period of time.  
2) Custom role with only logging.privateLogEntries.list permission won't let the auditor to access Log Explorer at all  
([https://cloud.google.com/logging/docs/access-control#console\\_permissions](https://cloud.google.com/logging/docs/access-control#console_permissions) - Minimal read-only access: logging.logEntries.list)  
upvoted 5 times

- mcaromit**

Most Recent

2 months, 2 weeks ago

B is correct  
upvoted 1 times
- [Removed]**

4 months ago

B is correct. Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.  
upvoted 1 times
- yuvi69**

4 months, 1 week ago

answer is B  
upvoted 1 times
- Sathya22**

4 months, 2 weeks ago

why here cloud storage is mentioned ? they are mentioning only access and why this is coming in the middle  
upvoted 2 times
- EABDAJA**

4 months, 3 weeks ago

B is correct  
upvoted 1 times
- GCP\_Student1**

5 months ago

B - Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.  
upvoted 2 times
- victory108**

5 months, 3 weeks ago

B - Assign the auditor the IAM role roles/logging.privateLogViewer. Direct the auditor to also review the logs for changes to Cloud IAM policy.  
upvoted 1 times
- Morgan91**

7 months, 2 weeks ago

A is correct  
roles/logging.privateLogViewer (Private Logs Viewer) includes roles/logging.viewer, plus the ability to read Access Transparency logs and Data Access audit logs. This role applies only to the \_Required and \_Default buckets.  
upvoted 1 times
- Bhagirathi**

7 months, 3 weeks ago

it is B .  
No other choices to make.  
upvoted 1 times
- Morgan91**

7 months, 4 weeks ago

for me b is the correct answer  
roles/logging.privateLogViewer (Private Logs Viewer) includes roles/logging.viewer, plus the ability to read Access Transparency logs and Data Access audit logs. This role applies only to the \_Required and \_Default buckets.  
upvoted 1 times
- Eshkrkrkr**

8 months, 3 weeks ago

I think it's B because question asks to audit all Cloud logs, not only Admin or Data logs - all the logs, that means that custom role with logging.privateLogEntries.list is not enough, what leads us to ask external auditor use Log explorer after we assign him/her a roles/logging.privateLogViewer role. Q does ask us to store log for any particular time period so there is no need in Cloud storage as well.  
upvoted 2 times
- Eshkrkrkr**

8 months, 2 weeks ago

To support my claim, why it's not D with custom role. [https://cloud.google.com/logging/docs/access-control#console\\_permissions](https://cloud.google.com/logging/docs/access-control#console_permissions)  
The following table lists the permissions needed to use the Logs Explorer.  
Minimal read-only access: logging.logEntries.list so if we create custom role with just logging.privateLogEntries.list permission the auditor won't be able to access Log Explorer.  
upvoted 2 times
- gcper**

8 months, 3 weeks ago

A  
  
roles/logging.privateLogViewer Private Logs Viewer Provides permissions of the Logs Viewer role and in addition, provides read-only access to log entries in private logs.  
  
roles/iam.securityReviewer Security Reviewer Provides permissions to list all resources and IAM policies on them.  
  
As stated in the documentation, privateLogViewer can not access Cloud IAM Policy.  
source: <https://cloud.google.com/iam/docs/understanding-roles>  
upvoted 1 times
- RockAJ**

9 months, 4 weeks ago

For me B  
upvoted 2 times
- rezavage**

10 months, 2 weeks ago

B is the correct answer because:

roles/logging.privateLogViewer (Private Logs Viewer) includes roles/logging.viewer, plus the ability to read Access Transparency logs and Data Access audit logs. This role applies only to the \_Required and \_Default buckets.

upvoted 1 times

 **nff** 10 months, 3 weeks ago

Answer is A - <https://cloud.google.com/logging/docs/audit>

upvoted 1 times

## Question #114

## Topic 1

You are managing several Google Cloud Platform (GCP) projects and need access to all logs for the past 60 days. You want to be able to explore and quickly analyze the log contents. You want to follow Google-recommended practices to obtain the combined logs for all projects. What should you do?

- A. Navigate to Stackdriver Logging and select resource.labels.project\_id="\*"
- B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.
- C. Create a Stackdriver Logging Export with a Sink destination to Cloud Storage. Create a lifecycle rule to delete objects after 60 days.
- D. Configure a Cloud Scheduler job to read from Stackdriver and store the logs in BigQuery. Configure the table expiration to 60 days.

**Correct Answer: B**

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

  **Verve** Highly Voted 12 months ago

Its B.

upvoted 16 times

  **[Removed]** Highly Voted 10 months, 2 weeks ago

The question is to view log past 60 days. B, c, D talks about deleting an object or truncation of table data

upvoted 9 times

  **[Removed]** 10 months, 2 weeks ago

Answer should be A

upvoted 2 times

  **[Removed]** 10 months, 2 weeks ago

Also A specifically talks about aggregation

upvoted 3 times

  **[Removed]** 10 months, 1 week ago

Also by default, you have a lot of flexibility when viewing logging in stack driver , to filter and query.

upvoted 2 times

  **xtian2900** 10 months ago

what about minimum retention is 30 days ? is it true ?

upvoted 3 times

  **[Removed]** 10 months ago

Ur correct so minimally is 30 for data access logs <https://cloud.google.com/logging/quotas> then B is the way to go.

upvoted 1 times

  **AD\_0525** Most Recent 1 month, 2 weeks ago



B is the correct one, option A does not give you the flexibility to analyze.

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago

B is correct as analysis of the log contents is a key requirement

upvoted 2 times

  **tifo16** 3 months, 1 week ago

Firstly i though that A was correct but when considering "default" retention period which is 30 days, i go for B.

<https://cloud.google.com/blog/products/it-ops/best-practices-for-working-with-google-cloud-audit-logging>

Log entries are held in Stackdriver Logging for a limited time known as the retention period. After that, the entries are deleted. To keep log entries longer, you need to export them outside of Stackdriver Logging by configuring log sinks.

upvoted 3 times

  **Crad** 3 months, 2 weeks ago

I think it's C.

If you set expiration date to 60days on a table then that table will be deleted after 60 days of creation.

Let's say you created some logs on 31st day - they'll be gone within 29 days and not 60.

upvoted 1 times

  **Crad** 3 months, 2 weeks ago

The output is in JSON so it's still possible to analyze the logs using BigQuery easily as an external source of data.

upvoted 1 times

🗨️ 👤 **[Removed]** 4 months ago

B is correct. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 1 times

🗨️ 👤 **GCP\_Student1** 4 months, 1 week ago

B. Create a Stackdriver Logging Export with a Sink destination to a BigQuery dataset. Configure the table expiration to 60 days.

upvoted 3 times

🗨️ 👤 **cloud\_\_guru** 4 months, 2 weeks ago

The answer is B because that's the simplest solution. In order to view AND analyze logs quickly as question asks, we need it to be in BigQuery

upvoted 1 times

🗨️ 👤 **EABDAJA** 4 months, 3 weeks ago

B is correct

upvoted 1 times

🗨️ 👤 **JKRowlings** 7 months, 2 weeks ago

Ans is B. <https://cloud.google.com/solutions/exporting-stackdriver-logging-for-security-and-access-analytics>

upvoted 2 times

🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago

It is B. that best fits.

upvoted 1 times

🗨️ 👤 **hiteshrup** 10 months, 3 weeks ago

Option B is my Answer as it has all require. Export Logs + Sink + BigQuery for Analytics + Partition Expiration Time. This is best solution for analytic solution.

For compliance scenario, option C is better choice

Option A is not best fit as minimum retention is 30 days for Data Access Audit logs in StackDriver.

Option D is not require as Sink will do same job as Cloud Scheduler.

upvoted 3 times

🗨️ 👤 **vlodia** 1 year ago

C better than B

upvoted 2 times

🗨️ 👤 **vlodia** 1 year ago

Well it seems B because "need to analyze the data"

upvoted 18 times

## Question #115

## Topic 1

You need to reduce GCP service costs for a division of your company using the fewest possible steps. You need to turn off all configured services in an existing GCP project. What should you do?

- A. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.
- B. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.
- C. 1. Verify that you are assigned the Organizational Administrator IAM role for this project. 2. Locate the project in the GCP console, enter the project ID and then click Shut down.
- D. 1. Verify that you are assigned the Organizational Administrators IAM role for this project. 2. Switch to the project in the GCP console, locate the resources and delete them.

**Correct Answer:** C

🗲️ 👤 **DarioFama23** Highly Voted 👍 1 year ago  
for me is A the correct answer  
upvoted 27 times

🗲️ 👤 **shafiqeee1** Highly Voted 👍 1 year ago  
A - I reproduced in my project  
upvoted 11 times

🗲️ 👤 **AD\_0525** Most Recent ⌚ 1 month, 2 weeks ago  
Answer should be A, For project shutdown you don't need any roles in org level. Just replicated in my free trial account, where I don't have any org but able to shutdown a project from console.  
upvoted 3 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago  
A is correct  
upvoted 2 times

🗲️ 👤 **noreen** 3 months, 2 weeks ago  
A Ref: <https://cloud.google.com/resource-manager/docs/organization-resource-management#delete-projects>  
upvoted 1 times

🗲️ 👤 **[Removed]** 3 months, 2 weeks ago  
The confusion is between A and C, because both can dump a project  
I Prefer C because of the wording  
A says 'click Shut down and then enter the project ID'  
C says 'enter the project ID and then click Shut down'

You cannot click and then enter something. Hence A should be false  
upvoted 1 times

🗲️ 👤 **FunkyTechnician** 2 months, 3 weeks ago  
Entering the project ID is for CONFIRMATION  
upvoted 2 times

🗲️ 👤 **Skiro** 3 months, 1 week ago  
After you click "Shut down" you should enter project ID for confirmation.  
upvoted 1 times

🗲️ 👤 **Tez1** 2 months ago  
In the GUI you actually enter the project Id, for confirmation, before shutting it down. Which is why A looks to be incorrect.  
upvoted 1 times

🗲️ 👤 **Rajusrinivasa** 4 months ago  
Ans) A  
To shut down a project using the Cloud Console:

Open the Settings page (found under IAM & admin) in the Google Cloud Console.

Open the Settings page

Click Select a project.

Select a project you want to delete, and click Open.

Click Shut down.

Enter the Project ID, then click Shut down.

ref: [https://cloud.google.com/resource-manager/docs/creating-managing-projects#console\\_4](https://cloud.google.com/resource-manager/docs/creating-managing-projects#console_4)

upvoted 1 times

🗲️ 👤 **[Removed]** 4 months ago

A is correct. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.

upvoted 1 times

🗲️ 👤 **yuvi69** 4 months, 1 week ago

answer is A

upvoted 1 times

🗲️ 👤 **ShakthiGCP** 4 months, 3 weeks ago

orgpolicy.policy.get  
resourcemanager.folders.get  
resourcemanager.folders.getIamPolicy  
resourcemanager.folders.list  
resourcemanager.folders.setIamPolicy  
resourcemanager.organizations.get  
resourcemanager.organizations.getIamPolicy  
resourcemanager.organizations.setIamPolicy  
resourcemanager.projects.get  
resourcemanager.projects.getIamPolicy  
resourcemanager.projects.list  
resourcemanager.projects.setIamPolicy

organization admin doesnt have a delete project permission to that role. Ans: A.

upvoted 5 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

A is not correct, how can you enter the project ID after shut down. C is correct. 1. Verify that you are assigned the Organizational Administrator IAM role for this project. 2. Locate the project in the GCP console, enter the project ID and then click Shut down.

upvoted 1 times

🗲️ 👤 **BenAji** 3 months, 1 week ago

you enter the ID after the shutdown of any project!!!. Its a way of re-confirming the exact project

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 5 months ago

A. 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.

upvoted 2 times

🗲️ 👤 **victory108** 5 months, 3 weeks ago

A - 1. Verify that you are assigned the Project Owners IAM role for this project. 2. Locate the project in the GCP console, click Shut down and then enter the project ID.

upvoted 1 times

🗲️ 👤 **rvgcp** 6 months ago

To me A is the right answer

upvoted 1 times

🗲️ 👤 **Liongeek** 7 months ago

Just for the record, despite the role the user should have to execute the shut down, when you are about to do it, you FIRST click Shut down and then enter the ID, you can try it. Only ans that fits is A.

upvoted 5 times

🗲️ 👤 **magistrum** 7 months ago

This just cements it

upvoted 1 times

🗲️ 👤 **Bhagirathi** 7 months, 3 weeks ago

It is equally fits both seems A & C works.

Same time, it is equally confusing ...

something to be clarified how Organization Admin IAM Role can do and not the Project Owner IAM Role?

upvoted 1 times

🗲️ 👤 **rramos96** 7 months ago

Only project owners can shut down or restore projects. ( <https://support.google.com/googleapi/answer/6251787> ) so the correct answer is [ A ]

upvoted 3 times



  **Sadagopan** 7 months ago

<https://cloud.google.com/resource-manager/docs/quickstart-organizations>  
Organization admin also can delete (shutdown) a project  
upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago

A or C - it adds to my understanding ..but at last I am confused than getting any clarity.  
upvoted 1 times

## Question #116

## Topic 1

You are configuring service accounts for an application that spans multiple projects. Virtual machines (VMs) running in the web-applications project need access to BigQuery datasets in crm-databases-proj. You want to follow Google-recommended practices to give access to the service account in the web-applications project. What should you do?

- A. Give "project owner" for web-applications appropriate roles to crm-databases-proj.
- B. Give "project owner" role to crm-databases-proj and the web-applications project.
- C. Give "project owner" role to crm-databases-proj and bigquery.dataViewer role to web-applications.
- D. Give bigquery.dataViewer role to crm-databases-proj and appropriate roles to web-applications.

**Correct Answer: C**

Reference:

<https://cloud.google.com/blog/products/gcp/best-practices-for-working-with-google-cloud-audit-logging>

  **ezat** Highly Voted 1 year ago

D cuz u just need read for DB at the other project  
upvoted 20 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

See the option correctly, as the web app needs access to the big query datasets we have to give access to the web app the data viewer role to only read the datasets! Hence, C  
upvoted 2 times

  **AnasAloan84** 7 months ago

Question didn't specify if the required access is Read only or more, its saying "access" which could be write permissions as well. I will go with C  
upvoted 1 times

  **AnasAloan84** 7 months ago

U r right, it D. why to give "project owner" as stated on C. correct answer is D  
upvoted 2 times

  **GCPACE2020** 1 month ago

but why giving bigquery.dataViewer to crm-databases-proj. we should give for web-application.  
upvoted 1 times

  **DarioFama23** 1 year ago

U re right, D is the correct answee  
upvoted 4 times

  **DarioFama23** Highly Voted 1 year ago

C is correct..  
upvoted 10 times

  **GCPACE2020** 1 month ago

But why giving project owner role to crm-databases-proj ?  
upvoted 1 times

  **AD\_0525** Most Recent 1 month, 2 weeks ago

D is correct, you just don't need any owner access in any of the projects.  
upvoted 2 times

  **Praveen\_N** 2 months ago

A and B are wrong as they talk about giving Project owner to web-applications which is unnecessary. If we go with D, crm-databases-proj will only have read-only permissions and it cannot write data into that VM so we cannot go with this one which rules that option. So only option left is C which is logical as project owner will have write permission to crm-databases-proj  
upvoted 3 times

  **arsh1916** 2 months, 1 week ago

D is answer  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

D is correct  
upvoted 1 times

  **ri\_unhou119** 2 months, 3 weeks ago

Why not A?  
I think A is correct  
upvoted 1 times

  **Jacky\_YO** 3 months, 4 weeks ago

my answer is D  
1.configuring service accounts for an application that spans multiple projects.  
2.web-applications project need access to BigQuery datasets in crm-databases-proj.

Access control examples  
[https://cloud.google.com/bigquery/docs/access-control-examples#read\\_access\\_to\\_data\\_in\\_a\\_different\\_project](https://cloud.google.com/bigquery/docs/access-control-examples#read_access_to_data_in_a_different_project)

Read access to data in a different project  
upvoted 3 times

  **Itw71** 4 months ago

VMS or Compute Engine in web-applications project need to access the BQ datasets in crm-databases-proj - so the initiation comes FROM web-applications project. This service account (call it SAWEB) would be created in web-applications project and would need roles in it's own project to operate the Compute Engine etc. But it would also need to have permission to access datasets in the crm-databases-proj. So you would add the service account SAWEB which you created in web-applications as an IAM user in the crm-databases-proj and in that project you would grant it the least permissions it needs to get the access. Very doubtful that it would need to be owner and you certainly wouldn't want a service account from one project to be an owner in another project. So the correct answer is D.



upvoted 2 times

  **[Removed]** 4 months ago

C is correct. Give  $\lambda$ project owner $\lambda$  role to crm-databases-proj and bigquery.dataViewer role to web-applications.  
upvoted 2 times

  **pondai** 4 months, 1 week ago

Why many guys select C.  
In this case. We create a service account in crm-databases-proj. And We can use this service account in our application. That all we need to do.  
I vote D  
upvoted 1 times

  **yuvi69** 4 months, 1 week ago

answer is D  
upvoted 1 times

  **sumanshu** 4 months, 1 week ago

In question it's written ==> project need access to BigQuery datasets in crm-databases-proj.  
So which type of access required its not mention ==> So answer should be "Appropriate roles to Database" . Is not ? which is in 'A'  
upvoted 2 times

  **GoCloud** 4 months, 1 week ago

Def D, read the question, the permission needs to be given to the service account.  
upvoted 3 times

  **user843983409** 4 months, 1 week ago

Yes, no need to have a project owner permission for crm-databases-proj .  
In case of D, service account gets the following permissions:  
For crm-databases-proj -> dataViewer = enough to access bigquery datasets (No need to give project owner permissions)  
For web-applications project -> appropriate permissions since service account is in the web-applications project.  
Hence, my vote for D  
upvoted 3 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

here, we create a service account for the VM's and ask the big query to grant the permission to access the data sets. Then, a big query will assign a data viewer role to the web applications. What will you gain if yoi give data viewer access to crm-databases-proj, so this means big query wants to access from the web app??? Hence, we can go with C.  
upvoted 1 times

  **Sskhan** 5 months ago

C Is correct.  
[https://cloud.google.com/bigquery/docs/access-control-examples#read\\_access\\_to\\_data\\_in\\_a\\_different\\_project](https://cloud.google.com/bigquery/docs/access-control-examples#read_access_to_data_in_a_different_project)

As D is talking about to give dataViewer access to crm-databases-proj which is not relevant because Bigquery is in crm-databases-proj project only. This dataset required to have access from web-applications project so web-applications should have bigquery.dataViewer role  
upvoted 2 times

  **hicham** 5 months ago

I vote for C too  
upvoted 2 times

  **ShakthiGCP** 5 months ago

Answer is C : Dont go with the Answer D because of high 'upvote'. read the answer C carefully and it will make sense. I was thinking Answer D but realized it is wrong. Going with C  
upvoted 2 times



An employee was terminated, but their access to Google Cloud Platform (GCP) was not removed until 2 weeks later. You need to find out this employee accessed any sensitive customer information after their termination. What should you do?

- A. View System Event Logs in Stackdriver. Search for the user's email as the principal.
- B. View System Event Logs in Stackdriver. Search for the service account associated with the user.
- C. View Data Access audit logs in Stackdriver. Search for the user's email as the principal.
- D. View the Admin Activity log in Stackdriver. Search for the service account associated with the user.

Correct Answer: B

- samvegas

Highly Voted

1 year ago

Answer = C, because the question is - "Accessing Sensitive information" specifically.

upvoted 27 times
- DarioFama23

Highly Voted

1 year ago

for me is C the correct answer

upvoted 15 times
- DarioFama23

1 year ago

<https://cloud.google.com/logging/docs/audit>

upvoted 2 times
- Jagadesh01

Most Recent

2 months ago

C - Specifically asking for Data access

upvoted 2 times
- mcaromit

2 months, 2 weeks ago

C is correct

upvoted 1 times
- [Removed]

4 months ago

C is correct. View Data Access audit logs in Stackdriver. Search for the user's email as the principal.

upvoted 1 times
- yuvi69

4 months, 1 week ago

answer is c

upvoted 1 times
- GCP\_Student1

4 months, 2 weeks ago

C. View Data Access audit logs in Stackdriver. Search for the user's email as the principal.

upvoted 2 times
- EABDAJA

4 months, 3 weeks ago

C is correct

upvoted 1 times
- victory108

4 months, 3 weeks ago

C - View Data Access audit logs in Stackdriver. Search for the user's email as the principal.

upvoted 2 times
- ShakthiGCP

4 months, 3 weeks ago

sorry, I take back my previous answer. Correct Answer : C , Reason: I tried the below in Log Explorer with Data Access Log enable. Was able to fetch the data using the below Query.  
resource.type="gcs\_bucket"  
logName="projects/mydemo-sample-proj-1001/logs/cloudaudit.googleapis.com%2Fdata\_access"  
protoPayload.authenticationInfo.principalEmail="sampleEmail@gmail.com" . Hope this helps.

upvoted 2 times
- JackGlemins



5 months ago



B & D are wrong. Service Account is not this case.



C is correct:  
Data Access audit logs  
Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.



D: are generated by Google systems, not users.



System Event audit logs  
System Event audit logs contain log entries for Google Cloud actions that modify the configuration of resources. System Event audit logs are generated by Google systems; they are not driven by direct user action.  
System Event audit logs are always written; you can't configure or disable them.  
upvoted 2 times



  **JackGlemins** 5 months ago  
<https://cloud.google.com/logging/docs/audit#data-access>  
upvoted 1 times



  **ShakthiGCP** 5 months ago  
Ans: A . Reason: They are asking for a User so it cannot be a Service Account. and the other reason is, you can search the System Event log with the Principle email in Query like below. . "authenticationInfo: { principalEmail: "useremail@gmail.com" }"  
upvoted 1 times

  **Joh3o8an** 5 months, 2 weeks ago  
Audit log is for Admin activity. The answer should be A  
upvoted 1 times

  **rvgcp** 6 months ago  
C is my choice  
upvoted 1 times

  **ladh** 7 months, 2 weeks ago  
C is answer  
upvoted 3 times



  **Bhagirathi** 8 months, 1 week ago  
what is finally to choose ?  
upvoted 1 times



  **ayj** 8 months, 2 weeks ago  
This hints at C  
[https://cloud.google.com/solutions/exporting-stackdriver-logging-for-security-and-access-analytics#which\\_users\\_accessed\\_the\\_data\\_in\\_the\\_accounts\\_table\\_last\\_month](https://cloud.google.com/solutions/exporting-stackdriver-logging-for-security-and-access-analytics#which_users_accessed_the_data_in_the_accounts_table_last_month)  
  
Which users accessed the data in the "accounts" table last month?  
The following query uses the Data Access Cloud Audit Logs logs to find the user account that most frequently queried the "accounts" table.  
Replace your-project-id with your project ID.  
upvoted 4 times



You need to create a custom IAM role for use with a GCP service. All permissions in the role must be suitable for production use. You also want to clearly share with your organization the status of the custom role. This will be the first version of the custom role. What should you do?



- A. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- B. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to BETA while testing the role permissions.
- C. Use permissions in your role that use the "testing" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.
- D. Use permissions in your role that use the "testing" support level for role permissions. Set the role stage to BETA while testing the role permissions.

**Correct Answer: C**

  **raksteer** Highly Voted 1 year ago  
You need a custom role with permissions supported in prod and you want to publish the status of the role.  
<https://cloud.google.com/iam/docs/custom-roles-permissions-support>  
SUPPORTED The permission is fully supported in custom roles.  
TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.  
NOT\_SUPPORTED The permission is not supported in custom roles.  
You can't use TESTING as it is not good for prod. And you need first version which should be ALPHA. Answer should be A.  
upvoted 43 times



  **passmepls** 11 months, 1 week ago  
good job  
upvoted 2 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (A):  
  
Testing and deploying  
Custom roles include a launch stage, which is stored in the stage property for the role. The launch stage is informational; it helps you keep track of whether each role is ready for widespread use.  
  
Each custom role can have one of the following launch stages:  
  
Launch stages  
ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.  
BETA The role has been tested on a limited basis, or it includes permissions for Google Cloud services or features that are not generally available.  
GA The role has been widely tested, and all of its permissions are for Google Cloud services or features that are generally available.  
upvoted 22 times

  **ESP\_SAP** 11 months, 1 week ago  
Correct Answer is (A): Continuation

Support levels for permissions in custom roles  
You can include many, but not all, Identity and Access Management (IAM) permissions in custom roles. Each permission has one of the following support levels:  
  
Support level Description  
SUPPORTED The permission is fully supported in custom roles.  
TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.  
NOT\_SUPPORTED The permission is not supported in custom roles.

The first version of the Custom Role is ALPHA then suitable to productions all permissions in "Supported"...  
upvoted 11 times

  **GCP\_Student1** 5 months ago  
ESP\_SAP  
There is a discrepancy between your first post and the second post. Compare these two sentences;  
1st POST - ALPHA The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

2nd POST - SUPPORTED The permission is fully supported in custom roles.

Are you still going to go with A ?  
upvoted 1 times

  **mcaromit**



Most Recent

 2 months, 2 weeks ago



A is correct  
upvoted 1 times

  **[Removed]** 4 months ago



A is correct. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.  
upvoted 1 times

  **yuvi69** 4 months, 1 week ago

answer is A. beacuse it contains SUPPORTED which we must see when creating custom roles and as it first version we must set it to ALPHA  
upvoted 1 times



  **cloud\_\_guru** 4 months, 2 weeks ago

Looks like A to be because the role is the first iteration and still under testing, we need to select 'Alpha' for the custom role. Additionally, we shouldn't use 'testing' permissions in the role after they said not to  
upvoted 1 times



  **SSunny** 4 months, 2 weeks ago

A:  
Supporting Points:  
a. The status of the custom rule is to be presented to the organization. And since this is yet in testing state, the custom rule state will be Alpha.  
b. The Support level, must be SUPPORTED so as to be used in production.  
If the support level is at TESTING, then in that case - the custom role must not be used as its results may change and unexpected behaviour may be expected.

SUPPORTED The permission is fully supported in custom roles.  
TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.  
NOT\_SUPPORTED The permission is not supported in custom roles.  
<https://cloud.google.com/iam/docs/custom-roles-permissions-support>  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct.  
upvoted 1 times

  **JackGlemins** 5 months ago

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

SUPPORTED or TESTING is by Google, and refer to the permissions (Not the custom role)

Hence C & D Is wrong: Not recommended for production use.  
TESTING The permission is being tested to check its compatibility with custom roles. You can include the permission in custom roles, but you might see unexpected behavior. Not recommended for production use.

Now A & B is tricky because is the stages of the Custom role: [https://cloud.google.com/iam/docs/understanding-custom-roles#testing\\_and\\_deploying](https://cloud.google.com/iam/docs/understanding-custom-roles#testing_and_deploying)



The Key word is: "This will be the first version of the custom role"

To me A is the correct:



ALPHA: The role is still being developed or tested, or it includes permissions for Google Cloud services or features that are not yet public. It is not ready for widespread use.

BETA: The role has been tested on a limited basis.



It is the first version it cant be tested yet. I think B is wrong.  
upvoted 1 times

  **victory108** 5 months, 3 weeks ago

A - Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.  
upvoted 1 times

  **DucSiu** 5 months, 3 weeks ago

A. Use permissions in your role that use the "supported" support level for role permissions. Set the role stage to ALPHA while testing the role permissions.  
upvoted 1 times

  **cRobert** 7 months, 1 week ago

Agree with ESP\_SAP  
upvoted 2 times



🗨️ 👤 **Bhagirathi** 8 months, 1 week ago

it is very confusing for anyone....what shall we finally choose?

upvoted 1 times

🗨️ 👤 **rvgcp** 6 months ago

A - First version is Alpha and since it is production, it has to be supported.

upvoted 1 times

🗨️ 👤 **ayj** 8 months, 2 weeks ago

A seems legit:

Must be supported as you shouldn't use roles in testing in production.

When you create a custom role, set its launch stage to ALPHA. Ask a few members of your organization to test the role. After you confirm that the custom role works correctly, change the launch stage to BETA or GA.

upvoted 2 times

🗨️ 👤 **Lieder** 8 months, 2 weeks ago

I go for A because:

1. role stage status ALPHA ist for testing purposes "The role is still being developed or tested..." - and reasonable as the question asks for a first version

<https://cloud.google.com/iam/docs/understanding-custom-roles>

2. The support levels for permissions in custom roles should be Supported as Testing is "Not recommended for production use."

<https://cloud.google.com/iam/docs/custom-roles-permissions-support>

upvoted 4 times

🗨️ 👤 **BabaMat** 9 months ago

Answer Is A

upvoted 2 times

🗨️ 👤 **RockAJ** 9 months, 4 weeks ago

Ans is A

upvoted 2 times

Your company has a large quantity of unstructured data in different file formats. You want to perform ETL transformations on the data. You need to make the data accessible on Google Cloud so it can be processed by a Dataflow job. What should you do?

- A. Upload the data to BigQuery using the bq command line tool.
- B. Upload the data to Cloud Storage using the gsutil command line tool.
- C. Upload the data into Cloud SQL using the import function in the console.
- D. Upload the data into Cloud Spanner using the import function in the console.

**Correct Answer:** *B*

Reference:

<https://cloud.google.com/solutions/performing-etl-from-relational-database-into-bigquery>

- DarioFama23** Highly Voted

1 year ago

B looks correct. Key work unstructured data

upvoted 18 times
- associatecloudexamuser** Most Recent

2 weeks, 3 days ago

B is correct. Since GCS is a datalake

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

B is correct

upvoted 1 times
- [Removed]**

4 months ago

B is correct. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 1 times
- cloud\_\_guru**

4 months, 2 weeks ago

B because it's unstructured file data and needs to be ETL transformed

upvoted 1 times
- GCP\_Student1**

5 months ago

B. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 2 times
- Bhagirathi**

7 months, 1 week ago

B 2000%

upvoted 2 times
- swatititame**

8 months, 1 week ago

B. Upload the data to Cloud Storage using the gsutil command line tool.

upvoted 1 times
- BabaMat**

9 months ago

Answer is B

upvoted 1 times
- hiteshrup**

10 months, 3 weeks ago

unstructured Data ... Wow that is key for finding a solution. All other than option B (Cloud Storage) is used for structure data.. So undoubtedly answer is B

upvoted 2 times
- stepkurniawan**

11 months, 1 week ago

B because "unstructured data" + "different file formats"

upvoted 2 times
- SSPC**

11 months, 2 weeks ago

B is correct. "UNSTRUCTURED DATA"

upvoted 2 times
- szakaria**

1 year ago

B for sure

upvoted 2 times

## Question #120

## Topic 1

You need to manage multiple Google Cloud projects in the fewest steps possible. You want to configure the Google Cloud SDK command line interface (CLI) so that you can easily manage multiple projects. What should you do?

- A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- B. 1. Create a configuration for each project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project
- C. 1. Use the default configuration for one project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.
- D. 1. Use the default configuration for one project you need to manage. 2. Use gcloud init to update the configuration values when you need to work with a non-default project.

**Correct Answer: D**

  **SSunny** Highly Voted 4 months, 2 weeks ago

A

Cloud SDK comes with a default configuration. To create multiple configurations, use gcloud config configurations create, and gcloud config configurations activate to switch between them.

<https://cloud.google.com/sdk/gcloud/reference/config/set>

upvoted 9 times

  **GCP\_Student1** Highly Voted 4 months, 2 weeks ago

A. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 6 times

  **arsh1916** Most Recent 2 months, 1 week ago

A 100%

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago

A is correct

upvoted 2 times

  **ri\_unhou119** 2 months, 2 weeks ago

Why not C?

I think C is correct.

[https://cloud.google.com/sdk/docs/configurations?hl=ja#default\\_configuration](https://cloud.google.com/sdk/docs/configurations?hl=ja#default_configuration)

upvoted 1 times

  **[Removed]** 4 months ago

A is correct. 1. Create a configuration for each project you need to manage. 2. Activate the appropriate configuration when you work with each of your assigned Google Cloud projects.

upvoted 3 times

  **user843983409** 4 months, 1 week ago

Why not c? question asks for fewest steps possible

upvoted 3 times

  **marialix87** 4 months, 1 week ago

tambien pienso que es la A

upvoted 2 times

  **cloud\_\_guru** 4 months, 2 weeks ago

A. The question is asking how do you configure? You need to create a separate config for each project and use that config to make further changes in CLI

upvoted 2 times

## Question #121

## Topic 1

Your managed instance group raised an alert stating that new instance creation has failed to create new instances. You need to maintain the number of running instances specified by the template to be able to process expected application traffic. What should you do?

- A. Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.
- B. Create an instance template that contains valid syntax that will be used by the instance group. Verify that the instance name and persistent disk name values are not the same in the template.
- C. Verify that the instance template being used by the instance group contains valid syntax. Delete any persistent disks with the same name as instance names. Set the disks.autoDelete property to true in the instance template.
- D. Delete the current instance template and replace it with a new instance template. Verify that the instance name and persistent disk name values are not the same in the template. Set the disks.autoDelete property to true in the instance template.

**Correct Answer: C**

Reference:

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances>

🗲️ 👤 **ESP\_SAP** Highly Voted 👍 11 months, 1 week ago

Correct Answer is (C):

Your instance template has set the disks.autoDelete option to false for boot persistent disks so that when a VM has been deleted (for example, because of autohealing), the persistent disk was not deleted. When the managed instance group attempted to recreate the VM with the same name, it ran into the same issue where a persistent disk already exists with the same name. Delete the existing persistent disk to resolve the immediate problem and update the instance template to set the disks.autoDelete to true if you would like boot persistent disks to be deleted alongside the instance

<https://cloud.google.com/compute/docs/instance-groups/creating-groups-of-managed-instances#troubleshooting>  
upvoted 26 times

🗲️ 👤 **magistrum** 7 months ago

Can't update instance templates, see below  
upvoted 6 times

🗲️ 👤 **JackGlemins** 5 months ago

[https://cloud.google.com/compute/docs/instance-templates#how\\_to\\_update\\_instance\\_templates](https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates)  
upvoted 5 times

🗲️ 👤 **TAvenger** Highly Voted 👍 5 months, 1 week ago

Really tricky question.

Ideal scenario would be

1. create new template, while creating ensure that in the new template disks.autoDelete=true, 3. delete existing persistent disks, 4. make rolling update ...

In order to switch to new template we need "Rolling update". Unfortunately, it is not mentioned.

With current options

C - not correct, we cannot update existing template

D - not correct, we cannot delete existing template when it is in use (just checked in GCP) (We need rolling update)

B - will not solve our problem without Rolling update

A - This is the only option (I know that it can be temporary) that will work without Rolling update according to

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>  
upvoted 10 times

🗲️ 👤 **ShakthiGCP** 4 months, 3 weeks ago

Well reasoned. I'm also going with A.  
upvoted 1 times

🗲️ 👤 **Blaaa** Most Recent 🕒 3 weeks, 5 days ago

<https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs>  
upvoted 2 times

🗲️ 👤 **Umesh09** 1 month, 3 weeks ago

D looks better option C is close but an existing instance template cannot be modified hence my answer would be D  
upvoted 2 times

🗲️ 👤 **Praveen\_N** 2 months ago

It should be A

[https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs#the\\_boot\\_disk\\_already\\_exists](https://cloud.google.com/compute/docs/troubleshooting/troubleshooting-migs#the_boot_disk_already_exists)

upvoted 2 times

🗲️ 👤 **herwig100** 2 months ago

I think B is correct.

You can't delete a disk when it is attached and you can't delete a template while it is used.

That only leaves B.

upvoted 1 times

🗲️ 👤 **arsh1916** 2 months, 1 week ago

C correct

upvoted 1 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

sorry B is correct...the PDs might be attached to other VMs so can't delete

upvoted 2 times

🗲️ 👤 **mcaromit** 2 months, 2 weeks ago

A is correct...instance templates can't be updated or deleted when in use & the instance-template already is in-use & failed to create instances as it has PDs with same names

upvoted 1 times

🗲️ 👤 **meh009** 3 months, 1 week ago

D for sure. You cannot update an existing instance template or change an instance template after you create it.

upvoted 3 times

🗲️ 👤 **Jacky\_YO** 3 months, 3 weeks ago

ANC : D

Always told us

"You cannot update an existing instance template or change an instance template after it has been created"

1.Instance templates

[https://cloud.google.com/compute/docs/instance-templates#how\\_to\\_update\\_instance\\_templates](https://cloud.google.com/compute/docs/instance-templates#how_to_update_instance_templates)

2.Creating a VM Instance from an Instance Template

[https://cloud.google.com/compute/docs/instances/create-vm-from-instance-template#gcloud\\_1](https://cloud.google.com/compute/docs/instances/create-vm-from-instance-template#gcloud_1)

3.Creating instance templates

<https://cloud.google.com/compute/docs/instance-templates/create-instance-templates>

upvoted 2 times

🗲️ 👤 **yuvarajmrec** 4 months ago

D is correct. we cannot update instance template.

upvoted 2 times

🗲️ 👤 **GCP\_Student1** 4 months, 2 weeks ago

A - Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.

upvoted 2 times

🗲️ 👤 **EABDAJA** 4 months, 3 weeks ago

A is correct

upvoted 1 times

🗲️ 👤 **victory108** 5 months, 3 weeks ago

A - Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.

upvoted 2 times

🗲️ 👤 **CloudPadawan** 6 months ago

Hi, in my opinion:

C and D are wrong because you cannot delete an instance template if a managed instance group references it.

(<https://cloud.google.com/compute/docs/instance-templates/get-list-delete-instance-templates>) neither update or modify

(<https://cloud.google.com/compute/docs/instance-templates>).

B Looks wrong too. To resolve the problem immediately you should delete the existing persistent disk, thing that B doesn't.

A Sounds better no? Create an instance template that contains valid syntax which will be used by the instance group. Delete any persistent disks with the same name as instance names.

upvoted 1 times

🗲️ 👤 **pca2b** 4 months ago

<https://cloud.google.com/compute/docs/instance-templates/get-list-delete-instance-templates>: see last para...older referenced templates can be deleted.

also see: above link (<https://stackoverflow.com/questions/41622489/google-compute-engine-how-to-change-instance-template>)

makes D viable

i.e. ideally should've been worded: create a new template for the MIG, and delete old template, but i'll buy that slip, then...make sure new template does not overlap with older VM/disk names.

and setting disks.autoDelete true, will avoid current predicament, of persistent disk from previous generation VM still exists, so cannot create a new VM.

I vote for D

upvoted 1 times

🗲️ 👤 **pca2b** 4 months ago

instance name vs persistent disk name:

instance name normally will be MIG name+few (4?)random-chars, but this is a MIG in which "you need to maintain the number of running instances specified by the template to be able to process expected application traffic", so it has fixed size i.e. no autoscaling, and looks like it has VMs with specific names, which checks out in link:

[https://cloud.google.com/compute/docs/instance-groups/working-with-managed-instances#adding\\_instances\\_with\\_specific\\_names](https://cloud.google.com/compute/docs/instance-groups/working-with-managed-instances#adding_instances_with_specific_names)

Then we just verify PD names in template are not same as these specific instance names.  
ensure disks.autoDelete is true.

and I'm back to D  
upvoted 1 times

  **pca2b** 4 months ago

ignore "instance name vs persistent disk name" comment above, i cant delete that rabbit-hole comment.

PD boot disk name is default same as instance name.

Current template didnt work, we just need to ensure there is no syntax error,  
then delete any loose PD's with same names as instances, and create a working template with disk.autoDelete true.

I change my vote to C  
upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

B 2000%

upvoted 2 times

  **ShakthiGCP** 4 months, 3 weeks ago

B is wrong. Please dont confuse everyone with 2000%. you can have both instance name and Disk with the same name. I would recommend everyone to math the probabilities. My Answer is : A .

upvoted 4 times

  **gcpengineer** 1 day, 23 hours ago

D is the ans

upvoted 1 times

## Question #122

## Topic 1

Your company is moving from an on-premises environment to Google Cloud. You have multiple development teams that use Cassandra environments as backend databases. They all need a development environment that is isolated from other Cassandra instances. You want to move to Google Cloud quickly and with minimal support effort. What should you do?

- A. 1. Build an instruction guide to install Cassandra on Google Cloud. 2. Make the instruction guide accessible to your developers.
- B. 1. Advise your developers to go to Cloud Marketplace. 2. Ask the developers to launch a Cassandra image for their development work.
- C. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Use the snapshot to create instances for your developers.
- D. 1. Build a Cassandra Compute Engine instance and take a snapshot of it. 2. Upload the snapshot to Cloud Storage and make it accessible to your developers. 3. Build instructions to create a Compute Engine instance from the snapshot so that developers can do it themselves.

**Correct Answer: D**

  **DarioFama23** Highly Voted 1 year ago

B is correct for me.. launch a solution from marketplace  
upvoted 19 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (B):

<https://medium.com/google-cloud/how-to-deploy-cassandra-and-connect-on-google-cloud-platform-with-a-few-clicks-11ee3d7001d1>  
upvoted 13 times

  **nightflyer** 7 months, 2 weeks ago

But we are moving from on premises to gcp  
upvoted 1 times

  **xtian2900** 10 months, 3 weeks ago

thanks, i always look for your insight  
upvoted 2 times

  **arsh1916** Most Recent 2 months, 1 week ago

B is correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago

B is correct  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

B is correct  
upvoted 2 times

  **GCP\_Student1** 5 months ago

B. 1. Advise your developers to go to Cloud Marketplace. 2. Ask the developers to launch a Cassandra image for their development work.  
upvoted 2 times

  **Bhagirathi** 7 months, 1 week ago

B 200%  
upvoted 1 times

  **tardao** 6 months ago

why only 200% and not 2000%?  
upvoted 21 times

  **ayj** 8 months, 2 weeks ago

B and D work, but B is minimum effort as per the question  
upvoted 2 times

  **nsibuea** 8 months, 3 weeks ago

D. Question says to move the database, not build it from scratch  
upvoted 1 times

  **kahina** 8 months, 4 weeks ago

B for me .  
upvoted 1 times

🗨️ 👤 **RockAJ** 9 months, 4 weeks ago

B also for me  
upvoted 2 times

🗨️ 👤 **Ale1973** 10 months, 3 weeks ago

Requirement 1: multiple Cassandra environments to deployment teams.  
Requirement 2: each Cassandra instance must be isolated from others.  
Requirement 3: You want to move to GCP quickly and with minimal support effort.  
My answer is B because A & S are not quickly and C implies a high support effort.  
upvoted 5 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

B is correct for me  
upvoted 3 times

🗨️ 👤 **szakaria** 1 year ago

Both B and D are correct. However, B is more expensive. I vote for D.  
upvoted 2 times

🗨️ 👤 **someoneinthecloud** 12 months ago

The question only asks for minimal effort, not cost, hence, B is correct. D involves a lot of work.  
upvoted 3 times

🗨️ 👤 **szakaria** 1 year ago

It says "You want to move to GCP quickly and with minimal support effort" . Hence, B is correct.  
upvoted 7 times



You have a Compute Engine instance hosting a production application. You want to receive an email if the instance consumes more than 90% of its CPU resources for more than 15 minutes. You want to use Google services. What should you do?

- A. 1. Create a consumer Gmail account. 2. Write a script that monitors the CPU usage. 3. When the CPU usage exceeds the threshold, have that script send an email using the Gmail account and smtp.gmail.com on port 25 as SMTP server.
- B. 1. Create a Stackdriver Workspace, and associate your Google Cloud Platform (GCP) project with it. 2. Create an Alerting Policy in Stackdriver that uses the threshold as a trigger condition. 3. Configure your email address in the notification channel.
- C. 1. Create a Stackdriver Workspace, and associate your GCP project with it. 2. Write a script that monitors the CPU usage and sends it as a custom metric to Stackdriver. 3. Create an uptime check for the instance in Stackdriver.
- D. 1. In Stackdriver Logging, create a logs-based metric to extract the CPU usage by using this regular expression: CPU Usage: ([0-9] {1,3})% 2. In Stackdriver Monitoring, create an Alerting Policy based on this metric. 3. Configure your email address in the notification channel.

Correct Answer: D

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (B):

Specifying conditions for alerting policies  
This page describes how to specify conditions for alerting policies.

The conditions for an alerting policy define what is monitored and when to trigger an alert. For example, suppose you want to define an alerting policy that emails you if the CPU utilization of a Compute Engine VM instance is above 80% for more than 3 minutes. You use the conditions dialog to specify that you want to monitor the CPU utilization of a Compute Engine VM instance, and that you want an alerting policy to trigger when that utilization is above 80% for 3 minutes.

<https://cloud.google.com/monitoring/alerts/ui-conditions-ga>  
upvoted 27 times

  **DarioFama23** Highly Voted 1 year ago



Maybe B, stackdriver monitoring  
upvoted 12 times

  **tanito83** Most Recent 1 month, 2 weeks ago



The correct answer is B. Please, modify it.  
upvoted 2 times

  **Umesh09** 2 months ago



Why B ? as soon as CPU his 90% it will send an alert right ? But the question states "90% for more than 15 min" so the answer should be D, can someone help in putting some thoughts between B and D which is correct ?  
upvoted 1 times

  **sanhoo** 1 month, 3 weeks ago

In the "for" menu you can select time <https://cloud.google.com/monitoring/alerts/ui-conditions-ga>. So B is correct  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago



B is correct  
upvoted 1 times

  **meh009** 3 months, 1 week ago

B for sure. plus also GCP are all about automation and ease and are not likely to expect you to create scripts for such simple but highly requested tasks.  
upvoted 1 times

  **passnow** 4 months ago

D is the correct answer, i work this every day  
upvoted 2 times

  **Hi2ALL** 4 months, 2 weeks ago

B is more realistic  
upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago

B is correct  
upvoted 1 times

🗨️ 👤 **GCP\_Student1** 5 months ago

B. 1. Create a Stackdriver Workspace, and associate your Google Cloud Platform (GCP) project with it. 2. Create an Alerting Policy in Stackdriver that uses the threshold as a trigger condition. 3. Configure your email address in the notification channel.

upvoted 2 times

🗨️ 👤 **DucSiu** 5 months, 3 weeks ago

why is not choose D?

upvoted 1 times

🗨️ 👤 **rvgcp** 6 months ago

B is my choice

upvoted 1 times

🗨️ 👤 **nony1992** 7 months, 1 week ago

Ans is B

upvoted 1 times

🗨️ 👤 **RockAJ** 9 months, 4 weeks ago

B for me

upvoted 1 times

🗨️ 👤 **[Removed]** 10 months, 3 weeks ago

I agree that B is the way to go.

Don't make a custom script. Just create an alert as ESP\_SAP describes. Using logs (D) is one way to do it, but that is not the main point of logs. D is a jury-rigged solution for something that Stackdriver Alerting policies already provide.

upvoted 1 times

🗨️ 👤 **phire** 1 year ago

B: GCP -> Monitoring -> Alerting -> Create Policy -> Set Conditions and Notifications

upvoted 8 times

🗨️ 👤 **szakaria** 1 year ago

D is correct because I just tried it now.

upvoted 1 times

🗨️ 👤 **ayj** 8 months, 2 weeks ago

but it talks about logging?

upvoted 1 times

🗨️ 👤 **DarioFama23** 1 year ago

It's B, in the section of condition trigger u can select threshold and minutes..

upvoted 3 times

- You have an application that uses Cloud Spanner as a backend database. The application has a very predictable traffic pattern. You want to automatically scale up or down the number of Spanner nodes depending on traffic. What should you do?
- A. Create a cron job that runs on a scheduled basis to review Cloud Monitoring metrics, and then resize the Spanner instance accordingly.
  - B. Create a Cloud Monitoring alerting policy to send an alert to oncall SRE emails when Cloud Spanner CPU exceeds the threshold. SREs would scale resources up or down accordingly.
  - C. Create a Cloud Monitoring alerting policy to send an alert to Google Cloud Support email when Cloud Spanner CPU exceeds your threshold. Google support would scale resources up or down accordingly.
  - D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

**Correct Answer:** D

- GCP\_Student1**

Highly Voted

4 months, 2 weeks ago

D. Create a Cloud Monitoring alerting policy to send an alert to webhook when Cloud Spanner CPU is over or under your threshold. Create a Cloud Function that listens to HTTP and resizes Spanner resources accordingly.

upvoted 10 times
- jcpls**

Most Recent

2 months ago

Option D. There's an official repository that does something similar to provide autoscaling to Cloud Spanner.  
<https://github.com/cloudspannerecosystem/autoscaler>

upvoted 2 times
- arsh1916**

2 months, 1 week ago

D no doubt

upvoted 2 times
- mcaromit**

2 months, 2 weeks ago

D is perfect

upvoted 2 times
- kopper2019**

3 months, 2 weeks ago

D the way to go

upvoted 1 times
- pondai**

4 months ago

D  
A cron job not what we want.  
B Call SRE to do this not automatic  
C Google don't help you do this.

upvoted 1 times
- SSunny**

4 months, 2 weeks ago

D  
Excerpt from - <https://cloud.google.com/spanner/docs/instances>  
Note: You can scale the number of nodes in your instance based on the Cloud Monitoring metrics on CPU or storage utilization in conjunction with Cloud Functions.

upvoted 3 times

## Question #125

## Topic 1

Your company publishes large files on an Apache web server that runs on a Compute Engine instance. The Apache web server is not the only application running in the project. You want to receive an email when the egress network costs for the server exceed 100 dollars for the current month as measured by Google Cloud.

What should you do?

- A. Set up a budget alert on the project with an amount of 100 dollars, a threshold of 100%, and notification type of `λemail.λ€`
- B. Set up a budget alert on the billing account with an amount of 100 dollars, a threshold of 100%, and notification type of `λemail.λ€`
- C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.
- D. Use the Cloud Logging Agent to export the Apache web server logs to Cloud Logging. Create a Cloud Function that uses BigQuery to parse the HTTP response log data in Cloud Logging for the current month and sends an email if the size of all HTTP responses, multiplied by current Google Cloud egress prices, totals over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.

**Correct Answer: A**

  **GCP\_Student1** Highly Voted 4 months, 2 weeks ago

C. Export the billing data to BigQuery. Create a Cloud Function that uses BigQuery to sum the egress network costs of the exported billing data for the Apache web server for the current month and sends an email if it is over 100 dollars. Schedule the Cloud Function using Cloud Scheduler to run hourly.

upvoted 11 times

  **MyName\_** Highly Voted 4 months ago

[C]

I think the keyword here is "as measured by Google Cloud". In Answer D you calculate the price yourself, in C you use the billing provided by GCP. Thus I think the Answer is C.

upvoted 6 times

  **Xumbegnows** Most Recent 1 month ago

How about B?

You can define the scope of the budget. For example, you can scope the budget to apply to the spend in an entire Cloud Billing account, or narrow the scope to one or more projects, and/or one or more products, and/or other budget filters applicable to your Cloud Billing account.

<https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 1 times

  **arsh1916** 2 months, 1 week ago

C, remember you need to filter.

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago

C is best option given that the web-server is not the only application in the project emitting traffic

upvoted 1 times

  **ApaMokus** 2 months, 4 weeks ago

It's C indeed.

A: Not correct as there are other services that can generate egress

B: A billing account is always linked on project level, so it contains other services

C: A bit complex, but this is the right choice

D: Again, not filtered for the Apache web server egress

upvoted 3 times

  **Ebaa** 1 month ago

A Cloud Billing account can be linked to one or more projects.

<https://cloud.google.com/billing/docs/concepts>

upvoted 1 times

  **djgodzilla** 1 month, 2 weeks ago

D: looks like it's already filtered => "Cloud Logging Agent to export the Apache web server logs to Cloud Logging" and parsed in BigQuery by the function every hour . C doesn't say that the billing data is exported every hour only the function runs each hour.

upvoted 1 times

  **Biju1** 3 months, 2 weeks ago

C is the answer



upvoted 1 times

  **pondai** 4 months ago

I Vote D

C Egress not only apache web,so I don't think C is correct.

upvoted 2 times

  **MaxImbrox** 4 months, 1 week ago

Answer is C :)

upvoted 1 times

  **marialix87** 4 months, 1 week ago

voy con D

upvoted 1 times

  **GCP\_user** 4 months, 2 weeks ago

D: There are more than 1 server running under this project and the question is "You want to receive an email when the egress network costs for the server exceed 100 dollars for the current month as measured by Google Cloud." the server, means this server, the web server, then check the http log for egress (http responses) traffic multiple rate make sense.

upvoted 2 times

  **raman23125** 4 months, 2 weeks ago

I will go with c

D : only exporting apache server logs, but there are other application too.

Still I'm confused



<https://blog.doit-intl.com/the-truth-behind-google-cloud-egress-traffic-6e8f57b5c2f8>

upvoted 1 times

  **djgodzilla** 1 month, 2 weeks ago



not the only application running in the project but presumably only one in the server.

upvoted 1 times

  **MQQ** 2 months ago

Not D, because the way mentioned in this blog needs VPC Flow logs not the Apache logs.

upvoted 1 times

  **lio123** 4 months, 2 weeks ago

i think C

upvoted 3 times

You have designed a solution on Google Cloud that uses multiple Google Cloud products. Your company has asked you to estimate the costs of the solution. You need to provide estimates for the monthly total cost. What should you do?

- A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.
- B. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Create a Google Sheet that summarizes the expected monthly costs for each product.
- C. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Navigate to the Billing Report page in the Cloud Console. Multiply the 1 week cost to determine the monthly costs.
- D. Provision the solution on Google Cloud. Leave the solution provisioned for 1 week. Use Cloud Monitoring to determine the provisioned and used resource amounts. Multiply the 1 week cost to determine the monthly costs.

Correct Answer: C

- GCP\_user**

Highly Voted

4 months, 2 weeks ago

Yes I agree with A. Makes more sense.

upvoted 13 times
- jcloud965**

Most Recent

1 week, 6 days ago

Correct is A. It is the main purpose of Pricing calculator. You have to check pricing details on products pages to understand what charges apply and how to size it.

B. It works bur it is more complicated than A.

C. It works but you need to wait one week and you pay for this. Moreoever, you won't be in production so somes charges may lack like net egress charges

D. No, because Cloud Monitoring can't be used for billing purpose

upvoted 3 times
- tanito83**

1 month, 2 weeks ago

The correct answer is A. Please, modify it.

upvoted 2 times
- arsh1916**

2 months, 1 week ago

A, is best option

upvoted 1 times
- mcaromit**

2 months, 2 weeks ago

A is correct

upvoted 1 times
- lio123**

4 months, 2 weeks ago

it is should be A

upvoted 1 times
- hasnaa**

4 months, 2 weeks ago

For me the correct answer is A

upvoted 2 times
- omunoz**

4 months, 2 weeks ago

Should be A –GCP best practices is to use Billing Calculator to estimate costs

upvoted 2 times
- GCP\_Student1**

4 months, 2 weeks ago

A. For each Google Cloud product in the solution, review the pricing details on the products pricing page. Use the pricing calculator to total the monthly costs for each Google Cloud product.

upvoted 4 times

## Question #127

## Topic 1

You have an application that receives SSL-encrypted TCP traffic on port 443. Clients for this application are located all over the world. You want to minimize latency for the clients. Which load balancing option should you use?

- A. HTTPS Load Balancer
- B. Network Load Balancer
- C. SSL Proxy Load Balancer
- D. Internal TCP/UDP Load Balancer. Add a firewall rule allowing ingress traffic from 0.0.0.0/0 on the target instances.

**Correct Answer: C**

Reference:

<https://cloud.google.com/load-balancing/docs/ssl>

  **Bituz** Highly Voted 1 year ago

SSL Proxy Load Balancing support for the following ports: 25, 43, 110, 143, 195, 443, 465, 587, 700, 993, 995, 1883, 3389, 5222, 5432, 5671, 5672, 5900, 5901, 6379, 8085, 8099, 9092, 9200, and 9300. When you use Google- managed SSL certificates with SSL Proxy Load Balancing, the frontend port for traffic must be 443 to enable the Google-managed SSL certificates to be provisioned and renewed.

upvoted 16 times

  **DarioFama23** Highly Voted 1 year ago

C is correct

upvoted 15 times

  **aa\_desh** Most Recent 1 month, 4 weeks ago

Answer is:B

he SSL proxy load balancer terminates TLS in locations that are distributed globally, so as to minimize latency between clients and the load balancer. If you require geographic control over where TLS is terminated, you should use Network Load Balancing instead, and terminate TLS on backends that are located in regions appropriate to your needs.

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 1 times

  **aa\_desh** 1 month, 4 weeks ago

The answer is B,

The SSL proxy load balancer terminates TLS in locations that are distributed globally, so as to minimize latency between clients and the load balancer. If you require geographic control over where TLS is terminated, you should use Network Load Balancing instead, and terminate TLS on backends that are located in regions appropriate to your needs.



<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 2 times

  **mcaromit** 2 months, 2 weeks ago

C is correct as we need a Global, SSL offload load-balancer

upvoted 2 times

  **yuvarajmrec** 4 months, 1 week ago

they gave 'TCP traffic and it is SSL-encrypted ' which means we must use SSL proxy loadbalancer.

see the flowchart in the link

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

upvoted 3 times

  **marialix87** 4 months, 1 week ago

C is correct

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

C is correct.

upvoted 1 times

  **GCP\_Student1** 5 months, 2 weeks ago

The answer is C

You have an application that receives SSL-encrypted TCP traffic on port 443. Clients for this application are located all over the world. You want to minimize latency for the clients. Which load balancing option should you use?

SSL Proxy Load Balancing is a reverse proxy load balancer that distributes SSL traffic coming from the internet to virtual machine (VM) instances in your Google Cloud VPC network.

When using SSL Proxy Load Balancing for your SSL traffic, user SSL (TLS) connections are terminated at the load balancing layer, and then proxied



to the closest available backend instances by using either SSL (recommended) or TCP. For the types of backends that are supported, see Backends.

With the Premium Tier, SSL Proxy Load Balancing can be configured as a global load balancing service. With Standard Tier, the SSL proxy load balancer handles load balancing regionally.

SSL Proxy Load Balancing is intended for non-HTTP(S) traffic. For HTTP(S) traffic, we recommend that you use HTTP(S) Load Balancing

<https://cloud.google.com/load-balancing/docs/ssl>

upvoted 6 times

  **lutoa** 5 months, 2 weeks ago

I think it must be C and maybe google have adjusted the name of load balancer to TCP proxy because these docs specifically mention:

“ TCP Proxy Load Balancing lets you use a single IP address for all users worldwide. The TCP proxy load balancer automatically routes traffic to the backends that are closest to the user.”

which is the only type that would improve LATENCY which is a key word in the question. If latency not mentioned then probably https would have been better.

<https://cloud.google.com/load-balancing/docs/tcp>

upvoted 1 times

  **andregrijp** 5 months, 3 weeks ago

The correct is C, because is SSL and doesn't mention HTTPs so looks like the traffic is HTTP over SSL. Also for SSL proxy you still need 443 port

2.3. Port Number

The first data that an HTTP server expects to receive from the client is the Request-Line production. The first data that a TLS server (and hence an HTTP/TLS server) expects to receive is the ClientHello. Consequently, common practice has been to run HTTP/TLS over a separate port in order to distinguish which protocol is being used. When HTTP/TLS is being run over a TCP/IP connection, the default port is 443. This does not preclude HTTP/TLS from being run over another transport. TLS only presumes a reliable connection-oriented data stream.

URL: <https://tools.ietf.org/html/rfc2818>

upvoted 1 times

  **AnasAloan84** 7 months ago

C is correct, if the question mentioned "HTTPS on 443" then we go with option A

upvoted 2 times

  **Bhagirathi** 7 months, 1 week ago

confusing.....go with C.

upvoted 1 times

  **saher677** 8 months ago

Answer is B,

If you require geographic control over where TLS is terminated (which is our scenario with clients located all over the world), you should use Google Cloud Network Load Balancing instead, and terminate TLS on backends that are located in regions appropriate to your needs. Ref:



<https://cloud.google.com/load-balancing/docs/https#control-tls-termination> Network Load Balancer is the right answer.

upvoted 3 times

  **swatititame** 8 months, 1 week ago



• C. SSL Proxy Load Balancer

upvoted 1 times

  **kahina** 8 months, 4 weeks ago

C IS CORRECT

upvoted 1 times

  **RockAJ** 9 months, 4 weeks ago

C for me

upvoted 1 times



## Question #128

## Topic 1

You have an application on a general-purpose Compute Engine instance that is experiencing excessive disk read throttling on its Zonal SSD Persistent Disk. The application primarily reads large files from disk. The disk size is currently 350 GB. You want to provide the maximum amount of throughput while minimizing costs.

What should you do?

- A. Increase the size of the disk to 1 TB.
- B. Increase the allocated CPU to the instance.
- C. Migrate to use a Local SSD on the instance.
- D. Migrate to use a Regional SSD on the instance.

**Correct Answer: C**

Reference:

<https://cloud.google.com/compute/docs/disks/performance>

  **DarioFama23** Highly Voted 1 year ago

C is correct, local SSD has more IOPS  
upvoted 23 times

  **TAvenger** 5 months, 1 week ago

Agree. This is also cheaper than having 350 Gb persistent SSD:  
Here are calculations (taken from GCP when creating instance)  
350 Gb SSD Persistent disk: 59.50\$/month, read IOPS: 10 500 with n1-standard-1  
1000 Gb SSD Persistent disk: 170.00\$/month, read IOPS: 15 000 with n1-standard-1  
375 Gb Local SSD (NVMe): 30.00\$/month, read IOPS: 170 000 with n1-standard-1  
upvoted 3 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (C):

Performance

Standard persistent disks are efficient and economical for handling sequential read/write operations, but they aren't optimized to handle high rates of random input/output operations per second (IOPS). If your apps require high rates of random IOPS, use SSD persistent disks. SSD persistent disks are designed for single-digit millisecond latencies. Observed latency is application specific.

<https://cloud.google.com/compute/docs/disks#performance>  
upvoted 10 times

  **jcloud965** Most Recent 1 week, 6 days ago



Local SSD disks are not used for the same usage than Persistent SSD.  
C can't be correct.

To increase disk performance, start with the following steps:  
1. Resize your persistent disks to increase the per-disk IOPS and throughput limits.  
2. Change the machine type and number of vCPUs on the instance to increase the per-instance IOPS and throughput limits.  
[https://cloud.google.com/compute/docs/disks/performance#optimize\\_disk\\_performance](https://cloud.google.com/compute/docs/disks/performance#optimize_disk_performance)

A & B are corrects but you have to add a lot of CPU (8 by 8) to achieve throughput increase. So the correct answer is A.  
upvoted 1 times

  **mcaromit** 2 months, 2 weeks ago


Since pd-ssd is already being used so adding more storage (1TB) won't make it cheaper but performance would definitely improve. So the choice is to be made on IOPS where Local SSD beats PD-SSD...hence C is correct  
upvoted 1 times

  **tifo16** 3 months, 1 week ago

I go for A.  
even Local SSD better meet performance and costing requirements BUT as the existing used disk is a ZONAL SSD, i think that this choice consider the durability concerns. Or local SSD is mainly for temporary usage. That's why increasing the existing zonal SSD size seems more appropriate for me  
upvoted 1 times

  **GCP\_Student1** 5 months ago

C. Migrate to use a Local SSD on the instance.  
upvoted 2 times

  **TAvenger** 5 months, 1 week ago

Here are calculations (taken from GCP when creating instance)

350 Gb SSD Persistent disk: 59.50\$/month, read IOPS: 10 500 with n1-standard-1


1000 Gb SSD Persistent disk: 170.00\$/month, read IOPS: 15 000 with n1-standard-1

375 Gb Local SSD (NVMe): 30.00\$/month, read IOPS: 170 000 with n1-standard-1

So switching to local SSD makes it cheaper and faster. Adding CPUs will make it more expensive than the old price.

If we don't need information persistence the answer is "C", Local SSD.


upvoted 2 times

  **gg\_robin** 5 months, 3 weeks ago

A

Local SSD:s are not persistent. Increasing the size of the existing disk increases the read IOPS from estimated 10500 to 15000.



upvoted 3 times

  **andregrijp** 5 months, 3 weeks ago

C is correct because according to this table, Local SSD have minimum capacity of 375GB, so will be ok for this task:

<https://cloud.google.com/compute/docs/disks>

upvoted 1 times

  **DucSiu** 5 months, 4 weeks ago

Local SSD is type more IOPS

upvoted 1 times

  **SimoneP** 5 months, 4 weeks ago

IMHO ANS is A

Guys, local SSD disks are not persistent

<https://cloud.google.com/compute/docs/disks/local-ssd>

Local SSDs are suitable only for temporary storage such as caches, processing space, or low value data. To store data that is not temporary or ephemeral in nature, use one of our durable storage options.

upvoted 4 times

  **devscorpio2001** 8 months, 1 week ago

How about B ? Increasing the CPU will help with IOPS .

upvoted 3 times

  **piotrekpal** 10 months, 2 weeks ago

Correct is "C" - Local disk cost ~33 USD (350 GB), 1TB Zonal Persistent Disk cost ~100 USD

upvoted 5 times

  **amalave** 10 months, 3 weeks ago

Hi guys, the correct answer is A.

It says, "You want to provide the maximum amount of throughput while minimizing costs." If we want the maximum amount of IOPS and keep the cost low, we can increase the disk size to 1TB to get 30K IOPS instead of the 10K with 350GB and the monthly price would be 5.8 USD.

A local SSD disk (375GB max. Per instance) would give us a greater amount of IOPS but its monthly price is 30 USD, the lowest price would be with a 3-year contract for 13.5 USD per month. Also, an SSD drive is not general purpose.

upvoted 4 times

  **rramos96** 7 months ago

Using GCP Pricing Calculator I get the following prices (us-central1 / n1-standard-n4: 97.09\$):

- [A]: Instance + 1 TB Zonal SSD Persistent Disk: 138.05\$

- [C]: Instance with Local SSD 375GB: 127.09\$



So, using Local SSD [C] is cheaper (+ better performance as @ESP\_SAP said).

upvoted 1 times

  **nightflyer** 7 months, 2 weeks ago

We go with the cheapest

upvoted 1 times

  **SSPC** 11 months, 2 weeks ago

I would go to the answer C

upvoted 1 times

  **vlodia** 1 year ago

D

Regional persistent disks have storage qualities that are similar to zonal persistent disks (standard and SSD). However, regional persistent disks provide durable storage and replication of data between two zones in the same region.

upvoted 2 times

  **DarioFama23** 1 year ago

Regional SSD is for HA hence we have disks across a region. For more IOPS throughput we need Local SSD

upvoted 6 times

Your Dataproc cluster runs in a single Virtual Private Cloud (VPC) network in a single subnet with range 172.16.20.128/25. There are no private IP addresses available in the VPC network. You want to add new VMs to communicate with your cluster using the minimum number of steps. What should you do?

- A. Modify the existing subnet range to 172.16.20.0/24.
- B. Create a new Secondary IP Range in the VPC and configure the VMs to use that range.
- C. Create a new VPC network for the VMs. Enable VPC Peering between the VMs’s VPC network and the Dataproc cluster VPC network.
- D. Create a new VPC network for the VMs with a subnet of 172.32.0.0/16. Enable VPC network Peering between the Dataproc VPC network and the VMs VPC network. Configure a custom Route exchange.

Correct Answer: A

- francisco\_guerra

Highly Voted

11 months, 3 weeks ago

I think is A,  
upvoted 25 times
- passmepls

11 months, 1 week ago

thank you  
upvoted 3 times
- ESP\_SAP

Highly Voted

11 months, 1 week ago

Correction.  
Correct Answers is (A):  
  
gcloud compute networks subnets expand-ip-range  
NAME  
gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork  
upvoted 17 times
- arsh1916

Most Recent

2 months, 1 week ago

A should be answer  
upvoted 2 times
- meh009

3 months, 1 week ago

A - you can modify existing subnet ranges to a larger size IF the new range encompasses the old range. In this case it will so A  
upvoted 2 times
- kopper2019

3 months, 2 weeks ago

A is the fastest  
upvoted 2 times
- GCP\_Student1

4 months, 1 week ago

A. Modify the existing subnet range to 172.16.20.0/24.  
upvoted 2 times
- raman23125

4 months, 2 weeks ago

answer a  
  
VPC networks do not have any IP address ranges associated with them. IP ranges are defined for the subnets. A network must have at least one subnet before you can use it. Auto mode VPC networks create subnets in each region automatically.  
upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

A is correct  
upvoted 1 times
- victory108

5 months, 3 weeks ago

A - Modify the existing subnet range to 172.16.20.0/24.  
upvoted 1 times
- andregrijp

5 months, 3 weeks ago

For minimum steps will just need to run this command:  
To expand the IP range of SUBNET to /24, run:  
gcloud compute networks subnets expand-ip-range SUBNET --region=REGION --prefix-length=24  
So A is better for this task  
upvoted 1 times

- 🗨️ 👤 **rvgcp** 6 months ago  
My choice is A.  
upvoted 1 times
- 🗨️ 👤 **gg\_robin** 6 months ago  
When it comes to A, I interpret the question as the dataproc cluster runs in a single subnet, not that the VPC only contains a single subnet.  
upvoted 1 times
- 🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago  
Guys, Anyone of you tried testing?  
if we chose A and if that works?  
why not C <  
upvoted 1 times
- 🗨️ 👤 **Bhagirathi** 8 months, 1 week ago  
you have confused everyone ...  
which one is really correct ? C ?? A ?? B ??  
upvoted 2 times
- 🗨️ 👤 **ladh** 7 months, 2 weeks ago  
u should try out. I tried out C and A.  
upvoted 1 times
- 🗨️ 👤 **Lieder** 8 months, 1 week ago  
I think A, too.  
But is there a typo or is it a red hering?  
172.16.20.0/24 is proposed in the answer. But shouldn't it read 172.16.20.128/24?  
upvoted 2 times
- 🗨️ 👤 **cuban123** 8 months, 1 week ago  
for /24 172.16.20.0/24 is the network. you gotta have full range that /24 offers  
upvoted 4 times
- 🗨️ 👤 **ritzheg** 8 months, 3 weeks ago  
A. you can expand the subnet, but you cannot shrink it. /25 gives you 128 IP, however extending it to /24 will increase it to 256.  
you can use this to calculate the cidr IP <https://www.ipaddressguide.com/cidr>  
upvoted 2 times
- 🗨️ 👤 **AZahid** 10 months ago  
Ans Is B  
upvoted 1 times

## Question #130

## Topic 1

You manage an App Engine Service that aggregates and visualizes data from BigQuery. The application is deployed with the default App Engine Service account.

The data that needs to be visualized resides in a different project managed by another team. You do not have access to this project, but you want your application to be able to read data from the BigQuery dataset. What should you do?

- A. Ask the other team to grant your default App Engine Service account the role of BigQuery Job User.
- B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer.
- C. In Cloud IAM of your project, ensure that the default App Engine service account has the role of BigQuery Data Viewer.
- D. In Cloud IAM of your project, grant a newly created service account from the other team the role of BigQuery Job User in your project.

**Correct Answer: B**

The Owner, Editor, and Viewer primitive roles include the BigQuery Admin (roles/bigquery.dataOwner), BigQuery Data Editor (roles/bigquery.dataEditor), and BigQuery Data Viewer (roles/bigquery.dataViewer) roles, respectively. This means the Owner, Editor, and Viewer primitive roles have BigQuery access as defined for the respective BigQuery roles.

Reference:

<https://cloud.google.com/bigquery/docs/access-control>

  **Hjameel** Highly Voted 11 months, 2 weeks ago

I think B is the answer  
upvoted 18 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (B):  
Sorry, I copied/pasted the the wrong statement.  
This is the proper explanation regarding to Big Query Data Viewer Role.  
The resource that you need to get access is in the other project.

roles/bigquery.dataViewer BigQuery Data Viewer  
When applied to a table or view, this role provides permissions to:

Read data and metadata from the table or view.  
This role cannot be applied to individual models or routines.

When applied to a dataset, this role provides permissions to:

Read the dataset's metadata and list tables in the dataset.  
Read data and metadata from the dataset's tables.  
When applied at the project or organization level, this role can also enumerate all datasets in the project. Additional roles, however, are necessary to allow the running of jobs.  
upvoted 9 times

  **[Removed]** 10 months, 1 week ago

A is correct, data viewer role does not allow you to execute query, that can be done if you are the user role.  
upvoted 2 times

  **[Removed]** 10 months, 1 week ago

The question states you want to aggregate and visualize data, that is run aggregate SQL on data before visualizing.  
upvoted 1 times

  **arsh1916** Most Recent 2 months, 1 week ago

B, simple  
upvoted 1 times

  **tavva\_prudhvi** 4 months ago

"you want your application to be able to read data from the BigQuery dataset", simply asks us to read the data from another project. So, dataViewer is sufficient for that. How can you execute the query without reading the data?  
upvoted 1 times

  **pca2b** 4 months ago

A  
<https://cloud.google.com/bigquery/docs/access-control>  
roles/bigquery.dataViewer: Additional roles, however, are necessary to allow the running of jobs.  
  
roles/bigquery.jobUser: Provides permissions to run jobs, including queries, within the project.

upvoted 1 times

  **pca2b** 4 months ago

not D: data is in the other project, not your project

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Does the jobUser "read" the data, and without able to read the data how can you query your data?

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Also, yes additional roles are necessary to run the jobs for Data Viewer role, hence we give the jobUser after that.

upvoted 1 times

  **JackGlemins** 5 months ago

I think A is right: the app need "aggregate and visualizes data".

Data Viewer just let you Read data, metadata and datasets.

BigQuery Job User Provides permissions to run jobs, including queries, within the project.

"bigquery.jobs.create"

resourcemanager.projects.get

resourcemanager.projects.list

<https://cloud.google.com/bigquery/docs/access-control>

upvoted 2 times

  **JackGlemins** 5 months ago

Adicional info: to Data Viewer Additional roles are necessary to allow the running of jobs.

upvoted 2 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Yes, you are right but what about the last line saying "the application should be able to read the data"??? jobUser doesn't have the access to read data. To query your data, you have to access the data first, then you can perform SQL queries as mentioned.

upvoted 1 times

  **GCP\_Student1** 5 months ago

B. Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer.

upvoted 3 times

  **nitinz** 5 months, 1 week ago

Correct answer is A, dataviewer can only view the data and list tables but can not run quires.

<https://cloud.google.com/bigquery/docs/access-control>

See the diff between the two roles. You app wont work if you go with data viewer.

upvoted 2 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

Just say me, how can you access the data without even reading it, does the job user provide you any functionality?

upvoted 1 times

  **victory108** 5 months, 2 weeks ago

B - Ask the other team to grant your default App Engine Service account the role of BigQuery Data Viewer.

upvoted 3 times

  **hiteshrup** 10 months, 3 weeks ago

Has to be B

upvoted 3 times

  **[Removed]** 10 months, 3 weeks ago

I agree that B is the way to go.

The job user role grants the ability to run queries, but it doesn't give the ability to read the data.

upvoted 3 times

  **pepepy** 10 months, 4 weeks ago

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork. Correct is B

upvoted 2 times

  **ESP\_SAP** 11 months, 1 week ago

Correct Answer is (B):

The resource that you need to get access is in the other project.

Correction.

Correct Answers is (A):

gcloud compute networks subnets expand-ip-range

NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork  
upvoted 1 times



## Question #131

## Topic 1

You need to create a copy of a custom Compute Engine virtual machine (VM) to facilitate an expected increase in application traffic due to a business acquisition.

What should you do?

- A. Create a Compute Engine snapshot of your base VM. Create your images from that snapshot.
- B. Create a Compute Engine snapshot of your base VM. Create your instances from that snapshot.
- C. Create a custom Compute Engine image from a snapshot. Create your images from that image.
- D. Create a custom Compute Engine image from a snapshot. Create your instances from that image.

**Correct Answer: D**

A custom image belongs only to your project. To create an instance with a custom image, you must first have a custom image.

Reference:

<https://cloud.google.com/compute/docs/instances/create-start-instance>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (D):

Preparing your instance for an image

You can create an image from a disk even while it is attached to a running VM instance. However, your image will be more reliable if you put the instance in a state that is easier for the image to capture. Use one of the following processes to prepare your boot disk for the image:

Stop the instance so that it can shut down and stop writing any data to the persistent disk.

If you can't stop your instance before you create the image, minimize the amount of writes to the disk and sync your file system.

Pause apps or operating system processes that write data to that persistent disk.

Run an app flush to disk if necessary. For example, MySQL has a FLUSH statement. Other apps might have similar processes.

Stop your apps from writing to your persistent disk.

Run `sudo sync`.

After you prepare the instance, create the image.

[https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images#prepare\\_instance\\_for\\_image](https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images#prepare_instance_for_image)  
upvoted 12 times

  **SSPC** Highly Voted 11 months, 2 weeks ago

I would go with the answer "D"

upvoted 7 times

  **Eshkrkrkr** 8 months, 3 weeks ago

D is wrong!

Creating a VM instance from a snapshot: If you backed up a boot persistent disk with a snapshot, you can use that snapshot to create an instance.

Keep in mind that if you plan to create many instances from the same boot disk snapshot, consider creating a custom image and creating instances from that image instead. - We need only one copy in the question! its B!

upvoted 2 times

  **saminlo** Most Recent 1 month, 1 week ago

B&D are both possible but to facilitate an increase in application traffic, likely multiple VM instances need to be created. You should use an image when creating multiple copies of VMs.

Hence D.

upvoted 1 times

  **arsh1916** 2 months, 1 week ago

D is answer

upvoted 1 times

  **pca2b** 4 months ago

B:

we just need to make 'a copy' of the VM, B works well for that.

not D: Had the question mentioned more copies, we would need to go the way of images...templates etc. D will work but not needed here.

upvoted 5 times

  **djgodzilla** 1 month, 2 weeks ago

custom images are better a fit if its for a new business workload you just acquired


upvoted 1 times



  **ShakthiGCP** 4 months, 4 weeks ago

Sorry... Going with Ans : D . just realised the Keyword 'Custom VM' in question. Which means you can perform only by Custom image and not by snapshot. Ans: D . <https://cloud.google.com/compute/docs/images>

upvoted 2 times

  **ShakthiGCP** 4 months, 4 weeks ago

Agreed. Ans is B. Well explained.

upvoted 2 times

  **ShakthiGCP** 4 months, 4 weeks ago

Both B and D works. but B is much faster to do.

upvoted 1 times

  **GCP\_Student1** 5 months ago

B. Create a Compute Engine snapshot of your base VM. Create your instances from that snapshot.

Since it is asking to create (a) copy of the VM instance, a single copy, then creating from snapshot is the right option. otherwise image is recommended.

upvoted 2 times

  **GCP\_Student1** 5 months ago

Creating a VM instance from a snapshot

If you backed up a boot persistent disk with a snapshot, you can use that snapshot to create a VM.

Keep in mind that if you plan to create more than one VM from the same boot disk snapshot, consider creating a custom image and creating VMs from that image instead. Boot disks created from custom images take less time to create than boot disks created from snapshots.

Creating a VM from a custom image

A custom image belongs only to your project. To create a VM with a custom image, you must first have a custom image.

upvoted 2 times

  **GCP\_Student1** 5 months ago

<https://cloud.google.com/compute/docs/instances/create-start-instance>

upvoted 2 times

  **GCP\_Student1** 5 months ago

An image is a complete backup of your server including all volumes.

A snapshot can be done from a specific volume (for example you have a server with a volume containing the OS and another one containing the application data, and you want to use different snapshot strategies on both volumes).

Creating snapshots of your volumes gives you total freedom which volumes you want to back up while images are convenient for full backups of your instance.

[https://www.scaleway.com/en/faq/what-is-the-difference-between-an-image-and-a-snapshot/#:~:text=An%20image%20is%20a%20complete,snapshot%20strategies%20on%20both%20volumes\).](https://www.scaleway.com/en/faq/what-is-the-difference-between-an-image-and-a-snapshot/#:~:text=An%20image%20is%20a%20complete,snapshot%20strategies%20on%20both%20volumes).)

upvoted 1 times

  **victory108** 5 months, 3 weeks ago

B - Create a Compute Engine snapshot of your base VM. Create your instances from that snapshot.

upvoted 3 times

  **DucSiu** 6 months ago

B is correct

upvoted 1 times

  **Pavan82** 6 months, 1 week ago

Correct Answer is (B): <https://geekflare.com/clone-google-cloud-vm/>

Two steps to copy/clone the VM, create a snapshot and then create an instance from that snapshot

upvoted 1 times

  **nherrerab** 6 months, 2 weeks ago

D is correct.

upvoted 1 times

  **Range2019** 7 months, 2 weeks ago

The correct is is B.

Both B & D can achieve the objective but the question said that you have create a copy to facilitate an expected increase in application traffic due to a business acquisition. That means that you CANNOT afford to shutdown your VM instance to create an custom image. Snapshot does not require a VM shutdown unlike custom image.

upvoted 4 times

  **wapry** 7 months ago

The image is created from snapshot so you don't need to shutdown the VM. In my opinion both B and D works but image is most suitable for new VM deploy because do not contain references of source VM. At the end, for me D is the correct one.

upvoted 2 times

  **magistrum** 7 months ago

D has one extra step than B, so B is simpler

upvoted 1 times

🗨️ 👤 **Morgan91** 7 months, 2 weeks ago

the virtual machine is already custom, the question says to make a copy of a custom VM.

B is correct

snapshot> VM

upvoted 1 times

🗨️ 👤 **Jignesh\_Gamdha** 8 months, 4 weeks ago

D is correct

<https://cloud.google.com/compute/docs/instances/create-start-instance#createsnapshot>

upvoted 1 times

🗨️ 👤 **Eshkrkrkr** 8 months, 3 weeks ago

D is wrong! Go read the link you paste here!

Creating a VM instance from a snapshot If you backed up a boot persistent disk with a snapshot, you can use that snapshot to create an instance.

Keep in mind that if you plan to create many instances from the same boot disk snapshot, consider creating a custom image and creating instances from that image instead. - We need only one copy in the question! its B!

upvoted 4 times

🗨️ 👤 **magistrum** 7 months ago

I guess wording makes a difference...if it said "You need to multiple copies..." D would have been more appropriate

upvoted 1 times

## Question #132

## Topic 1

You have deployed an application on a single Compute Engine instance. The application writes logs to disk. Users start reporting errors with the application. You want to diagnose the problem. What should you do?



- A. Navigate to Cloud Logging and view the application logs.
- B. Connect to the instance's serial console and read the application logs.
- C. Configure a Health Check on the instance and set a Low Healthy Threshold value.
- D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging.

**Correct Answer: D**



  **hiteshrup** Highly Voted 10 months, 3 weeks ago



Answer: D



App logs can't be visible to Cloud Logging until we install Cloud Logging Agent on GCE  
upvoted 14 times



  **hiteshrup** 10 months, 3 weeks ago  
Continuation of reasoning.

If Problem statement is not having this statement "The application writes logs to disk", then we might assume that application is writing logs on Cloud Logging with google-fluentd agent API library. However, problem statement is clearly mentioned that logs are writing down on disk, we need agent installed on GCE to fetch those logs from disk to Cloud Logging. If that is not desirable, then option B is left  
upvoted 3 times

  **hiteshrup** 10 months, 3 weeks ago  
(Correction) Answer is A after rethinking and doing some research by focusing words "App Engine", which has by default enabled Request Logs which has App logs on each request and those logs are enabled for Cloud Logging ..  
[https://cloud.google.com/appengine/docs/standard/python/logs#request\\_logs\\_vs\\_application\\_logs](https://cloud.google.com/appengine/docs/standard/python/logs#request_logs_vs_application_logs)  
upvoted 2 times

  **Eshkrkrkr** 8 months, 3 weeks ago  
Wrong! Request Logs has the LIST of App logs and ONLY associated with that request! Read the links you provide!  
upvoted 1 times

  **cuban123** 8 months, 1 week ago  
you must still install the agent:  
  
[https://cloud.google.com/error-reporting/docs/setup/compute-engine#using\\_logging](https://cloud.google.com/error-reporting/docs/setup/compute-engine#using_logging)  
upvoted 1 times



  **rezavage** 10 months, 2 weeks ago  
Cloud logging without agent only works for App engine as you stated . but the question is about the compute engine which has to be equipped first with Logging Agent in order to write logs into Cloud Logging. so based your assumption the correct answer is "D"  
upvoted 3 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (D):

In its default configuration, the Logging agent streams logs from common third-party applications and system software to Logging; review the list of default logs. You can configure the agent to stream additional logs; go to Configuring the Logging agent for details on agent configuration and operation.

It is a best practice to run the Logging agent on all your VM instances. The agent runs under both Linux and Windows. To install the Logging agent, go to Installing the agent.

<https://cloud.google.com/logging/docs/agent>  
upvoted 11 times

  **ESP\_SAP** 11 months, 1 week ago  
CORRECTION.  
Correct Answer is (A):

Activity logging is enabled by default for all Compute Engine projects.

You can see your project's activity logs through the Logs Viewer in the Google Cloud Console:

In the Cloud Console, go to the Logging page.  
Go to the Logging page

When in the Logs Viewer, select and filter your resource type from the first drop-down list.  
 From the All logs drop-down list, select `compute.googleapis.com/activity_log` to see Compute Engine activity logs.  
[https://cloud.google.com/compute/docs/logging/activity-logs#viewing\\_logs](https://cloud.google.com/compute/docs/logging/activity-logs#viewing_logs)

Besides:

Activity logs are provided as part of the Cloud Logging service. For more information about Logging in general, read the Cloud Logging documentation.

<https://cloud.google.com/compute/docs/logging/activity-logs>

upvoted 5 times

  **mexblood1** 11 months ago

Activity Logs do not include 2rd party application logs. Activity logs are more related to operations and changes in the infrastructure. This question is tricky, I think it's either D or B, because if it's only an application on a single instance, you can connect to the instance and read the application logs directly and you save the cost of logging agent.

upvoted 9 times

  **mexblood1** 11 months ago

Maybe I was assuming serial console is the same than system console, technically I guess they're not the same, hence I guess D will be my chosen answer.

upvoted 1 times

  **ESP\_SAP** 11 months, 1 week ago

Additional information about VM Image for AWS EC2:

The Logging agent streams logs from your VM instances and from selected third-party software packages to Cloud Logging. It is a best practice to run the Logging agent on all your VM instances.

The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances. The agent runs under both Linux and Windows.

If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page.

upvoted 2 times

  **magistrum** 7 months ago

This points to D then

upvoted 2 times

  **sapguru** 11 months, 1 week ago

Cloud logging enabled by default for compute engine

upvoted 1 times

  **hchafloque** Most Recent 1 month, 2 weeks ago

Could Cloud logging read log app files saved on disk or create new one for analysis?

upvoted 1 times

  **meh009** 3 months, 1 week ago

Answer is A

Reasoning behind this is because Activity logs (now audit logs) are on by default so no need to install/configure.

<https://cloud.google.com/compute/docs/logging/audit-logging>

upvoted 2 times

  **tavva\_prudhvi** 3 months, 3 weeks ago

It's D, clearly says in this article <https://cloud.google.com/logging/docs/agent/installation>

"The VM images for Compute Engine and Amazon Elastic Compute Cloud (EC2) don't include the Logging agent, so you must complete these steps to install it on those instances, and If your VMs are running in Google Kubernetes Engine or App Engine, the agent is already included in the VM image, so you can skip this page."

upvoted 1 times

  **pca2b** 4 months ago

B

application is already deployed on GCE.

writes logs to disk.

there are errors being reported.

would've been nice to have logging agent installed. but we're not there.

we need to troubleshoot current errors, so need to get on serial console (SSH works) and check the app logs written to disk.

default logs sent to Cloud Logging wont be of much use in app troubleshooting, we can follow best practices for future, but for now, read logs from disk --> B

upvoted 3 times

  **bnkt13** 4 months, 1 week ago

but if the application writes logs to disk, we need to look into de vm to see the logs. ans B

upvoted 2 times

  **SSunny** 4 months, 2 weeks ago

D - Compute Engine require Logging agent to be installed.  
 Except from link- <https://cloud.google.com/run/docs/logging>

A note about logging agents

If you've used Cloud Logging with certain Google Cloud products, such as Compute Engine, you may have used Cloud Logging logging agents. Cloud Run does not use logging agents because it has built-in support for log collection.

upvoted 2 times

  **EABDAJA** 4 months, 3 weeks ago



D is correct

upvoted 1 times

  **ShakthiGCP** 4 months, 4 weeks ago

Ans: D . Always best to have Logging Agent installed in the applications.

upvoted 3 times

  **ShakthiGCP** 4 months, 4 weeks ago

fluent-logger that we install is a Logging Agent . <https://cloud.google.com/error-reporting/docs/setup/compute-engine>

upvoted 1 times

  **ShakthiGCP** 4 months, 3 weeks ago

Ans:D looks correct. but i got a confusion after reading this... <https://cloud.google.com/logging/docs/agent/default-logs> ...The Logging agent is pre-configured to send logs from VM instances to Cloud Logging. If this is the case, then VM instance can send logs to Cloud Logging. Then Ans: A . anyone thinking about the same ?

upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

What you say here again confirms D being the right answer, What is your confusion then?

With A you can't see the logs, the Agent has to be installed first.

upvoted 1 times

  **GCP\_Student1** 5 months ago

D. Install and configure the Cloud Logging Agent and view the logs from Cloud Logging.

upvoted 2 times

  **register829** 5 months, 2 weeks ago

D


[https://cloud.google.com/logging/docs/agent/configuration#third-party\\_application\\_log\\_input\\_configuration](https://cloud.google.com/logging/docs/agent/configuration#third-party_application_log_input_configuration)

upvoted 1 times

  **victory108** 5 months, 3 weeks ago

D - Install and configure the Cloud Logging Agent and view the logs from Cloud Logging.

upvoted 1 times

  **boss lady** 5 months, 4 weeks ago

if the application runs on app engine or gke then no need of installing if its compute engine you have to install the agent.

upvoted 2 times

  **rvgcp** 6 months ago

correct answer is D. Agent needs to be installed to view the application logs.

upvoted 2 times

  **guid1984** 6 months ago

Correct Answer is D, monitoring and logging are not enabled by default. Just tested it by creating a VM and navigating to monitoring section, which says No logging and No monitoring set up. You need to run below command in order to set up logging agent on VM

`gcloud beta compute ssh instance-1 --project=<PROJECT_ID> --zone=<YOUR SUBNET ZONE> --command="curl -sSO`

`https://dl.google.com/cloudagents`

`/add-logging-agent-repo.sh && sudo bash add-logging-agent-repo.sh && sudo apt-get update && sudo apt-get install google-fluentd && sudo`

`apt-get install -y google-fluentd-catch-all-co`

`nfig && sudo service google-fluentd start"`

upvoted 3 times

  **Alexou** 6 months, 2 weeks ago

Hey, so what answer is it ? Do we need to install the Logging Agent ?

upvoted 1 times



An application generates daily reports in a Compute Engine virtual machine (VM). The VM is in the project corp-iot-insights. Your team operates only in the project corp-aggregate-reports and needs a copy of the daily exports in the bucket corp-aggregate-reports-storage. You want to configure access so that the daily reports from the VM are available in the bucket corp-aggregate-reports-storage and use as few steps as possible while following Google-recommended practices. What should you do?


- A. Move both projects under the same folder.
- B. Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.
- C. Create a Shared VPC network between both projects. Grant the VM Service Account the role Storage Object Creator on corp-iot-insights.
- D. Make corp-aggregate-reports-storage public and create a folder with a pseudo-randomized suffix name. Share the folder with the IoT team.

**Correct Answer: A**

Reference:

<https://cloud.google.com/billing/docs/onboarding-checklist>

-   **ESP\_SAP**

Highly Voted 

 11 months, 1 week ago



Correct Answer is (B):

Predefined roles



The following table describes Identity and Access Management (IAM) roles that are associated with Cloud Storage and lists the permissions that are contained in each role. Unless otherwise noted, these roles can be applied either to entire projects or specific buckets.


Storage Object Creator (roles/storage.objectCreator) Allows users to create objects. Does not give permission to view, delete, or overwrite objects.

<https://cloud.google.com/storage/docs/access-control/iam-roles#standard-roles>

upvoted 19 times
-   **ESP\_SAP** 11 months, 1 week ago



Basically, you are giving the permissions to the VM Service Account to create a copy of the daily report on the bucket that the other team has access.

upvoted 3 times
-   **francisco\_guerra**



Highly Voted 

 11 months, 3 weeks ago

i think is B



upvoted 13 times
-   **francisco\_guerra** 11 months, 3 weeks ago


Object creator cant see object so i think is D

upvoted 1 times
-   **lxgywil** 2 months, 3 weeks ago

VM doesn't need to see the ojects - just to create them. It's B:



The VM is located in project "corp-iot-insights" - give its SA the Storage Object Creator role for bucket "corp-aggregate-reports-storage" that is located in project "corp-aggregate-reports", where your team operates.

upvoted 3 times
-   **arsh1916**



Most Recent 

 2 months, 1 week ago



B, assign access is less step

upvoted 1 times
-   **pradr** 4 months, 2 weeks ago



From stackoverflow: Bucket names are globally unique, so your app will refer to an existing bucket in another project in the same way that it refers to buckets in its own project. Hence the shared VPC is not required to access the bucket. Just the IAM role.

upvoted 2 times
-   **EABDAJA** 4 months, 3 weeks ago

B is correct

upvoted 1 times
-   **GCP\_Student1** 5 months ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 2 times
-   **victory108** 5 months, 2 weeks ago

B - Grant the VM Service Account the role Storage Object Creator on corp-aggregate-reports-storage.

upvoted 1 times

🗨️ 👤 **Bhagirathi** 8 months, 1 week ago

why not A ?

upvoted 1 times

🗨️ 👤 **magistrum** 7 months ago

lots of steps vs B

upvoted 1 times

🗨️ 👤 **RockAJ** 9 months, 4 weeks ago

For me is B

upvoted 1 times

🗨️ 👤 **shan316** 10 months ago

Why the answer can't be C since the communication is between two different projects. If there is no VPC peering then there won't be any communication at the first place. What you guys think?

upvoted 1 times

🗨️ 👤 **Examan1** 5 months, 3 weeks ago

The VM is in the project corp-iot-insights. Your team operates only in the project corp-aggregate-reports and needs a copy of the daily exports in the bucket corp-aggregate-reports-storage. Hence you need Storage Object Creator on the bucket?

upvoted 1 times

🗨️ 👤 **Eshkrkrkr** 8 months, 3 weeks ago

We think that you are digging in the wrong direction. So you've shared a network, what next? How does C helps to solve the gestion problem?

upvoted 2 times

🗨️ 👤 **Ale1973** 10 months, 3 weeks ago

You are right, but to solve the scenario only is required that VM could write the report on the bucket.

upvoted 1 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

I would go with the answer B

upvoted 3 times

🗨️ 👤 **Hjameel** 11 months, 2 weeks ago

I think its B, all you need from the VM is to create a copy in the bucket where the users reside.

upvoted 4 times



## Question #134

## Topic 1

You built an application on your development laptop that uses Google Cloud services. Your application uses Application Default Credentials for authentication and works fine on your development laptop. You want to migrate this application to a Compute Engine virtual machine (VM) and set up authentication using Google- recommended practices and minimal changes. What should you do?

- A. Assign appropriate access for Google services to the service account used by the Compute Engine VM.
- B. Create a service account with appropriate access for Google services, and configure the application to use this account.
- C. Store credentials for service accounts with appropriate access for Google services in a config file, and deploy this config file with your application.
- D. Store credentials for your user account with appropriate access for Google services in a config file, and deploy this config file with your application.

**Correct Answer: B**

Reference:

<https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (B):

Best practices

In general, Google recommends that each instance that needs to call a Google API should run as a service account with the minimum permissions necessary for that instance to do its job. In practice, this means you should configure service accounts for your instances with the following process:

Create a new service account rather than using the Compute Engine default service account.

Grant IAM roles to that service account for only the resources that it needs.

Configure the instance to run as that service account.

Grant the instance the <https://www.googleapis.com/auth/cloud-platform> scope to allow full access to all Google Cloud APIs, so that the IAM permissions of the instance are completely determined by the IAM roles of the service account.

Avoid granting more access than necessary and regularly check your service account permissions to make sure they are up-to-date.

[https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best\\_practices](https://cloud.google.com/compute/docs/access/create-enable-service-accounts-for-instances#best_practices)

upvoted 31 times

  **cRobert** 7 months, 1 week ago

From your quote:



Configure the "instance" to run as that service account.

From answer B:

and configure the "application" to use this account.

You don't add service accounts to applications, ans A

upvoted 11 times

  **lxgywil** 2 months, 3 weeks ago

When you use a GCP service within your app (code), you have to use its client libraries. When you instantiate a client with client libraries you can pass it a Service Account key, which will define on behalf of which SA the client will be acting. That's how you can configure your app to use a particular service account.



E.g. [https://cloud.google.com/storage/docs/reference/libraries#using\\_the\\_client\\_library](https://cloud.google.com/storage/docs/reference/libraries#using_the_client_library)

upvoted 1 times

  **magistrum** 7 months ago

wording is the clue :)

upvoted 1 times

  **TAvenger** 5 months, 1 week ago

It's dirty play with words... All understand that we need custom SA, grant required permissions and attach this SA to the VM...

Why Google does this?

upvoted 2 times

  **filco72** Highly Voted 11 months, 3 weeks ago

I would choose: A. Assign appropriate access for Google services to the service account used by the Compute Engine VM. as there is no need to create a new service account.



























upvoted 10 times

  **Hjameel** 11 months, 2 weeks ago

I agree, there is no need to create a new service account

upvoted 6 times



-   **gcpengineer** Most Recent 1 day, 19 hours ago  
question asked for minimal effort and B says configure the app, which rules out B.  
A is the answer  
upvoted 2 times
-   **pYWORLD** 2 days, 20 hours ago  
You don't need to create another SA for this application. If you will put the application inside of the Compute Engine instance then it will use the SA from the Compute Engine SA.  
  
Answer from my perspective is: A  
upvoted 1 times
-   **associatecloudexamuser** 2 weeks, 3 days ago  
I think the correct answer is A. We can supply a service account with required permissions to access GCP resources while creating a new VM instance. Configuring application to use service account means, we again use GOOGLE\_APPLICATION\_CREDENTIALS of the service account. Since VM instance is already inside GCP, that is not necessary.  
upvoted 1 times
-   **arsh1916** 2 months, 1 week ago  
B, assign access to service account  
upvoted 1 times
-   **Biju1** 3 months, 1 week ago  
A is correct. 100 %. No need to create new account  
upvoted 4 times
-   **pradr** 4 months, 1 week ago  
Source: <https://cloud.google.com/iam/docs/service-accounts>  
A service account is a special kind of account used by an application or a virtual machine (VM) instance, not a person. Applications use service accounts to make authorized API calls, authorized as either the service account itself, or as Google Workspace or Cloud Identity users through domain-wide delegation.  
  
So I think this means you can attach a service account to an application  
  
Also,  
[https://cloud.google.com/iam/docs/service-accounts#application\\_default\\_credentials](https://cloud.google.com/iam/docs/service-accounts#application_default_credentials)  
  
Application Default Credentials is a tool that Google Cloud Client Libraries use to automatically discover service account credentials. You can specify a service account key in an environment variable, and Application Default Credentials automatically uses that service account key.  
  
Which I guess means that we could allow the application to use the specific key of the new service account that was created. And the best practice is to create a new service account so as to isolate the application work from the VM's service account which comes with editor role. So in my opinion it would be (B)  
upvoted 4 times
-   **pca2b** 4 months ago  
B, I agree  
upvoted 1 times
-   **Koen1978** 4 months, 3 weeks ago  
Not B. A service account is not added to an application but to an instance  
upvoted 1 times
-   **tavva\_prudhvi** 3 months, 2 weeks ago  
It can be used by any application to access other services, but NOT A PERSON!!! That's the main reason we use SA for app-to-app communication like making API calls.  
upvoted 1 times
-   **GCP\_Student1** 5 months ago  
B. Create a service account with appropriate access for Google services, and configure the application to use this account.  
upvoted 2 times
-   **victory108** 5 months, 3 weeks ago  
A - Assign appropriate access for Google services to the service account used by the Compute Engine VM.  
upvoted 2 times
-   **rramos96** 7 months ago  
I think that those answers are a bit confusing.  
Best practices recommend to create a new service account (1) with minimum required roles (2) and use it in the VM (3).  
  
- [B]: Creates a new SA (1), set the minimum roles (2) but it configures the app instead of VM.  
- [A]: We can assume that we have created a new SA (1) (or not), set the minimum roles (2) and use it in the VM (3).  
  
So I think that the best answer is [A], having a new SA (3/3) or using default SA (2/3), as it fills 3/3 or 2/3 of best practices requirements.  
upvoted 2 times
-   **vlodia** 10 months ago

Answer is B.

Passing credentials manually

If your application runs outside Google Cloud environments that provide a default service account, you need to manually create one. Then you can create one or more service account keys, which are credentials associated with the service account. Service account keys can then be manually passed to your application.

<https://cloud.google.com/docs/authentication/production#manually>

upvoted 2 times

  **Sheeda** 11 months, 1 week ago

The question says that it is using ADC. So ADC is also used with a service account.

<https://www.jhanley.com/google-cloud-application-default-credentials/>

Question says, using minimal changes... so why should we create new service account ??

See this para from the site

[https://cloud.google.com/iam/docs/service-accounts#application\\_default\\_credentials](https://cloud.google.com/iam/docs/service-accounts#application_default_credentials)

Application default credentials are a mechanism to make it easy to use service accounts when operating inside and outside Google Cloud, as well as across multiple Google Cloud projects. The most common use case is testing code on a local machine, and then moving to a development project in Google Cloud, and then moving to a production project in Google Cloud. Using Application Default Credentials ensures that the service account works seamlessly; when testing on your local machine, it uses a locally-stored service account key, but when running on Compute Engine, it uses the project's default Compute Engine service account. See Application Default Credentials for more information.

upvoted 1 times

  **Eshkrkrkr** 8 months, 3 weeks ago

Sheeda, read what you paste here: but when running on Compute Engine, it uses the PROJECT'S default Compute Engine service account, so it's refers to a customly created service account in dev project, not general Compute Engine default service account. Again its B! Oh my god, ppl, just read the damn docs you paste here!

upvoted 6 times

  **Examan1** 5 months ago

Using Application Default Credentials ensures that the service account works seamlessly. When testing on your local machine, it uses a locally-stored service account key, but when running on Compute Engine, it uses the project's default Compute Engine service account. So we have to provide access to the service account used by the compute engine VM and not the service account used by the application.

upvoted 1 times

  **[Removed]** 10 months, 3 weeks ago

I agree with you on this. Why create a new service account when you can just use the service account that the VM is currently using. Note that in the answer choice it does not say that the VM is using the default service account, but that it is using a service account.

Moreover, according to ESP\_SAP's notes, while you should create a new service account and only assign relevant permissions, it also says that the applications running on the VM will use whatever service account the VM has assigned. There is no configuration needed in this scenario, unlike what's implied by option B.

upvoted 2 times

  **mexblood1** 11 months ago


It also says using Google recommended practices, so I guess that leaves B as the correct answer.

upvoted 1 times

  **francisco\_guerra** 11 months, 3 weeks ago

I think is b because you dont need to add credentials to the default service account that any other has access.

upvoted 3 times

  **SSPC** 11 months, 2 weeks ago

I agree with you, because it don't need to add credentials to the default

upvoted 1 times

  **tanito83** 1 month, 2 weeks ago

I suppose, you mean A. It's name the default account.

upvoted 1 times

You need to create a Compute Engine instance in a new project that doesn't exist yet. What should you do?

- A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.
- B. Enable the Compute Engine API in the Cloud Console, use the Cloud SDK to create the instance, and then use the --project flag to specify a new project.
- C. Using the Cloud SDK, create the new instance, and use the --project flag to specify the new project. Answer yes when prompted by Cloud SDK to enable the Compute Engine API.
- D. Enable the Compute Engine API in the Cloud Console. Go to the Compute Engine section of the Console to create a new instance, and look for the Create In A New Project option in the creation form.

Correct Answer: B

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (A):

Quickstart: Creating a New Instance Using the Command Line

Before you begin

1. In the Cloud Console, on the project selector page, select or create a Cloud project.
2. Make sure that billing is enabled for your Google Cloud project. Learn how to confirm billing is enabled for your project.

To use the gcloud command-line tool for this quickstart, you must first install and initialize the Cloud SDK:

1. Download and install the Cloud SDK using the instructions given on Installing Google Cloud SDK.
2. Initialize the SDK using the instructions given on Initializing Cloud SDK.

To use gcloud in Cloud Shell for this quickstart, first activate Cloud Shell using the instructions given on Starting Cloud Shell.

<https://cloud.google.com/ai-platform/deep-learning-vm/docs/quickstart-cli#before-you-begin>

upvoted 24 times

  **filco72** Highly Voted 11 months, 3 weeks ago

I would choose A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

as first I need to create a project. Instance creation cannot automatically create a project.

upvoted 11 times

  **arsh1916** Most Recent 2 months, 1 week ago



A is correct

upvoted 2 times

  **kopper2019** 3 months, 1 week ago



A the way to go

upvoted 1 times

  **pondai** 4 months, 1 week ago

Vote A

upvoted 1 times

  **Devgela** 4 months, 2 weeks ago

A for me

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

A is correct

upvoted 1 times

  **neerajgoyal** 4 months, 3 weeks ago

A is the right answer

upvoted 1 times

  **GCP\_Student1** 5 months ago

A. Using the Cloud SDK, create a new project, enable the Compute Engine API in that project, and then create the instance specifying your new project.

upvoted 2 times

  **JackGlemins** 5 months, 1 week ago

I think is A: <https://cloud.google.com/sdk/gcloud/reference/projects/create>

upvoted 1 times

🗉 👤 **lutoa** 5 months, 2 weeks ago  
yep has to be A, new project must be created first  
upvoted 1 times

🗉 👤 **RockAJ** 9 months, 4 weeks ago  
A for me  
upvoted 1 times

🗉 👤 **hiteshrup** 10 months, 3 weeks ago  
Answer: A

Project creation is always a first step as we need to link billing account before creating any new resources. So Answer is A  
upvoted 5 times

🗉 👤 **stepkurniawan** 11 months, 1 week ago  
2 things that you need to create VM:  
1. a project  
2. Compute API  
So answer A is correct  
upvoted 2 times

🗉 👤 **Hjameel** 11 months, 2 weeks ago  
Answer B, first you need to create a project and then create the instance.  
upvoted 1 times

🗉 👤 **Hjameel** 11 months, 2 weeks ago  
I mean A, Project first then the instance creation  
upvoted 4 times

## Question #136

## Topic 1

Your company runs one batch process in an on-premises server that takes around 30 hours to complete. The task runs monthly, can be performed offline, and must be restarted if interrupted. You want to migrate this workload to the cloud while minimizing cost. What should you do?

- A. Migrate the workload to a Compute Engine Preemptible VM.
- B. Migrate the workload to a Google Kubernetes Engine cluster with Preemptible nodes.
- C. Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.
- D. Create an Instance Template with Preemptible VMs On. Create a Managed Instance Group from the template and adjust Target CPU Utilization. Migrate the workload.

**Correct Answer: B**

Reference:

<https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool>

  **juliandm** Highly Voted 11 months, 1 week ago



i understand preemptible as a no-go because of "must be restarted if interrupted" here meaning "starting from scratch" . So C seems right  
upvoted 25 times

  **jcloud965** 1 week, 6 days ago

I agree, C.  
you won't run 30 hours job on preemptible instances that can be stopped at any time and can't run more than 24 hours.  
If the job could be splitted, then preemptible VM is an option.  
upvoted 1 times

  **stepkurniawan** Highly Voted 11 months, 1 week ago

Preemptible will be perfect for a batch job that takes less than 24 hours. But it's not in this case.  
upvoted 9 times

  **Linus11** 2 months, 4 weeks ago

What if it is a managed group of Pre emptible instances like in D. If one instance stops, another instance will take over.

I choose D.  
upvoted 3 times

  **sanhoo** 1 month, 3 weeks ago

is there an option to specify Pre emptible instances while creating template? I couldn't find that. If so then D can't be true  
upvoted 1 times

  **djgodzilla** 1 month, 2 weeks ago

Yes under management> Availability policy > preemptibility ON/OFF  
upvoted 2 times

  **AD\_0525** Most Recent 1 month, 1 week ago

must be restarted implies, the job is restartable. So valid options now A and D.  
Now considering there is nothing in the questions on load sharing or the resource utilization, I will definitely go for A.  
upvoted 1 times

  **AD\_0525** 1 month, 2 weeks ago

Batch, restartable, cost effective - It has to be preemptible.  
MIG is cheaper to K8 with preemptible instances. So will go with option D  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago

I selected A.  
upvoted 1 times

  **A\_ksh\_a\_Y** 2 months, 1 week ago



Here it clearly states :  
1) task runs monthly,  
2) can be performed offline,  
and  
3) must be restarted if interrupted

Pre-emptible serves this purpose because :

- Preemptible VMs are priced lower
- Since job can be offline and it can be re-started (points 2 and 3 above) with pre-emptible, it would resume where it last stopped with new pre-emptible instances.

Hence Answer seems 'A'

upvoted 2 times

  **lxgywil** 2 months, 3 weeks ago

To me it seems the correct answer is "A".

"Batch processing jobs can run on preemptible instances. If some of those instances stop during processing, the job slows but does not completely stop. Preemptible instances complete your batch processing tasks without placing additional workload on your existing instances and without requiring you to pay full price for additional normal instances."

Since MIG with preemptible VMs can create or add new preemptible instances, option "C" can be more costly.

Ref: <https://cloud.google.com/compute/docs/instances/preemptible>

upvoted 2 times

  **Jacky\_YO** 3 months, 3 weeks ago

ANS : C

1. runs one batch process in an on-premises server
2. hat takes around 30 hours to complete. The task runs monthly
3. can be performed offline
4. must be restarted if interrupted
5. minimizing cost

#Compute Engine always stops preemptible instances after they run for 24 hours.

Preemptible VM instances

<https://cloud.google.com/compute/docs/instances/preemptible#testing-preemption-settings>

upvoted 6 times

  **yuvarajmrec** 4 months ago



i think it's D because preemptible VM's are 80% cheaper when compared to normal VM . And moreover using instance group , it will restart preemptible VM when they shut down

upvoted 1 times

  **tavva\_prudhvi** 4 months ago

but, they last only upto 24hrs what about that, and after restarting it it needs to be activated from the scratch.

upvoted 1 times

  **Linus11** 2 months, 4 weeks ago

Batch processes dont need to start from scratch, after restart. It can restart, from the previous point where it stopped.'

The very mentioning of "Batch process and Restart" indicates Preemptibles play here. not compute enginer. 30 hour continuous execution by a single Vm not needed here.

upvoted 1 times

  **Devgela** 4 months, 2 weeks ago


C for me

upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago

C is correct

upvoted 1 times

  **Koen1978** 4 months, 3 weeks ago


Answer is D. It's the only way to restart the batch job automatically

upvoted 2 times

  **tavva\_prudhvi** 3 months, 3 weeks ago

30 hours bro? what about that?

upvoted 1 times

  **Linus11** 2 months, 4 weeks ago



Batch processes dont need to start from scratch, after restart. It can restart, from the previous point where it stopped.'

upvoted 2 times

  **GCP\_Student1** 5 months ago

C - Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.

upvoted 2 times

  **nitinz** 5 months, 1 week ago


Guys read the question. If stopped need to restart \*AND\* it takes 30 hours. Preemptible will stop in 24 hours and restart will happen - meaning JOB will never finish. Be it VM or K8s. The answer is C.

upvoted 5 times

  **jmfmeireles** 5 months, 3 weeks ago

Why not D? MIGs are free, as you can see on the following link: <https://cloud.google.com/compute/docs/instance-groups#pricing>

"There is no additional charge for using managed or unmanaged instance groups. You are charged based on the resources that your group uses."  
upvoted 1 times

  **victory108** 5 months, 3 weeks ago

C - Migrate the workload to a Compute Engine VM. Start and stop the instance as needed.  
upvoted 1 times

  **CloudPadawan** 6 months, 1 week ago

<https://cloud.google.com/compute/docs/instances/preemptible#testing-preemption-settings>  
Dots 4, google says that Preemptible instances can't live migrate to a regular VM instance, or be set to automatically restart..  
I will choose C  
upvoted 1 times



You are developing a new application and are looking for a Jenkins installation to build and deploy your source code. You want to automate the installation as quickly and easily as possible. What should you do?

- A. Deploy Jenkins through the Google Cloud Marketplace.
- B. Create a new Compute Engine instance. Run the Jenkins executable.
- C. Create a new Kubernetes Engine cluster. Create a deployment for the Jenkins image.
- D. Create an instance template with the Jenkins executable. Create a managed instance group with this template.

**Correct Answer:** C

Reference:

<https://cloud.google.com/solutions/jenkins-on-kubernetes-engine>

  **filco72** Highly Voted 11 months, 3 weeks ago

I would choose A. Deploy Jenkins through the Google Cloud Marketplace.  
as this is a well known opportunity on the GCP Marketplace  
upvoted 16 times

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (A):

Installing Jenkins  
In this section, you use Cloud Marketplace to provision a Jenkins instance. You customize this instance to use the agent image you created in the previous section.

Go to the Cloud Marketplace solution for Jenkins.

Click Launch on Compute Engine.

Change the Machine Type field to 4 vCPUs 15 GB Memory, n1-standard-4.

Machine type selection for Jenkins deployment.

Click Deploy and wait for your Jenkins instance to finish being provisioned. When it is finished, you will see:

Jenkins has been deployed.


[https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine#installing\\_jenkins](https://cloud.google.com/solutions/using-jenkins-for-distributed-builds-on-compute-engine#installing_jenkins)  
upvoted 10 times

  **vinaybattula** Most Recent 1 month, 3 weeks ago

Option A is the quickest way.  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago

A is fastest way  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago



A is the Answer  
upvoted 1 times

  **GCP\_Student1** 5 months ago


A. Deploy Jenkins through the Google Cloud Marketplace.  
upvoted 1 times

  **nitinz** 5 months, 1 week ago

A all the way. Fast  
upvoted 1 times













  **Bhagirathi** 8 months, 1 week ago

should be A  
upvoted 1 times

  **bosslady** 5 months, 4 weeks ago

how many percent?  
upvoted 5 times



-   **berserkersyco** 11 months ago  
For me it looks 'D'. Here we are talking about automating the process.  
upvoted 1 times
-   **[Removed]** 10 months, 2 weeks ago  
Jenkins is a single tool for the environment you don't need a swarm of that. As said by Ale1973 we are not looking to automate the Vm creation.  
upvoted 1 times
-   **Ale1973** 10 months, 3 weeks ago  
We are not talking about automating all the deployment process, only the Jenkin's installation.  
upvoted 2 times
-   **stepkurniawan** 11 months, 1 week ago  
I'm a simple guy, I see quick and easy, I will search for marketplace directly  
upvoted 6 times
-   **SSPC** 11 months, 2 weeks ago  
A is the best option I think  
upvoted 2 times
-   **Hjameel** 11 months, 2 weeks ago  
A is the best answer.  
upvoted 4 times

## Question #138

## Topic 1

You have downloaded and installed the gcloud command line interface (CLI) and have authenticated with your Google Account. Most of your Compute Engine instances in your project run in the europe-west1-d zone. You want to avoid having to specify this zone with each CLI command when managing these instances.

What should you do?

- A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.
- B. In the Settings page for Compute Engine under Default location, set the zone to europe-west1-d.
- C. In the CLI installation directory, create a file called default.conf containing zone=europe-west1-d.
- D. Create a Metadata entry on the Compute Engine page with key compute/zone and value europe-west1-d.

**Correct Answer: C**

Reference:

<https://cloud.google.com/compute/docs/gcloud-compute>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (A):

Change your default zone and region in the metadata server

Note: This only applies to the default configuration.

You can change the default zone and region in your metadata server by making a request to the metadata server. For example:

```
gcloud compute project-info add-metadata \
--metadata google-compute-default-region=europe-west1,google-compute-default-zone=europe-west1-b
```

The gcloud command-line tool only picks up on new default zone and region changes after you rerun the gcloud init command. After updating your default metadata, run gcloud init to reinitialize your default configuration.

[https://cloud.google.com/compute/docs/gcloud-compute#change\\_your\\_default\\_zone\\_and\\_region\\_in\\_the\\_metadata\\_server](https://cloud.google.com/compute/docs/gcloud-compute#change_your_default_zone_and_region_in_the_metadata_server)  
upvoted 16 times

  **Examan1** 5 months ago

Using gcloud config you can set the zone in your active configuration only. This setting does not apply to other gcloud configurations and does not become the default for the project.

Ref: <https://cloud.google.com/sdk/gcloud/reference/config/set>

So I believe correct answer is B as per <https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>

In the Cloud Console, go to the Settings page.

From the Zone drop-down menu, select a default zone.

upvoted 1 times

  **tavva\_prudhvi** 3 months, 3 weeks ago

bro, it mentioned going into the console settings, not the compute engine settings!

To change your default region or zone:

In the Cloud Console, go to the Settings page.

Go to the Settings page

From the Region drop-down menu, select a default region.

From the Zone drop-down menu, select a default zone.

upvoted 1 times

  **jcloud965** 1 week, 6 days ago

This setting in the Cloud Console won't be taken into account for gcloud on your active config

upvoted 1 times

  **xtian2900** 10 months, 2 weeks ago

does your comment imply that the answer is D ? i'm confused

upvoted 1 times

  **SSPC** Highly Voted 11 months, 2 weeks ago

I would go with the answer A

upvoted 9 times

- 🗲️ 👤 **jcloud965** Most Recent 1 week, 6 days ago  
A is correct.  
"gcloud set config compute/zone" will be applied only for the active config BUT we assume there is only one project here.  
upvoted 1 times
- 🗲️ 👤 **jcloud965** 1 week, 6 days ago  
A is correct.  
"gcloud set config compute/zone" will be applied only for the active config BUT we assume there is only one project here.  
upvoted 1 times
- 🗲️ 👤 **JieHeng** 1 month ago  
Should be A, "You have downloaded and installed the gcloud command line interface (CLI)", now refer to [https://cloud.google.com/compute/docs/gcloud-compute#set\\_default\\_zone\\_and\\_region\\_in\\_your\\_local\\_client](https://cloud.google.com/compute/docs/gcloud-compute#set_default_zone_and_region_in_your_local_client) on how to "Set default region and zone in your local client"  
upvoted 1 times
- 🗲️ 👤 **JH86** 1 month, 2 weeks ago  
The answer SHOULD BE B. i have looked at this questions multiple times and has initially selected A due to the high upvote and the google doc links provided.  
  
A - you can perform a gcloud command to set a default zone but the catch is here - the answer has indicated gcloud config instead of GLOUD COMPUTE. the command that allows you to set default region and zone is GLOUD COMPUTE.  
<https://cloud.google.com/compute/docs/gcloud-compute>  
  
B - the google doc provides 2 methods to changing default region/zone. One is using the gcloud command (gcloud compute) OR using the console method, by changing the settings in the cloud console.  
<https://cloud.google.com/compute/docs/regions-zones/changing-default-zone-region#console>  
upvoted 1 times
- 🗲️ 👤 **vinaybattula** 1 month, 3 weeks ago  
we can set the default zone using command -- gcloud config set compute/region europe-west1-b.  
Option A is correct.  
upvoted 1 times
- 🗲️ 👤 **arsh1916** 2 months, 1 week ago  
A, asking about CLI command  
upvoted 1 times
- 🗲️ 👤 **arsh1916** 2 months, 3 weeks ago  
Ans C, since this will help you to avoid mentioning every time default zone.  
upvoted 1 times
- 🗲️ 👤 **arsh1916** 2 months, 3 weeks ago  
It should be C, bcz it's saying everytime when you you work with CLI command need to avoid mentioning zone.  
upvoted 1 times
- 🗲️ 👤 **kopper2019** 3 months, 1 week ago  
It's A since it is talking about CLI  
upvoted 1 times
- 🗲️ 👤 **Jacky\_YO** 3 months, 3 weeks ago  
ANS : A  
  
A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.  
upvoted 1 times
- 🗲️ 👤 **GCP\_Student1** 5 months ago  
A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.  
upvoted 2 times
- 🗲️ 👤 **andregrijp** 5 months, 3 weeks ago  
gcloud config set compute/region <ZONE> so looks like is A  
upvoted 1 times
- 🗲️ 👤 **DucSiu** 5 months, 4 weeks ago  
A. Set the europe-west1-d zone as the default zone using the gcloud config subcommand.  
upvoted 2 times
- 🗲️ 👤 **Bhagirathi** 7 months, 1 week ago  
A 200%  
upvoted 2 times
- 🗲️ 👤 **Faraz2** 7 months, 3 weeks ago  
Answer is A  
upvoted 1 times

## Question #139

## Topic 1

The core business of your company is to rent out construction equipment at large scale. All the equipment that is being rented out has been equipped with multiple sensors that send event information every few seconds. These signals can vary from engine status, distance traveled, fuel level, and more. Customers are billed based on the consumption monitored by these sensors. You expect high throughput  $\lambda$  up to thousands of events per hour per device  $\lambda$  and need to retrieve consistent data based on the time of the event. Storing and retrieving individual signals should be atomic. What should you do?

- A. Create a file in Cloud Storage per device and append new data to that file.
- B. Create a file in Cloud Filestore per device and append new data to that file.
- C. Ingest the data into Datastore. Store data in an entity group based on the device.
- D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.

**Correct Answer: D**

  **hiteshrup** Highly Voted  10 months, 3 weeks ago

Answer: D

Keyword need to look for

- "High Throughput",
- "Consistent",
- "Property based data insert/fetch like engine status, distance traveled, fuel level, and more." which can be designed in column,
- "Large Scale Customer Base + Each Customer has multiple sensor which send event in seconds" This will go for petabytes situation,
- Export data based on the time of the event.
- Atomic

o BigTable will fit all requirement.

o DataStore is not fully Atomic

o CloudStorage is not a option where we can export data based on time of event. We need another solution to do that

o Firestore can be used with MobileSDK.

So go with Option D: Big Table

upvoted 20 times

  **[Removed]** 10 months, 1 week ago

Its File store not firestore. But still, the argument is correct here as file store will not be automatic like cloud store that no SQL.



If it would be Firestore I would have gone with it, then big data for throughout v/s cost.

upvoted 1 times

  **Hjameel** Highly Voted  11 months, 2 weeks ago

D is the best answer , Cloud Bigtable

upvoted 6 times

  **har\_riy** 5 months, 2 weeks ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 4 times

  **CompteBidon91** Most Recent  2 weeks, 4 days ago

D,

Keyword : IOT, Database, large scale = Big Table

upvoted 1 times

  **vinaybattula** 1 month, 3 weeks ago

D is correct. If Data is huge, unstructured and related time Bigtable is best option

upvoted 1 times

  **arsh1916** 2 months, 1 week ago



D is ans 100%

upvoted 1 times

  **GCP\_Student1** 5 months ago

D. Ingest the data into Cloud Bigtable. Create a row key based on the event timestamp.



upvoted 3 times

  **har\_riy** 5 months, 2 weeks ago

Simple analogy.

Information every few seconds --> Time Series --> Big Table

upvoted 3 times

  **har\_riy** 5 months, 2 weeks ago



Simple analogy.  
Inforamtion every few seconds --> Time Series --> Big Table  
upvoted 1 times

  **Bhagirathi** 7 months, 1 week ago

D 200%  
upvoted 3 times

  **stepkurniawan** 11 months, 1 week ago

Cloud BigTable if it is required for something BIG and Atomic...  
upvoted 4 times

  **SSPC** 11 months, 1 week ago

I would go with the answer "D". <https://galvarado.com.mx/post/comparaci%C3%B3n-de-bases-de-datos-en-google-cloud-datastore-vs-bigtable-vs-cloud-sql-vs-spanner-vs-bigquery/>  
upvoted 3 times

You are asked to set up application performance monitoring on Google Cloud projects A, B, and C as a single pane of glass. You want to monitor CPU, memory, and disk. What should you do?

- A. Enable API and then share charts from project A, B, and C.
- B. Enable API and then give the metrics.reader role to projects A, B, and C.
- C. Enable API and then use default dashboards to view all projects in sequence.
- D. Enable API, create a workspace under project A, and then add projects B and C.

Correct Answer: C

- jlclaude

Highly Voted

11 months, 2 weeks ago

D. workspaces is made for monitoring multiple projects.

upvoted 22 times
- Hjameel

Highly Voted

11 months, 2 weeks ago

D , Workspace to monitor multiple projects.

upvoted 8 times
- Khoka

8 months, 2 weeks ago

<https://cloud.google.com/monitoring/workspaces>

upvoted 1 times
- vinaybattula

Most Recent

1 month, 3 weeks ago

Option D is correct. For monitoring first thing we should create workspace for one project and then can add more projects to the sme workspace for monitoring.

upvoted 1 times
- arsh1916

2 months, 1 week ago

D i selected

upvoted 1 times
- EABDAJA

4 months, 3 weeks ago

D is correct

upvoted 1 times
- GCP\_Student1

5 months ago

D. Enable API, create a workspace under project A, and then add project B and C.

upvoted 2 times
- hicham

5 months, 1 week ago

I choose D. we always need Workspace to monitor several projects

upvoted 1 times
- victory108

5 months, 3 weeks ago

D - Enable API, create a workspace under project A, and then add project B and C.

upvoted 2 times
- DucSiu

5 months, 4 weeks ago

D. Enable API, create a workspace under project A, and then add project B and C.

upvoted 1 times
- sarahf

7 months, 2 weeks ago

It is weird that none of the alternatives follow the Google best practice as stated both in <https://cloud.google.com/monitoring/workspaces#monitored-projects> and <https://cloud.google.com/monitoring/workspaces/create#multi-project-ws>

"If you plan to use a Workspace to monitor more than the host project, then the best practice is to use a new, empty Google Cloud project to host the Workspace and then to add the projects and AWS accounts you want to monitor to your Workspace."



"Create a Google Cloud project to be the Workspace host project. This project should be empty of resources. That is, it doesn't have any VM instances, Cloud Spanner tables, or other resources."



Both state that you should create a new empty project to host a Workspace, which none of the alternatives say. Or maybe I am not understanding alternative C correctly. Alternative D works but it is not best practice.



upvoted 3 times
- magistrum

6 months, 4 weeks ago

I guess the question didn't ask for best practice  
upvoted 1 times

  **Bhagirathi** 7 months, 3 weeks ago  
I am choosing D.  
upvoted 1 times

  **Bhagirathi** 8 months, 1 week ago  
how it can be C?  
I must be D...  
upvoted 1 times

  **dat987** 9 months ago  
D is correct? i think C  
upvoted 2 times

You created several resources in multiple Google Cloud projects. All projects are linked to different billing accounts. To better estimate future charges, you want to have a single visual representation of all costs incurred. You want to include new cost data as soon as possible. What should you do?

- A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.
- B. Visit the Cost Table page to get a CSV export and visualize it using Data Studio.
- C. Fill all resources in the Pricing Calculator to get an estimate of the monthly cost.
- D. Use the Reports view in the Cloud Billing Console to view the desired cost information.

**Correct Answer:** A

Reference:  
<https://cloud.google.com/billing/docs/how-to/visualize-data>

**Zakaullah**

Highly Voted

11 months, 1 week ago

Highly recommended website for exam prep.just passed the exam.  
upvoted 20 times

**IRobot**

11 months, 1 week ago

Hope the best for my exam tomorrow  
upvoted 3 times

**[Removed]**

10 months, 3 weeks ago

how did it go for you? I sit mine tomorrow.  
upvoted 2 times

**naveedpk00**

Most Recent

11 months, 1 week ago

Salam Zaka, Can you please share correct answers with me? Please sent it naveedpk00@hotmail.com Much Appreciated.  
upvoted 1 times

**arsh1916**

Most Recent

2 months, 1 week ago

A is best answer  
upvoted 4 times

**EABDAJA**

4 months, 3 weeks ago

A is correct  
upvoted 3 times

**GCP\_Student1**

5 months ago

A. Configure Billing Data Export to BigQuery and visualize the data in Data Studio.  
upvoted 4 times

**nitinz**

5 months, 1 week ago

A, Any time you see question about billing and has BQ - 99.99% of time correct answer is that has BQ in it.  
upvoted 1 times

**victory108**

5 months, 3 weeks ago

A - Configure Billing Data Export to BigQuery and visualize the data in Data Studio.  
upvoted 2 times

**sarahf**

7 months, 2 weeks ago

I was thinking A first.  
  
<https://cloud.google.com/billing/docs/how-to/export-data-bigquery>  
  
"Cloud Billing export to BigQuery enables you to export detailed Google Cloud billing data (such as usage, cost estimates, and pricing data) automatically throughout the day to a BigQuery dataset that you specify."  
  
It also says that:  
"Be aware that your BigQuery dataset only reflects Google Cloud billing data incurred from the date you set up Cloud Billing export, and after. That is, Google Cloud billing data is not added retroactively, so you won't see Cloud Billing data from before you enable export."  
  
But I am leaning towards alternative D after reading this page: <https://cloud.google.com/billing/docs/how-to/reports>.  
  
You don't have to export/import to see data and you can see costs from different projects. All you need is the permission: `billing.accounts.getSpendingInformation`. It can forecast future costs too. According to the video at the same page, BiqQuery exports to Data Studio is for deeper cost analysis.



upvoted 4 times

  **rramos96** 7 months ago

As we can read in the 1st paragraph at <https://cloud.google.com/billing/docs/how-to/reports>, Reports View from Cloud Billing Console shows info about a single Billing Account.

This question says:

- "all projects are linked to different billing accounts".
- "we want to have a single visual representation".

So, in my opinion, the best answer is [A] (BQ+data studio)

upvoted 6 times

  **don\_v** 8 months ago

I guess it's A.



I'm not sure how to get that info for multiple billing accounts with the billing report as they say "All projects are linked to different billing accounts."

upvoted 1 times

  **Bhagirathi** 8 months, 1 week ago

A or D which one to choose ?

upvoted 2 times

  **Lieder** 8 months, 2 weeks ago

Why not make use of Cloud Billing console? In the reports you get a dashboard with all the information needed (also forecasts) - across all billing accounts. And it is up to date - new resources are billed immediately. Everything is in place. No additional effort - like creating a dashboard in Data Studio.

I would go for D.

upvoted 3 times

  **[Removed]** 10 months, 2 weeks ago

I am also favour A, where you can use visualise the cost. But exporting and visualizing will show past data. We are talking about "estimating the further charges", that can be done by AI on big query else the closest will be pricing calculator so C. But since we don't know about the variable that will add on the cost that is you, for example, your data traffic is going to increase we can still go with A as answer here.

upvoted 1 times

  **stepkurniawan** 11 months, 1 week ago

I think it is A.



You can combine all to a single BigQuery using service account, and then push it to DataStudio for analytics

upvoted 3 times

  **Hjameel** 11 months, 2 weeks ago



The key is "To better estimate future charges" , A is the best answer, although exporting the billing account to BQ would take one day to start populating data, also this is the google best practice.

upvoted 3 times

  **SSPC** 11 months, 2 weeks ago

Yes, "A" is correct. Thanks

upvoted 3 times

  **SSPC** 11 months, 2 weeks ago



With Data Studio, can you see the cost of different billing accounts?


upvoted 1 times

Your company has workloads running on Compute Engine and on-premises. The Google Cloud Virtual Private Cloud (VPC) is connected to your WAN over a Virtual Private Network (VPN). You need to deploy a new Compute Engine instance and ensure that no public Internet traffic can be routed to it. What should you do?

- A. Create the instance without a public IP address.
- B. Create the instance with Private Google Access enabled.
- C. Create a deny-all egress firewall rule on the VPC network.
- D. Create a route on the VPC to route all traffic to the instance over the VPN tunnel.

**Correct Answer: B**  
Get private access to Google services, such as storage, big data, analytics, or machine learning, without having to give your service a public IP address.  
Reference:  
<https://cloud.google.com/vpc>

-   **[Removed]**

Highly Voted 



 10 months, 3 weeks ago


A for sure

B - this allows internal communicaiton, but does nothing to limit public traffic

C - deny all is nice, but it's for egress -- we're looking for ingress



D - this is way to invasive and doesn't explicitly address the issue of preventing public internet traffic from reaching your instance -- if it does, someone let me know how.


upvoted 13 times
-   **MohammedGhouse**

Highly Voted 

 11 months, 3 weeks ago



A: answer looks right

upvoted 8 times
-   **arsh1916**



Most Recent 

 2 months, 1 week ago



A is correct

upvoted 2 times
-   **Devgela** 4 months, 2 weeks ago



A for me

upvoted 2 times
-   **GCP\_Student1** 5 months ago



A. Create the instance without a public IP address.

upvoted 2 times
-   **victory108** 5 months, 3 weeks ago



A - Create the instance without a public IP address.

upvoted 1 times
-   **DucSiu** 5 months, 4 weeks ago



A. Create the instance without a public IP address.

upvoted 1 times
-   **guid1984** 6 months ago

Seems like a multi-select option questions in this case its A+D



upvoted 1 times
-   **Bhagirathi** 8 months, 1 week ago

solution given is B and why B also explained - why not B?

upvoted 1 times
-   **JackGlemins** 5 months ago

Because Private Google Access is for reach Google APIs and services with external/public IP addresses. The question dont talk about it.

<https://cloud.google.com/vpc/docs/private-google-access>

upvoted 1 times
-   **Bhagirathi** 8 months, 1 week ago

A and or D ? which one to choose?

upvoted 2 times

🗲️ 👤 **swatititame** 8 months, 3 weeks ago

- A. Create the instance without a public IP address.

I am using this approach for my prod as well as non prod environment's.

upvoted 2 times

🗲️ 👤 **JCGO** 9 months, 2 weeks ago

IMO this is a multiple choise question. A+D. simple.

upvoted 2 times

🗲️ 👤 **RockAJ** 9 months, 3 weeks ago

A for me

upvoted 1 times

🗲️ 👤 **[Removed]** 10 months, 2 weeks ago

The question is very confusing as to where you are hosting a new CE instance is not clear.

If it is on-prem then nothing new has to be done. But if we are talking about new instance on GCP then what it does? , does it reach out to other CE on PREM or only access GCP services. Please help me out understand the context.

upvoted 1 times

🗲️ 👤 **[Removed]** 10 months, 1 week ago

Will take back my comment. We are talking about a set of machine on PREM and in GCP you want to add to CE that means will be GCP. So avoiding public IP will prevent any access to the machine from the internet and will belong to the VPC.

upvoted 3 times

🗲️ 👤 **Ale1973** 10 months, 3 weeks ago

My answer is A.

When you are creating the instance you could configure external IP to none, to avoid external Internet access.

upvoted 4 times

🗲️ 👤 **stepkurniawan** 11 months, 1 week ago

Do you need a public IP to connect to the VPN?

How about D?

upvoted 3 times

🗲️ 👤 **magistrum** 6 months, 4 weeks ago

If you put a public IP on the machine, then public internet traffic can route through it over the vpn tunnel

upvoted 1 times

🗲️ 👤 **sapguru** 11 months, 1 week ago

C can also be considered

upvoted 1 times

🗲️ 👤 **Chapman** 9 months, 3 weeks ago

But traffic can still route to it

upvoted 2 times

🗲️ 👤 **mexblood1** 11 months ago



C is not an option because it says egress instead of ingress.

upvoted 1 times

Your team maintains the infrastructure for your organization. The current infrastructure requires changes. You need to share your proposed changes with the rest of the team. You want to follow Google's recommended best practices. What should you do?

- A. Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket.
- B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories.
- C. Apply the changes in a development environment, run `gcloud compute instances list`, and then save the output in a shared Storage bucket.
- D. Apply the changes in a development environment, run `gcloud compute instances list`, and then save the output in Cloud Source Repositories.

Correct Answer: B

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (A):



Connecting to Cloud Storage buckets



Cloud Storage is a flexible, scalable, and durable storage option for your virtual machine instances. You can read and write files to Cloud Storage buckets from almost anywhere, so you can use buckets as common storage between your instances, App Engine, your on-premises systems, and other cloud services.



<https://cloud.google.com/compute/docs/disks/gcs-buckets>



Why not (B)?  
Caution



Cloud Source Repositories are intended to store only the source code for your app and not user or personal data. Don't store any Core App Engine Customer Data (as defined in your License Agreement) in Cloud Source Repositories.  
<https://cloud.google.com/source-repositories/docs/features>  
upvoted 15 times



  **gcpengineer** 1 day, 13 hours ago  
B is the ans  
upvoted 1 times



  **pYWORLD** 2 days, 20 hours ago  
I agree with what are you saying, but the problem that you know how the deployment manager template looks? Is `jinja/yaml` file that means that are source code, so better to put them inside of an repository.  
  
So, for my perspective I will go with the B.  
upvoted 1 times



  **stepkurniawan** 11 months, 1 week ago  
you store the sensitive data NOT in the instance template, that is the current best practice. But you need version control like GIT or Google's GIT (Cloud Source Repo) to backup your code somehow and able to roll back if needed.  
upvoted 7 times

  **JohnnieWalker** 4 weeks, 1 day ago  
B is the answer. Deployment Manager Template can be written in either Jinja or Python, this is Infrastructure as Code (IaC) we are talking about here, same as AWS Cloudformation, or Terraform. Therefore, they should be stored on a git repository such as Google Cloud Source Repositories.  
upvoted 1 times

  **magistrum** 6 months, 4 weeks ago  
Look at my post above, cloud repo is for code, not templates  
upvoted 1 times

  **[Removed]** 10 months, 1 week ago  
You can do all thing you are mentioning in the cloud store also. Ethically answer is cloud store as you are not dealing with a source file but a template. Again here the argument can go that config is also a part of the source so B answers. To make life easy let's call this template/config file as the proposal file, so the best way to share will be cloud store.  
upvoted 1 times

  **SSPC** Highly Voted 11 months, 2 weeks ago  
B is correct. <https://cloud.google.com/source-repositories/docs/features>  
upvoted 12 times

  **magistrum** 6 months, 4 weeks ago  
I don't see how you can do this when I tried creating:  
Add code to your repository

info

Your repository is currently empty. Add some code using a selected method and then refresh your browser. Contents added to this repository can take some time to show up in search results. Learn more.

Select an option to push code to your repository:

Push code from a local Git repository

Clone your repository to a local Git repository

upvoted 2 times

  **AmitKM** 11 months, 1 week ago

Using Cloud Storage Repos, you can add comments and describe your changes to the team. Hence this might be a better option.

upvoted 3 times

  **JieHeng** Most Recent 1 month ago

A is correct, from the doc - [https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally#hosting\\_external\\_templates](https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally#hosting_external_templates)

"You can host external templates in Google Cloud Storage or in a publicly-accessible location, such as GitHub."

Why NOT B - Cloud Source Repo is not a "publicly-accessible location", check this - [https://cloud.google.com/source-repositories/docs/configure-access-control#granting\\_member\\_access](https://cloud.google.com/source-repositories/docs/configure-access-control#granting_member_access)

upvoted 1 times

  **arsh1916** 2 months, 1 week ago

A is answer

upvoted 1 times

  **ri\_unhou119** 2 months, 2 weeks ago

I think A is correct.

[https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally#hosting\\_external\\_templates](https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally#hosting_external_templates)

You can host external templates in Google Cloud Storage or in a publicly-accessible location, such as GitHub. As long as there is a URL to the file and you have access to the templates, Deployment Manager will be able to download it. If the file uses a schema to import other templates, you must also host those templates in the same base URL.

upvoted 1 times

  **pondai** 4 months ago

B is correct

Cloud Source Repositories provides fully featured, private Git repositories hosted on Google Cloud.

Now we see what deploymentmanager template doc say

<https://cloud.google.com/deployment-manager/docs/configuration/templates/create-basic-template?hl=zh-tw>

For examples of templates that you can use in your own deployments, see the Deployment Manager GitHub repository.

So I think B is better way than A

upvoted 3 times

  **pca2b** 4 months ago

B: DM is Infra as Code, and use CSR for versioning and working/collaborating in a team.

A will work, but not good practice for code repository

upvoted 2 times

  **GCP\_Student1** 5 months ago

A. Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket.

upvoted 1 times

  **TAvenger** 5 months, 1 week ago

According to this doc <https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally> you can easily use/apply your templates externally.

Source code is really good for versioning, but looks like this is not the case.

Answer is "A"

upvoted 2 times

  **lutoa** 5 months, 2 weeks ago

A is correct. The question asks for infrastructure changes NOT code changes which means the code repository answer is a trick answer.

I have seen variations of this question where the correct answer is to use deployment manager to preview the changes and store those in cloud storage.

In this case, A is the closest answer.

upvoted 2 times

  **victory108** 5 months, 3 weeks ago

B - Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories.

upvoted 1 times

  **victory108** 4 months, 2 weeks ago

A is correct. Use Deployment Manager templates to describe the proposed changes and store them in a Cloud Storage bucket.

upvoted 1 times

  **tavva\_prudhvi** 3 months, 2 weeks ago

What made you change your decision after a month, it would have better if you could keep some additional articles supporting your statement.

upvoted 2 times

🗨️ 👤 **sarahf** 7 months, 3 weeks ago

Not that this link made it more clear:

<https://cloud.google.com/deployment-manager/docs/configuration/templates/hosting-templates-externally>

upvoted 2 times

🗨️ 👤 **tavva\_prudhvi** 3 months, 2 weeks ago

I reckon the article is not inclined to the question. Yes, I have seen the articles saying to store your templates externally via Cloud Storage buckets, but here it's different right, we are looking for an configuration file. May be B, suits better?

upvoted 1 times

🗨️ 👤 **Bhagirathi** 7 months, 3 weeks ago

Referring all the discussions....

A or B ? what to finally choose.

upvoted 1 times

🗨️ 👤 **Bhagirathi** 8 months, 1 week ago

B seems correct...anyone tested/tried.

upvoted 1 times

🗨️ 👤 **swatititame** 8 months, 1 week ago

- B. Use Deployment Manager templates to describe the proposed changes and store them in Cloud Source Repositories.

upvoted 1 times

🗨️ 👤 **ting\_66** 9 months, 2 weeks ago

B, the point is "proposed change". with code repo, you can comment and reply with a change and compare the changes. Option A, Cloud Storage doesn't have that feature.

upvoted 2 times

🗨️ 👤 **RockAJ** 9 months, 3 weeks ago

B for me

upvoted 1 times

🗨️ 👤 **hiteshrup** 10 months, 3 weeks ago

i will go with Option B.

Cloud Deployment manager template is also Infrastructure code and valid for versioning for all change proposed. So this should go with Repository than Bucket.

upvoted 3 times

You have a Compute Engine instance hosting an application used between 9 AM and 6 PM on weekdays. You want to back up this instance daily for disaster recovery purposes. You want to keep the backups for 30 days. You want the Google-recommended solution with the least management overhead and the least number of services. What should you do?

- A. 1. Update your instances' metadata to add the following value: snapshot"schedule: 0 1 \* \* \* 2. Update your instances' metadata to add the following value: snapshot"retention: 30
- B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: - Schedule frequency: Daily - Start time: 1:00 AM 2:00 "€ AM - Autodelete snapshots after: 30 days
- C. 1. Create a Cloud Function that creates a snapshot of your instance's disk. 2. Create a Cloud Function that deletes snapshots that are older than 30 days. 3. Use Cloud Scheduler to trigger both Cloud Functions daily at 1:00 AM.
- D. 1. Create a bash script in the instance that copies the content of the disk to Cloud Storage. 2. Create a bash script in the instance that deletes data older than 30 days in the backup Cloud Storage bucket. 3. Configure the instance's crontab to execute these scripts daily at 1:00 AM.

Correct Answer: B

- ESP\_SAP

Highly Voted

11 months, 1 week ago

Correct Answer is (B):

Creating scheduled snapshots for persistent disk  
This document describes how to create a snapshot schedule to regularly and automatically back up your zonal and regional persistent disks. Use snapshot schedules as a best practice to back up your Compute Engine workloads. After creating a snapshot schedule, you can apply it to one or more persistent disks.

<https://cloud.google.com/compute/docs/disks/scheduled-snapshots>  
upvoted 26 times
- [Removed]

10 months, 3 weeks ago

Definitely B.

With something like this, you should not have to write any custom scripts, custom functions, or cron jobs. This is google's way of saying 'hey, we've already built that stuff in to our snapshot schedules feature.

upvoted 5 times
- gcpengineer

Most Recent

1 day, 12 hours ago

why not C?

upvoted 1 times
- arsh1916

2 months, 1 week ago

B is correct

upvoted 2 times
- Hi2ALL

4 months, 2 weeks ago

B is more realistic approach

upvoted 3 times
- GCP\_user

4 months, 2 weeks ago

B is the best option so far. However just wonder this: Schedule frequency: Daily " Start time: 1:00 AM " 2:00 AM" Autodelete snapshots: after 30 days; For Saturday and Sunday it will be a waste of resource to create snapshots since the instance is running during weekdays.

upvoted 2 times
- GCP\_Student1

5 months ago

B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 2 times
- DucSiu

6 months ago

It's B

upvoted 1 times
- Bhagirathi

8 months, 1 week ago

B for sure, any doubt?

upvoted 1 times



🗨️ 👤 **swatititame** 8 months, 1 week ago

• B. 1. In the Cloud Console, go to the Compute Engine Disks page and select your instance's disk. 2. In the Snapshot Schedule section, select Create Schedule and configure the following parameters: "" Schedule frequency: Daily "" Start time: 1:00 AM "" 2:00 AM "" Autodelete snapshots after 30 days

upvoted 1 times

🗨️ 👤 **hiteshrup** 10 months, 3 weeks ago

Has to B...

- scripting which means overhead and maintenance So Option D neglected.
- Cloud function, adding another service which not require.
- Metadata, I am not sure if meta data have something to define as cron job schedule. So not be an option.

upvoted 1 times

🗨️ 👤 **stepkurniawan** 11 months, 1 week ago

Either A or B, I am not sure.

I have tested B and it works, but I dont know the name of the key to do A.

upvoted 1 times

🗨️ 👤 **juliandm** 11 months, 1 week ago

B is right <https://www.cloudbooklet.com/backup-google-cloud-vm-automatically-with-snapshot-schedules>

upvoted 2 times

🗨️ 👤 **baonguyen** 11 months, 1 week ago

The Ans is B.

Checked on my GCP account.

upvoted 4 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

I would go with the answer "D" but I don't totally sure

upvoted 1 times

🗨️ 👤 **[Removed]** 10 months, 2 weeks ago

You need to write the script and test it that is a big step.

upvoted 1 times

🗨️ 👤 **[Removed]** 10 months, 2 weeks ago

So will go with B

upvoted 2 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

B?? I think is "D". [https://cloud.google.com/solutions/dr-scenarios-planning-guide#design\\_for\\_end-to-end\\_recovery](https://cloud.google.com/solutions/dr-scenarios-planning-guide#design_for_end-to-end_recovery)

upvoted 1 times

🗨️ 👤 **Eshkrkrkr** 8 months, 3 weeks ago

It's B - question asks for Google-recommended and with the least management overhead what makes D the wrong answer.

upvoted 3 times

🗨️ 👤 **francisco\_guerra** 11 months, 2 weeks ago

Yes maybe A or D

upvoted 1 times



## Question #145

## Topic 1

Your existing application running in Google Kubernetes Engine (GKE) consists of multiple pods running on four GKE n1-standard2 nodes. You need to deploy additional pods requiring n2-highmem16 nodes without any downtime. What should you do?

- A. Use gcloud container clusters upgrade. Deploy the new services.
- B. Create a new Node Pool and specify machine type n2-highmem16. Deploy the new pods.
- C. Create a new cluster with n2-highmem16 nodes. Redeploy the pods and delete the old cluster.
- D. Create a new cluster with both n1-standard2 and n2-highmem16 nodes. Redeploy the pods and delete the old cluster.

**Correct Answer: B**

  **GCP\_Student1** Highly Voted  4 months, 2 weeks ago

B is correct answer, read below form google docs;

This tutorial demonstrates how to migrate workloads running on a Google Kubernetes Engine (GKE) cluster to a new set of nodes within the same cluster without incurring downtime for your application. Such a migration can be useful if you want to migrate your workloads to nodes with a different machine type.

#### Background

A node pool is a subset of machines that all have the same configuration, including machine type (CPU and memory) authorization scopes. Node pools represent a subset of nodes within a cluster; a container cluster can contain one or more node pools.

When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.

To migrate your workloads without incurring downtime, you need to:

Mark the existing node pool as unschedulable.  
Drain the workloads running on the existing node pool.  
Delete the existing node pool.

[https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool#creating\\_a\\_node\\_pool\\_with\\_large\\_machine\\_type](https://cloud.google.com/kubernetes-engine/docs/tutorials/migrating-node-pool#creating_a_node_pool_with_large_machine_type)  
upvoted 10 times


  **arsh1916** Most Recent  2 months, 1 week ago

B is correct  
upvoted 2 times



  **Jacky\_YO** 3 months, 3 weeks ago

ANS : B

1. The title did not say to delete four GKE n1.  
upvoted 2 times

  **pondai** 4 months, 1 week ago

B, You need to create new node pool for cluster  
upvoted 2 times

  **dunhill** 4 months, 2 weeks ago

I guess it's B. I couldn't find resize parameter under cluster upgrade. C and D are incorrect because it's no need to create new cluster.  
upvoted 2 times

  **GCP\_Student1** 4 months, 2 weeks ago

A. Use gcloud container clusters upgrade. Deploy the new services.  
upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

I take it back, the correct answer is "B"  
B. Create a new Node Pool and specify machine type n2-highmem16. Deploy the new pods.  
upvoted 2 times

  **Jamaal\_a** 4 months, 2 weeks ago

Answer is B - When you need to change the machine profile of your Compute Engine cluster, you can create a new node pool and then migrate your workloads over to the new node pool.  
upvoted 3 times

You have an application that uses Cloud Spanner as a database backend to keep current state information about users. Cloud Bigtable logs all events triggered by users. You export Cloud Spanner data to Cloud Storage during daily backups. One of your analysts asks you to join data from Cloud Spanner and Cloud Bigtable for specific users. You want to complete this ad hoc request as efficiently as possible. What should you do?

- A. Create a dataflow job that copies data from Cloud Bigtable and Cloud Storage for specific users.
- B. Create a dataflow job that copies data from Cloud Bigtable and Cloud Spanner for specific users.
- C. Create a Cloud Dataproc cluster that runs a Spark job to extract data from Cloud Bigtable and Cloud Storage for specific users.
- D. Create two separate BigQuery external tables on Cloud Storage and Cloud Bigtable. Use the BigQuery console to join these tables through user fields, and apply appropriate filters.

Correct Answer: B

- AmitKM**

Highly Voted

11 months, 2 weeks ago

I think it should be D. <https://cloud.google.com/bigquery/external-data-sources>

upvoted 19 times
- SSPC**

11 months, 1 week ago

The question says: " Join data from Cloud Spanner and Cloud Bigtable for specific users" You can see the Google documentation in the link <https://cloud.google.com/spanner/docs/export>

upvoted 3 times
- Eshkrkrkr**

8 months, 3 weeks ago

Oh my god, SSPC read you your links!  
The process uses Dataflow and exports data to a folder in a Cloud Storage bucket. The resulting folder contains a set of Avro files and JSON manifest files. And what next? I will tell - next you read below: Compute Engine: Before running your export job, you must set up initial quotas for Recommended starting values are:

CPUs: 200  
In-use IP addresses: 200  
Standard persistent disk: 50 TB  
Still think its A?

upvoted 2 times
- punjabishiva123**

8 months ago

Hi,

Could u pls mail me all the correct answer to [shivaasingh1104@gmail.com](mailto:shivaasingh1104@gmail.com),

Really appreciated  
Anyone pls help

upvoted 1 times
- ash500**

7 months ago

can you pls also share the correct answers with me

upvoted 1 times
- kp0916**

5 months ago

Does anyone have list of correct answers,let me know

upvoted 1 times
- ESP\_SAP**

Highly Voted

11 months, 1 week ago

Correct Answer is (D):

Introduction to external data sources  
This page provides an overview of querying data stored outside of BigQuery.

<https://cloud.google.com/bigquery/external-data-sources>


upvoted 11 times
- [Removed]**

10 months, 2 weeks ago

As per your comment D is the answer.  
I also agree.  
But can BigQurey read backed up data? , as we have backup data on Cloud storage, did not get any evidence in the link you shared.

upvoted 1 times



 **ESP\_SAP** 11 months, 1 week ago

BigQuery offers support for querying data directly from:

Bigtable  
Cloud Storage  
Google Drive  
Cloud SQL (beta)  
upvoted 3 times

  **djgodzilla** 1 month, 2 weeks ago

but here we're not talking about joining Cloud Storage and Cloud Bigtable external tables.  
the join happens between a distributed relational database (Spanner) and key-value NoSQL Database (BigTable) . how's converting Spanner to cloud storage an implicit and trivial step.  
upvoted 1 times


  **djgodzilla** 1 month, 2 weeks ago

"The Cloud Spanner to Cloud Storage Text template is a batch pipeline that reads in data from a Cloud Spanner table, optionally transforms the data via a JavaScript User Defined Function (UDF) that you provide, and writes it to Cloud Storage as CSV text files."  
<https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#cloudspannertogcstext>



"The Dataflow connector for Cloud Spanner lets you read data from and write data to Cloud Spanner in a Dataflow pipeline"  
<https://cloud.google.com/spanner/docs/dataflow-connector>  
upvoted 1 times

  **BalSan** Most Recent 2 weeks, 4 days ago

One of the key word here is "ad hoc" request.... It doesn't make sense to create Dataflow Job for an adhoc request, ruling out A & B. C doesn't make sense for this use case and hence the answer is D  
upvoted 1 times

  **arsh1916** 2 months, 1 week ago



D is best option  
upvoted 1 times

  **arsh1916** 2 months, 3 weeks ago

Ans B, simple they are asking about Cloud bigtable & spanner  
upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

A. Create a dataflow job that copies data from Cloud Bigtable and Cloud Storage for specific users.  
upvoted 1 times

  **Amu89** 2 months, 4 weeks ago

Can you please confirm why you choose Option A over D  
upvoted 1 times

  **EABDAJA** 4 months, 3 weeks ago



D is correct  
upvoted 1 times

  **GCP\_Student1** 5 months ago

D. Create two separate BigQuery external tables on Cloud Storage and Cloud Bigtable. Use the BigQuery console to join these tables through user fields, and apply appropriate filters.  
upvoted 1 times

  **jamesET209** 7 months, 2 weeks ago

Answer will be D  
upvoted 1 times

  **Bhagirathi** 8 months, 1 week ago

D for for sure. it fits well to the ask.  
upvoted 2 times

  **swatititame** 8 months, 1 week ago



You have an application that uses Cloud Spanner as a database backend to keep current state information about users. Cloud Bigtable logs all events triggered by users. You export Cloud Spanner data to Cloud Storage during daily backups. One of your analysts asks you to join data from Cloud Spanner and Cloud Bigtable for specific users. You want to complete this ad hoc request as efficiently as possible. What should you do?  
upvoted 1 times

  **cpd** 9 months, 3 weeks ago

BigQuery supports:  
Cloud Bigtable  
Cloud Storage  
Google Drive  
Cloud SQL

D


upvoted 2 times

  **RockAJ** 9 months, 3 weeks ago


D for me  
upvoted 2 times

  **[Removed]** 10 months, 2 weeks ago

B option does what the analyst is asking for.  
D is trying to merge data that is backed from cloud storage same is A and C,  
from A, C, D options D is the best if we know BigQuery External Source can be configured to read Cloud Spanner data backed (that can if read a back up file)  
upvoted 1 times

  **[Removed]** 10 months, 1 week ago



Take back my comment, D is the solution. B only copes but does not talk about joining.  
upvoted 2 times

  **stepkurniawan** 11 months, 1 week ago

I cant decide between A or B  
upvoted 1 times

  **technicalguru** 11 months, 1 week ago

D is the correct answer! Agreed to AmitKM  
upvoted 2 times

  **SSPC** 11 months, 2 weeks ago

What do you think the correct answer is A?  
upvoted 1 times

You are hosting an application from Compute Engine virtual machines (VMs) in us-central1-a. You want to adjust your design to support the failure of a single Compute Engine zone, eliminate downtime, and minimize cost. What should you do?

- A. Create Compute Engine resources in us-central1-b. Balance the load across both us-central1-a and us-central1-b.
- B. Create a Managed Instance Group and specify us-central1-a as the zone. Configure the Health Check with a short Health Interval.
- C. Create an HTTP(S) Load Balancer. Create one or more global forwarding rules to direct traffic to your VMs.
- D. Perform regular backups of your application. Create a Cloud Monitoring Alert and be notified if your application becomes unavailable. Restore from backups when notified.

Correct Answer: A

- GCP\_Student1

Highly Voted

4 months, 2 weeks ago

A. Create Compute Engine resources in us-central1-b. Balance the load across both us-central1-a and us-central1-b.  
upvoted 5 times
- arsh1916

Most Recent

2 months, 1 week ago

A is best option  
upvoted 4 times
- mj98

3 months, 2 weeks ago

Can someone explain how A?  
upvoted 2 times
- tifo16

3 months, 1 week ago

in order to remediate to the problem of single point of failure, we have to replicate VMs within multiple zones. Only A choice consider this concern  
upvoted 7 times
- nzexamtopics

3 months, 3 weeks ago

A? Really? how?  
upvoted 1 times
- NARWAL

3 months, 3 weeks ago

A is correct.  
upvoted 2 times
- victory108

4 months, 2 weeks ago

A - Create Compute Engine resources in us-central1-b. Balance the load across both us-central1-a and us-central1-b.  
upvoted 2 times

## Question #148

## Topic 1

A colleague handed over a Google Cloud Platform project for you to maintain. As part of a security checkup, you want to review who has been granted the Project Owner role. What should you do?

- A. In the console, validate which SSH keys have been stored as project-wide keys.
- B. Navigate to Identity-Aware Proxy and check the permissions for these resources.
- C. Enable Audit Logs on the IAM & admin page for all resources, and validate the results.
- D. Use the command `gcloud projects get-iam-policy` to view the current role assignments.

**Correct Answer: A**

Reference:

<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (D):

A simple approach would be to use the command flags available when listing all the IAM policy for a given project. For instance, the following command:

```
`gcloud projects get-iam-policy $PROJECT_ID --flatten="bindings[].members" --format="table(bindings.members)" --filter="bindings.role:roles/owner"`
```

outputs all the users and service accounts associated with the role 'roles/owner' in the project in question.



<https://groups.google.com/g/google-cloud-dev/c/Z6sZs7TvygQ?pli=1>

upvoted 23 times

  **MohammedGhouse** Highly Voted 11 months, 3 weeks ago

D: is the answer

upvoted 11 times

  **SSPC** 11 months, 2 weeks ago

D is the correct.

upvoted 2 times

  **yurstev** 7 months, 1 week ago

D IS THE ANSWER

upvoted 3 times

  **arsh1916** Most Recent 2 months, 1 week ago

D is correct

upvoted 2 times

  **GCP\_Student1** 5 months ago

D. Use the command `gcloud projects get-iam-policy` to view the current role assignments.

upvoted 3 times

  **Bhagirathi** 7 months, 1 week ago

D 200%

upvoted 2 times

  **Bhagirathi** 8 months, 1 week ago

anyone will be confused - solution says one answer

same time, all you guys have different choices here. what to take from this ?

upvoted 1 times

  **swatititame** 8 months, 1 week ago

D. Use the command `gcloud projects get-iam-policy` to view the current role assignments.

upvoted 1 times

  **adeyemi5700** 9 months, 2 weeks ago

C is the likely answer. With D you see the current users with permission. With audit log you see those with prev. and current permission.

upvoted 1 times

  **[Removed]** 10 months, 2 weeks ago

D is the answer but the command is wrong  
gcloud projects get-iam-policy  
upvoted 2 times

  **stepkurniawan** 11 months, 1 week ago

Why C is not the answer?  
upvoted 1 times

  **[Removed]** 10 months, 2 weeks ago

You need to wait until all of them access to determine who has access.  
upvoted 1 times

  **Eshkrkrkr** 8 months, 3 weeks ago

No, you don't. With the command you'll see the current list of Owners (and all the other assignments as well).  
upvoted 3 times

  **francisco\_guerra** 11 months ago

Because D is a easy way to do the same  
upvoted 2 times

## Question #149

## Topic 1

You are running multiple VPC-native Google Kubernetes Engine clusters in the same subnet. The IPs available for the nodes are exhausted, and you want to ensure that the clusters can grow in nodes when needed. What should you do?

- A. Create a new subnet in the same region as the subnet being used.
- B. Add an alias IP range to the subnet used by the GKE clusters.
- C. Create a new VPC, and set up VPC peering with the existing VPC.
- D. Expand the CIDR range of the relevant subnet for the cluster.

**Correct Answer: D**

🗨️ **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (D):

gcloud compute networks subnets expand-ip-range  
NAME

gcloud compute networks subnets expand-ip-range - expand the IP range of a Compute Engine subnetwork

<https://cloud.google.com/sdk/gcloud/reference/compute/networks/subnets/expand-ip-range>

upvoted 16 times

🗨️ **magistrum** 6 months, 4 weeks ago

Ok D it is, here's the GKE specific documentation

<https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips>

Every subnet must have a primary IP address range. You can expand the primary IP address range at any time, even when Google Cloud resources use the subnet; however, you cannot shrink or change a subnet's primary IP address scheme after the subnet has been created. The first two and last two IP addresses of a primary IP address range are reserved by Google Cloud.

upvoted 3 times

🗨️ **MohammedGhouse** Highly Voted 11 months, 3 weeks ago

D: is the answer

upvoted 11 times

🗨️ **SSPC** 11 months, 2 weeks ago

I agree with you. [https://cloud.google.com/vpc/docs/configure-alias-ip-ranges#gcloud\\_1](https://cloud.google.com/vpc/docs/configure-alias-ip-ranges#gcloud_1)

upvoted 2 times

🗨️ **GCP\_Student1** Most Recent 5 months ago

This might help

Node limiting ranges

The maximum number of Pods and Services for a given GKE cluster is limited by the size of the cluster's secondary ranges. The maximum number of nodes in the cluster is limited by the size of the cluster's subnet's primary IP address range and the cluster's Pod address range.

The Cloud Console shows error messages like the following to indicate that either the subnet's primary IP address range or the cluster's Pod IP address range (the subnet's secondary IP address range for Pods) has been exhausted:

Instance [node name] creation failed: IP space of [cluster subnet] is exhausted

Note: Secondary subnets are not visible in Cloud Console. If you can't find the [cluster subnet] reported by the above error message it means that the error is caused by IP exhaustion in a secondary subnet. In this case check the secondary ranges of the primary subnet.

[https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips#node\\_limiters](https://cloud.google.com/kubernetes-engine/docs/concepts/alias-ips#node_limiters)

upvoted 5 times

🗨️ **GCP\_Student1** 4 months, 2 weeks ago

By the way the answer is;

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 3 times

🗨️ **Ozymandiox** 7 months ago

UHmmm, 1 question. The description of the problem says that the ip's are EXHAUSTED. So, no more IP's available in this subnet.



It also states that we're having a multi-VPC environment... as allways we should not interpret, just take the questions literally.

IF we do not know the actual size of the deployment it cna be ANY size, adn if IP's are EXHAUSTED, it should BE, BIG as Galactic sized or so....

With all this I wonder if the right answer it is not C...





upvoted 2 times

  **Bhagirathi** 8 months, 1 week ago

D best option to think here.

upvoted 2 times

  **hicham** 5 months, 1 week ago

totaly agree

upvoted 1 times

  **swatititame** 8 months, 1 week ago

D. Expand the CIDR range of the relevant subnet for the cluster.

upvoted 1 times

## Question #150

## Topic 1

You have a batch workload that runs every night and uses a large number of virtual machines (VMs). It is fault-tolerant and can tolerate some of the VMs being terminated. The current cost of VMs is too high. What should you do?

- A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.
- B. Run a test using simulated maintenance events. If the test is successful, use N1 Standard VMs when running future jobs.
- C. Run a test using a managed instance group. If the test is successful, use N1 Standard VMs in the managed instance group when running future jobs.
- D. Run a test using N1 standard VMs instead of N2. If the test is successful, use N1 Standard VMs when running future jobs.

**Correct Answer:** B

Reference:

<https://cloud.google.com/compute/vm-instance-pricing>

  **ESP\_SAP** Highly Voted  11 months, 1 week ago

Correct Answer is (A):

Creating and starting a preemptible VM instance

This page explains how to create and use a preemptible virtual machine (VM) instance. A preemptible instance is an instance you can create and run at a much lower price than normal instances. However, Compute Engine might terminate (preempt) these instances if it requires access to those resources for other tasks. Preemptible instances will always terminate after 24 hours. To learn more about preemptible instances, read the preemptible instances documentation.

Preemptible instances are recommended only for fault-tolerant applications that can withstand instance preemptions. Make sure your application can handle preemptions before you decide to create a preemptible instance. To understand the risks and value of preemptible instances, read the preemptible instances documentation.

<https://cloud.google.com/compute/docs/instances/create-start-preemptible-instance>

upvoted 19 times

  **MohammedGhouse** Highly Voted  11 months, 3 weeks ago

A: is the answer

upvoted 13 times

  **juliandm** 11 months, 1 week ago



What about a mixture of preemptible N1 and normal N1 instances? i can't believe just having preemptible is a good practice

upvoted 1 times

  **Ale1973** 10 months, 3 weeks ago

Good point, in real-world your solution, is the best. For this scenario, the answer is A.

upvoted 7 times

  **SSPC** 11 months, 2 weeks ago

"A" is correct

upvoted 3 times

  **AD\_0525** Most Recent  1 month, 2 weeks ago

Keyword- Fault tolerant, so answer should be pre emptible VMs, option A

upvoted 1 times

  **GCP\_Student1** 5 months ago

A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs

upvoted 3 times

  **swatititame** 8 months, 1 week ago



• A. Run a test using simulated maintenance events. If the test is successful, use preemptible N1 Standard VMs when running future jobs.

upvoted 2 times

You are working with a user to set up an application in a new VPC behind a firewall. The user is concerned about data egress. You want to configure the fewest open egress ports. What should you do?

- A. Set up a low-priority (65534) rule that blocks all egress and a high-priority rule (1000) that allows only the appropriate ports.
- B. Set up a high-priority (1000) rule that pairs both ingress and egress ports.
- C. Set up a high-priority (1000) rule that blocks all egress and a low-priority (65534) rule that allows only the appropriate ports.
- D. Set up a high-priority (1000) rule to allow the appropriate ports.

Correct Answer: C

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (A):



Implied rules  
Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:



Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see Internet access requirements.

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.

[https://cloud.google.com/vpc/docs/firewalls#default\\_firewall\\_rules](https://cloud.google.com/vpc/docs/firewalls#default_firewall_rules)  
upvoted 14 times



  **MohammedGhouse** Highly Voted 11 months, 3 weeks ago  
A: is the answer  
upvoted 10 times



  **tanito83** Most Recent 1 month, 2 weeks ago  
The correct answer is A. Please, modify it.  
upvoted 3 times



  **kopper2019** 3 months, 1 week ago  
A, lower is high priority  
Implied rules  
Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console:



Implied allow egress rule. An egress rule whose action is allow, destination is 0.0.0.0/0, and priority is the lowest possible (65535) lets any instance send traffic to any destination, except for traffic blocked by Google Cloud. A higher priority firewall rule may restrict outbound access. Internet access is allowed if no other firewall rules deny outbound traffic and if the instance has an external IP address or uses a Cloud NAT instance. For more information, see Internet access requirements.

Implied deny ingress rule. An ingress rule whose action is deny, source is 0.0.0.0/0, and priority is the lowest possible (65535) protects all instances by blocking incoming connections to them. A higher priority rule might allow incoming access. The default network includes some additional rules that override this one, allowing certain types of incoming connections.  
upvoted 2 times

  **Devgela** 4 months, 2 weeks ago  
A for me  
upvoted 1 times



  **EABDAJA** 4 months, 3 weeks ago  
A is correct.  
upvoted 1 times

  **GCP\_Student1** 5 months ago  
A. Set up a low-priority (65534) rule that blocks all egress and a high-priority rule (1000) that allows only the appropriate ports.  
upvoted 2 times

  **Bhagirathi** 8 months, 1 week ago  
A sounds OK - anyone tested it in your project?  
upvoted 1 times

  **swatititame** 8 months, 1 week ago

- A. Set up a low-priority (65534) rule that blocks all egress and a high-priority rule (1000) that allows only the appropriate ports.  
upvoted 1 times

  **RockAJ** 9 months, 3 weeks ago

A for me

upvoted 2 times

  **[Removed]** 10 months, 1 week ago

A as per the priority working of firewall.

[https://cloud.google.com/vpc/docs/firewalls#priority\\_order\\_for\\_firewall\\_rules](https://cloud.google.com/vpc/docs/firewalls#priority_order_for_firewall_rules)

upvoted 2 times

  **[Removed]** 10 months, 1 week ago

"The highest priority rule applicable to a target for a given type of traffic takes precedence. Target specificity does not matter. For example, a higher priority ingress rule for certain ports and protocols intended for all targets overrides a similarly defined rule for the same ports and protocols intended for specific targets."

That is if I do C that is high rule for blocking all , and low for allowing some then still will be blocked for all.

So A is better.

upvoted 1 times



  **Hjameel** 11 months, 2 weeks ago

Every VPC network has two implied firewall rules. These rules exist, but are not shown in the Cloud Console: Implied allow egress rule, Implied deny ingress rule

if a higher priority firewall rule may restrict outbound access,

A is the answer.

upvoted 3 times

  **SSPC** 11 months, 2 weeks ago

A is correct

upvoted 5 times

Your company runs its Linux workloads on Compute Engine instances. Your company will be working with a new operations partner that does not use Google Accounts. You need to grant access to the instances to your operations partner so they can maintain the installed tooling. What should you do?

- A. Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User.
- B. Tag all the instances with the same network tag. Create a firewall rule in the VPC to grant TCP access on port 22 for traffic from the operations partner to instances with the network tag.
- C. Set up Cloud VPN between your Google Cloud VPC and the internal network of the operations partner.
- D. Ask the operations partner to generate SSH key pairs, and add the public keys to the VM instances.

Correct Answer: B

Reference:

<https://cloud.google.com/vpc/docs/firewalls>

- Bhagirathi

Highly Voted

8 months, 1 week ago

full of confusions for any reader....  
You guys all say A, B, C & D but which one is correct ?  
upvoted 12 times
- yc25744

2 weeks ago

nothing  
upvoted 1 times
- kulikBro

Highly Voted

3 months, 2 weeks ago

A - <https://cloud.google.com/iap/docs/external-identities>  
upvoted 10 times
- YAS007

Most Recent

1 day, 23 hours ago

for me its D, but I'm not sur :  
[https://cloud.google.com/compute/docs/access#granting\\_users\\_ssh\\_access\\_to\\_vm\\_instances](https://cloud.google.com/compute/docs/access#granting_users_ssh_access_to_vm_instances)  
upvoted 1 times
- umashankar\_a

4 weeks ago

Option A - IAP  
By default, IAP uses Google identities and IAM. By leveraging Identity Platform instead, you can authenticate users with a wide range of external identity providers, such as:  
  
Email/password  
OAuth (Google, Facebook, Twitter, GitHub, Microsoft, etc.)  
SAML  
OIDC  
Phone number  
Custom  
Anonymous  
upvoted 2 times
- jcloud965

1 week, 6 days ago


How to grant Cloud IAP Tunnel User role without Google Accounts ?  
upvoted 2 times
- JieHeng

1 month ago

I would go for D after reading [https://cloud.google.com/compute/docs/instances/access-overview#managing\\_user\\_access](https://cloud.google.com/compute/docs/instances/access-overview#managing_user_access) even though it might have risk as indicated in the doc - <https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>  
  
B & C - are merely about network connectivity, not access  
  
A - <https://cloud.google.com/iap/docs/concepts-overview> If you need to create Google Accounts for your existing users, you can use Google Cloud Directory Sync to synchronize with your Active Directory or LDAP server. But then <https://cloud.google.com/iap/docs/enabling-compute-howto#iap-access> mention that it requires google account.  
upvoted 3 times
- nana1995

1 month, 3 weeks ago

D ... idk.... i didn't find anything that said you can not use this party ssh key and add it to your VM  
<https://cloud.google.com/compute/docs/instances/adding-removing-ssh-keys>  
upvoted 1 times
-

 **arsh1916** 2 months, 1 week ago

I selected B

upvoted 1 times



  **meh009** 3 months, 1 week ago

Ans: D

Not really sure I can add much more to the already confusing discussions here however D for me seems the most likely one considering what I beleive the Question is trying to get at. By Default, within the Default VPC, SSH ingress is permitted (default-allow-ssh).

So we can (I guess) assume that network access portion is already permitted and the answer is in more in relation to auth

upvoted 1 times

  **noreen** 3 months, 2 weeks ago

. By combining IAP and Identity Platform, you can authenticate users with a wide range of identity providers (such as OAuth, SAML, OIDC, and more), instead of just Google accounts.

upvoted 4 times

  **pca2b** 4 months ago

really messy one...

not A: ops partner does not use Google accounts, so cant use IAP

not B: assumes all GCE VMs have external IPs for exposing to network partner (seems a waste), doesnt discuss Linux OS access



not C: doesnt discuss next step of Linux access

not D: assumes all GCE instances have external IPs and exposed to network, accessible by ops network

iterate...



C is least distasteful, gives all GCE instance internal IP access to ops partner, we can worry about next step of Linux access then

upvoted 4 times

  **JH86** 1 month, 2 weeks ago

can use IAP for external identities - <https://cloud.google.com/iap/docs/enable-external-identities>

upvoted 3 times

  **pondai** 4 months, 1 week ago



Answer D

A) Need a google account

B)Can't access to Linux

C)Can't access to Linux

upvoted 4 times

  **JH86** 1 month, 2 weeks ago

do not need google account - <https://cloud.google.com/iap/docs/enable-external-identities>

upvoted 2 times

  **jcloud965** 1 week, 6 days ago

You do need to grant Cloud IAP Tunnel User role. Second part of the answer

upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

C. Set up Cloud VPN between your Google Cloud VPC and the internal network of the operations partner.

upvoted 1 times

  **tavva\_prudhvi** 3 months, 4 weeks ago



It's possible, but i think its time taking compared to D.

upvoted 1 times

  **kopper2019** 3 months, 1 week ago

yes but for D, you need a bastion host, all VMs with public IPs or a Cloud VPN

upvoted 2 times

  **Amu89** 2 months, 4 weeks ago

Can you please confirm the answer ?

upvoted 1 times

  **ShakthiGCP** 4 months, 4 weeks ago

IAM needs Google sign in to Authenticate. Whenever you talk about on-premises connectivity , think about Cloud VPN. Ans : C . this will be the initial step to enable the connectivity to Instance. - "Cloud VPN lets you connect your existing network to your Google Cloud network by using an IPsec connection to a VPN gateway device. This allows direct routing of traffic from your premises to the private IP interfaces of Compute Engine instances. Traffic is encrypted as it transits over public links to Google."

upvoted 1 times

  **ShakthiGCP** 4 months, 3 weeks ago









I dont know why we over think with all those crazy words and alternatives over their..it is just to confuse us. I'm taking back my previous answer.

Ans is B . It is simple and no unnecessary complications. It is so straight forward to use the network Tags and give permission to the partner to SSH into the instance. I tried myself and it is working and it is simple.

upvoted 1 times

  **tavva\_prudhvi** 3 months, 4 weeks ago

May i know why it isn't D, i reckon it does the same!  
upvoted 2 times

-   **nitinz** 5 months, 1 week ago  
Tricky question, Answer is A.  
IAP TCP forwarding allows you to establish an encrypted tunnel over which you can forward SSH, RDP, and other traffic to VM instances.  
<https://cloud.google.com/iap/docs/using-tcp-forwarding>  
upvoted 3 times
-   **victory108** 5 months, 3 weeks ago  
A - Enable Cloud IAP for the Compute Engine instances, and add the operations partner as a Cloud IAP Tunnel User.  
upvoted 1 times
-   **INASR** 6 months ago  
Answer is D  
upvoted 2 times
-   **zanaad231** 7 months, 1 week ago  
D for me  
upvoted 3 times



You have created a code snippet that should be triggered whenever a new file is uploaded to a Cloud Storage bucket. You want to deploy this code snippet. What should you do?


- A. Use App Engine and configure Cloud Scheduler to trigger the application using Pub/Sub.
- B. Use Cloud Functions and configure the bucket as a trigger resource.
- C. Use Google Kubernetes Engine and configure a CronJob to trigger the application using Pub/Sub.
- D. Use Dataflow as a batch job, and configure the bucket as a data source.

**Correct Answer: A**

Reference:

<https://cloud.google.com/scheduler/docs/tut-pub-sub>

-   **ESP\_SAP**

Highly Voted 

 11 months, 1 week ago

Correct Answer is (B):

Google Cloud Storage Triggers

Cloud Functions can respond to change notifications emerging from Google Cloud Storage. These notifications can be configured to trigger in response to various events inside a bucket—object creation, deletion, archiving and metadata updates.

Note: Cloud Functions can only be triggered by Cloud Storage buckets in the same Google Cloud Platform project.

Event types

Cloud Storage events used by Cloud Functions are based on Cloud Pub/Sub Notifications for Google Cloud Storage and can be configured in a similar way.

Supported trigger type values are:

google.storage.object.finalize

google.storage.object.delete

google.storage.object.archive



google.storage.object.metadataUpdate


Object Finalize

Trigger type value: google.storage.object.finalize

This event is sent when a new object is created (or an existing object is overwritten, and a new generation of that object is created) in the bucket.



[https://cloud.google.com/functions/docs/calling/storage#event\\_types](https://cloud.google.com/functions/docs/calling/storage#event_types)

upvoted 20 times
-   **francisco\_guerra**



Highly Voted 

 11 months, 3 weeks ago



The answer is B


upvoted 14 times
-   **SSPC** 11 months, 2 weeks ago

Sure B? Please you could share the link with the Google documentation

upvoted 1 times
-   **Ale1973** 10 months, 3 weeks ago

<https://cloud.google.com/functions/docs/calling/storage>

upvoted 3 times
-   **Zimcruza**



Most Recent 

 1 month, 1 week ago



Question asks:

You want to deploy this code snippet. What should you do?

To me, none of the answers is relevant to DEPLOYMENT - they all are about how you get the trigger to run ...

upvoted 1 times
-   **GoCloud** 4 months, 1 week ago

vote for B

upvoted 1 times
-   **marialix87** 4 months, 1 week ago

I think is B

upvoted 2 times



- 🗨️ 👤 **GCP\_Student1** 4 months, 2 weeks ago  
B. Use Cloud Functions and configure the bucket as a trigger resource.  
upvoted 2 times
- 🗨️ 👤 **lutoa** 5 months, 2 weeks ago  
Answer is B, app engine is for applications, in this case it's just a code snippet which cloud functions is more suitable  
upvoted 2 times
- 🗨️ 👤 **victory108** 5 months, 3 weeks ago  
B - Use Cloud Functions and configure the bucket as a trigger resource.  
upvoted 2 times
- 🗨️ 👤 **Bhagirathi** 8 months, 1 week ago  
Most of you suggest B.  
  
someone chosed A - I will choose B  
upvoted 1 times
- 🗨️ 👤 **swatititame** 8 months, 1 week ago  
• B. Use Cloud Functions and configure the bucket as a trigger resource.  
upvoted 2 times
- 🗨️ 👤 **ritzheg** 8 months, 3 weeks ago  
B.  
<https://cloud.google.com/functions/docs/tutorials/storage#objectives>  
upvoted 2 times
- 🗨️ 👤 **JJ\_ME** 9 months, 3 weeks ago  
I think A.  
You can send notifications to any Pub/Sub topic in any project for which you have sufficient permissions.  
<https://cloud.google.com/storage/docs/pubsub-notifications#overview>  
  
Cloud Functions - Note that your bucket must reside in the same project as Cloud Functions  
[https://cloud.google.com/storage/docs/pubsub-notifications#other\\_notification\\_options](https://cloud.google.com/storage/docs/pubsub-notifications#other_notification_options)  
  
Without knowing if the Storage bucket is in the same project as Cloud Functions, one must use Pub/Sub as the preferred option.  
upvoted 1 times
- 🗨️ 👤 **RockAJ** 9 months, 3 weeks ago  
B for me  
upvoted 2 times
- 🗨️ 👤 **technicalguru** 11 months, 1 week ago  
Agreed with B  
upvoted 1 times
- 🗨️ 👤 **SSPC** 11 months, 2 weeks ago  
Yes B is correct  
upvoted 2 times



You have been asked to set up Object Lifecycle Management for objects stored in storage buckets. The objects are written once and accessed frequently for 30 days. After 30 days, the objects are not read again unless there is a special need. The objects should be kept for three years, and you need to minimize cost.  
What should you do?

- A. Set up a policy that uses Nearline storage for 30 days and then moves to Archive storage for three years.
- B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.
- C. Set up a policy that uses Nearline storage for 30 days, then moves the Coldline for one year, and then moves to Archive storage for two years.
- D. Set up a policy that uses Standard storage for 30 days, then moves to Coldline for one year, and then moves to Archive storage for two years.

**Correct Answer: A**

Reference:



[https://books.google.com.pk/books?id=q0nhDwAAQBAJ&pg=PA52&lpg=PA52&dq=Set+up+a+policy+that+uses+Nearline+storage+for+30+days+and+then+moves+to+Archive+storage+for+three+years.&source=bl&ots=kYLZN1ymA8&sig=ACfU3U2XLmzQ39cmPDwjfWxRbNtDNLc\\_6g&hl=en&sa=X&ved=2ahUKEwjZmefOpr7qAhVzQkEAHTUgASYQ6AEwAHoECAoQAQ#v=onepage&q=Set%20up%20a%20policy%20that%20uses%20Nearline%20storage%20for%2030%20days%20and%20then%20moves%20to%20Archive%20storage%20for%20three%20years.&f=false](https://books.google.com.pk/books?id=q0nhDwAAQBAJ&pg=PA52&lpg=PA52&dq=Set+up+a+policy+that+uses+Nearline+storage+for+30+days+and+then+moves+to+Archive+storage+for+three+years.&source=bl&ots=kYLZN1ymA8&sig=ACfU3U2XLmzQ39cmPDwjfWxRbNtDNLc_6g&hl=en&sa=X&ved=2ahUKEwjZmefOpr7qAhVzQkEAHTUgASYQ6AEwAHoECAoQAQ#v=onepage&q=Set%20up%20a%20policy%20that%20uses%20Nearline%20storage%20for%2030%20days%20and%20then%20moves%20to%20Archive%20storage%20for%20three%20years.&f=false)



  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (B):



The key to understand the requirement is : "The objects are written once and accessed frequently for 30 days"  
Standard Storage  
Standard Storage is best for data that is frequently accessed ("hot" data) and/or stored for only brief periods of time.



Archive Storage  
Archive Storage is the lowest-cost, highly durable storage service for data archiving, online backup, and disaster recovery. Unlike the "coldest" storage services offered by other Cloud providers, your data is available within milliseconds, not hours or days. Archive Storage is the best choice for data that you plan to access less than once a year.


<https://cloud.google.com/storage/docs/storage-classes#standard>  
upvoted 29 times



  **naveedpk00** 10 months, 2 weeks ago  
What if we chose option D to minimize the cost as asked in the question? What do you think?  
upvoted 1 times

  **gcper** 9 months, 2 weeks ago  
It doesn't minimize the costs. Check the costs of coldline vs archival  
upvoted 2 times

  **SSPC** Highly Voted 11 months, 3 weeks ago  
I think the correct one is B. Because Nearline has a 30-day minimum storage duration.  
<https://cloud.google.com/storage/docs/storage-classes>  
upvoted 12 times

  **pepepy** 11 months, 3 weeks ago  
The object should be kept for three years, and you need to minimize cost, after 30 days it will be moved to archive, ans A  
upvoted 1 times

  **pepepy** 11 months, 3 weeks ago  
Sorry you are right accessed frequently for 30 days, its B  
upvoted 6 times

  **AD\_0525** Most Recent 1 month, 2 weeks ago  
Answer should be D. object will be frequently access for first 30 days. cost of data access from standard storage is less than Nearline. Then after 30 days the objects can be moved to archive.  
upvoted 1 times

  **AD\_0525** 1 month, 2 weeks ago

Typo, answer should e B. Description remains the same.

upvoted 1 times

  **jahnu** 3 months ago



My Answer is B. why because Nearline Storage is a low-cost, highly durable storage service for storing infrequently accessed data. so we use standard storage for frequent access.

upvoted 1 times

  **kopper2019** 3 months, 1 week ago



B, Standard since they frequently access the for 30 days and after that rarely accessed and you need to reduce cost so Archive

upvoted 2 times

  **pondai** 4 months, 2 weeks ago

B is my answer after 30 day no read again ,so put in archive is the way to min cost

upvoted 1 times

  **Devgela** 4 months, 2 weeks ago

B for sure

upvoted 1 times

  **neerajgoyal** 4 months, 3 weeks ago

I will go with B

upvoted 1 times

  **GCP\_Student1** 5 months ago

B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.

upvoted 2 times

  **JackGlemins** 5 months ago



I think B is the correct. <https://cloud.google.com/storage/docs/storage-classes>

Standard Storage is best for data that is frequently accessed ("hot" data) and/or stored for only brief periods of time.

Nearline Storage is ideal for data you plan to read or modify on average once per month or less.

Archive Storage is the lowest-cost, highly durable storage service for data archiving

upvoted 1 times

  **Faraz2** 7 months, 3 weeks ago

bhai B hy

upvoted 3 times

  **swatititame** 8 months, 1 week ago

B. Set up a policy that uses Standard storage for 30 days and then moves to Archive storage for three years.

upvoted 1 times

  **adeyemi5700** 9 months, 2 weeks ago

The correct answer is D. The cost would be really high after 3years when you decide to move to archive.

upvoted 1 times

  **JJ\_ME** 9 months, 3 weeks ago

I think A.

The question is about minimizing costs.

The objects are written ONCE and accessed frequently for 30 days.

Nearline has a minimum storage duration of 30 days, so it can be used in this scenario. Nearline storage costs are half the price of Standard.

Multiply that by 30 days and that is quite a bit of savings.

Costs for accessing objects in Standard and Nearline are the same.

Writing the object ONCE in Nearline will higher than in Standard.

Overall, the pricing is better for Nearline to Standard.

[https://cloud.google.com/storage/docs/storage-classes#available\\_storage\\_classes](https://cloud.google.com/storage/docs/storage-classes#available_storage_classes)


<https://cloud.google.com/storage/pricing>

upvoted 1 times

  **Rom0817** 1 month, 1 week ago

take note of "accessed frequently for 30 days" so standard storage to archive. B

upvoted 1 times

  **RockAJ** 9 months, 3 weeks ago

B for me

upvoted 1 times

  **muk5658** 10 months, 1 week ago

Undoubtedly its B , as it clear asks we need to use the file frequently for 30 days which means standard and after 30 days we dont need it. So archive class is the best.

upvoted 2 times



 **[Removed]** 10 months, 2 weeks ago

Its D,

1) standard storage as it is frequently accessed for 30 days . So we can have option B and D

2) It can be accessed after that so if it is required that if kept in nearline storage the breakage in retention changes are less compared to Coldline or Archival. So D

upvoted 2 times

  **[Removed]** 10 months, 1 week ago

Take back my previous comment.

Read frequently for 30 days so standard storage. After that not accessed just to archive so archival storage

Answers B

upvoted 2 times

You are storing sensitive information in a Cloud Storage bucket. For legal reasons, you need to be able to record all requests that read any of the stored data. You want to make sure you comply with these requirements. What should you do?

- A. Enable the Identity Aware Proxy API on the project.
- B. Scan the bucket using the Data Loss Prevention API.
- C. Allow only a single Service Account access to read the data.
- D. Enable Data Access audit logs for the Cloud Storage API.

**Correct Answer: D**  
Reference:  
<https://cloud.google.com/storage/docs/audit-logs>

- francisco\_guerra**

Highly Voted

11 months, 3 weeks ago

D is the correct one  
upvoted 15 times
- SSPC**

11 months, 2 weeks ago

Yes D is the correct  
upvoted 5 times
- ESP\_SAP**

Highly Voted

11 months, 1 week ago

Correct Answer is (D):  
  
Logged information  
Within Cloud Audit Logs, there are two types of logs:  
  
Admin Activity logs: Entries for operations that modify the configuration or metadata of a project, bucket, or object.  
  
Data Access logs: Entries for operations that modify objects or read a project, bucket, or object. There are several sub-types of data access logs:  
  
ADMIN\_READ: Entries for operations that read the configuration or metadata of a project, bucket, or object.  
  
DATA\_READ: Entries for operations that read an object.  
  
DATA\_WRITE: Entries for operations that create or modify an object.  
  
<https://cloud.google.com/storage/docs/audit-logs#types>  
upvoted 9 times
- WakandaF**

Most Recent

2 months, 1 week ago

seems that B is the right!  
  
Cloud Data Loss Prevention (DLP) helps you to understand and manage such sensitive data. It provides fast, scalable classification and redaction for sensitive data elements. Using the Data Loss Prevention API and Cloud Functions, you can automatically scan this data before it is uploaded to the shared storage bucket.  
upvoted 1 times
- YAS007**

1 day, 22 hours ago

the question doesn't ask you to manage or understand sensitive data :  
" you need to be able to record all requests that read any of the stored data"  
upvoted 1 times
- victory108**

4 months, 2 weeks ago

D - Enable Data Access audit logs for the Cloud Storage API.  
upvoted 1 times
- EABDAJA**

4 months, 3 weeks ago



D is correct  
upvoted 1 times
- GCP\_Student1**

5 months ago

D. Enable Data Access audit logs for the Cloud Storage API.  
upvoted 2 times
- swatititime**



8 months, 1 week ago

- D. Enable Data Access audit logs for the Cloud Storage API.upvoted 1 times

  **RockAJ** 9 months, 3 weeks ago

Ans is D

upvoted 1 times

  **pepepy** 11 months, 2 weeks ago

agree with D

upvoted 2 times

## Question #156

## Topic 1

You are the team lead of a group of 10 developers. You provided each developer with an individual Google Cloud Project that they can use as their personal sandbox to experiment with different Google Cloud solutions. You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment. What should you do?

- A. Create a single budget for all projects and configure budget alerts on this budget.
- B. Create a separate billing account per sandbox project and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per billing account.
- C. Create a budget per project and configure budget alerts on all of these budgets.
- D. Create a single billing account for all sandbox projects and enable BigQuery billing exports. Create a Data Studio dashboard to plot the spending per project.

**Correct Answer: C**

Reference:

<https://cloud.google.com/billing/docs/how-to/budgets>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago

Correct Answer is (C):

Set budgets and budget alerts

Overview

Avoid surprises on your bill by creating Cloud Billing budgets to monitor all of your Google Cloud charges in one place. A budget enables you to track your actual Google Cloud spend against your planned spend. After you've set a budget amount, you set budget alert threshold rules that are used to trigger email notifications. Budget alert emails help you stay informed about how your spend is tracking against your budget.

2. Set budget scope

Set the budget Scope and then click Next.

In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all the projects in the Cloud Billing account, choose Select all.

<https://cloud.google.com/billing/docs/how-to/budgets#budget-scope>

upvoted 23 times

  **Hjameel** Highly Voted 11 months, 2 weeks ago

I think C is the best answer.

upvoted 7 times

  **pondai** Most Recent 4 months, 2 weeks ago


Does anyone knows Data Studio can be alert to email?If it can't I'll pick C

upvoted 2 times

  **GCP\_Student1** 5 months ago

C. Create a budget per project and configure budget alerts on all of these budgets.

upvoted 2 times

  **andregrip** 5 months, 3 weeks ago

I believe is A because of this...

Projects: In the Projects field, select one or more projects that you want to apply the budget alert to. To apply the budget alert to all of the projects in the Cloud Billing account, choose Select all.

Some costs are not related to a project, such as the costs of subscriptions or Support costs.

In the budget's project scope, in the list of projects you can filter on, [Charges not specific to a project] is not an option you can select.

If you choose Select all, then the costs in all projects, including Charges not specific to a project, are included in the budget and cost trend chart cost calculations.

If you select one or more projects - but not all projects - then the Charges not specific to a project are not included in the budget and cost trend chart cost calculations.

You can view your costs that are not related to a project in the billing reports. Using the projects filter in the reports page, you can select and view [Charges not specific to a project].

URL: <https://cloud.google.com/billing/docs/how-to/budgets>

upvoted 2 times

  **Bhagirathi** 8 months, 1 week ago

which one correct?

upvoted 1 times

  **swatititame** 8 months, 1 week ago

- C. Create a budget per project and configure budget alerts on all of these budgets.

upvoted 1 times

  **xtian2900** 10 months, 2 weeks ago

is C, not A

with A, i guess if you create a single budget for all projects, together they can easily beat the \$500 mark and you need to know if "one deveoper" did it.

so one budget per project is the solution.

upvoted 2 times

  **[Removed]** 10 months, 2 weeks ago

Yes C is correct, I would have gone with B and D as billing export is the crucial element for billing, but both the option don't talk about notifying about spending.

Also, you don't want to combine the billing for all as each one can spend up to 500 so it will be better if they are individual so A is out.

upvoted 1 times

  **pepepy** 10 months, 4 weeks ago



You want to be notified if any of the developers are spending above \$500 per month on their sandbox environment, so Answer is C

upvoted 2 times

  **tothecloud** 11 months, 2 weeks ago



C is correct

upvoted 4 times

  **SSPC** 11 months, 2 weeks ago

Do you think "C" is correct?

upvoted 2 times

  **SSPC** 11 months, 2 weeks ago

A budget per project?

upvoted 1 times

  **Ale1973** 10 months, 3 weeks ago

Yes, when you are creating a budget, the budget can be scoped to focus on a specific set of resources (Projects, Products or Labels)

upvoted 2 times



Question #157

Topic 1



You are deploying a production application on Compute Engine. You want to prevent anyone from accidentally destroying the instance by clicking the wrong button. What should you do?

- A. Disable the flag `--delete-boot-disk-when-instance-is-deleted`.
- B. Enable delete protection on the instance.
- C. Disable Automatic restart on the instance.
- D. Enable Preemptibility on the instance.

**Correct Answer: A**

Reference:

<https://googlecloudplatform.uservoice.com/forums/302595-compute-engine/suggestions/14227521-set-delete-boot-disk-when-instance-is-deleted-to>

  **ESP\_SAP** Highly Voted 11 months, 1 week ago  
Correct Answer is (B):



Preventing Accidental VM Deletion  
This document describes how to protect specific VM instances from deletion by setting the `deletionProtection` property on an Instance resource. To learn more about VM instances, read the [Instances documentation](#).



As part of your workload, there might be certain VM instances that are critical to running your application or services, such as an instance running a SQL server, a server used as a license manager, and so on. These VM instances might need to stay running indefinitely so you need a way to protect these VMs from being deleted.



By setting the `deletionProtection` flag, a VM instance can be protected from accidental deletion. If a user attempts to delete a VM instance for which you have set the `deletionProtection` flag, the request fails. Only a user that has been granted a role with `compute.instances.create` permission can reset the flag to allow the resource to be deleted.



<https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>  
upvoted 18 times



  **MohammedGhouse** Highly Voted 11 months, 3 weeks ago  
"B" is the answer  
upvoted 10 times



  **NARWAL** Most Recent 3 months, 3 weeks ago  
B should be the answer.  
upvoted 1 times



  **GCP\_Student1** 5 months, 1 week ago  
B. Enable delete protection on the instance.  
upvoted 2 times

  **Lomy** 5 months, 1 week ago  
B  
The ans is B  
upvoted 1 times

  **RockAJ** 9 months, 3 weeks ago  
B for me  
upvoted 2 times

  **muk5658** 10 months, 1 week ago  
Correct Answer is 'B'  
upvoted 3 times

  **SSPC** 11 months, 2 weeks ago  
B is correct answer. <https://cloud.google.com/compute/docs/instances/preventing-accidental-vm-deletion>  
upvoted 5 times

  **francisco\_guerra** 11 months, 3 weeks ago  
I think is B  
upvoted 3 times

## Question #158

## Topic 1

Your company uses a large number of Google Cloud services centralized in a single project. All teams have specific projects for testing and development. The DevOps team needs access to all of the production services in order to perform their job. You want to prevent Google Cloud product changes from broadening their permissions in the future. You want to follow Google-recommended practices. What should you do?

- A. Grant all members of the DevOps team the role of Project Editor on the organization level.
- B. Grant all members of the DevOps team the role of Project Editor on the production project.
- C. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the production project.
- D. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the organization level.

Correct Answer: B

ESP\_SAP Highly Voted 11 months, 1 week ago

Correct Answer is (C):

Understanding IAM custom roles

Key Point: Custom roles enable you to enforce the principle of least privilege, ensuring that the user and service accounts in your organization have only the permissions essential to performing their intended functions.

Basic concepts

Custom roles are user-defined, and allow you to bundle one or more supported permissions to meet your specific needs. Custom roles are not maintained by Google; when new permissions, features, or services are added to Google Cloud, your custom roles will not be updated automatically.

When you create a custom role, you must choose an organization or project to create it in. You can then grant the custom role on the organization or project, as well as any resources within that organization or project.

[https://cloud.google.com/iam/docs/understanding-custom-roles#basic\\_concepts](https://cloud.google.com/iam/docs/understanding-custom-roles#basic_concepts)  
upvoted 22 times

anony1234567 Highly Voted 10 months, 3 weeks ago

Question looks incorrect. In 1st line they are saying one project and in the second line they are saying multiple projects. I don't know if i understood that correctly..  
upvoted 8 times

omunoz 4 months, 2 weeks ago

I believe in the second line they are referring to company projects and no GCP projects... but yes, is confuse..  
upvoted 1 times

rjperdigao 5 months, 1 week ago

please look at the following link:  
[https://console.cloud.google.com/getting-started/enterprise;step=RESOURCE\\_MANAGEMENT](https://console.cloud.google.com/getting-started/enterprise;step=RESOURCE_MANAGEMENT)  
Imagine you have only a shared project for development, testing and production  
upvoted 1 times

jahnu Most Recent 3 months ago

Answer C: custom role enable privileges.  
upvoted 2 times

GCP\_Student1 4 months, 2 weeks ago

C. Create a custom role that combines the required permissions. Grant the DevOps team the custom role on the production project.  
upvoted 2 times

swatititame 8 months, 2 weeks ago

A single folder contains multiple projects. In the below line, it should be a single organization or folder. If we open the GCP console we can see all the services without selecting the project.  
Your company uses a large number of Google Cloud services centralized in a single project.  
upvoted 1 times

SSPC 11 months, 2 weeks ago

C better yes  
upvoted 7 times

SSPC 11 months, 2 weeks ago

What do you think C or D?  
upvoted 1 times

🗨️ 👤 **[Removed]** 10 months, 3 weeks ago

I think it's C because the question says "The DevOps team needs access to all of the production services in order to perform their job" and so to give them the custom permissions across the whole organization would be too permissive for what they need.

upvoted 2 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

"You want to prevent Google Cloud product changes from broadening their permissions in the future." then CUSTOM ROLE

upvoted 7 times

🗨️ 👤 **francisco\_guerra** 11 months, 3 weeks ago

Maybe is C because is only one project

upvoted 4 times

🗨️ 👤 **SSPC** 11 months, 2 weeks ago

I think that the answer is D, because It says "DevOps team"

upvoted 1 times

You are building an application that processes data files uploaded from thousands of suppliers. Your primary goals for the application are data security and the expiration of aged data. You need to design the application to:

- ⌘ Restrict access so that suppliers can access only their own data.
- ⌘ Give suppliers write access to data only for 30 minutes.
- ⌘ Delete data that is over 45 days old.

You have a very short development cycle, and you need to make sure that the application requires minimal maintenance. Which two strategies should you use?

(Choose two.)

- A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.
- B. Use signed URLs to allow suppliers limited time access to store their objects.
- C. Set up an SFTP server for your application, and create a separate user for each supplier.
- D. Build a Cloud function that triggers a timer of 45 days to delete objects that have expired.
- E. Develop a script that loops through all Cloud Storage buckets and deletes any buckets that are older than 45 days.

Correct Answer: AE

- ESP\_SAP

Highly Voted

11 months, 1 week ago

Correct Answers are: (AB):

(A) Object Lifecycle Management  
Delete  
The Delete action deletes an object when the object meets all conditions specified in the lifecycle rule.

Exception: In buckets with Object Versioning enabled, deleting the live version of an object causes it to become a noncurrent version, while deleting a noncurrent version deletes that version permanently.  
<https://cloud.google.com/storage/docs/lifecycle#delete>

(B) Signed URLs  
This page provides an overview of signed URLs, which you use to give time-limited resource access to anyone in possession of the URL, regardless of whether they have a Google account

<https://cloud.google.com/storage/docs/access-control/signed-urls>  
upvoted 20 times
- francisco\_guerra

Highly Voted

11 months, 3 weeks ago

AB is the answer

upvoted 12 times
- jahnu

Most Recent

3 months ago

My Answer: A B

upvoted 2 times
- GCP\_Student1

4 months, 2 weeks ago

A. Build a lifecycle policy to delete Cloud Storage objects after 45 days.

B. Use signed URLs to allow suppliers limited time access to store their objects.

upvoted 2 times
- Bhagirathi

7 months, 1 week ago

A & B helps.

upvoted 1 times
- tuta

7 months ago

did you pass exam ?

upvoted 3 times
- RockAJ

9 months, 3 weeks ago

for me A and B

upvoted 2 times
- [Removed]

10 months, 3 weeks ago

100% A and B.

A. lifecycle policies are designed for situations just like these. It can cleanly schedule deletion of objects that are past 45 days old.

B. signed URLs also can allow the user to 1) access exclusively their own data and 2) access it for a set period of time, in this case 30 minutes  
C. an SFTP server? Really? This is GCP, not circa 2000. Plus it's high maintenance.  
D. Ok, though it's better than C, it's still clunky. Why not use a pre-built lifecycle policy? Plus this doesn't tell us how objects are marked as expired anyways. Sounds like extra work.  
E. Now this one is as bad as C. All that looping isn't sustainable as the dataset grows. Plus it's overly complex, and doesn't even mention how the script gets invoked in the first place. Another can of worms.

A and B for sure.

upvoted 6 times

  **stepkurniawan** 11 months, 1 week ago



But if we implement A: Storage lifecycle -> it will affect all the files, including the files that are not created by the supplier...

upvoted 1 times

  **tothecloud** 11 months, 2 weeks ago

AB is correct

upvoted 3 times

  **SSPC** 11 months, 2 weeks ago



A, E are correct. <https://cloud.google.com/storage/docs/lifecycle>

upvoted 4 times

  **pepepy** 11 months, 2 weeks ago

It says "Give suppliers write access to data only for 30 minutes"

upvoted 1 times

  **SSPC** 11 months, 2 weeks ago

True. Do you think A and B are the answers?



upvoted 1 times

  **pepepy** 11 months, 2 weeks ago

Use signed URLs to allow suppliers limited time access to store their objects.

so A-B

upvoted 4 times

  **SSPC** 11 months, 2 weeks ago

Thank you!

upvoted 2 times

## Question #160

## Topic 1

Your company wants to standardize the creation and management of multiple Google Cloud resources using Infrastructure as Code. You want to minimize the amount of repetitive code needed to manage the environment. What should you do?

- A. Develop templates for the environment using Cloud Deployment Manager.
- B. Use curl in a terminal to send a REST request to the relevant Google API for each individual resource.
- C. Use the Cloud Console interface to provision and manage all related resources.
- D. Create a bash script that contains all requirement steps as gcloud commands.

**Correct Answer: A**

Reference:

<https://cloud.google.com/deployment-manager/docs/fundamentals>  
(see templates)

  **jmgf** Highly Voted 4 months, 2 weeks ago

A

You can use Google Cloud Deployment Manager to create a set of Google Cloud resources and manage them as a unit, called a deployment. For example, if your team's development environment needs two virtual machines (VMs) and a BigQuery database, you can define these resources in a configuration file, and use Deployment Manager to create, change, or delete these resources. You can make the configuration file part of your team's code repository, so that anyone can create the same environment with consistent results.

<https://cloud.google.com/deployment-manager/docs/quickstart>  
upvoted 15 times



  **GCP\_Student1** Highly Voted 4 months, 2 weeks ago

A. Develop templates for the environment using Cloud Deployment Manager.  
upvoted 6 times

  **TAvenger** Most Recent 4 months, 2 weeks ago

A.

According to this document <https://cloud.google.com/solutions/infrastructure-as-code>  
IaC (Infrastructure as code) tools for Google Cloud:  
Deployment Manager, Terraform, Puppet, Chef ...  
upvoted 1 times

  **lio123** 4 months, 2 weeks ago



I think A


upvoted 1 times

You are performing a monthly security check of your Google Cloud environment and want to know who has access to view data stored in your Google Cloud Project. What should you?



- A. Enable Audit Logs for all APIs that are related to data storage.
- B. Review the IAM permissions for any role that allows for data access.
- C. Review the Identity-Aware Proxy settings for each resource.
- D. Create a Data Loss Prevention job.


**Correct Answer:** *B*  
Reference:  
<https://cloud.google.com/compute/docs/access>

-   **GCP\_Student1**



Highly Voted 



 4 months, 2 weeks ago



B. Review the IAM permissions for any role that allows for data access.  
upvoted 13 times
-   **AD\_0525**



Most Recent 

 1 month, 1 week ago

B is the correct option, option A will enable logging from the current time and not useful for past purpose.  
upvoted 3 times
-   **kernel1973** 1 month, 1 week ago

A. Enable Audit Logs for all APIs that are related to data storage.  
upvoted 1 times
-   **kopper2019** 3 months, 1 week ago

B the way to go  
upvoted 3 times
-   **lio123** 4 months, 2 weeks ago

it is B  
upvoted 1 times
-   **greatsam321** 4 months, 2 weeks ago

B seems to be the right answer  
upvoted 2 times

## Question #162

## Topic 1

Your company has embraced a hybrid cloud strategy where some of the applications are deployed on Google Cloud. A Virtual Private Network (VPN) tunnel connects your Virtual Private Cloud (VPC) in Google Cloud with your company's on-premises network. Multiple applications in Google Cloud need to connect to an on-premises database server, and you want to avoid having to change the IP configuration in all of your applications when the IP of the database changes.

What should you do?

- A. Configure Cloud NAT for all subnets of your VPC to be used when egressing from the VM instances.
- B. Create a private zone on Cloud DNS, and configure the applications with the DNS name.
- C. Configure the IP of the database as custom metadata for each instance, and query the metadata server.
- D. Query the Compute Engine internal DNS from the applications to retrieve the IP of the database.

**Correct Answer: A**

  **kopper2019** Highly Voted 3 months, 1 week ago

B,

Forwarding zones

Cloud DNS forwarding zones let you configure target name servers for specific private zones. Using a forwarding zone is one way to implement outbound DNS forwarding from your VPC network.

A Cloud DNS forwarding zone is a special type of Cloud DNS private zone. Instead of creating records within the zone, you specify a set of forwarding targets. Each forwarding target is an IP address of a DNS server, located in your VPC network, or in an on-premises network connected to your VPC network by Cloud VPN or Cloud Interconnect.



A does not apply, that is to provide internet access to resources

C, does not apply

D, I don't get it



so B

upvoted 8 times

  **meh009** 3 months, 1 week ago

Agreed, It's B although I chose A initially. After some careful consideration and understanding how Cloud NAT works, I'm sticking with B <https://cloud.google.com/nat/docs/overview>



upvoted 1 times

  **meh009** 3 months, 1 week ago

Further clarification:

"On-premises clients can resolve records in private zones, forwarding zones, and peering zones for which the VPC network has been authorized. On-premises clients use Cloud VPN or Cloud Interconnect to connect to the VPC network."

upvoted 1 times

  **djgodzilla** 1 month, 2 weeks ago

this is talking about On-premises client resolving nodes outside their network . the question is about how would the application tier within the VPC would resolve the database server . you're confusing the resolution direction my friend

upvoted 1 times

  **djgodzilla** 1 month, 2 weeks ago

It is still B , but it's rather outbound forward that's needed here :

DNS outbound Forwarding :

- Set up outbound forwarding private zones to query on-premises servers (On-prem Authoritative Zone: corp.example.com)
- In Cloud Router , add a custom route advertisement for GCP DNS proxies range 35.199.192.0/19 to the on-premises environment.
- Make sure inbound DNS traffic from 35.199.192.0/19 is allowed on on-prem firewall
- Cloud Router should be learning on-prem network route from On-prem Router

[https://youtu.be/OH\\_Jw8NhEGU?t=1283](https://youtu.be/OH_Jw8NhEGU?t=1283)

[https://cloud.google.com/dns/docs/best-practices#use\\_forwarding\\_zones\\_to\\_query\\_on-premises\\_servers](https://cloud.google.com/dns/docs/best-practices#use_forwarding_zones_to_query_on-premises_servers)

upvoted 2 times

  **gloisv** Most Recent 1 day, 17 hours ago

IT's D, because:

A) Cloud NAT direction will be from the cloud resources to the on prem, an the DB is on prem (It will not work if the IP of the database on prem changes, and you have an VPN you should traffic your data trough it).

B) If you create a private zone and configure the applications, if your servers doesnt use the dns will not work.

C) I think is not suitable. I guess you should re-deploy all your Apps with the new conf if the address change.

D) If you have an A record ip of your DB HOST(wich is on prem) in Engine internal DNS and if it changes, you can update the registry quickly to change it to the new IP address, so it will be the best option for sure.

upvoted 1 times

  **iloveexam** 2 days, 17 hours ago



I choose D.

It means setting up a compute engine as internal DNS that contains a domain that resolves IP address of the database.

upvoted 1 times

🗲️ 👤 **AkshayKalbhor** 1 month ago

Options A, B, C do not make sense to me, I think there is a typo in option D:

D: Query the [Compute Engine] internal DNS from the applications to retrieve the IP of the database.

D: Query the ON-PREMISE internal DNS from the applications to retrieve the IP of the database.

upvoted 2 times

🗲️ 👤 **crazyaboutazure** 1 month ago

Must be A

When you create a Cloud NAT gateway, you can choose to have the gateway automatically allocate regional external IP addresses. Alternatively, you can manually assign a fixed number of regional external IP addresses to the gateway. For details about each method, see NAT IP addresses.

You can configure the number of source ports that each Cloud NAT gateway reserves to each VM for which it should provide NAT services. The VMs for which NAT should be provided are determined by the subnet IP address ranges that the gateway is configured to serve.

upvoted 1 times

🗲️ 👤 **crazyaboutazure** 1 month ago

Sorry it has to be B

<https://gcloud.devoteam.com/blog/google-cloud-platform-dns-forwarding-big-thing-enterprises>

upvoted 2 times

🗲️ 👤 **AD\_0525** 1 month, 2 weeks ago

Answer is A

upvoted 1 times

🗲️ 👤 **mj98** 3 months, 2 weeks ago

Answer is A. Since Cloud DNS uses DNS forwarding. The Cloud DNS will forward it to a single IP specified. So when IP address changes it's not functional for the scenario given here

upvoted 2 times

🗲️ 👤 **Jacky\_YO** 3 months, 3 weeks ago

ANS : A

Cloud DNS use "DNS forwarding" , not configuration DNS Name.

<https://cloud.google.com/dns/docs/best-practices>

upvoted 2 times

🗲️ 👤 **coul** 4 months ago

B doesn't look right. If you have DNS name it will be resolved to IP after all.

upvoted 1 times

🗲️ 👤 **pondai** 4 months ago

[https://cloud.google.com/dns/docs/best-practices#best\\_practices\\_for\\_dns\\_forwarding\\_zones\\_and\\_server\\_policies](https://cloud.google.com/dns/docs/best-practices#best_practices_for_dns_forwarding_zones_and_server_policies)

Cloud DNS offers DNS forwarding zones and DNS server policies to allow lookups of DNS names between your on-premises and Google Cloud environment. You have multiple options for configuring DNS forwarding. The following section lists best practices for hybrid DNS setup. These best practices are illustrated in the Reference architectures for hybrid DNS.

So I think B is correct

upvoted 4 times

🗲️ 👤 **victory108** 4 months, 2 weeks ago

B - Create a private zone on Cloud DNS, and configure the applications with the DNS name.

upvoted 1 times

🗲️ 👤 **GCP\_Student1** 4 months, 2 weeks ago

A. Configure Cloud NAT for all subnets of your VPC to be used when egressing from the VM instances.

upvoted 3 times

🗲️ 👤 **JDoutthere** 4 months, 2 weeks ago

having read some more... seems after all A is better using Cloud Nat.

upvoted 1 times

🗲️ 👤 **TAvenger** 4 months, 2 weeks ago

I've rechecked google documentation, <https://cloud.google.com/nat/docs/overview>

NAT lets Google Cloud virtual machine (VM) instances without external IP addresses and private Google Kubernetes Engine (GKE) clusters send outbound packets to the internet and receive any corresponding established inbound response packets.

There is nothing about hybrid connections or VPNs etc...

In our case it is clearly mentioned that we have hybrid, VPN, VPC ... and the only thing we need is not to be dependent on IP change.

From google documentation private zone on Cloud DNS ("B") will help us to solve this issue

<https://cloud.google.com/dns/docs/best-practices>

upvoted 5 times

🗲️ 👤 **TAvenger** 4 months, 2 weeks ago

Can you provide link to google cloud docs?

According to <https://cloud.google.com/dns/docs/best-practices> it should be "B"

upvoted 4 times

  **TAvenger** 4 months, 2 weeks ago



I'd choose

B. Create a private zone on Cloud DNS, and configure the applications with the DNS name.

Such examples are in google docs:

<https://cloud.google.com/dns/docs/best-practices>

upvoted 4 times

  **dunhill** 4 months, 2 weeks ago

I guess B. Create a private zone on Cloud DNS, and configure the applications with the DNS name.

Since the applications are on Google Cloud and using Cloud DNS to map IP and name should be a good idea no matter how change the IP of the database.

upvoted 2 times

  **JDoutthere** 4 months, 2 weeks ago

I'd use C

upvoted 2 times

## Question #163

## Topic 1

You have developed a containerized web application that will serve internal colleagues during business hours. You want to ensure that no costs are incurred outside of the hours the application is used. You have just created a new Google Cloud project and want to deploy the application. What should you do?

- A. Deploy the container on Cloud Run for Anthos, and set the minimum number of instances to zero.
- B. Deploy the container on Cloud Run (fully managed), and set the minimum number of instances to zero.
- C. Deploy the container on App Engine flexible environment with autoscaling, and set the value min\_instances to zero in the app.yaml.
- D. Deploy the container on App Engine flexible environment with manual scaling, and set the value instances to zero in the app.yaml.

**Correct Answer: C**

  **crysk** Highly Voted 4 months, 2 weeks ago

I think that is B the correct answer, because Cloud Run can scale to 0:  
<https://cloud.google.com/run/docs/about-instance-autoscaling>  
 And App Engine Flexible can't scale to 0, the minimum instance number is 1:  
[https://cloud.google.com/appengine/docs/the-appengine-environments#comparing\\_high-level\\_features](https://cloud.google.com/appengine/docs/the-appengine-environments#comparing_high-level_features)  
 upvoted 13 times

  **pca2b** Highly Voted 4 months ago

B:  
 not A because Anthos is an add-on to GKE clusters, 'new project' means we dont have a GKE cluster to work with  
[https://cloud.google.com/kuberun/docs/architecture-overview#components\\_in\\_the\\_default\\_installation](https://cloud.google.com/kuberun/docs/architecture-overview#components_in_the_default_installation)  
 upvoted 6 times

  **umashankar\_a** Most Recent 4 weeks ago

It's Option B... Folks !!  
 "By default, container instances have min-instances turned off, with a setting of 0"  
 This statement means - By default min instances is set to Zero,unless you want to change it and keep your payload warm for reducing the latency.  
 App Engine Flexible documentation clearly stated that min value cannot be less than 1. Crysk is absolutely right and that leaves us with option B - Cloud Run (Fully Managed.)  
 Please read through the doc supplied by Crysk.  
 upvoted 1 times

  **AkshayKalbhor** 1 month ago

B IS WRONG,  
 Docs state the below line.  
 "By default, container instances have min-instances turned off, with a setting of 0"  
 which means that the "minimum instances" value cannot have a value of '0'.  
  
 Reference : <https://cloud.google.com/run/docs/configuring/min-instances>  
  
 C IS CORRECT,  
 As we can deploy containerized web application and we can set the "min\_instances" value to zero in the app.xml as mentioned in the answer.  
 upvoted 1 times

  **AD\_0525** 1 month, 2 weeks ago

App engine Flexible does not support basic scaling i.e. it can't scale to 0. So answer should be B  
 upvoted 1 times

  **jahnu** 3 months ago


Ans: B Your workloads are automatically scaled out or in to zero depending on the traffic to your app. u can pay only when app is running.  
 upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

B. Deploy the container on Cloud Run (fully managed), and set the minimum number of instances to zero  
 upvoted 2 times

  **TAvenger** 4 months, 2 weeks ago

App Engine Flex has at least 1 instance running.  
 Anthos will charge additional costs, so  
 the answer is  
 "B" CloudRun  
 upvoted 3 times

  **ravixkumar** 4 months, 2 weeks ago

Ans: B  
 Your workloads are automatically scaled out or in to zero depending on the traffic to your app. You only pay when your app is running, billed to the nearest 100 milliseconds.

upvoted 4 times

Question #164

Topic 1

You have experimented with Google Cloud using your own credit card and expensed the costs to your company. Your company wants to streamline the billing process and charge the costs of your projects to their monthly invoice. What should you do?

- A. Grant the financial team the IAM role of “Billing Account User” on the billing account linked to your credit card.
- B. Set up BigQuery billing export and grant your financial department IAM access to query the data.
- C. Create a ticket with Google Billing Support to ask them to send the invoice to your company.
- D. Change the billing account of your projects to the billing account of your company.

Correct Answer: D

- j\_mrn**

Highly Voted

4 months, 2 weeks ago

1000% Ans D

upvoted 13 times
- arsh1916**

Most Recent

2 months, 1 week ago

D is correct

upvoted 2 times
- sivre**

4 months, 1 week ago

Does the user have access to the company billing account? Not sure of D, I think more A.

To change the Cloud Billing account for a project, you need to be able to move a project from one Cloud Billing account to another. To accomplish this task, you need permissions adequate to unlink the project from the existing Cloud Billing account AND to link the project to the target Cloud Billing account.  
Roles with adequate permissions to perform this task: Project Owner or Project Billing Manager on the project, AND Billing Account Administrator or Billing Account User for the target Cloud Billing account

[https://cloud.google.com/billing/docs/how-to/modify-project#change\\_the\\_billing\\_account\\_for\\_a\\_project](https://cloud.google.com/billing/docs/how-to/modify-project#change_the_billing_account_for_a_project)

upvoted 1 times
- GoCloud**

4 months, 1 week ago

A would be ok if assigning Project Billing Manager IAM role to the finance team.

So the answer will be D.

upvoted 1 times
- muneebarshad**

3 weeks, 1 day ago

"Billing Account User" Role Link projects to billing accounts. Since It does not say anything about user permissions , lets assume if user does not have sufficient permission then Granting the financial team the ""Billing Account User" " role will do the job

upvoted 1 times
- GCP\_Student1**

4 months, 2 weeks ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 4 times
- TAvenger**

4 months, 2 weeks ago

D. Change the billing account of your projects to the billing account of your company.

upvoted 2 times
- lio123**

4 months, 2 weeks ago

yes it is D

upvoted 2 times

## Question #165

## Topic 1

You are running a data warehouse on BigQuery. A partner company is offering a recommendation engine based on the data in your data warehouse. The partner company is also running their application on Google Cloud. They manage the resources in their own project, but they need access to the BigQuery dataset in your project. You want to provide the partner company with access to the dataset. What should you do?

- A. Create a Service Account in your own project, and grant this Service Account access to BigQuery in your project.
- B. Create a Service Account in your own project, and ask the partner to grant this Service Account access to BigQuery in their project.
- C. Ask the partner to create a Service Account in their project, and have them give the Service Account access to BigQuery in their project.
- D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.**

**Correct Answer:** D

  **GCP\_Student1** Highly Voted 4 months, 2 weeks ago

D. Ask the partner to create a Service Account in their project, and grant their Service Account access to the BigQuery dataset in your project.  
upvoted 14 times

  **pondai** Highly Voted 4 months ago



BigQuery is in our project,so we need to create a service account and grant it access BigQuery role.That can make partner company to use this account to use it to access our project's BigQuery.So I vote A.  
upvoted 7 times

  **tavva\_prudhvi** 3 months, 4 weeks ago

See, the ones who want our access needs to create a service account(in our case it's the partner company), then we give access to the service account with the user permissions. Clearly, D says the same thing!  
upvoted 4 times

  **AD\_0525** Most Recent 1 month, 1 week ago

Answer should be D, as the other company project needs access in your project.  
upvoted 3 times

  **EmreA** 2 months, 2 weeks ago



It's A  
upvoted 2 times

  **kopper2019** 3 months, 1 week ago



D is the answer  
upvoted 4 times

  **TAvenger** 4 months, 2 weeks ago

I think it is "D" but I have not found such usecase when you share dataset with another organization via service account  
upvoted 2 times

  **lio123** 4 months, 2 weeks ago

it is D  
upvoted 3 times

  **nshtmnt** 4 months, 2 weeks ago

D looks correct to me.  
upvoted 4 times

Your web application has been running successfully on Cloud Run for Anthos. You want to evaluate an updated version of the application with a specific percentage of your production users (canary deployment). What should you do?

- A. Create a new service with the new version of the application. Split traffic between this version and the version that is currently running.
- B. Create** a new revision with the new version of the application. Split traffic between this version and the version that is currently running.
- C. Create a new service with the new version of the application. Add HTTP Load Balancer in front of both services.
- D. Create a new revision with the new version of the application. Add HTTP Load Balancer in front of both revisions.

**Correct Answer: A**

- crysk**

Highly Voted

4 months, 2 weeks ago

In my opinion correct answer is B:  
[https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration?utm\\_campaign=CDR\\_ahm\\_aap-severless\\_cloud-run-faq\\_&utm\\_source=external&utm\\_medium=web](https://cloud.google.com/run/docs/rollouts-rollbacks-traffic-migration?utm_campaign=CDR_ahm_aap-severless_cloud-run-faq_&utm_source=external&utm_medium=web)

Cloud Run can split traffic between revisions  
upvoted 18 times
- TAvenger**

4 months, 2 weeks ago

The google doc link is incorrect. You need to specify CloudRun for Anthos  
<https://cloud.google.com/kuberun/docs/rollouts-rollbacks-traffic-migration>

Anyway principles for CloudRun and CloundRun for Anthos are the same. Traffic can be split between multiple revisions.  
The answer is "B"  
upvoted 2 times
- GCP\_Student1**

Highly Voted

4 months, 2 weeks ago

B. Create a new revision with the new version of the application. Split traffic between this version and the version that is currently running.  
upvoted 6 times
- AkshayKalbhor**

Most Recent

1 month ago

Answer should be B.  
upvoted 2 times
- RishiAbhishek**

1 month, 2 weeks ago

Canary deployments are a method of releasing software to a subset of users or servers. The plan is to deliver the update to a small selection of servers first, test it, and then roll it out to the other servers.

So it's a Revision or new update or a new version, not a service.  
upvoted 1 times
- nana1995**

1 month, 3 weeks ago

Its B  
<https://cloud.google.com/run/docs/managing/revisions>  
upvoted 1 times
- Ashii**

2 months, 1 week ago

Its B, revision or version, not service.  
upvoted 1 times

## Question #167

## Topic 1

Your company developed a mobile game that is deployed on Google Cloud. Gamers are connecting to the game with their personal phones over the Internet. The game sends UDP packets to update the servers about the gamers' actions while they are playing in multiplayer mode. Your game backend can scale over multiple virtual machines (VMs), and you want to expose the VMs over a single IP address. What should you do?

- A. Configure an SSL Proxy load balancer in front of the application servers.
- B. Configure an Internal UDP load balancer in front of the application servers.
- C. Configure an External HTTP(s) load balancer in front of the application servers.
- D. Configure an External Network load balancer in front of the application servers.**

**Correct Answer: A**

Reference:

<https://cloud.google.com/solutions/connecting-securely>

  **kopper2019** Highly Voted 3 months, 1 week ago

Answer is D, cell phones are sending UDP packets and the only that can receive that type of traffic is a External Network TCP/UDP  
<https://cloud.google.com/load-balancing/docs/network>  
upvoted 11 times

  **associatecloudexamuser** Most Recent 2 weeks, 2 days ago

Oh. Missed "Internal" term. So it should be External Load Balancer.  
Answer is D.  
upvoted 3 times

  **JH86** 1 month, 2 weeks ago

Answer is D. there are so many confusion here, from B,C or D. For myself im eliminating all options except B,D due to the traffic type. which leaves me with B or D. Then next the traffic source either external or internal which in this case is an external traffic from the internet, therefore my final answer is D.

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>  
upvoted 3 times

  **curious\_kitty9** 2 months, 1 week ago

Answer is D, Reason- UDP and preserve client ID  
upvoted 1 times

  **jahnu** 3 months ago

Ans: B check in <https://cloud.google.com/load-balancing/docs/internal/multiple-forwarding-rules-same-ip>  
upvoted 2 times

  **mj98** 3 months, 2 weeks ago

D is the correct answer. External Network Load Balancing does provide a single IP. Also it allows UDP traffic too  
upvoted 4 times

  **tavva\_prudhvi** 3 months, 4 weeks ago

Now, as we have UDP packets it's either External Network TCP/UDP Load balancing or Internal TCP/UDP Load balancing. But, at the end, it also mentions to "expose the VMs over a single IP address" then i have seen the official documentation which clearly mentions the point.  
<https://cloud.google.com/load-balancing/docs/https>  
upvoted 3 times

  **yuvarajmrec** 4 months ago

answer is D  
upvoted 3 times

  **victory108** 4 months, 2 weeks ago

C - Configure an External HTTP(s) load balancer in front of the application servers.  
Network Load balancer is regional and not global.  
upvoted 3 times

  **tavva\_prudhvi** 3 months, 4 weeks ago

Yeah, but what about the UDP packets mentioned here, will you choose D then?  
upvoted 2 times

  **victory108** 3 months, 3 weeks ago

No, requirement is to expose the VMs over a single IP address. UDP packets just sent to update the servers, it's not asked to serve UDP traffic.  
upvoted 2 times



🗳️ 👤 **jcolds** 1 month, 4 weeks ago

Here, <https://cloud.google.com/load-balancing/images/choose-lb.svg> specifies Internet-to-GCP traffic, which means incoming UDP traffic. Also, it's possible to expose the VMs with a single external IP address: [https://cloud.google.com/load-balancing/docs/features#ip\\_addresses](https://cloud.google.com/load-balancing/docs/features#ip_addresses).

upvoted 3 times

🗳️ 👤 **GCP\_Student1** 4 months, 2 weeks ago

C. Configure an External HTTP(s) load balancer in front of the application servers.

upvoted 3 times

🗳️ 👤 **TAvenger** 4 months, 2 weeks ago

Read some theory! Or try to google!

Here is the link

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>

UDP traffic is possible with 2 options:

- External Network TCP/UDP Load balancing

- Internal TCP/UDP Load balancing

For LoadBalancer to be available outside of GCP we will need External Network TCP/UDP LoadBalancer.

"D" is correct answer!

upvoted 4 times

🗳️ 👤 **GCP\_Student1** 4 months, 2 weeks ago

It will make life easier if we understand the question correctly.

1. I think the load balancer behavior would not change for the type of external traffic received, be it TCP/UDP etc. It will distribute it to the appropriate recipient at the backend.

2. This question asks only one thing. (To expose the VMs on a SINGLE IP Address). This is not doable with External Network TCP/UDP Load Balancer.

upvoted 3 times

🗳️ 👤 **pca2b** 4 months ago

you want to expose the VMs over a single IP address. What should you do?

...just looking at backends, they will need an internal TCP/UDP LB, to provide a single IP for multiple VMs.

Game users will use an external NLB also, but that is not the point of the question.

so I vote for B

upvoted 2 times

🗳️ 👤 **tavva\_prudhvi** 3 months, 4 weeks ago

Sorry, can you mention the supporting articles stating the same?

upvoted 1 times

🗳️ 👤 **pca2b** 3 months, 3 weeks ago

the game backends need to be consolidated behind 1x IP address. messages are coming to backend in UDP, so a UDP internal LB is needed. i.e. B. Please follow the flowchart: <https://cloud.google.com/load-balancing/images/choose-lb.svg>

Not to confuse with several discussions here about external LB. there may be a need for that too...but the question is asking specifically about the backends behind ONE IP address. That has to be TCP/UDP internal LB, imho.

upvoted 1 times

🗳️ 👤 **iri\_gcp** 4 months, 2 weeks ago

Correct is D - External Network load balancer. <https://cloud.google.com/load-balancing/docs/network>

upvoted 3 times

🗳️ 👤 **GCP\_Student1** 4 months, 2 weeks ago

The main part of the question is in quotes

"you want to expose the VMs over a SINGLE IP address."

This is possible only with HTTPs load balancing. The HTTPs Load Balancer terminate all incoming traffic at load balancer point and establish new connection to VMs in the backend. This way the outside world can only see a SINGLE IP address of the load balancer.

upvoted 5 times

🗳️ 👤 **JieHeng** 1 month ago

I thought External Network load balancer can also have a single static external IP address? [https://cloud.google.com/load-balancing/docs/network/setting-up-network-backend-service#testing\\_the\\_load\\_balancer](https://cloud.google.com/load-balancing/docs/network/setting-up-network-backend-service#testing_the_load_balancer)

upvoted 1 times

🗳️ 👤 **yuvarajmrec** 4 months ago

HTTP loadbalancer does not terminate incoming traffic at loadbalancer. it just balances the traffic among backend resources

upvoted 1 times

🗳️ 👤 **TAvenger** 4 months, 2 weeks ago

The answer is "D" !!!

Read some theory! Or try to google!

Here is the link. Check what types of LoadBalancers support UDP

<https://cloud.google.com/load-balancing/docs/choosing-load-balancer>



upvoted 1 times

🗳️ 👤 **GCP\_Student1** 4 months, 2 weeks ago



Follow this link for more info on this;

<https://cloud.google.com/load-balancing/docs/network>  
upvoted 3 times

  **alexin2** 4 months, 2 weeks ago



Answer is D as it states it that gamers need to send "UDP packets" - Therefore external network loadbalancer.  
See: <https://cloud.google.com/load-balancing/docs/choosing-load-balancer>  
upvoted 3 times

  **GCP\_Student1** 4 months, 2 weeks ago


How about this sentence in question ?  
"you want to expose the VMs over a SINGLE IP address."  
upvoted 4 times

  **TAvenger** 4 months, 2 weeks ago

External Network TCP/UDP LoadBalancer will provide you single IP address. What's the problem?  
upvoted 3 times

  **pdbeta** 4 months, 2 weeks ago

Answer is D. because UPD Packet  
upvoted 3 times

  **raman23125** 4 months, 2 weeks ago

answer is d  
For UDP traffic, use:  
Network Load Balancing  
Internal TCP/UDP Load Balancing  
upvoted 4 times

  **Jamaal\_a** 4 months, 2 weeks ago

Answer is C - Google Cloud HTTP(S) Load Balancing is a global, proxy-based Layer 7 load balancer that enables you to run and scale your services worldwide behind a single external IP address.  
upvoted 3 times

## Question #168

## Topic 1

You are working for a hospital that stores its medical images in an on-premises data room. The hospital wants to use Cloud Storage for archival storage of these images. The hospital wants an automated process to upload any new medical images to Cloud Storage. You need to design and implement a solution. What should you do?

- A. Create a Pub/Sub topic, and enable a Cloud Storage trigger for the Pub/Sub topic. Create an application that sends all medical images to the Pub/Sub topic.
- B. Deploy a Dataflow job from the batch template, `gsutil` Datastore to Cloud Storage. Schedule the batch job on the desired interval.
- C. Create a script** that uses the `gsutil` command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.
- D. In the Cloud Console, go to Cloud Storage. Upload the relevant images to the appropriate bucket.

**Correct Answer: C**

  **GCP\_Student1** Highly Voted 4 months, 2 weeks ago

C. Create a script that uses the `gsutil` command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.

upvoted 10 times

  **TAvenger** Highly Voted 4 months, 2 weeks ago

From the question the key point is "upload ANY NEW medical images to Cloud Storage". So we are not interested in old images. That's why we need some trigger that will upload images. I think option "A" with PubSub is the best

upvoted 6 times

  **pca2b** 4 months ago

Pub/Sub will be good for all future files in in-prem data-storage.

we want to sync all + new, so a local on-prem server running a cron job (not GCE CronJob) to run `gsutil` to transfer files to Cloud Storage would work.

I vote for C

upvoted 2 times

  **yvinisiupacuando** 2 months, 3 weeks ago

Sorry you are wrong, the question clearly indicates "The hospital wants an automated process to upload ANY NEW medical images to Cloud Storage." It does not mention the need to upload the original stock of images, only the new ones. Then I think the right answer must be A, as you said "Pub/sub will be good for all future files in prem data-storage" which is exactly what the questions is pointing to.

upvoted 3 times

  **dunhill** 4 months, 2 weeks ago

I am not sure but the question also mentions that "wants to use Cloud Storage for archival storage of these images". It can create an application that sends all medical images to storage and no need via PubSub?

upvoted 1 times

  **jackdbd** Most Recent 3 days, 22 hours ago

I would go with C.

Here is my rationale for excluding option A.

Pub/Sub accepts a maximum of 1,000 messages in a batch, and the size of a batch can not exceed 10 megabytes. A single CT scan can be 35 MB, so unless we chunk it up in many smaller pieces and publish them as smaller messages, using a Pub/Sub topic is not a viable alternative.

upvoted 1 times

  **associatecloudexamuser** 2 weeks, 2 days ago

Answer is C. `gsutil rsync <source_location> <destination_location>`. This can sync content with Google cloud storage locations

upvoted 1 times

  **JieHeng** 1 month ago

I would go with C

Not A, don't think you can send image files to Pub/Sub. Technically you can do so by converting image to some binary text, but then we don't know the size of the image and there is a limitation on message size. Not recommended.

Not B – there is only this template "Datastore to Cloud Storage Text", as the name implies it is for text, <https://cloud.google.com/dataflow/docs/guides/templates/provided-batch#datastore-to-cloud-storage-text>, and it reads from datastore which is definitely not where the medical images are stored, from the question "... stores its medical images in an on-premises data room".

Not D – it's not automated

upvoted 1 times

- 🗒️ 👤 **Umesh09** 2 months, 1 week ago  
I am really confuse the question states to upload new images to cloud why do we need the new images to be uploaded first in on premises datastore and then syn it using gsutil  
upvoted 1 times
- 🗒️ 👤 **arsh1916** 2 months, 1 week ago  
C 100%  
upvoted 1 times
- 🗒️ 👤 **[Removed]** 3 months, 2 weeks ago  
<https://stackoverflow.com/questions/37662416/how-to-sync-a-local-folder-with-a-folder-in-a-google-cloud-platform-bucket>  
Answer is C  
gsutil rsync + cron job  
upvoted 4 times
- 🗒️ 👤 **victory108** 3 months, 3 weeks ago  
Need to upload new images to Cloud Storage. Option A doesn't say anything about that.  
C is correct - Create a script that uses the gsutil command line interface to synchronize the on-premises storage with Cloud Storage. Schedule the script as a cron job.  
upvoted 3 times
- 🗒️ 👤 **GoCloud** 4 months, 1 week ago  
Cron jobs only support Http, pub/sub and app engine http target types.  
  
The answer is B.  
upvoted 2 times
- 🗒️ 👤 **GCP\_user** 4 months, 2 weeks ago  
Just don't know how to use gsutil command to sync on premises storage and how they connect to each other (On premises storage and cloud storage.) B may be possible however need to verify.  
upvoted 2 times
- 🗒️ 👤 **kopper2019** 3 months, 1 week ago  
you install it in a Windows o linux machine for example and a cron job to run let's say every 30 min  
<https://cloud.google.com/storage/docs/gsutil/commands/cp>  
upvoted 1 times
- 🗒️ 👤 **ace\_student** 4 months, 1 week ago  
Dataflow is stream and batch data processing. In this scenario we are not processing any data, we are just uploading images to GCS. I would go for option C  
upvoted 2 times
- 🗒️ 👤 **raman23125** 4 months, 2 weeks ago  
answer C  
gsutil rsync + cron job  
upvoted 1 times

Your auditor wants to view your organization's use of data in Google Cloud. The auditor is most interested in auditing who accessed data in Cloud Storage buckets. You need to help the auditor access the data they need. What should you do?

- A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.
- B. Assign the appropriate permissions, and then create a Data Studio report on Admin Activity Audit Logs.
- C. Assign the appropriate permissions, and the use Cloud Monitoring to review metrics.
- D. Use the export logs API to provide the Admin Activity Audit Logs in the format they want.

**Correct Answer:** D

Reference:

<https://cloud.google.com/storage/docs/audit-logging>

- iri\_gcp**

Highly Voted

4 months, 2 weeks ago

It should be A.  
Data access log are not enabled by default due to the fact that it incurs costs.  
So you need to enable it first.  
And then you can filter it in the log viewer  
upvoted 16 times
- GCP\_Student1**

Highly Voted

4 months, 2 weeks ago

A. Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.  
upvoted 6 times
- learnazureportal**

Most Recent

1 month, 4 weeks ago

based on the <https://cloud.google.com/storage/docs/audit-logging> link, Answer is D.  
upvoted 4 times
- JieHeng**

1 month ago

from the link, it's the Data Access logs that contain "Entries for operations that read the configuration or metadata of a project, bucket, or object." and "Entries for operations that read an object."  
Admin Activity logs do not have these.  
upvoted 3 times
- arsh1916**

2 months, 1 week ago

A is correct option  
upvoted 1 times
- jahnu**

3 months ago

Ans: A Turn on Data Access Logs for the buckets they want to audit, and then build a query in the log viewer that filters on Cloud Storage.  
upvoted 2 times
- mj98**

3 months, 2 weeks ago

A is the correct option  
upvoted 1 times
- relaxm**

4 months ago

Excuse me Should be D  
upvoted 2 times
- GCPACE2020**

4 weeks, 1 day ago

you are excused, but still answer is A  
upvoted 2 times
- mj98**

3 months, 2 weeks ago


Wrong. It's A  
upvoted 1 times
- tavva\_prudhvi**

3 months, 4 weeks ago


Initially, you mentioned it B, I don't know why and now you say it's D again. it would be better if you mention few supporting articles w.r.t the option you choose. Then, we can discuss it.  
upvoted 2 times
- relaxm**

4 months ago

In my opinion B  
upvoted 1 times


-  **GCP\_user** 4 months, 2 weeks ago

Agree with option: A

upvoted 1 times
-  **TAvenger** 4 months, 2 weeks ago

"A"

This is the only option that allows to read Data access logs

upvoted 1 times
-  **JDoutthere** 4 months, 2 weeks ago

I'd go for A

upvoted 1 times


Question #170


Topic 1

You received a JSON file that contained a private key of a Service Account in order to get access to several resources in a Google Cloud project. You downloaded and installed the Cloud SDK and want to use this private key for authentication and authorization when performing gcloud commands. What should you do?

- A. Use the command gcloud auth login and point it to the private key.
- B. Use the command gcloud auth activate-service-account and point it to the private key.
- C. Place the private key file in the installation directory of the Cloud SDK and rename it to `credentials.json`.
- D. Place the private key file in your home directory and rename it to `GOOGLE_APPLICATION_CREDENTIALS`.

**Correct Answer:** B  
Reference:  
<https://cloud.google.com/sdk/docs/authorizing>

-  **GCP\_Student1**

Highly Voted 


 4 months, 2 weeks ago


B. Use the command gcloud auth activate-service-account and point it to the private key.

Authorizing with a service account

gcloud auth activate-service-account authorizes access using a service account. As with gcloud init and gcloud auth login, this command saves the service account credentials to the local system on successful completion and sets the specified account as the active account in your Cloud SDK configuration.

[https://cloud.google.com/sdk/docs/authorizing#authorizing\\_with\\_a\\_service\\_account](https://cloud.google.com/sdk/docs/authorizing#authorizing_with_a_service_account)

upvoted 18 times
-  **TAvenger**

Highly Voted 

 4 months, 2 weeks ago

B.

gcloud auth activate-service-account --help

NAME)

gcloud auth activate-service-account - authorize access to Google Cloud Platform with a service account

SYNOPSIS


gcloud auth activate-service-account [ACCOUNT] --key-file=KEY\_FILE


[--password-file=PASSWORD\_FILE | --prompt-for-password]

[GLOUD\_WIDE\_FLAG ...]

DESCRIPTION

To allow gcloud (and other tools in Cloud SDK) to use service account credentials to make requests, use this command to import these credentials from a file that contains a private authorization key, and activate them for use in gcloud. gcloud auth activate-service-account serves the same function as gcloud auth login but uses a service account rather than Google user credentials.

upvoted 8 times
-  **ravixkumar**

Most Recent 

 4 months, 2 weeks ago

Ans : B

gcloud auth activate-service-account --key-file=/test-service-account.json

upvoted 6 times

## Question #171

## Topic 1

You are working with a Cloud SQL MySQL database at your company. You need to retain a month-end copy of the database for three years for audit purposes.

What should you do?

- A. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.
- B. Save the automatic first-of-the-month backup for three years. Store the backup file in an Archive class Cloud Storage bucket.
- C. Set up an on-demand backup for the first of the month. Write the backup to an Archive class Cloud Storage bucket.
- D. Convert the automatic first-of-the-month backup to an export file. Write the export file to a Coldline class Cloud Storage bucket.

**Correct Answer: B**

  **TAvenger** Highly Voted 4 months, 2 weeks ago

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

not B: Automatic backups are made EVERY SINGLE DAY. You can set only the number of backups up to 365. Also you cannot choose your Archival storage as destination

not C: You cannot setup "on-demand" backup. User would have to make backups manually every month. Also you cannot choose your Archival storage as destination

not D: You cannot convert backup to export file. Also Coldline class is less cost-effective than Archival class.

The only option left is "A"

You can set up your job with any date/time schedule. You can export file to any storage with any storage class.

upvoted 20 times

  **djgodzilla** 1 month, 2 weeks ago

from the same link :

Can I export a backup?

No, you can't export a backup. You can only export instance data. See Exporting data from Cloud SQL to a dump in Cloud storage.

upvoted 2 times

  **JieHeng** Most Recent 1 month ago

First need to understand backup vs export, two different concepts. - <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

A – yes, you can export data from Cloud SQL to Cloud Storage- <https://cloud.google.com/sql/docs/mysql/import-export/exporting#cloud-sql>

Not B, C, D – be it automatic or on-demand backup, according to the doc "No, you can't export a backup. You can only export instance data."

upvoted 2 times

  **arsh1916** 2 months, 1 week ago

B is correct


upvoted 1 times

  **jahnu** 3 months ago

Ans: B Backups are managed by Cloud SQL according to retention policies, and are stored separately from the Cloud SQL instance. so take the backup and store in Bucket.

check: <https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 1 times

  **kopper2019** 3 months, 1 week ago

Answer is A, there is no way to take the backups made by SQL-MYSQL and send them to Cloud Storage so you have to make an export... Default backup only allows 365 retention points so you need to make export I mean program exports and send them to Cloud Datastore so A.

Backups are managed by Cloud SQL according to retention policies, and are stored separately from the Cloud SQL instance. Cloud SQL backups differ from an export uploaded to Cloud Storage, where you manage the lifecycle. Backups encompass the entire database. Exports can select specific contents.

<https://cloud.google.com/sql/docs/mysql/backup-recovery/backups>

upvoted 1 times

  **GCP\_Student1** 4 months, 2 weeks ago

A. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.

upvoted 2 times

  **victory108** 4 months, 2 weeks ago

B- Save the automatic first-of-the-month backup for three years. Store the backup file in an Archive class Cloud Storage bucket.

upvoted 3 times

  **victory108** 4 months ago

A is correct one. Set up an export job for the first of the month. Write the export file to an Archive class Cloud Storage bucket.

upvoted 1 times



You are monitoring an application and receive user feedback that a specific error is spiking. You notice that the error is caused by a Service Account having insufficient permissions. You are able to solve the problem but want to be notified if the problem recurs. What should you do?


- A. In the Log Viewer, filter the logs on severity `ERROR` and the name of the Service Account.
- B. Create a sink to BigQuery to export all the logs. Create a Data Studio dashboard on the exported logs.
- C. Create a custom log-based metric for the specific error to be used in an Alerting Policy.
- D. Grant Project Owner access to the Service Account.

**Correct Answer: A**

Reference:



<https://cloud.google.com/logging/docs/view/advanced-queries>


-   **GCP\_Student1**

Highly Voted 

 4 months, 2 weeks ago



C. Create a custom log-based metrics for the specific error to be used in an Alerting Policy.


upvoted 11 times
-   **greatsam321**

Highly Voted 

 4 months, 2 weeks ago



C seems to be the right answer.

upvoted 8 times
-   **AD\_0525**



Most Recent 

 1 month, 1 week ago



You want to be aerted next time, so only option C meets that criteria.

upvoted 2 times
-   **Enzo** 1 month, 4 weeks ago

C is correct : You are able to solve the problem but want to be notified if the problem recurs.

upvoted 3 times
-   **arsh1916** 2 months, 1 week ago

A is correct

upvoted 1 times
-   **kopper2019** 3 months, 1 week ago

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?



A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.

B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.



C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team’s internal data warehouse.

D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.



it's A

upvoted 2 times
-   **lxgywil** 2 months, 3 weeks ago



The question is saying "want to be notified if the problem recurs", I don't see how A meets that requirement.

upvoted 1 times
-   **Petza** 3 months ago



A. "Every time the dashboard is refreshed, it pulls new data from the view, which in turn dynamically reflects the latest data in BigQuery". Data Science on the Google Cloud Platform: Implementing End-to-End Real-Time (C)

upvoted 2 times
-   **mj98** 3 months ago

I think it's D. Can anyone confirm?

upvoted 2 times
-   **kopper2019** 3 months, 1 week ago

Question 178

upvoted 1 times
-   **kopper2019** 3 months, 1 week ago

You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

A. Create and deploy a Custom Resource Definition per microservice.



- B. Create and deploy a Docker Compose File.
- C. Create and deploy a Job per microservice.
- D. Create and deploy a Deployment per microservice.


I think is D

upvoted 3 times

  **AmineHM** 3 months, 1 week ago

I think B

upvoted 1 times

  **AmineHM** 3 months, 1 week ago

It's D, yes !

upvoted 1 times

  **noreen** 3 months ago


its c; Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a YAML file to configure your application's services. Then, with a single command, you create and start all the services from your configuration.

upvoted 1 times

  **noreen** 3 months ago

b \* B. Create and deploy a Docker Compose File.

upvoted 1 times

  **kopper2019** 3 months, 1 week ago

You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.
- B. Provision Compute Engine instances with GPUs attached.
- C. Provision Compute Engine instances with local SSDs attached.
- D. Provision Compute Engine instances with M1 machine type.

Answer is D type M,

upvoted 3 times

  **AmineHM** 3 months, 1 week ago

it's D

upvoted 1 times

  **kopper2019** 3 months, 1 week ago

hey guys go Question number 1 for New Question from 173 to 179

upvoted 3 times

  **AmineHM** 3 months, 1 week ago

I don't find 178 and question for 179

upvoted 1 times

  **NARWAL** 3 months, 3 weeks ago

based on the "You are able to solve the problem but want to be notified if the problem recurs." , we definitely need to create an alerting policy. So Answer should be C.

<https://cloud.google.com/logging/docs/logs-based-metrics/charts-and-alerts>



upvoted 3 times

  **TAvenger** 4 months, 2 weeks ago

C.

Just checked in GC Console. Created custom log-based metric and set up an alert.



upvoted 2 times

  **jackycc** 4 months, 2 weeks ago

Answer is C.

<https://cloud.google.com/logging/docs/logs-based-metrics/charts-and-alerts>

upvoted 2 times

  **ravixkumar** 4 months, 2 weeks ago

Ans : C

upvoted 5 times

## Question #173

## Topic 1

You are developing a financial trading application that will be used globally. Data is stored and queried using a relational structure, and clients from all over the world should get the exact identical state of the data. The application will be deployed in multiple regions to provide the lowest latency to end users. You need to select a storage option for the application data while minimizing latency. What should you do?

- A. Use Cloud Bigtable for data storage.
- B. Use Cloud SQL for data storage.
- C. Use Cloud Spanner for data storage.
- D. Use Firestore for data storage.

**Correct Answer:** C

Reference:

<https://cloud.google.com/solutions/best-practices-compute-engine-region-selection>

🗨️ 👤 **associatecloudexamuser** 2 weeks, 2 days ago

Yes. C is right answer  
upvoted 2 times

🗨️ 👤 **JieHeng** 1 month ago

C, Cloud Spanner, keywords are globally, relational structure and lastly "clients from all over the world should get the exact identical state of the data" which implies strong consistency is needed.  
upvoted 3 times

🗨️ 👤 **Fidget\_** 2 months, 4 weeks ago

C) Spanner - Global, low latency, relational  
upvoted 2 times

🗨️ 👤 **Rightsaidfred** 2 months, 4 weeks ago

Yes C, Cloud Spanner  
upvoted 1 times

🗨️ 👤 **norrec9** 3 months ago

C is the answer  
upvoted 1 times

🗨️ 👤 **Biju1** 3 months ago

C Cloud Spanner to store relational data  
upvoted 1 times

## Question #174

## Topic 1

You are about to deploy a new Enterprise Resource Planning (ERP) system on Google Cloud. The application holds the full database in-memory for fast data access, and you need to configure the most appropriate resources on Google Cloud for this application. What should you do?

- A. Provision preemptible Compute Engine instances.
- B. Provision Compute Engine instances with GPUs attached.
- C. Provision Compute Engine instances with local SSDs attached.
- D. Provision** Compute Engine instances with M1 machine type.

**Correct Answer: C**

Reference:

<https://cloud.google.com/compute/docs/disks/local-ssd>

  **Rightsaidfred** Highly Voted 2 months, 4 weeks ago

Yes D, M1 Machine types for ERP i.e. SAP-HANA:  
<https://cloud.google.com/compute/docs/machine-types>  
upvoted 12 times

  **jackdbd** Most Recent 4 days, 15 hours ago

Note that VM instances m1-megamem-96 are both from the M1 family AND can have local SSDs attached to them.  
[https://cloud.google.com/compute/docs/memory-optimized-machines#m1\\_vms](https://cloud.google.com/compute/docs/memory-optimized-machines#m1_vms)  
upvoted 1 times

  **associatecloudexamuser** 2 weeks, 2 days ago



Answer is D.  
Applications of Memory optimized VMs are,  
1. Medium-large in-memory databases such as SAP HANA  
2. In-memory databases and in-memory analytics  
3. Microsoft SQL Server and similar databases  
upvoted 1 times

  **JieHeng** 1 month ago

D, "The application holds the full database in-memory for fast data access", so it'll be more appropriate to use memory-optimized machine types -  
[https://cloud.google.com/compute/docs/machine-types#m1\\_machine\\_types](https://cloud.google.com/compute/docs/machine-types#m1_machine_types)  
upvoted 2 times

  **AD\_0525** 1 month, 1 week ago

Here the question is about data access not work load type, hence local SSDs are best option. So it is C  
upvoted 1 times

  **sanhoo** 1 month, 3 weeks ago

C: All the options use compute engine. A is incorrect because it uses preemptible instances, which only runs for max 24 hours. Hence A is not an option. Then look at the key "fast data access". Local SSD is the fastest  
upvoted 2 times

  **arsh1916** 2 months, 1 week ago

D 100%  
upvoted 1 times

  **Ashii** 2 months, 2 weeks ago

ERP solutions - M1/M2. Its D  
upvoted 1 times


  **norrec9** 3 months ago


D is the answer  
upvoted 1 times


You have developed an application that consists of multiple microservices, with each microservice packaged in its own Docker container image. You want to deploy the entire application on Google Kubernetes Engine so that each microservice can be scaled individually. What should you do?

- A. Create and deploy a Custom Resource Definition per microservice.
- B. Create and deploy a Docker Compose File.
- C. Create and deploy a Job per microservice.
- D. Create and deploy a Deployment per microservice.


Correct Answer: D


- 


 **Kollipara**

Highly Voted 


2 months, 4 weeks ago


D is the correct answer  
upvoted 12 times
- 


 **norrec9**

Highly Voted 


3 months ago


D is the answer  
upvoted 5 times
- 

 **associatecloudexamuser**


Most Recent 


2 weeks, 2 days ago

Yes. D is correct. Can deploy each service through  
kubectl apply -f <deployment\_config.yaml>  
upvoted 1 times
- 

 **pacman\_user**

2 weeks, 3 days ago

D is the correct answer  
upvoted 1 times
- 

 **AD\_0525**



1 month, 1 week ago


D is correct  
upvoted 2 times

You will have several applications running on different Compute Engine instances in the same project. You want to specify at a more granular level the service account each instance uses when calling Google Cloud APIs. What should you do?



- A. When creating the instances, specify a Service Account for each instance.
- B. When creating the instances, assign the name of each Service Account as instance metadata.
- C. After starting the instances, use gcloud compute instances update to specify a Service Account for each instance.
- D. After starting the instances, use gcloud compute instances update to assign the name of the relevant Service Account as instance metadata.


Correct Answer: C

-   **GoCloud**



Highly Voted 


 2 months, 3 weeks ago

A .  
upvoted 12 times
-   **JieHeng**



Highly Voted 



 1 month ago

A, when you create an instance using the gcloud command-line tool or the Google Cloud Console, you can specify which service account the instance uses when calling Google Cloud APIs - [https://cloud.google.com/compute/docs/access/service-accounts#associating\\_a\\_service\\_account\\_to\\_an\\_instance](https://cloud.google.com/compute/docs/access/service-accounts#associating_a_service_account_to_an_instance)  
upvoted 5 times
-   **AD\_0525**

Most Recent 

 1 month, 2 weeks ago

A should be correct  
upvoted 3 times
-   **norrec9** 3 months ago

A is the answer  
upvoted 3 times
-   **Biju1** 3 months ago

It should be A  
upvoted 2 times

You are creating an application that will run on Google Kubernetes Engine. You have identified MongoDB as the most suitable database system for your application and want to deploy a managed MongoDB environment that provides a support SLA. What should you do?

- A. Create a Cloud Bigtable cluster, and use the HBase API.
- B. Deploy MongoDB Atlas from the Google Cloud Marketplace.
- C. Download a MongoDB installation package, and run it on Compute Engine instances.
- D. Download a MongoDB installation package, and run it on a Managed Instance Group.

Correct Answer: C

- arsh1916**

Highly Voted

2 months, 1 week ago

Simple it's B  
upvoted 9 times
- lxgywil**

Highly Voted

2 months, 3 weeks ago

MongoDB Atlas is actually managed and supported by third-party service providers.  
  
<https://console.cloud.google.com/marketplace/details/gc-launcher-for-mongodb-atlas/mongodb-atlas>  
upvoted 5 times
- lxgywil**

2 months, 2 weeks ago

I think that's it. The answer is B  
upvoted 1 times
- AD\_0525**

Most Recent

1 month, 2 weeks ago

Answer should be B  
upvoted 1 times
- GoCloud**

2 months, 3 weeks ago

B .  
upvoted 1 times
- norrec9**

3 months ago

Answer is b  
upvoted 2 times
- Biju1**

3 months ago

it should be B  
upvoted 1 times

You are managing a project for the Business Intelligence (BI) department in your company. A data pipeline ingests data into BigQuery via streaming. You want the users in the BI department to be able to run the custom SQL queries against the latest data in BigQuery. What should you do?

- A. Create a Data Studio dashboard that uses the related BigQuery tables as a source and give the BI team view access to the Data Studio dashboard.
- B. Create a Service Account for the BI team and distribute a new private key to each member of the BI team.
- C. Use Cloud Scheduler to schedule a batch Dataflow job to copy the data from BigQuery to the BI team’s internal data warehouse.
- D. Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.

Correct Answer: C

- ApaMokus

Highly Voted

2 months, 3 weeks ago

D is correct  
roles/bigquery.user  
When applied to a dataset, this role provides the ability to read the dataset's metadata and list tables in the dataset.  
When applied to a project, this role also provides the ability to run jobs, including queries, within the project. A member with this role can enumerate their own jobs, cancel their own jobs, and enumerate datasets within a project. Additionally, allows the creation of new datasets within the project; the creator is granted the BigQuery Data Owner role (roles/bigquery.dataOwner) on these new datasets.  
upvoted 11 times
- AD\_0525

Most Recent

1 month, 2 weeks ago

D should be the answer  
upvoted 3 times
- victory108

2 months, 3 weeks ago

D - Assign the IAM role of BigQuery User to a Google Group that contains the members of the BI team.  
upvoted 3 times
- norrec9

3 months ago

D is the answer  
upvoted 2 times
- Biju1

3 months ago

D is the answer  
upvoted 2 times

## Question #179

## Topic 1

Your company is moving its entire workload to Compute Engine. Some servers should be accessible through the Internet, and other servers should only be accessible over the internal network. All servers need to be able to talk to each other over specific ports and protocols. The current on-premises network relies on a demilitarized zone (DMZ) for the public servers and a Local Area Network (LAN) for the private servers. You need to design the networking infrastructure on Google Cloud to match these requirements. What should you do?

- A. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- B. 1. Create a single VPC with a subnet for the DMZ and a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.
- C. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ.
- D. 1. Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public egress traffic for the DMZ.

**Correct Answer:** A

  **yvinisiupacuando** Highly Voted 2 months, 3 weeks ago

A is the Right answer. You can discard B and C because they lack the need of creating Network Peering to communicate the DMZ VPC with the LAN VPC (LAN VPC is not exposed to public so they need to communicate via private addresses which cannot be achieved with 2 VPCs without Network Peering). Plus, you can discard B, as you don't need to enable the egress traffic, you always need to enable the ingress traffic as this is never enabled by default.

upvoted 10 times

  **Alela** 2 months ago

A is wrong. You don't need to set up firewall rules between subnets of the same VPC. C is the answer

upvoted 6 times

  **gcpengineer** 1 day, 2 hours ago

You need fw rules

upvoted 1 times

  **Ashii** 2 months ago

C is Create a VPC with a subnet for the DMZ and another VPC with a subnet for the LAN. 2. Set up firewall rules to open up relevant traffic between the DMZ and the LAN subnets, and another firewall rule to allow public ingress traffic for the DMZ. Without peering 2 VPC's how this this be done ?

upvoted 1 times

  **associatecloudexamuser** Most Recent 2 weeks, 2 days ago

Yes. Correct answer is A. No need to complicate the setup by creating two different VPC networks.

upvoted 1 times

  **perdigiorno** 2 weeks, 6 days ago

Passed the test today. About 80% of the questions are here.

upvoted 3 times

  **associatecloudexamuser** 2 weeks, 2 days ago

Congratulations!

upvoted 1 times

  **shayanahmed** 3 weeks, 4 days ago

All questions are still valid. I cleared my paper yesterday (shayan18@live.com)

upvoted 2 times

  **associatecloudexamuser** 2 weeks, 2 days ago

Congratulations!

upvoted 1 times

  **JieHeng** 1 month ago

Should be A.

Not B, allow public egress traffic for DMZ won't help anything (also by default there is already this allow public egress traffic rule)

Not C & D, Network Peering is needed to allow internal IP address connectivity across two Virtual Private Cloud (VPC)

<https://cloud.google.com/vpc/docs/vpc-peering>

upvoted 3 times



🗲️ 👤 **AD\_0525** 1 month, 1 week ago

C should be the correct answer. Once you put both set of VMs in a single VPC, either all will be exposed to external traffic or none will be exposed.  
upvoted 1 times

🗲️ 👤 **Tez1** 1 month, 3 weeks ago

as we are looking to emulate the existing onsite environment - having 2 separate networks with firewalls in-between the networks better matches a DMZ and internal network setup.

Therefore my choice is C.

upvoted 1 times

🗲️ 👤 **zsdfaq** 1 month, 3 weeks ago

whats the ans?

upvoted 1 times

🗲️ 👤 **MQQ** 2 months ago

vote for C

upvoted 2 times

🗲️ 👤 **norrec9** 3 months ago

A is the answer

upvoted 2 times

🗲️ 👤 **Biju1** 3 months ago

A is correct

upvoted 2 times